



- Expert Verified, Online, **Free**.

Custom View Settings

Question #48

A company maintains multiple web and mobile applications. Each application uses custom in-house identity providers as well as social identity providers.

You need to implement single sign-on (SSO) for all the applications.

What should you do?

- A. Use Azure Active Directory B2C (Azure AD B2C) with custom policies.
- B. Use Azure Active Directory B2B (Azure AD B2B) and enable external collaboration.
- C. Use Azure Active Directory B2C (Azure AD B2C) with user flows.
- D. Use Azure Active Directory B2B (Azure AD B2B).

Correct Answer: B

You can add Google as an identity provider for B2B guest users.

Federation with SAML/WS-Fed identity providers for guest users.

Make sure your organization's external collaboration settings are configured such that you're allowed to invite guests.

Note 1: As a user who is assigned any of the limited administrator directory roles, you can use the Azure portal to invite B2B collaboration users. You can invite guest users to the directory, to a group, or to an application. After you invite a user through any of these methods, the invited user's account is added to Azure

Active Directory (Azure AD), with a user type of Guest.

Note 2: Direct federation in Azure Active Directory is now referred to as SAML/WS-Fed identity provider (IdP) federation.

Reference:

<https://docs.microsoft.com/en-us/azure/active-directory/external-identities/google-federation> <https://docs.microsoft.com/en-us/azure/active-directory/external-identities/add-users-administrator>

Community vote distribution

A (82%) C (18%)

  **tcybu**  1 year, 7 months ago

The provided answer is wrong!

Based on this article: <https://docs.microsoft.com/en-us/azure/active-directory/external-identities/external-collaboration-settings-configure>
External collaboration settings let you specify what roles in your organization can invite external users for B2B collaboration. These settings also include options for allowing or blocking specific domains, and options for restricting what external guest users can see in your Azure AD directory.

So, you use B2B external collaboration to invite guests into your Azure AD tenant.

I vote for Custom Policies. Both Custom Policies and User Flows support external identity providers, but because of required custom in-house providers support, I'd choose Custom Policies over the User Flows - see <https://docs.microsoft.com/en-us/azure/active-directory-b2c/user-flow-overview#comparing-user-flows-and-custom-policies>

upvoted 15 times

  **ShoaibAnwar**  1 month, 2 weeks ago

Selected Answer: A

Azure AD B2C is designed specifically for customer-facing applications where you need to support a variety of identity providers, including social identities (like Facebook, Google, etc.) and custom (or in-house) identity providers. Custom policies (also known as Identity Experience Framework policies) in Azure AD B2C allow for deep customization of the authentication experience, enabling complex scenarios including integrating with various identity providers, custom user attributes, and multifactor authentication flows. This option provides the flexibility needed to implement SSO while accommodating the diverse identity requirements described.

upvoted 2 times

  **Ciupaz** 5 months, 1 week ago

Selected Answer: A

The best solution for implementing SSO for multiple web and mobile applications with a variety of identity providers is to use Azure AD B2C with custom policies.

upvoted 1 times

  **adilkhan** 1 year ago

A is correct

upvoted 3 times

  **JOSEEVILLASMIL** 1 year ago

Given answer is correct, remember that the question says "company maintains many applications" and "those application have their own in house login providers" so there is a business with other business, is B2B, no B2C.

And B2B with external collaboration says:

"You can also enable self-service sign-up user flows to let external users sign up for apps or resources themselves."

<https://learn.microsoft.com/en-us/azure/active-directory/external-identities/what-is-b2b>

upvoted 3 times

 **adilkhan** 1 year, 2 months ago

Should be A A company maintains multiple web and mobile applications. Each application uses custom in-house identity providers as well as social identity providers.

upvoted 2 times

 **mitu7** 1 year, 3 months ago

Selected Answer: C

"C. Use Azure Active Directory B2C (Azure AD B2C) with user flows."

Azure AD B2C is a cloud-based identity management service that enables you to customize and control how users sign up, sign in, and manage their profiles when interacting with your applications. With Azure AD B2C, you can create user flows that define the steps in the authentication and authorization process for your applications, and you can use these user flows to enable SSO across your applications.

User flows allow you to customize the authentication experience for your users, including specifying the identity providers that you want to use, such as custom in-house identity providers or social identity providers. You can also customize the user interface, branding, and policies that apply to each user flow.

References: <https://learn.microsoft.com/en-us/azure/active-directory-b2c/user-flow-overview>

upvoted 3 times

 **guchao2000** 1 year, 2 months ago

It's A

The message is "... uses custom in-house identity providers", this requires Azure AD B2C with custom policies.

upvoted 1 times

 **Yumi21** 1 year, 5 months ago

For Azure AD B2C check this: <https://learn.microsoft.com/en-us/azure/active-directory-b2c/overview>

upvoted 1 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: A

A - correct answer

upvoted 4 times

 **RochaG2** 1 year, 6 months ago

Selected Answer: A

A for sure

upvoted 4 times

 **finnishr** 1 year, 7 months ago

Selected Answer: A

A is correct 100%

upvoted 3 times

 **Knightie** 1 year, 7 months ago

The Answer seems Correct. B2C seems for external totally, so no mention of the existing AAD authentication to the app.

<https://docs.microsoft.com/en-us/azure/active-directory/external-identities/external-identities-overview?bc=%2Fazure%2Factive-directory-b2c%2Fbread%2Ftoc.json&toc=%2Fazure%2Factive-directory-b2c%2FTOC.json>

upvoted 1 times

 **jeanfmc** 1 year, 7 months ago

A

<https://docs.microsoft.com/en-us/azure/active-directory-b2c/custom-policy-reference-sso>

upvoted 4 times

Question #49

You develop a Python application for image rendering that uses GPU resources to optimize rendering processes. You deploy the application to an Azure Container Instances (ACI) Linux container. The application requires a secret value to be passed when the container is started. The value must only be accessed from within the container. You need to pass the secret value. What are two possible ways to achieve this goal? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. Create an environment variable Set the `secureValue` property to the secret value.
- B. Add the secret value to the container image. Use a managed identity.
- C. Add the secret value to the application code Set the container startup command.
- D. Add the secret value to an Azure Blob storage account. Generate a SAS token.
- E. Mount a secret volume containing the secret value in a secrets file.

Correct Answer: AE

A: Secure environment variables -

Another method (another than a secret volume) for providing sensitive information to containers (including Windows containers) is through the use of secure environment variables.

E: Use a secret volume to supply sensitive information to the containers in a container group. The secret volume stores your secrets in files within the volume, accessible by the containers in the container group. By storing secrets in a secret volume, you can avoid adding sensitive data like SSH keys or database credentials to your application code.

Reference:

<https://docs.microsoft.com/en-us/azure/container-instances/container-instances-volume-secret>

Community vote distribution

AE (100%)

✉  **osagiestar**  1 year, 6 months ago

Selected Answer: AE

A and E

upvoted 9 times

✉  **TonyMel** 1 year ago

correct, in 2023Mar24, score: 904/1000.

upvoted 2 times

✉  **elequiel**  1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 7 times

✉  **BaoNguyen2411**  8 months, 3 weeks ago

got this question on 06-29-2023

upvoted 4 times

✉  **lorenaizzo** 1 year, 7 months ago

Selected Answer: AE

Secure values

Objects with secure values are intended to hold sensitive information like passwords or keys for your application. Using secure values for environment variables is both safer and more flexible than including it in your container's image. Another option is to use secret volumes, described in Mount a secret volume in Azure Container Instances.....

<https://docs.microsoft.com/en-us/azure/container-instances/container-instances-environment-variables>

upvoted 6 times

Question #50

You are developing a user portal for a company.

You need to create a report for the portal that lists information about employees who are subject matter experts for a specific topic. You must ensure that administrators have full control and consent over the data.

Which technology should you use?

- A. Microsoft Graph data connect
- B. Microsoft Graph API
- C. Microsoft Graph connectors

Correct Answer: A

Data Connect grants a more granular control and consent model: you can manage data, see who is accessing it, and request specific properties of an entity. This enhances the Microsoft Graph model, which grants or denies applications access to entire entities.

Microsoft Graph Data Connect augments Microsoft Graph's transactional model with an intelligent way to access rich data at scale. The data covers how workers communicate, collaborate, and manage their time across all the applications and services in Microsoft 365.

Incorrect:

Not B: The Microsoft Graph API is a RESTful web API that enables you to access Microsoft Cloud service resources. After you register your app and get authentication tokens for a user or service, you can make requests to the Microsoft Graph API.

A simplistic definition of a Graph API is an API that models the data in terms of nodes and edges (objects and relationships) and allows the client to interact with multiple nodes in a single request.

Not C: Microsoft Graph connectors, your organization can index third-party data so that it appears in Microsoft Search results.

With Microsoft Graph connectors, your organization can index third-party data so that it appears in Microsoft Search results.

Reference:

<https://docs.microsoft.com/en-us/graph/data-connect-concept-overview>

Community vote distribution

A (69%)

B (31%)

✉  **nvtienanh**  1 year, 6 months ago

Selected Answer: A

<https://learn.microsoft.com/en-us/graph/overview#access-microsoft-graph-data-at-scale>: text=Find%20you%20the%20closest%20expert%20on%20a%20topic%20to%20you%20in%20your%20organization

upvoted 14 times

✉  **coffecold**  1 year, 5 months ago

Big Microsoft Brother is watching you with some kind of shady tool : 'Data Connect'

upvoted 7 times

✉  **FeriAZ**  2 months ago

Selected Answer: A

Microsoft Graph data connect is a tool that allows large-scale data extraction from the Microsoft Graph, providing a secure method to access Microsoft 365 data at scale. It is designed to work with Azure Data Factory, enabling administrators to create pipelines that transfer Microsoft 365 data into Azure storage for analysis and reporting. Data connect supports granular consent and governance controls, allowing administrators to have full control over the data, including what data is accessed and how it is used. This makes it suitable for scenarios where large volumes of data need to be analyzed or processed while ensuring that administrative consent and control are maintained.

upvoted 2 times

✉  **LJNZ2049** 2 months, 3 weeks ago

Selected Answer: B

I think the question is talking more about access control rather than data processing. Prefer B would be the right answer

upvoted 2 times

✉  **FeriAZ** 3 months ago

Selected Answer: A

A. Microsoft Graph data connect

Microsoft Graph data connect is designed to handle complex data processing tasks while providing strong data governance and administrator control. It allows administrators to consent to exactly what data is shared and ensures that large-scale data transfers are handled securely and efficiently.

It is more suitable for scenarios where bulk data processing is required, and administrator oversight on data access is a priority, aligning well with your need for creating detailed reports while maintaining strict control over data access and use.

upvoted 1 times

✉ **CarlosTheBoldest** 4 months, 1 week ago

I got this question on my exam, 2023Dec, go with what I remember was the most voted answer. Score 902, most of the questions were here, slightly different on wording because the Azure Ad <-> Entra Id change. Case was City Power & Light. Good luck!
Important tip, you have access to microsoft learn during the exam!

upvoted 5 times

✉ **Ciupaz** 5 months, 1 week ago

Selected Answer: B

The Microsoft Graph API is the best technology to use for creating a report that lists information about employees who are subject matter experts for a specific topic. It provides a comprehensive set of APIs that allow you to retrieve, create, update, and delete data across Microsoft 365 services. Additionally, the Microsoft Graph API provides a number of security features that can be used to ensure that administrators have full control and consent over the data.

upvoted 3 times

✉ **Morgist** 7 months ago

Selected Answer: A

Microsoft Graph Data Connect provides access to many users or groups so it meets the requirements

upvoted 1 times

✉ **oreakinodidi** 9 months ago

Selected Answer: B

The Microsoft Graph API provides a comprehensive set of REST APIs that allows you to access and manipulate a wide range of data and resources within Microsoft 365 and Azure Active Directory. It offers powerful querying capabilities and supports various operations to retrieve and filter employee information based on specific criteria such as topic expertise.

While Microsoft Graph data connect and Microsoft Graph connectors also have their use cases, they are not directly applicable to the requirement of creating a report for the user portal. Microsoft Graph data connect is focused on data integration and export scenarios, while Microsoft Graph connectors are used for connecting external data sources to Microsoft Search.

Therefore, for the given scenario, the Microsoft Graph API is the most suitable option to retrieve and present the required employee information in the user portal report.

upvoted 1 times

✉ **DarthVee** 11 months, 1 week ago

Selected Answer: B

In summary, Microsoft Graph API provides the necessary security mechanisms to ensure that administrators have full control and consent over the data accessed by your application.

upvoted 3 times

✉ **karthikwarrior** 1 year, 1 month ago

The people who are commenting the right ans have paid subscription (Contributor role) as bec i am one of them and you can notice, comments are reduced from others who doesnt have access to this role. they cant access this content..

upvoted 4 times

✉ **macobuzi** 7 months, 4 weeks ago

Which is sad since we might not have enough information to decide the right answer!

upvoted 1 times

✉ **lorenaizzo** 1 year, 7 months ago

Selected Answer: A

<https://docs.microsoft.com/en-us/graph/overview>

upvoted 2 times

Question #51

HOTSPOT -

You are a developer building a web site using a web app. The web site stores configuration data in Azure App Configuration.

Access to Azure App Configuration has been configured to use the identity of the web app for authentication. Security requirements specify that no other authentication systems must be used.

You need to load configuration data from Azure App Configuration.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
Host .CreateDefaultBuilder(args)
    .ConfigureWebHostDefaults(wb =>
{
    wb.ConfigureAppConfiguration((hc, config) =>
    {
        var settings = config.Build();
        config. (options =>
{
            AddAzureKeyVault
            DefaultAzureCredential
            ChainedTokenCredential
            ManagedIdentityCredential
            AddAzureAppConfiguration
        }

        options.Connect(new Uri(settings["AppConfig:Endpoint"]),
            new (()));
            AddAzureKeyVault
            DefaultAzureCredential
            ChainedTokenCredential
            ManagedIdentityCredential
            AddAzureAppConfiguration
        });

    });
})
```

Answer Area

```

Host .CreateDefaultBuilder(args)
    .ConfigureWebHostDefaults(wb =>
    {

        wb.ConfigureAppConfiguration((hc, config) =>
        {
            var settings = config.Build();
            config. (options =>
AddAzureKeyVault
DefaultAzureCredential
ChainedTokenCredential
ManagedIdentityCredential
AddAzureAppConfiguration

            options.Connect(new Uri(settings["AppConfig:Endpoint"]),
                new (());
AddAzureKeyVault
DefaultAzureCredential
ChainedTokenCredential
ManagedIdentityCredential
AddAzureAppConfiguration

        });
    })
}

```

Correct Answer:

Box 1: AddAzureAppConfiguration -

Load data from App Configuration, code example:

```

public static IHostBuilder CreateHostBuilder(string[] args) =>
    Host.CreateDefaultBuilder(args)
        .ConfigureWebHostDefaults(webBuilder =>
            webBuilder.ConfigureAppConfiguration((hostingContext, config) =>
            {
                var settings = config.Build();
                config.AddAzureAppConfiguration(options =>
                {

```

Etc.

Box 2: ManagedIdentityCredential

Use managed identities to access App Configuration

If you want to use a user-assigned managed identity, be sure to specify the clientId when creating the ManagedIdentityCredential.

```

config.AddAzureAppConfiguration(options =>
{
    options.Connect(new Uri(settings["AppConfig:Endpoint"]), new ManagedIdentityCredential("<your_clientId>"))
});

```

Full code sample:

```

public static IHostBuilder CreateHostBuilder(string[] args) =>
    Host.CreateDefaultBuilder(args)
        .ConfigureWebHostDefaults(webBuilder =>
            webBuilder.ConfigureAppConfiguration((hostingContext, config) =>
            {
                var settings = config.Build();
                config.AddAzureAppConfiguration(options =>
                    options.Connect(new Uri(settings["AppConfig:Endpoint"]), new ManagedIdentityCredential());
                )
            .UseStartup<Startup>());

```

Reference:

<https://docs.microsoft.com/en-us/azure/azure-app-configuration/howto-integrate-azure-managed-service-identity?tabs=core5x&pivots=framework-dotnet>

✉  **coffecold**  1 year, 6 months ago

Answer is right.

"Security requirements specify that no other authentication systems must be used"
So you cannot use DefaultAzureCredential, risk of other authentication systems being used.
otherwise you could have used DefaultAzureCredential for ease

upvoted 25 times

✉  **serpevi**  1 year, 7 months ago

Got this in 09/22 , went with the answers given, score 927.

upvoted 15 times

✉  **SSR999**  3 months, 1 week ago

<https://learn.microsoft.com/en-us/azure/azure-app-configuration/howto-integrate-azure-managed-service-identity?tabs=core6x&pivots=framework-dotnet#use-a-managed-identity>

upvoted 1 times

✉  **junkz** 5 months, 2 weeks ago

second answer is wrong -><https://learn.microsoft.com/en-us/answers/questions/987373/azure-app-configuration-label>

upvoted 1 times

✉  **Weam** 5 months ago

However in the question requirement that there should be no other authentication methods tried so we should use ManagedIdentityCredential in specific since we are using web app

upvoted 1 times

✉  **OPT_001122** 1 year, 4 months ago

Ans is correct

AddAzureAppConfiguration
ManagedIdentityCredential

upvoted 6 times

Question #52

HOTSPOT -

You are building an application that stores sensitive customer data in Azure Blob storage. The data must be encrypted with a key that is unique for each customer.

If the encryption key has been corrupted it must not be used for encryption.

You need to ensure that the blob is encrypted.

How should you complete the code segment? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
Uri blobUri = ... ; TokenCredential c = ...
byte[] key = ...; string verify = ...
var x = new
    AesManaged(key)
    AsnEncodedData(key)
    CustomerProvidedKey(key)
    BlobContainerEncryptionScopeOptions { DefaultEncryptionScope = key }

if ( ) {
    x.IV == verify
    x.RawData == verify
    x.Encryption == verify
    x.PreventEncryptionScopeOverride == verify
}

var o = new BlobClientOptions()
{
    = x
    Version
    Transport
    EncryptionScope
    CustomerProvidedKey
};

var blobClient = new BlobClient(blobUri, c, o);
}
```

Answer Area

```
Uri blobUri = ... ; TokenCredential c = ...
byte[] key = ...; string verify = ...
var x = new
    AesManaged(key)
    AsnEncodedData(key)
    CustomerProvidedKey(key) CustomerProvidedKey(key)
    BlobContainerEncryptionScopeOptions { DefaultEncryptionScope = key }

if ( ) {
    x.IV == verify
    x.RawData == verify
    x.Encryption == verify x.Encryption == verify
    x.PreventEncryptionScopeOverride == verify
}

var o = new BlobClientOptions()
{
    = x
    Version
    Transport
    EncryptionScope
    CustomerProvidedKey CustomerProvidedKey
};

var blobClient = new BlobClient(blobUri, c, o);
}
```

Correct Answer:

Box 1: CustomerProvidedKey(key)

The data must be encrypted with a key that is unique for each customer.

Sample code:

```
async static Task UploadBlobWithClientKey(Uri blobUri,
```

```

Stream data,
byte[] key,
string keySha256)
{
// Create a new customer-provided key.
// Key must be AES-256.
var cpk = new CustomerProvidedKey(key);

```

Box 2: Encryption -

CustomerProvidedKey.EncryptionKey Property

Sample code continued:

```

// Check the key's encryption hash.
if (cpk.EncryptionKeyHash != keySha256)
{
throw new InvalidOperationException("The encryption key is corrupted.");
}

```

Box 3: CustomerProvidedKey -

Sample code continued:

// Specify the customer-provided key on the options for the client.

BlobClientOptions options = new BlobClientOptions()

{

CustomerProvidedKey = cpk -

};

// Create the client object with options specified.

BlobClient blobClient = new BlobClient(

blobUri,

new DefaultAzureCredential(),

options);

Incorrect:

* Version - Gets the BlobClientOptions.ServiceVersion of the service API used when making requests.

Transport - The HttpPipelineTransport to be used for this client.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-customer-provided-key>

✉  **Praks13**  1 year, 7 months ago

This looks correct!

<https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-customer-provided-key>

upvoted 16 times

✉  **gmishra88**  1 year, 6 months ago

Shouldn't it be EncryptionKeyHash based on the link provided?

upvoted 11 times

✉  **Ciupaz**  5 months, 1 week ago

Where are talking about costumer, so:

1) CustomerProvidedKey(key)

We have to check the encryption, so:

2) x.Encryption == verify

Always about customer, so:

3) CustomerProvidedKey

upvoted 1 times

✉  **NPE_** 5 months, 3 weeks ago

AZ204 should be renamed as Azure .NET/c# Developer

upvoted 3 times

✉  **NightshadeRC** 8 months, 3 weeks ago

Had this question in today's exam: 2023-07-26

Also the exam did have the value as EncryptedKeyHash

upvoted 7 times

✉  **Elbanna** 6 months, 3 weeks ago

Does All exam questions from here?

upvoted 1 times

✉  **willchenxa** 1 year, 7 months ago

The second box probably should be EncryptionKey, CustomerProvidedKey doesn't even has Encryption property, <https://docs.microsoft.com/en-us/dotnet/api/azure.storage.blobs.models.customerprovidedkey?view=azure-dotnet>

upvoted 3 times

✉  **coffecold** 1 year, 6 months ago

Must be either EncryptionKey or EncryptionKeyHash. But why verify the EncryptionKey while it is given ? I would suggest the hash.

upvoted 2 times

✉  **coffecold** 1 year, 6 months ago

Base64 encoded string of the AES256 encryption key's SHA256 hash.

upvoted 1 times

Question #53

HOTSPOT

You are developing a web application that uses the Microsoft Identity platform for user and resource authentication. The web application called several REST APIs.

You are implementing various authentication and authorization flows for the web application.

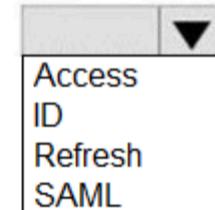
You need to validate the claims in the authentication token.

Which token type should you use? To answer, select the appropriate options in the answer area.

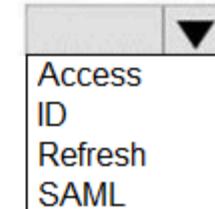
NOTE: Each correct selection is worth one point.

Answer Area**Requirement****Token type**

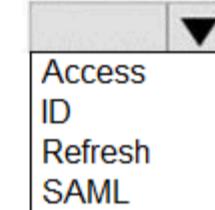
Identify users for the application by using a JWT token that contains claims.



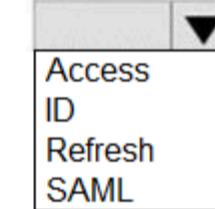
Identify the permissions granted to APIs by using a JWT token that contains claims.



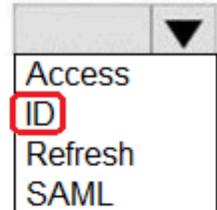
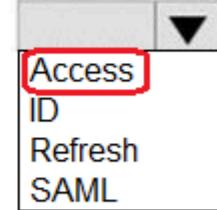
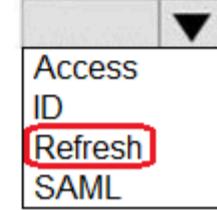
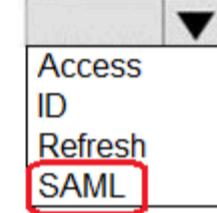
Provide the web application with long-term access to resources on behalf of users without requiring interaction with those users.



Provide XML representations of claims that can be consumed by applications that use WS-Federation.



Answer Area

	Requirement	Token type
	Identify users for the application by using a JWT token that contains claims.	
Correct Answer:	Identify the permissions granted to APIs by using a JWT token that contains claims.	
	Provide the web application with long-term access to resources on behalf of users without requiring interaction with those users.	
	Provide XML representations of claims that can be consumed by applications that use WS-Federation.	

✉  **Ciupaz**  5 months, 1 week ago

How to remember:

- 1) Identify users -> ID (ID of the user)
- 2) Permissions -> Access (how to access? with permissions)
- 3) Without interactions -> Refresh (with refresh, there is no interaction)
- 4) Provide XML -> SAML (similar in writing)

upvoted 22 times

✉  **Rabi3** 1 month, 4 weeks ago

Thanks

upvoted 1 times

✉  **JH81**  9 months, 3 weeks ago

Got this on 6/28/2023 and passed with 850. Answer is correct.

upvoted 5 times

✉  **SachinV**  5 months ago

Answer is correct, got this one in Nov182023, scored 962

upvoted 3 times

✉  **warchoon** 1 year, 1 month ago

looks correct

<https://learn.microsoft.com/en-us/azure/active-directory/develop/active-directory-configurable-token-lifetimes#token-lifetime-policies-for-access-saml-and-id-tokens>

<https://learn.microsoft.com/en-us/azure/active-directory/develop/refresh-tokens#refresh-token-lifetime>

upvoted 5 times

✉  **NombreFalso** 1 year, 2 months ago

Suggested answer makes sense

upvoted 2 times

✉  **mabdo** 1 year, 2 months ago

was on exam 02/23

upvoted 2 times

✉  **g2000** 1 year, 3 months ago

correct

<https://learn.microsoft.com/en-us/azure/active-directory/develop/access-tokens>

upvoted 2 times

Question #54

Topic 4

HOTSPOT

You are developing a content management application for technical manuals. The application is deployed as an Azure Static Web app.

Authenticated users can view pages under/manuals but only contributors can access the page /manuals/new.html.

You need to configure the routing for the web app.

How should you complete the configuration? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
"routes": [
  {
    "route": "/manuals*",
    "allowedRoles": [
      "contributors",
      "authenticated"
    ],
    "allowedAuthentications": [
      "aad"
    ]
  },
  {
    "route": "/manuals/new.html",
    "allowedRoles": [
      "contributors"
    ],
    "allowedAuthentications": [
      "aad"
    ]
  }
]
```

Answer Area

```

"routes": [
  {
    "route": "/manuals*",
    "methods": [
      "contributors",
      "authenticated"
    ],
    "allowedRoles": [
      "/manuals*",
      "contributors"
    ],
    "allowedAuthLevel": [
      "/manuals*",
      "contributors",
      "authenticated"
    ],
    "onSuccess": [
      "/manuals/new.html"
    ]
  }
]
  
```

Correct Answer:

```

"routes": [
  {
    "route": "/manuals*",
    "methods": [
      "contributors",
      "authenticated"
    ],
    "allowedRoles": [
      "/manuals*",
      "contributors",
      "authenticated"
    ],
    "onSuccess": [
      "/manuals/new.html"
    ]
  }
]
  
```

2fa  1 year, 3 months ago

Correct.

"Rule evaluation stops at the first match. A match occurs when the route property and a value in the methods array (if specified) match the request. Each request can match at most one rule." So the first rule must be for the /manuals/new.html page.

<https://learn.microsoft.com/en-us/azure/static-web-apps/configuration#routes>

upvoted 27 times

juancar  9 months, 2 weeks ago

This was on the exam (July 2023). Went with proposed. Scored 917

upvoted 6 times

unravel  7 months ago

Rule evaluation stops at the first match. If authenticated came first all users including contributors would have access to /manuals/new.html page which goes against the requirement

upvoted 2 times

rockstar89 8 months, 3 weeks ago

got it exam 31/07/2023. correct

upvoted 2 times

plusJoyed 11 months ago

couldn't it just have easily have been manuals/authenticated then html/contributors?

upvoted 4 times

CarlosTheBoldest 4 months, 1 week ago

Nope, cause it takes the first rule

upvoted 1 times

Question #55

Topic 4

You are developing a web application that uses the Microsoft identity platform for user and resource authentication. The web application calls several REST APIs.

A REST API call must read the user's calendar. The web application requires permission to send an email as the user.

You need to authorize the web application and the API.

Which parameter should you use?

- A. tenant
- B. code_challenge
- C. state
- D. client_id
- E. scope

Correct Answer: E

Community vote distribution

E (100%)

✉  **alexein74**  1 year, 2 months ago

Selected Answer: E

E. scope

Scope is the parameter used to specify the permissions that the web application is requesting to access the user's resources, such as the calendar and the ability to send an email as the user. The scope parameter is typically included in the API call or in the authorization request sent to the Microsoft identity platform.

upvoted 14 times

✉  **g2000**  1 year, 3 months ago

right

<https://learn.microsoft.com/en-us/azure/active-directory/develop/scopes-oidc>

upvoted 9 times

✉  **FeriAZ**  3 months ago

Selected Answer: E

scope

The scope parameter in OAuth 2.0 specifies the permissions that the application is requesting. These permissions can be related to user delegated permissions (like reading a user's calendar) or application permissions (like sending an email on behalf of a user).

Relevance to Authorization: This is the key parameter for defining the specific permissions your web application requires to access Microsoft Graph or other APIs.

upvoted 1 times

✉  **NightshadeRC** 8 months, 3 weeks ago

Had this question in today's exam: 2023-07-26

upvoted 4 times

Question #56

HOTSPOT

You develop and deploy a web app to Azure App service. The web app allows users to authenticate by using social identity providers through the Azure B2C service. All user profile information is stored in Azure B2C.

You must update the web app to display common user properties from Azure B2C to include the following information:

- Email address
- Job title
- First name
- Last name
- Office location

You need to implement the user properties in the web app.

Which code library and API should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Requirement	Value
API to access user properties	<input type="checkbox"/> Microsoft Graph <input type="checkbox"/> Azure AD Graph <input type="checkbox"/> Azure Key Vault <input type="checkbox"/> Azure AD entitlement management
Code library to interface to Azure AD B2C	<input type="checkbox"/> Microsoft Authentication Library (MSAL) <input type="checkbox"/> Microsoft Azure Key Vault SDK <input type="checkbox"/> Azure Identity library

Requirement	Value
API to access user properties	<input checked="" type="checkbox"/> Microsoft Graph <input type="checkbox"/> Azure AD Graph <input type="checkbox"/> Azure Key Vault <input type="checkbox"/> Azure AD entitlement management
Code library to interface to Azure AD B2C	<input checked="" type="checkbox"/> Microsoft Authentication Library (MSAL) <input type="checkbox"/> Microsoft Azure Key Vault SDK <input type="checkbox"/> Azure Identity library

 **g2000**  1 year, 3 months ago

right..

<https://learn.microsoft.com/en-us/azure/active-directory-b2c/microsoft-graph-operations>

<https://learn.microsoft.com/en-us/azure/active-directory/develop/msal-net-aad-b2c-considerations>

upvoted 9 times

 **JH81**  9 months, 3 weeks ago

Got this on 6/28/2023 and passed with 850. Graph, MSL is correct.

upvoted 5 times

 **rockstar89**  8 months, 3 weeks ago

Received on 31/07/2023

upvoted 2 times

 **uncledana** 10 months ago

Correct

- Microsoft Graph
- MSAL

upvoted 1 times

 **Vmwarevirtual** 10 months, 3 weeks ago

Received on 27 May 2023

- Microsoft Graph
- MSAL

upvoted 3 times

Question #57

HOTSPOT

You develop and deploy the following staticwebapp.config.json file to the app_location value specified in the workflow file of an Azure Static Web app:

```
{  
  "routes": [  
    {  
      "route": "/api/*",  
      "methods": ["GET"],  
      "allowedRoles": ["registeredusers"]  
    },  
    {  
      "route": "/api/*",  
      "methods": ["PUT", "POST", "PATCH", "DELETE"],  
      "allowedRoles": ["administrator"]  
    },  
    {  
      "route": "/login",  
      "rewrite": "./.auth/login/github"  
    },  
    {  
      "route": "./.auth/login/twitter",  
      "statusCode": 404  
    },  
    {  
      "route": "/logout",  
      "redirect": "./.auth/logout"  
    }  
,  
  "navigationFallback": {  
    "rewrite": "index.html",  
    "exclude": ["/images/*.{png,jpg,gif}", "/css/*"]  
  },  
  "responseOverrides": {  
    "400": {  
      "rewrite": "/invalid-invitation-error.html"  
    },  
    "401": {  
      "redirect": "./.auth/login/aad",  
      "statusCode": 302  
    },  
    "403": {  
      "rewrite": "/forbidden.html"  
    },  
    "404": {  
      "rewrite": "/404.html"  
    }  
,  
  "mimeType": {  
    ".json": "text/json"  
  }  
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements

	Yes	No
Unauthenticated users are challenged to authenticate with GitHub.	<input type="radio"/>	<input type="radio"/>
A non-existent file in the <code>/images/</code> folder will generate a 404 response code.	<input type="radio"/>	<input type="radio"/>
HTTP GET method requests from authenticated users in the role named registeredusers are sent to the API folder.	<input type="radio"/>	<input type="radio"/>
Authenticated users that are not in the role named registeredusers and unauthenticated users are served a 401 HTTP error when accessing the API folder.	<input type="radio"/>	<input type="radio"/>

Answer Area		
Statements	Yes	No
Unauthenticated users are challenged to authenticate with GitHub.	<input type="radio"/>	<input checked="" type="checkbox"/>
A non-existent file in the <code>/images/</code> folder will generate a 404 response code.	<input type="radio"/>	<input checked="" type="checkbox"/>
HTTP GET method requests from authenticated users in the role named registeredusers are sent to the API folder.	<input checked="" type="checkbox"/>	<input type="radio"/>
Authenticated users that are not in the role named registeredusers and unauthenticated users are served a 401 HTTP error when accessing the API folder.	<input type="radio"/>	<input checked="" type="checkbox"/>

✉  **mabdo** Highly Voted  1 year, 2 months ago

1) Y. login with github
 2) Y. true because it is excluded from returning index.html in navigation fallback, so the other option is a 404 not found
 3) Y. straight forward from first routes object
 4) N. from <https://learn.microsoft.com/en-us/azure/static-web-apps/configuration>
 401 is for requesting when unauthenticated but here you are authenticated and so the error should be 403, because you are authenticated but don't have authorization (due to not having the registeredusers role)
 upvoted 15 times

✉  **oCaio** 8 months, 1 week ago

Initially I thought YNYN (since only 3 types of images are excluded in `/images/`) because:

<https://learn.microsoft.com/en-us/azure/static-web-apps/configuration>

"In the following example, requests for certain routes in the `/images` folder and all files in the `/css` folder are excluded from returning the fallback file."

JSON

```
{
  "navigationFallback": {
    "rewrite": "/index.html",
    "exclude": ["/images/*.{png,jpg,gif}", "/css/*"]
  }
}
```

However I changed to YYYN, because of

<https://learn.microsoft.com/en-us/azure/static-web-apps/configuration#example-configuration-file>

So this depends on which Microsoft article they are trying to test us.

upvoted 2 times

✉  **SachinV** 5 months ago

Agreed looks correct, in Nov182023, scored 962

upvoted 2 times

✉  **aboezebal** 1 year, 1 month ago

1 should be no.

Unauthenticated users get a 401 - Unauthorized response and are redirected by the responseOverrides to `"/.auth/login/aad"`. I believe that is not GitHub. You will only be redirected to GitHub if the an unauthenticated user routes to `"/login"`.

upvoted 26 times

✉  **surprise0011** 1 year ago

received 2023-04-17 went with N,Y,Y,N , score 926

i have doubts about second one - I don't know that status code will persist or not. We set new page 404.html without status code like in example with 401 where status code is set to 302. Needs to be verified, good luck

upvoted 12 times

✉  **adilkhan** 12 months ago

did you get most questions from exam topics?
upvoted 1 times

✉  **drkpnsn**  8 months ago

It should be:

1. N

In the example written here: <https://learn.microsoft.com/en-us/azure/static-web-apps/configuration#example-configuration-file>

- If you go to the login page, you redirect to GitHub, just like in the question.

- If the request got a 401 error (Unauthorized), it redirects the user to the login route (with 302 response), which then redirects to the GitHub login.

But in this question, if you check what happens when the request got a 401 error (Unauthorized), it redirects to the AAD route (with 302 response), not to GitHub. So this part of the question and the example are different.

2. Y

That part of the question seems the same as the example above.

A non-existent file in the /images/ folder -> A 404 error.

3. Y

That part of the question seems the same as the example above.

GET requests from authenticated users in the registeredusers role are sent to the API. Authenticated users not in the registeredusers role and unauthenticated users are served a 401 error.

4. N

It overrides the 401 error to the 302 and redirects the user to the AAD URL.

upvoted 12 times

✉  **Emil_Topics**  4 months, 3 weeks ago

I don't get why everybody states that the third option should be yes. How I read the statement is that users in that role automatically get redirected to the API path. But that is not what the configuration states. Only that users in that role have access to it. Not that they are redirected.

upvoted 2 times

✉  **oskx2** 1 month ago

Agreed. The users are NOT sent to the API folder. They are ALLOWED to the api folder.

upvoted 1 times

✉  **JH81** 9 months, 3 weeks ago

Got this on 6/28/2023 and passed with 850. I went with Y,N,Y,N. No idea if that is correct.

upvoted 1 times

✉  **BikashKatwal** 10 months, 1 week ago

<https://learn.microsoft.com/en-us/azure/static-web-apps/configuration#example-configuration-file>

I think based on the above config

1. Y - Unauthenticated users are challenged to authenticate with GitHub.

2. Y: A non-existent file in the /images/ folder results in a 404 error

3. Y: GET requests from authenticated users in the registeredusers role are sent to the API.

4. N: Authenticated users not in the administrator role and unauthenticated users are served a 401 error.

upvoted 2 times

✉  **nekkilodeon** 12 months ago

N,Y,Y,N

<https://learn.microsoft.com/en-us/azure/static-web-apps/configuration#example-configuration-file>

upvoted 6 times

✉  **AndySmith** 5 months, 3 weeks ago

Almost identical example from Microsoft with good explanation. No doubts about N Y Y N.

upvoted 1 times

✉  **adilkhan** 1 year ago

N,Y,Y,N IS CORRECT

upvoted 4 times

✉  **DP_Bhatt** 1 year ago

No - Unauthorized users receive a 401 response code, which subsequently redirects to "/.auth/login/aad" with a status code 302 (see responseOverrides). You will only be redirected to GitHub if the unauthenticated user routes to "/login".

Yes - The files in the path "/images/*" with extensions png, jpg and gif are excluded from the fallback which means if the files with the mentioned extensions do not exist in the path, an 404 will be returned. The index.html fallback will be returned only if the file has an extension other than the excluded extensions. For example, if the file with name /images/xyz.svg will return index.html because the .svg extension is not excluded. But if the file is /images/unknown.jpg and the file is not exist on the path will return 404.

Yes -

No - In this case, if the user does not have registeredusers role, the 403 status code will be returned for the authenticated users.

upvoted 5 times

✉  **DP_Bhatt** 1 year ago

No - Unauthorized users receive a 401 response code, which subsequently redirects to "./.auth/login/aad" with a status code 302 (see responseOverrides). You will only be redirected to GitHub if the unauthenticated user routes to "/login". Yes - The files in the path "/images/*" with extensions png, jpg and gif are excluded from the fallback which means if the files with the mentioned extensions do not exist in the path, an 404 will be returned. The index.html fallback will be returned only if the file has an extension other than the excluded extensions. For example, if the file with name /images/xyz.svg will return index.html because the .svg extension is not excluded. But if the file is /images/unknown.jpg and the file is not exist on the path will return 404.

Yes -

No - In this case, if the user does not have registeredusers role, the 403 status code will be returned for the authenticated users.

upvoted 2 times

 **Kluk_Kluk** 1 year, 1 month ago

I saw the same configuration in this link. All answers are Yes.

<https://learn.microsoft.com/en-us/azure/static-web-apps/configuration>

upvoted 2 times

 **courtguard** 1 year ago

the last one is No, as per your link.

If they are not in role registeredusers they are served 403 error, and if they are unauthenticated they are redirected to /login

upvoted 1 times

 **erieri** 1 year, 1 month ago

Haha, really different opinions about this one... really tricky, but I will give it a try:

1) N . Unauthenticated = 401 which is redirected to .../login/aad (with 302 status). See aboezebal's comment.

2) Y. See rpatsev's comment below. Always returns 404 even though some are rewritten with index.html

3) Y.

4) N. Authenticated users but not in correct role = 403. Unauthenticated user, see question.

upvoted 1 times

 **lixd** 1 year, 1 month ago

Not sure 100% the answers, but all the cases are described and covered in <https://learn.microsoft.com/en-us/azure/static-web-apps/configuration#example-configuration-file>

So, the answers with comments from the web-site

1. (Y) /login : Unauthenticated users are challenged to authenticate with GitHub.

2. (Y) An non-existent file in the /images/ folder : A 404 error.

3. (Y) /api/admin : GET requests from authenticated users in the registeredusers role are sent to the API.

4. (Y) /api/admin : POST, PUT, PATCH, and DELETE requests from authenticated users in the administrator role are sent to the API. Authenticated users not in the administrator role and unauthenticated users are served a 401 error.

upvoted 2 times

 **adilkhan** 1 year, 1 month ago

Y Y Y N is the correct answer

upvoted 2 times

 **baroo1** 1 year, 2 months ago

The option 2 is tricky.

If file does not exist in /images/. Considering below files does not exist:

- /images/Image.png > /images/*.png > Excluded > 404.

- /images/Image.csv > /images/*.csv > Included > Redirect to index.html.

As I can easily find at least 1 scenario where it does not return a 404 > Answer is NO!

upvoted 2 times

 **rpatsev** 1 year, 1 month ago

- /images/Image.csv > /images/*.csv > Included > Rewrite (not Redirect) to index.html.

"navigationFallback" mode is "rewrite" (not "redirect") it means the body will be rewritten, but the statusCode will still be 404.

upvoted 3 times

 **mabdo** 1 year, 2 months ago

was on exam 02/23

upvoted 1 times

 **sbutettino** 1 year, 1 month ago

can u tell us which answers are right?

upvoted 1 times

 **Osip21** 1 year, 2 months ago

Y, N, Y, Y

upvoted 2 times

 **TaoLu** 1 year, 2 months ago

Y

N

N

Y

upvoted 1 times

 **eugene1234** 1 year, 2 months ago

1. there is no redirect to /login route in response overrides, 4. 401 is redirected to aad, so no 302 response code; IMHO the answer to this specific JSON is all NO

upvoted 1 times

Question #58

Topic 4

You develop and deploy an Azure App Service web app named App1. You create a new Azure Key Vault named Vault1. You import several API keys, passwords, certificates, and cryptographic keys into Vault1.

You need to grant App1 access to Vault1 and automatically rotate credentials. Credentials must not be stored in code.

What should you do?

- A. Enable App Service authentication for App1. Assign a custom RBAC role to Vault1.
- B. Add a TLS/SSL binding to App1.
- C. Upload a self-signed client certificate to Vault1. Update App1 to use the client certificate.
- D. Assign a managed identity to App1.

Correct Answer: D

Community vote distribution

D (100%)

 **alexein74**  1 year, 2 months ago

Selected Answer: D

D. Assign a managed identity to App1.

To grant App1 access to Vault1 and automatically rotate credentials without storing them in code, you should assign a managed identity to App1. Managed identities for Azure resources enable Azure services to authenticate to other Azure resources without needing to manage the authentication details.

After you enable a managed identity for App1, you can grant the identity access to Vault1 and use Azure Key Vault's built-in rotation feature to automatically rotate the credentials.

Additionally, you can use Azure Key Vault's built-in rotation feature to automatically rotate the credentials.

upvoted 15 times

 **warchoon** 1 year, 1 month ago

Don't put ChatGPT shit here, please

upvoted 9 times

 **Ciupaz** 6 months, 3 weeks ago

In this case ChatGPT got the right answer.

upvoted 1 times

 **lyaq** 6 months, 3 weeks ago

Not a reliable source. Not accurate 100%.

upvoted 1 times

 **adilkhan**  1 year, 2 months ago

D is correct <https://learn.microsoft.com/en-us/azure/active-directory/fundamentals/service-accounts-managed-identities>

upvoted 9 times

 **FeriAZ**  2 months, 3 weeks ago

Selected Answer: D

Managed identities provide an identity for Azure services, allowing them to authenticate to and interact with other Azure services without needing credentials in the code. By assigning a managed identity to App1 and setting the appropriate access policies in Vault1, App1 can securely access the keys, secrets, and certificates.

upvoted 1 times

 **130nk3r5** 4 months ago

Selected Answer: D

Answer is correct

upvoted 1 times

 **PKUSYD** 7 months, 1 week ago

D is correct

upvoted 1 times

Question #59

You are developing a Java application to be deployed in Azure. The application stores sensitive data in Azure Cosmos DB.

You need to configure Always Encrypted to encrypt the sensitive data inside the application.

What should you do first?

- A. Create a new container to include an encryption policy with the JSON properties to be encrypted.
- B. Create a customer-managed key (CMK) and store the key in a new Azure Key Vault instance.
- C. Create a data encryption key (DEK) by using the Azure Cosmos DB SDK and store the key in Azure Cosmos DB.
- D. Create an Azure AD managed identity and assign the identity to a new Azure Key Vault instance.

Correct Answer: B

Community vote distribution

B (100%)

✉  **alexein74**  1 year, 2 months ago

Selected Answer: B

B. Create a customer-managed key (CMK) and store the key in a new Azure Key Vault instance.

In order to configure Always Encrypted for the Java application, you need to first create a customer-managed key (CMK) using Azure Key Vault. This key will be used to encrypt and decrypt the sensitive data stored in Cosmos DB.

After creating the key, you should store it in a new Azure Key Vault instance, which will be used to manage and secure the key.

Once you have the key stored in Key vault, you can use Azure Cosmos DB SDK to encrypt the sensitive data and store it in Cosmos DB.

It's important to note that Always Encrypted in Azure Cosmos DB is a client-side encryption feature that encrypts and decrypts sensitive data inside the application and not in the Azure Cosmos DB service.

upvoted 14 times

✉  **warchoon** 1 year, 1 month ago

CMK does not encrypt the data. It wraps the DEK. But CMK is produced first.

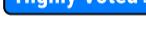
So the answer is correct, but the explanation is usual ChatGPT shit.

upvoted 3 times

✉  **warchoon** 1 year, 1 month ago

Don't post ChatGPT answers here. They are not always correct, even if they look logical.

upvoted 6 times

✉  **Samuelegaghagh**  1 year, 3 months ago

Selected Answer: B

<https://learn.microsoft.com/en-us/azure/cosmos-db/how-to-always-encrypted?tabs=dotnet>

upvoted 9 times

✉  **warchoon** 1 year, 1 month ago

=> <https://learn.microsoft.com/en-us/azure/cosmos-db/how-to-always-encrypted?tabs=dotnet#setup-azure-key-vault>

upvoted 2 times

✉  **Weam**  5 months ago

Selected Answer: B

Steps for always encrypted based on this link: <https://learn.microsoft.com/en-us/azure/cosmos-db/how-to-always-encrypted?tabs=dotnet>
1) create customer managed key stored in azure key vault

2) create a data encryption key from the SDK and wrap it with the key vault customer managed key

3) create a container with encryption policy

4) read and write encryption data

upvoted 1 times

✉  **dom271219** 6 months, 2 weeks ago

A of course. Why do you talk about Azure Key Vault ? The statement of the question does not address that.

upvoted 2 times

✉  **JH81** 9 months, 3 weeks ago

Selected Answer: B

Got this on 6/28/2023 and passed with 850. Answer is correct.

upvoted 3 times

 **Tom111** 1 year, 1 month ago

As the question asks what to do first I would select answer A.

"This policy must be provided when the container is created and it is immutable. In the current release, you can't update the encryption policy."

=> <https://learn.microsoft.com/en-us/azure/cosmos-db/how-to-always-encrypted?tabs=dotnet#encryption-policy>

upvoted 1 times

 **CarlosTheBoldest** 5 months ago

From your link: "Similar to an indexing policy, an encryption policy is a container-level specification describing how JSON properties should be encrypted. This policy must be provided when the container is created and it is immutable. In the current release, you can't update the encryption policy." No one cares about json properties but for the sensitive data stored on Cosmos DB, so A is a non sense

upvoted 1 times

Question #60

HOTSPOT

You develop a web app that interacts with Azure Active Directory (Azure AD) groups by using Microsoft Graph.

You build a web page that shows all Azure AD groups that are not of the type 'Unified'.

You need to build the Microsoft Graph query for the page.

How should you complete the query? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

https://graph.microsoft.com/v1.0/groups?

filter
search
contain

=

groupTypes/any(s:s ne 'Unified')
not groupTypes/contains('Unified')
not groupTypes/any(s:s eq 'Unified')
groupTypes/contains('Unified') eq false

&\$

\$top=true
\$count=true
\$filter=nested
\$consistencylevel=eventual

Correct Answer:

Answer Area

https://graph.microsoft.com/v1.0/groups?

filter
search
contain

=

groupTypes/any(s:s ne 'Unified')
not groupTypes/contains('Unified')
not groupTypes/any(s:s eq 'Unified')
groupTypes/contains('Unified') eq false

&\$

\$top=true
 \$count=true
\$filter=nested
\$consistencylevel=eventual

 **mmdex**  1 year, 1 month ago

~/groups?\$filter=NOT groupTypes/any(c:c eq 'Unified')&\$count=true

This exact example is mentioned in Microsoft documentation:

<https://learn.microsoft.com/en-us/graph/filter-query-parameter?tabs=http#examples-using-the-filter-query-operator>

upvoted 34 times

 **c_mazzotta** 1 year ago

Agree. if you try: "\$filter=groupTypes/any(s:s ne 'Unified')&\$count=true", you get: "Using 'ne' inside 'any' clause is invalid. Suggestion: consider using the 'not' operator instead."

upvoted 1 times

 **warchoon** 1 year, 1 month ago

=> [https://learn.microsoft.com/en-us/graph/filter-query-parameter?tabs=http#examples-using-the-filter-query-operator:~:text=~/groups%3F%24filter%3DNOT%20groupTypes/any\(c%3Ac%20eq%20%27Unified%27\)%26%24count%3Dtrue](https://learn.microsoft.com/en-us/graph/filter-query-parameter?tabs=http#examples-using-the-filter-query-operator:~:text=~/groups%3F%24filter%3DNOT%20groupTypes/any(c%3Ac%20eq%20%27Unified%27)%26%24count%3Dtrue)

upvoted 3 times

 **CarlosTheBoldest** 4 months, 2 weeks ago

Not sure since then but now it supports the ne parameter too:

"As Microsoft Entra ID continues to deliver more capabilities and improvements in stability, availability, and performance, Microsoft Graph also continues to evolve and scale to efficiently access the data. One way is through Microsoft Graph's increasing support for advanced query capabilities on various Microsoft Entra ID objects, also called directory objects, and their properties. For example, the addition of not (not), not

equals (ne), and ends with (endsWith) operators on the \$filter query parameter."

<https://learn.microsoft.com/en-us/graph/aad-advanced-queries?tabs=http#count-of-a-collection-in-a-filter-expression>

upvoted 4 times

✉ **g2000** Highly Voted 1 year, 3 months ago

it is correct

<https://learn.microsoft.com/en-us/azure/cosmos-db/how-to-always-encrypted?tabs=dotnet>

upvoted 7 times

✉ **tcybu** 1 year, 3 months ago

Agreed, it's correct, but here're the links with explanation

<https://learn.microsoft.com/en-us/graph/filter-query-parameter?tabs=http#any-operator>

<https://learn.microsoft.com/en-us/graph/aad-advanced-queries>

upvoted 2 times

✉ **Maskoo** Most Recent 2 months, 3 weeks ago

It's D

The 'ne' is not supported by default, we need count=true AND consistency level set to eventual which is not possible with this question.

upvoted 1 times

✉ **130nk3r5** 4 months ago

ne (Not Equals) is now supported.

upvoted 2 times

✉ **Acelexout** 6 months, 3 weeks ago

From documentation (<https://learn.microsoft.com/en-us/graph/filter-query-parameter?tabs=http#examples-using-the-filter-query-operator>):

List all non-Microsoft 365 groups in an organization.

GET ~/groups?\$filter=NOT groupTypes/any(c:c eq 'Unified')&\$count=true*

upvoted 3 times

✉ **BaoNguyen2411** 8 months, 3 weeks ago

got this question on 06/29/2023

upvoted 4 times

✉ **BaoNguyen2411** 8 months, 3 weeks ago

Got this answer on 06-29-2023

upvoted 3 times

✉ **Dianahu** 8 months, 3 weeks ago

Just try it yourself here: <https://developer.microsoft.com/en-us/graph/graph-explorer>

[https://graph.microsoft.com/v1.0/groups?\\$filter=groupTypes/any\(s:s ne 'Unified'\)&\\$count=true](https://graph.microsoft.com/v1.0/groups?$filter=groupTypes/any(s:s ne 'Unified')&$count=true) -> 400

[https://graph.microsoft.com/v1.0/groups?\\$filter=not groupTypes/contains\('Unified'\)&\\$count=true](https://graph.microsoft.com/v1.0/groups?$filter=not groupTypes/contains('Unified')&$count=true) -> 400

[https://graph.microsoft.com/v1.0/groups?\\$filter=not groupTypes/any\(s:s eq 'Unified'\)&\\$count=true](https://graph.microsoft.com/v1.0/groups?$filter=not groupTypes/any(s:s eq 'Unified')&$count=true) -> 200 with ConsistencyLevel: eventual

[https://graph.microsoft.com/v1.0/groups?\\$filter=groupTypes/contains\('Unified'\) eq false&\\$count=true](https://graph.microsoft.com/v1.0/groups?$filter=groupTypes/contains('Unified') eq false&$count=true) -> 400

upvoted 2 times

✉ **nekkilodeon** 12 months ago

From Microsoft documentation <https://learn.microsoft.com/en-us/graph/filter-query-parameter?tabs=http#examples-using-the-filter-query-operator>

List all non-Microsoft 365 groups in an organization. GET ~/groups?\$filter=NOT groupTypes/any(c:c eq 'Unified')&\$count=true*

upvoted 1 times

✉ **Heighte** 1 year ago

any idea why \$count=true though?

upvoted 4 times

✉ **Tomn82** 1 year ago

Answer ~/groups?\$filter=NOT groupTypes/any(c:c eq 'Unified')&\$count=true

Option 2: Use the \$filter query parameter with the ne operator. This request is not supported by default because the ne operator is only supported in advanced queries. Therefore, you must add the ConsistencyLevel header set to eventual and use the \$count=true query string.

<https://learn.microsoft.com/en-us/graph/aad-advanced-queries?tabs=http>

upvoted 1 times

✉ **baroo1** 1 year, 2 months ago

Answer: /groups?\$filter=not groupTypes/any(s:s eq 'Unified')&\$count=true

It also requires an extra header

ConsistencyLevel: eventual

"ne" is not supported.

message": "Using 'ne' inside 'any' clause is invalid. Suggestion: consider using the 'not' operator instead. For example: \$filter=not groupTypes/any(x:x eq 'Unified')

upvoted 2 times

✉ **vmakharashvili** 1 year, 1 month ago

I don't think so. I see if you want to use negative this way, you have to set parenthesis after not. like this:
`/groups$filter=not(groupTypes/any(s:s eq 'Unified'))`

upvoted 1 times

 **Net_IT** 1 year, 2 months ago

IMO it should be `filter=NOT groupTypes/any(s:s eq 'Unified')&$count=true`.

Source: <https://learn.microsoft.com/en-us/azure/cosmos-db/how-to-always-encrypted?tabs=dotnet>

In advanced querries you can also use the ne operator but then the consistencylevel has to be put to eventual and I'm not sure that's the case here?

upvoted 1 times

 **eugene1234** 1 year, 2 months ago

The answer is wrong! 1. filter should be `$filter`, and the expression should be `NOT groupTypes/any(s:s ne 'Unified')`, we want records that are NOT unified, see examples at <https://learn.microsoft.com/en-us/graph/filter-query-parameter?tabs=http>

upvoted 2 times

Question #61

DRAG DROP

You are developing an Azure solution.

You need to develop code to access a secret stored in Azure Key Vault.

How should you complete the code segment? To answer, drag the appropriate code segments to the correct location. Each code segment may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Code segments

DefaultAzureCredential
ClientSecretCredential
CloudClients
SecretClient

Answer area

```
def get_secret():

    var1 = os.environ.get("KEY_VAULT_URI")

    var2 =  ( var1 ,  )

    var3 = var2.get_secret("secret" )

    return "secret value {}" .format(var3.value)
```

Correct Answer:

Answer area

```
def get_secret():

    var1 = os.environ.get("KEY_VAULT_URI")

    var2 =  ( var1 ,  DefaultAzureCredential () )

    var3 = var2.get_secret("secret" )

    return "secret value {}" .format(var3.value)
```

✉  **g2000**  1 year, 3 months ago

it's correct

<https://learn.microsoft.com/en-us/azure/cosmos-db/how-to-always-encrypted?tabs=dotnet>

upvoted 13 times

✉  **JustHereToLearn** 1 year, 2 months ago

Looks good but I think this URL should be more accurate <https://learn.microsoft.com/en-us/python/api/azure-keyvault-secrets/azure.keyvault.secrets.secretclient?view=azure-python#examples>

upvoted 15 times

✉  **JH81**  9 months, 3 weeks ago

Got this on 6/28/2023 and passed with 850. Correct answer.

upvoted 6 times

Question #62

HOTSPOT

You are a developer building a web site using a web app. The web site stores configuration data in Azure App Configuration.

Access to Azure App Configuration has been configured to use the identity of the web app for authentication. Security requirements specify that no other authentication systems must be used.

You need to load configuration data from Azure App Configuration.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

```
from azure.identity import DefaultAzureCredential  
ChainedTokenCredential  
ManagedIdentityCredential  
AddAzureAppConfiguration  
  
from azure.app configuration import DefaultAzureCredential  
ChainedTokenCredential  
ManagedIdentityCredential  
AddAzureAppConfiguration  
  
credential = DefaultAzureCredential  
ChainedTokenCredential  
ManagedIdentityCredential  
AddAzureAppConfiguration ()  
  
client = DefaultAzureCredential  
ChainedTokenCredential  
ManagedIdentityCredential  
AddAzureAppConfiguration (base_url="...", credential=credential)
```

Correct Answer:

```
from azure.identity import DefaultAzureCredential ChainedTokenCredential  
ManagedIdentityCredential  
AddAzureAppConfiguration  
  
from azure.app configuration import DefaultAzureCredential  
ChainedTokenCredential  
ManagedIdentityCredential  
AddAzureAppConfiguration  
  
credential = DefaultAzureCredential ChainedTokenCredential  
ManagedIdentityCredential  
AddAzureAppConfiguration ()  
  
client = DefaultAzureCredential  
ChainedTokenCredential  
ManagedIdentityCredential  
AddAzureAppConfiguration (base_url="...", credential=credential)
```

 **Taschiloge**  1 year, 3 months ago

I think it should be ManagedIdentityCredential:

ManagedIdentityCredential: attempts to authenticate using a managed identity that is assigned to the deployment environment (if any).

DefaultAzureCredential can not be, since only one type of authentication is allowed.

ChainedTokenCredential could work if there is only one authentication type is specified in code.

Source: <https://yourazurecoach.com/2020/08/13/managed-identity-simplified-with-the-new-azure-net-sdks/>

upvoted 18 times

✉️  **tarek0811** 8 months, 3 weeks ago

chatgpt generated code with ManagedIdentityCredential

upvoted 1 times

✉️  **adilkhan** 1 year, 2 months ago

from azure.identity import DefaultAzureCredential

from azure.appconfiguration import AzureAppConfigurationClient

credential = DefaultAzureCredential()

client = AzureAppConfigurationClient(base_url="your_endpoint_url", credential=credential)

upvoted 7 times

✉️  **SharpZx** 1 year, 1 month ago

I think it should be DefaultAzureCredential. Check DefaultAzureCredential section at

<https://learn.microsoft.com/en-us/python/api/overview/azure/identity-readme?view=azure-python>

DefaultAzureCredential is appropriate for most applications which will run in the Azure Cloud because it combines common production credentials with development credentials

Managed Identity - If the application is deployed to an Azure host with Managed Identity enabled, DefaultAzureCredential will authenticate with it.

upvoted 2 times

✉️  **JamieS** 1 year ago

I disagree, if they used options (<https://docs.microsoft.com/en-us/dotnet/api/azure.identity.defaultazurecredentialoptions?view=azure-dotnet-preview>) they could exclude unwanted auth methods, but it's not in the code. Otherwise "It is not always clear what authentication method got executed".

Question states that "configured to use the identity" so I think it's ManagedIdentityCredential in place of azure default.

upvoted 4 times

✉️  **SharpZx** 1 year ago

JamieS you are right. It should be ManagedIdentityCredential. I overlooked .. "Security requirements specify that no other authentication systems must be used".

upvoted 3 times

✉️  **EricPerezVillar**  11 months, 2 weeks ago

Answer is correct. It is here:

<https://learn.microsoft.com/en-us/python/api/overview/azure/appconfiguration-readme?view=azure-python# create-a-client>

upvoted 16 times

✉️  **MikeAWS** 6 months ago

great catch! thanks

upvoted 1 times

✉️  **130nk3r5**  4 months ago

ManagedIdentityCredential > AddAzureAppconfig > ManagedIdentityCredential > AdAzureAppConf

upvoted 2 times

✉️  **Weam** 5 months ago

ManagedIdentityCredential according to the following link : <https://learn.microsoft.com/en-us/python/api/azure-identity/azure.identity.managedidentitycredential?view=azure-python>

We use ManagedIdentityCredential with apps or resources that support ManagedIdentityCredential.

upvoted 1 times

✉️  **Ciupaz** 6 months, 3 weeks ago

This question is not updated with the latest Microsoft docs.

The 2nd and 4th answers are not correct, and are not present in DDL.

Here the right Python code:

```
from azure.identity import DefaultAzureCredential
from azure.appconfiguration import AzureAppConfigurationClient
```

```
credential = DefaultAzureCredential()
```

```
client = AzureAppConfigurationClient(base_url="your_endpoint_url", credential=credential)
```

<https://learn.microsoft.com/en-us/python/api/overview/azure/appconfiguration-readme?view=azure-python# create-a-client>

upvoted 2 times

macobuzi 7 months, 4 weeks ago

And why the hell did Microsoft decide to include Python code in the exam?
upvoted 4 times

EliteAllen 8 months, 1 week ago

To load configuration data from Azure App Configuration using the identity of the web app for authentication, you would typically use the ManagedIdentityCredential class for authentication. Here's how you should complete the code:

```
from azure.identity import ManagedIdentityCredential
from azure.appconfiguration import AzureAppConfigurationClient

credential = ManagedIdentityCredential()

client = AzureAppConfigurationClient(base_url="your_endpoint_url", credential=credential)
```

upvoted 7 times

nickk6425 9 months, 3 weeks ago

Should be ManagedIdentity because "...configured to use the identity of the web app for authentication. Security requirements specify that no other authentication systems must be used."

upvoted 4 times

nickk6425 9 months, 3 weeks ago

DefaultAzureCredentials will try different authentication, ie. it will use AZ, Visual Studio, ManagedIdentity, and so on until one succeeds
upvoted 3 times

Dats1987 1 year ago

Note that the DefaultAzureCredential is used to authenticate using the identity of the web app. This credential provider tries multiple authentication methods (e.g. environment variables, Azure Managed Identity, Azure CLI) until it finds a suitable one. In this way, you can ensure that no other authentication systems are used

Above ans are correct.

upvoted 1 times

Dilmurod 1 year ago

It should be ManagedIdentityCredential.

Managed identity authentication is supported via either the DefaultAzureCredential or the ManagedIdentityCredential directly for the following Azure services:

Example:

```
from azure.identity import ManagedIdentityCredential
from azure.keyvault.secrets import SecretClient

credential = ManagedIdentityCredential()
client = SecretClient("https://my-vault.vault.azure.net", credential)
```

Reference: https://github.com/Azure/azure-sdk-for-python/tree/azure-appconfiguration_1.4.0/sdk/identity/azure-identity

upvoted 2 times

naivecoder786 1 year, 1 month ago

Correct one is ManagedIdentityCredential 110% Sure !

upvoted 3 times

[Removed] 1 year, 1 month ago

Can you explain please?

upvoted 1 times

karthikwarrior 1 year, 1 month ago

What is the right answer then? everyone thrown some kitchen sink.. pls keep it simple.

upvoted 4 times

MikeAWS 6 months ago

```
from azure.identity import DefaultAzureCredential
from azure.appconfiguration import AzureAppConfigurationClient

credential = DefaultAzureCredential()
```

```
client = AzureAppConfigurationClient(base_url="your_endpoint_url", credential=credential)
```

<https://learn.microsoft.com/en-us/python/api/overview/azure/appconfiguration-readme?view=azure-python# create-a-client>
upvoted 1 times

adilkhan 1 year, 2 months ago

The options are wrong <https://learn.microsoft.com/en-us/python/api/overview/azure/appconfiguration-readme?view=azure-python>

```
from azure.identity import DefaultAzureCredential
from azure.appconfiguration import AzureAppConfigurationClient

credential = DefaultAzureCredential()
```

```
client = AzureAppConfigurationClient(base_url="your_endpoint_url", credential=credential)
```

upvoted 3 times

✉ **ajayasa** 1 year, 2 months ago

@adilkhan

can you please provide right answers. your comment starts with answers being wrong and you are providing the same answers mentioned
upvoted 6 times

✉ **maiwufsiorn** 1 year, 2 months ago

No, he is not. Read it again: AzureAppConfigurationClient is different from AddAzureAppConfiguration.

upvoted 2 times

✉ **warchoon** 1 year, 1 month ago

It looks like typo. AddAzureAppConfiguration is a method, not object.

upvoted 1 times

✉ **g2000** 1 year, 3 months ago

DefaultAzureCredential is wrong because it tries different authentication methods.

<https://yourazurecoach.com/2020/08/13/managed-identity-simplified-with-the-new-azure-net-sdks/>

upvoted 4 times

Question #63

Topic 4

You are developing several microservices to deploy to a new Azure Kubernetes Service cluster. The microservices manage data stored in Azure Cosmos DB and Azure Blob storage. The data is secured by using customer-managed keys stored in Azure Key Vault.

You must automate key rotation for all Azure Key Vault keys and allow for manual key rotation. Keys must rotate every three months. Notifications of expiring keys must be sent before key expiry.

You need to configure key rotation and enable key expiry notifications.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Create and configure a new Azure Event Grid instance.
- B. Configure Azure Key Vault alerts.
- C. Create and assign an Azure Key Vault access policy.
- D. Create and configure a key rotation policy during key creation.

Correct Answer: AC

Community vote distribution



✉ **abcdx** Highly Voted 12 months ago

Selected Answer: AD

A&D for sure!

<https://learn.microsoft.com/en-us/azure/key-vault/keys/how-to-configure-key-rotation>

upvoted 20 times

✉ **halfway** Highly Voted 1 year ago

Selected Answer: AD

Key rotation policy and Event Grid notification

upvoted 6 times

✉ **Ciupaz** Most Recent 3 months, 1 week ago

Selected Answer: BD

For me, B and D for sure.

upvoted 1 times

✉ **130nk3r5** 4 months ago

Selected Answer: BD

B. Configure Azure Key Vault alerts. To receive notifications of expiring keys, you need to configure Azure Key Vault alerts. You can set up alerts for key expiration events, which will notify you before the key expires.

D. Create and configure a key rotation policy during key creation. To automate key rotation, you need to create and configure a key rotation policy when creating the keys in Azure Key Vault. You can set the rotation interval to three months, as required, and also allow for manual key rotation.

upvoted 1 times

✉ **macobuzi** 7 months, 4 weeks ago

Selected Answer: AD

We can use the Key Rotation Policy in Azure Key Vault combined with Event Grid to trigger sending notification when a secret in the key vault is about to expire.

<https://learn.microsoft.com/en-us/azure/key-vault/general/event-grid-tutorial>

upvoted 2 times

✉ **nekkilodeon** 12 months ago

Selected Answer: BC

B & C are correct

Alerts for notifications and Access policies for storage access to existing keys

upvoted 1 times

✉ **CarlosTheBoldest** 5 months ago

I thought as you and I was wrong :)

"As you start to scale your service, the number of requests sent to your key vault will rise. This rise has a potential to increase the latency of your requests. In extreme cases, it can cause your requests to be throttled and affect the performance of your service. You also need to know if your key vault is sending an unusual number of error codes, so you can quickly handle any problems with an access policy or firewall configuration." So the KV alert is used to raise alerts when it begins to send too many errors or receive too many petitions

<https://learn.microsoft.com/en-us/azure/key-vault/general/alert>

upvoted 1 times

✉  **paunski7** 12 months ago

B & D

To configure key rotation and enable key expiry notifications for Azure Key Vault, you should perform the following two actions:

B. Configure Azure Key Vault alerts: Configure alerts for when a key is expiring, so that you can receive notifications before the key expiry. You can configure these alerts in Azure Key Vault using Azure Monitor, which sends an email or a webhook notification to a recipient or service when the key is about to expire. You can specify the notification threshold in terms of days, so you can receive alerts, for example, seven days before the key expiry.

C. Create and assign an Azure Key Vault access policy: Create and assign an access policy for your Azure Key Vault that allows your microservices to perform key rotations manually, and automated key rotations using a key rotation script. You can create and assign access policies for Azure Key Vault through the Azure portal or the Azure CLI.

These actions enable you to configure key rotation and notifications for key expiry.

upvoted 4 times

✉  **paunski7** 12 months ago

Options A and D are not relevant to the solution. Azure Event Grid is a service that allows you to react to events in Azure services by routing them to different endpoints, but it is not required for key rotation or notifications for Azure Key Vault. Key rotation policies are not created during key creation, but rather they are created and applied to existing keys.

upvoted 2 times

✉  **smariussorin** 8 months, 1 week ago

<https://learn.microsoft.com/en-us/azure/key-vault/general/event-grid-tutorial>

check documentation: ". When one of the secrets in the key vault is about to expire (defined as 30 days before expiration date), Event Grid is notified of the status change and makes an HTTP POST to the endpoint. "

is A & D

upvoted 1 times

✉  **dy0917** 11 months, 1 week ago

Key vault alerts use to monitor health issues.

After you start to use Azure Key Vault to store your production secrets, it's important to monitor the health of your key vault to make sure that your service operates as intended.

upvoted 1 times

Question #64

Topic 4

You are developing a web application that uses the Microsoft identity platform to authenticate users and resources. The web application calls several REST APIs.

The APIs require an access token from the Microsoft identity platform.

You need to request a token.

Which three properties should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Redirect URI/URL
- B. Application ID
- C. Application name
- D. Application secret
- E. Supported account type

Correct Answer: ABD

Community vote distribution

ABD (100%)

✉  **AbidooKing**  9 months, 2 weeks ago

Answer seems correct!

<https://learn.microsoft.com/en-us/graph/auth-v2-user?tabs=http#3-request-an-access-token>
upvoted 10 times

✉  **Tralalaaz204**  3 months, 1 week ago

Selected Answer: ABD

correct

upvoted 1 times

✉  **CarlosTheBoldest** 4 months, 1 week ago

I got this question on my exam, 2023Dec, go with what I remember was the most voted answer. Score 902, most of the questions were here, slightly different on wording because the Azure Ad <-> Entra Id change. Case was City Power & Light. Good luck!

Important tip, you have access to microsoft learn during the exam!

upvoted 2 times

✉  **nardk** 7 months, 1 week ago

Selected Answer: ABD

correct

upvoted 2 times

✉  **Roel1** 9 months, 2 weeks ago

Seems plausible

upvoted 3 times

✉  **Steffexx** 9 months, 2 weeks ago

Answer A B D is correct

upvoted 1 times

Question #65

HOTSPOT

You are developing an application that uses Azure Storage to store customer data. The data must only be decrypted by the customer and the customer must be provided a script to rotate keys.

You need to provide a script to rotate keys to the customer.

How should you complete the command? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
$h = $(az keyvault show --hsm-name ... --query "properties.hsmUri"  
$x = az keyvault list-versions --name ""  
key  
secret  
recover  
certificate  
--vault-name "" --query "[0].kid"  
az storage account update  
--name ... \  
--resource-group ... \  
--encryption-key-name ... \  
--encryption-key-version $x \  
--encryption-key-source Microsoft.Secret  
Microsoft.Storage  
Microsoft.Keyvault  
Microsoft.Certificate  
--encryption-key-vault $h
```

Answer Area

```
$h = $(az keyvault show --hsm-name ... --query "properties.hsmUri"  
$x = az keyvault list-versions --name ""  
key  
secret  
recover  
certificate  
--vault-name "" --query "[0].kid"  
az storage account update  
--name ... \  
--resource-group ... \  
--encryption-key-name ... \  
--encryption-key-version $x \  
--encryption-key-source Microsoft.Secret  
Microsoft.Storage  
Microsoft.Keyvault  
Microsoft.Certificate  
--encryption-key-vault $h
```

Correct Answer:

✉  **FakerJani**  7 months, 1 week ago

Answers is correct.

<https://learn.microsoft.com/en-us/cli/azure/storage/account?view=azure-cli-latest#az-storage-account-update>

<https://learn.microsoft.com/en-us/cli/azure/keyvault/key?view=azure-cli-latest>

upvoted 5 times

✉  **130nk3r5** 3 months, 2 weeks ago

Got this today.

Went with answer here.

Score 927

upvoted 1 times

✉  **Momook**  7 months, 2 weeks ago

maybe the answer should be , secret and Microsoft.Keyvault ??

upvoted 2 times

✉  **zeugo** 4 months, 2 weeks ago

Would really help if you provided documentation to back up your answer.

upvoted 2 times

✉  **sekelsenmat** 2 months, 2 weeks ago

az in the next line is being called with "--encryption-key-name" which in the docs says it is for "The name of the KeyVault key."

So clearly it cannot be a secret, it has to be a key we are talking about here.

upvoted 1 times

✉  **BaoNguyen2411** 8 months, 3 weeks ago

got this question on 29/06/2023

upvoted 3 times

✉  **CloudlessCloud** 9 months, 1 week ago

Ty! (^///^)

upvoted 1 times

Question #66

Topic 4

You are developing several Azure API Management (APIM) hosted APIs.

You must transform the APIs to hide private backend information and obscure the technology stack used to implement the backend processing.

You need to protect all APIs.

What should you do?

- A. Configure and apply a new inbound policy scoped to a product.
- B. Configure and apply a new outbound policy scoped to the operation.
- C. Configure and apply a new outbound policy scoped to global.
- D. Configure and apply a new backend policy scoped to global.

Correct Answer: B

Community vote distribution



✉️  **Ciupaz** 3 months, 3 weeks ago

Selected Answer: C

All APIs -> global
upvoted 2 times

✉️  **130nk3r5** 4 months ago

Selected Answer: C

C. Configure and apply a new outbound policy scoped to global.

By configuring and applying a global outbound policy, you can transform the responses from all APIs in your Azure API Management instance. This will help you hide private backend information and obscure the technology stack used for backend processing across all APIs.

upvoted 1 times

✉️  **Tralalaaz204** 3 months, 1 week ago

I agree with you.

As far as I understand in the document, the operation scope is related only to one specific API

<https://learn.microsoft.com/en-us/azure/api-management/api-management-howto-policies#scopes>

upvoted 1 times

✉️  **JoaoLoop10** 4 months, 4 weeks ago

Selected Answer: B

Maybe is B: <https://learn.microsoft.com/en-us/azure/api-management/transform-api>

upvoted 2 times

✉️  **kakajos** 7 months ago

Selected Answer C:

Global scope is for All APIs in your API Management Instance. And since we need to protect all the APIs the, we want an outbound policy with a global scope. Ans

upvoted 2 times

✉️  **katrang** 7 months, 1 week ago

Selected Answer: C

Add outbound policy -> Set headers and the scope should be global (transform the APIs - no mention of specific operation)

upvoted 2 times

✉️  **devex** 7 months, 2 weeks ago

Selected Answer: C

C: You need to protect all APIs ("All operations")

upvoted 2 times

✉️  **tempacc4nk** 7 months, 3 weeks ago

Selected Answer: C

To safeguard private backend information and mask the underlying technology stack utilized in the backend processes of your Azure API Management (APIM) hosted APIs, implement and apply a fresh global-scoped outbound policy. This policy will uniformly apply the necessary

transformations and concealments to the responses of all APIs within your APIM instance, ensuring a consistent approach to securing backend details and technology stack disclosure. So C

upvoted 4 times

 **namlus** 7 months, 3 weeks ago

Its 'C'. Check <https://learn.microsoft.com/en-us/azure/api-management/set-edit-policies?tabs=form>

upvoted 4 times

 **lyggwtwtxnxhznebw** 8 months, 1 week ago

It should be an outbound policy. The articles says you shoyld select 'All operations', so either you take that as 'global' and go for C, or there is a typo in B and it should have said 'All operations'

upvoted 3 times

 **smariussorin** 8 months, 1 week ago

C, is a global policy. Check <https://learn.microsoft.com/en-us/azure/api-management/transform-api>.

"Select Demo Conference API > Design > All operations."

upvoted 2 times

 **Firo** 8 months, 3 weeks ago

Selected Answer: B

B is correct refer to below link

<https://learn.microsoft.com/en-us/azure/api-management/transform-api>

upvoted 4 times

 **macobuzi** 7 months, 4 weeks ago

Why not C?

upvoted 1 times

Question #67

HOTSPOT

You are developing an Azure Function App named App1. You also plan to use cross-origin requests (CORS).

You have the following requirements:

- App1 functions must securely access an Azure Blob Storage account.
- Access to the Azure Blob Storage account must not require the provisioning or rotation of secrets.
- JavaScript code running in a browser on an external host must not be allowed to interact with the function.

You need to implement App1.

Which configuration should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Requirement

Azure Blob Storage access

Configuration value

- Client secret credentials
User-assigned managed identity
System-assigned managed identity

Disallow access from other domains

- Configure CORS allowed origins to *
Configure CORS allowed origins to none
Configure CORS allowed origins to disable

Answer Area

Requirement

Azure Blob Storage access

Configuration value

- Client secret credentials
User-assigned managed identity
System-assigned managed identity

Correct Answer:

Disallow access from other domains

- Configure CORS allowed origins to *
Configure CORS allowed origins to none
Configure CORS allowed origins to disable

 **SSR999** 3 months, 1 week ago

is this question is part of any case study
upvoted 2 times

 **130nk3r5** 4 months ago

Answer is correct.

upvoted 2 times

✉ **130nk3r5** 4 months ago

System Managed Identity: System managed identities are tied to your Azure service and are automatically cleaned up when the resource is deleted. They are easier to use if your application only needs to authenticate to services that support Azure AD authentication.

Configure CORS allowed origins to none: To disallow JavaScript code running in a browser on an external host from interacting with the function, you should not include that host in your CORS policy. Configuring CORS allowed origins to none will prevent any domain from accessing your function app.

upvoted 3 times

✉ **Weam** 5 months ago

Correct answer is :

System managed Identity

Configure allowed origins to disable according to this link:

<https://learn.microsoft.com/en-us/cli/azure/functionapp/cors?view=azure-cli-latest#az-functionapp-cors-credentials>

upvoted 3 times

✉ **SachinV** 5 months ago

As the requirement says “Plan to Use CORS”, the correct response should be “Configure cors allowed origins *”. However, I got confused on the 2nd question label asking “Disallow access from another domain”. If we don’t set CORS, then by default, CORS is disabled for other domains.

upvoted 2 times

Question #68

HOTSPOT

You develop a containerized application. The application must be deployed to an existing Azure Kubernetes Service (AKS) cluster from an Azure Container Registry (ACR) instance. You use the Azure command-line interface (Azure CLI) to deploy the application image to AKS.

Images must be pulled from the registry. You must be able to view all registries within the current Azure subscription. Authentication must be managed by Microsoft Entra ID and removed when the registry is deleted. The solution must use the principle of least privilege.

You need to configure authentication to the registry.

Which authentication configuration should you use? To answer, select the appropriate configuration values in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Authentication

Registry authentication method

Configuration Value

Service principal
Repository-scoped access token
User-assigned managed identity
System-assigned managed identity

Registry Azure role-based access control (Azure RBAC) role

Reader
AcrPush
AcrPull
Contributor

Answer Area

Authentication

Registry authentication method

Configuration Value

Service principal
Repository-scoped access token
User-assigned managed identity
System-assigned managed identity

Correct Answer:

Registry Azure role-based access control (Azure RBAC) role

Reader
AcrPush
AcrPull
Contributor

 **zixys** 1 month, 3 weeks ago

System-assigned managed identity

Reader: The Reader has permissions to view the registry and pull images.

upvoted 4 times

 **b69794f** 1 month, 2 weeks ago

Correct <https://learn.microsoft.com/en-us/azure/container-registry/container-registry-roles?tabs=azure-cli>

upvoted 1 times

 **didi2087** 1 month, 3 weeks ago

Registry Azure RBAC Role: should be reader . to meet the requirements

"Images must be pulled from the registry." "You must be able to view all registries within the current Azure subscription."

Check the table :Role/Permission

<https://learn.microsoft.com/en-us/azure/container-registry/container-registry-roles?tabs=azure-cli#pull-image>

upvoted 3 times

 **FeriAZ** 2 months, 3 weeks ago

correct.

Registry Authentication Method: A System-assigned Managed Identity. It's tied to the AKS service and automatically managed by Azure, aligning with the requirement for authentication to be managed by Microsoft Entra ID and removed when the registry is deleted. It adheres to the principle of least privilege as it's specific to the AKS resource.

Registry Azure RBAC Role: The AcrPull role is the most fitting. It provides just enough permission to pull images from ACR, aligning with the principle of least privilege and meeting the requirement of the deployment process.

upvoted 2 times

Question #69

Topic 4

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. When you are ready to answer a question, click the Question button to return to the question.

Background -

Munson's Pickles and Preserves Farm is an agricultural cooperative corporation based in Washington, US, with farms located across the United States. The company supports agricultural production resources by distributing seeds, fertilizers, chemicals, fuel, and farm machinery to the farms.

Current Environment -

The company is migrating all applications from an on-premises datacenter to Microsoft Azure. Applications support distributors, farmers, and internal company staff.

Corporate website -

- The company hosts a public website located at <http://www.munsonspicklesandpreservesfarm.com>. The site supports farmers and distributors who request agricultural production resources.

Farms -

- The company created a new customer tenant in the Microsoft Entra admin center to support authentication and authorization for applications.

Distributors -

- Distributors integrate their applications with data that is accessible by using APIs hosted at <http://www.munsonspicklesandpreservesfarm.com/api> to receive and update resource data.

Requirements -

The application components must meet the following requirements:

Corporate website -

- The site must be migrated to Azure App Service.
- Costs must be minimized when hosting in Azure.

- Applications must automatically scale independent of the compute resources.
- All code changes must be validated by internal staff before release to production.
- File transfer speeds must improve, and webpage-load performance must increase.
- All site settings must be centrally stored, secured without using secrets, and encrypted at rest and in transit.
- A queue-based load leveling pattern must be implemented by using Azure Service Bus queues to support high volumes of website agricultural production resource requests.

Farms -

- Farmers must authenticate to applications by using Microsoft Entra ID.

Distributors -

- The company must track a custom telemetry value with each API call and monitor performance of all APIs.
- API telemetry values must be charted to evaluate variations and trends for resource data.

Internal staff -

- App and API updates must be validated before release to production.
- Staff must be able to select a link to direct them back to the production app when validating an app or API update.
- Staff profile photos and email must be displayed on the website once they authenticate to applications by using their Microsoft Entra ID.

Security -

- All web communications must be secured by using TLS/HTTPS.
- Web content must be restricted by country/region to support corporate compliance standards.
- The principle of least privilege must be applied when providing any user rights or process access rights.
- Managed identities for Azure resources must be used to authenticate services that support Microsoft Entra ID authentication.

Issues -

Corporate website -

- Farmers report HTTP 503 errors at the same time as internal staff report that CPU and memory usage are high.
- Distributors report HTTP 502 errors at the same time as internal staff report that average response times and networking traffic are high.
- Internal staff report webpage load sizes are large and take a long time to load.
- Developers receive authentication errors to Service Bus when they debug locally.

Distributors -

- Many API telemetry values are sent in a short period of time. Telemetry traffic, data costs, and storage costs must be reduced while preserving a statistically correct analysis of the data points sent by the APIs.

You need to implement farmer authentication.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Add the shared access signature (SAS) token to the app.
- B. Create a shared access signature (SAS) token.
- C. Create a user flow.
- D. Add the app to the user flow.
- E. Register the app in Microsoft Entra ID.

Correct Answer: CDE*Community vote distribution***CDE (100%)****👤** **FeriAZ** 2 months, 3 weeks ago**Selected Answer: CDE**

Create a user flow (Option C): This step is necessary to define how farmers will authenticate and possibly manage their profiles.
Add the app to the user flow (Option D): Essential for integrating the user flow with the application, ensuring the app utilizes the defined authentication process.

Register the app in Microsoft Entra ID (Option E): A fundamental step to allow the application to use Microsoft Entra ID for authentication.

upvoted 2 times

👤 **AzDeveloper** 3 months ago**Selected Answer: CDE**

<https://learn.microsoft.com/en-us/azure/active-directory-b2c/tutorial-create-user-flows?pivots=b2c-user-flow>

upvoted 1 times

👤 **Ciupaz** 3 months, 1 week ago**Selected Answer: CDE**

Correct.

By performing these 3 actions, you will effectively implement farmer authentication using Microsoft Entra ID.

upvoted 2 times

Question #70

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. When you are ready to answer a question, click the Question button to return to the question.

Background -

Munson's Pickles and Preserves Farm is an agricultural cooperative corporation based in Washington, US, with farms located across the United States. The company supports agricultural production resources by distributing seeds, fertilizers, chemicals, fuel, and farm machinery to the farms.

Current Environment -

The company is migrating all applications from an on-premises datacenter to Microsoft Azure. Applications support distributors, farmers, and internal company staff.

Corporate website -

- The company hosts a public website located at <http://www.munsonspicklesandpreservesfarm.com>. The site supports farmers and distributors who request agricultural production resources.

Farms -

- The company created a new customer tenant in the Microsoft Entra admin center to support authentication and authorization for applications.

Distributors -

- Distributors integrate their applications with data that is accessible by using APIs hosted at <http://www.munsonspicklesandpreservesfarm.com/api> to receive and update resource data.

Requirements -

The application components must meet the following requirements:

Corporate website -

- The site must be migrated to Azure App Service.
- Costs must be minimized when hosting in Azure.

- Applications must automatically scale independent of the compute resources.
- All code changes must be validated by internal staff before release to production.
- File transfer speeds must improve, and webpage-load performance must increase.
- All site settings must be centrally stored, secured without using secrets, and encrypted at rest and in transit.
- A queue-based load leveling pattern must be implemented by using Azure Service Bus queues to support high volumes of website agricultural production resource requests.

Farms -

- Farmers must authenticate to applications by using Microsoft Entra ID.

Distributors -

- The company must track a custom telemetry value with each API call and monitor performance of all APIs.
- API telemetry values must be charted to evaluate variations and trends for resource data.

Internal staff -

- App and API updates must be validated before release to production.
- Staff must be able to select a link to direct them back to the production app when validating an app or API update.
- Staff profile photos and email must be displayed on the website once they authenticate to applications by using their Microsoft Entra ID.

Security -

- All web communications must be secured by using TLS/HTTPS.
- Web content must be restricted by country/region to support corporate compliance standards.
- The principle of least privilege must be applied when providing any user rights or process access rights.
- Managed identities for Azure resources must be used to authenticate services that support Microsoft Entra ID authentication.

Issues -

Corporate website -

- Farmers report HTTP 503 errors at the same time as internal staff report that CPU and memory usage are high.
- Distributors report HTTP 502 errors at the same time as internal staff report that average response times and networking traffic are high.
- Internal staff report webpage load sizes are large and take a long time to load.
- Developers receive authentication errors to Service Bus when they debug locally.

Distributors -

- Many API telemetry values are sent in a short period of time. Telemetry traffic, data costs, and storage costs must be reduced while preserving a statistically correct analysis of the data points sent by the APIs.

You need to secure the corporate website to meet the security requirements.

What should you do?

- A. Create an Azure Cache for Redis instance. Update the code to support the cache.
- B. Create an Azure Content Delivery Network profile and endpoint. Configure the endpoint.
- C. Create an App Service instance with a standard plan. Configure the custom domain with a TLS/SSL certificate.
- D. Create an Azure Application Gateway with a Web Application Firewall (WAF). Configure end-to-end TLS encryption and the WAF.

Correct Answer: C

Community vote distribution

✉ **odinpodin** 1 month, 3 weeks ago

how can c be correct?

upvoted 1 times

✉ **FeriAZ** 2 months, 3 weeks ago

Selected Answer: D

Azure Application Gateway is a web traffic load balancer that enables you to manage traffic to your web applications. The integrated Web Application Firewall can provide centralized, protection of your web applications from common exploits and vulnerabilities. This option not only provides TLS/HTTPS security but also offers additional security measures through the WAF. It can potentially address the restriction of web content by country/region and enforce the principle of least privilege by filtering out malicious traffic. Based on the security requirements specified:

Securing all web communications with TLS/HTTPS.

Restricting web content by country/region.

Applying the principle of least privilege.

upvoted 2 times

✉ **AzDeveloper** 3 months ago

Selected Answer: D

<https://learn.microsoft.com/en-us/entra/identity/app-proxy/application-proxy-application-gateway-waf>

upvoted 1 times

✉ **Jedi** 3 months ago

There is also the requirement for "Web content must be restricted by country/region to support corporate compliance standards."

Thinking we either need CDN: <https://learn.microsoft.com/en-us/azure/cdn/cdn-restrict-access-by-country-region>

Or WAF: <https://learn.microsoft.com/en-us/azure/web-application-firewall/ag/geomatch-custom-rules>

upvoted 1 times

✉ **AzDeveloper** 3 months, 1 week ago

Answer C is correct

upvoted 1 times

✉ **AzDeveloper** 3 months ago

Answer is D

upvoted 1 times

Question #71

HOTSPOT

-

Case study

-

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

-

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. When you are ready to answer a question, click the Question button to return to the question.

Background

-

Munson's Pickles and Preserves Farm is an agricultural cooperative corporation based in Washington, US, with farms located across the United States. The company supports agricultural production resources by distributing seeds, fertilizers, chemicals, fuel, and farm machinery to the farms.

Current Environment

-

The company is migrating all applications from an on-premises datacenter to Microsoft Azure. Applications support distributors, farmers, and internal company staff.

Corporate website

-

- The company hosts a public website located at <http://www.munsonspicklesandpreservesfarm.com>. The site supports farmers and distributors who request agricultural production resources.

Farms

-

- The company created a new customer tenant in the Microsoft Entra admin center to support authentication and authorization for applications.

Distributors

-

- Distributors integrate their applications with data that is accessible by using APIs hosted at

<http://www.munsonspicklesandpreservesfarm.com/api> to receive and update resource data.

Requirements

-

The application components must meet the following requirements:

Corporate website

- The site must be migrated to Azure App Service.
- Costs must be minimized when hosting in Azure.
- Applications must automatically scale independent of the compute resources.
- All code changes must be validated by internal staff before release to production.
- File transfer speeds must improve, and webpage-load performance must increase.
- All site settings must be centrally stored, secured without using secrets, and encrypted at rest and in transit.
- A queue-based load leveling pattern must be implemented by using Azure Service Bus queues to support high volumes of website agricultural production resource requests.

Farms

-

- Farmers must authenticate to applications by using Microsoft Entra ID.

Distributors

-

- The company must track a custom telemetry value with each API call and monitor performance of all APIs.
- API telemetry values must be charted to evaluate variations and trends for resource data.

Internal staff

-

- App and API updates must be validated before release to production.
- Staff must be able to select a link to direct them back to the production app when validating an app or API update.
- Staff profile photos and email must be displayed on the website once they authenticate to applications by using their Microsoft Entra ID.

Security

-

- All web communications must be secured by using TLS/HTTPS.
- Web content must be restricted by country/region to support corporate compliance standards.
- The principle of least privilege must be applied when providing any user rights or process access rights.
- Managed identities for Azure resources must be used to authenticate services that support Microsoft Entra ID authentication.

Issues

-

Corporate website

-

- Farmers report HTTP 503 errors at the same time as internal staff report that CPU and memory usage are high.
- Distributors report HTTP 502 errors at the same time as internal staff report that average response times and networking traffic are high.
- Internal staff report webpage load sizes are large and take a long time to load.
- Developers receive authentication errors to Service Bus when they debug locally.

Distributors

- Many API telemetry values are sent in a short period of time. Telemetry traffic, data costs, and storage costs must be reduced while preserving a statistically correct analysis of the data points sent by the APIs.

You need to display the profile photo and email for signed-in internal staff on the website.

Which Microsoft Graph configuration should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Microsoft Graph configuration

Configuration setting	Configuration value
Endpoint	<input type="checkbox"/> /v1.0/me <input type="checkbox"/> /v1.0/me/people/?\$search=me <input type="checkbox"/> /v1.0/users?\$select=profilePhoto,mail
Permission	<input type="checkbox"/> User.Read <input type="checkbox"/> User.Export.All <input type="checkbox"/> User.ReadWrite <input type="checkbox"/> User.ManageIdentities.All

Microsoft Graph configuration

Configuration setting	Configuration value
Endpoint	<input type="checkbox"/> /v1.0/me <input type="checkbox"/> /v1.0/me/people/?\$search=me <input type="checkbox"/> /v1.0/users?\$select=profilePhoto,mail
Permission	<input type="checkbox"/> User.Read <input type="checkbox"/> User.Export.All <input type="checkbox"/> User.ReadWrite <input type="checkbox"/> User.ManageIdentities.All

Correct Answer:

Configuration setting	Configuration value
Endpoint	<input checked="" type="checkbox"/> /v1.0/me <input type="checkbox"/> /v1.0/me/people/?\$search=me <input type="checkbox"/> /v1.0/users?\$select=profilePhoto,mail
Permission	<input checked="" type="checkbox"/> User.Read <input type="checkbox"/> User.Export.All <input type="checkbox"/> User.ReadWrite <input type="checkbox"/> User.ManageIdentities.All

 **raymond_abcd** 2 months, 2 weeks ago

given answer is correct:

<https://learn.microsoft.com/en-us/graph/api/resources/profile-example?view=graph-rest-beta&viewFallbackFrom=graph-rest-1.0>

upvoted 2 times

 **misinho7** 2 months, 3 weeks ago

Can Anyone help me and tell me where I can find all question related to Farmers? Topics and questions number? Thank you very much !

upvoted 1 times

 **onlyforheros** 1 month, 1 week ago

Use search button

upvoted 1 times

 **FeriAZ** 2 months, 3 weeks ago

correct.

/v1.0/me: This endpoint is used to access the profile information of the currently signed-in user. It is the most direct way to get the user's own profile details, including their email.

User.Read: This permission allows an app to read the user's profile. It would typically allow access to basic profile information, including email address.

upvoted 2 times

✉ **emysa341** 3 months ago

v1.0/users?\$select=profilePhoto,mail is correccct

<https://learn.microsoft.com/en-us/graph/api/user-get?view=graph-rest-1.0&tabs=http#:~:text=to%20your%20query-,%24select%3DdisplayName%2CgivenName%2CpostalCode%2Cidentities,-Request>

upvoted 1 times

✉ **AzDeveloper** 3 months, 1 week ago

Answer is:

1. /v1.0/me

2. User.Read

profilePhoto is not valid name in graph api.

<https://learn.microsoft.com/en-us/graph/api/profilephoto-get?view=graph-rest-1.0&tabs=http>

upvoted 4 times

✉ **Swekker** 3 months, 1 week ago

Answer is:

Endpoint: /v1.0/users?\$select=profilePhoto,mail

Permission: User.ManageIdentities.All

1. You cannot use the /me endpoint because you need to list all the employees on the website.

2. Second /me endpoint wont return the user photo.

upvoted 4 times

✉ **de8a1d8** 1 month, 2 weeks ago

From what I understand, you only need to display the profile of the signed-in user, not all of the existing users at once. So /me should be enough.

upvoted 1 times

✉ **azure1282** 3 months, 1 week ago

I don't think it is users endpoint as we need the id or userPrincipalName in the endpoint e.g.

[https://graph.microsoft.com/v1.0/users/{id|userPrincipalName}?\\$select=id,employeeOrgData](https://graph.microsoft.com/v1.0/users/{id|userPrincipalName}?$select=id,employeeOrgData)

upvoted 1 times

✉ **Tralalaaz204** 3 months, 1 week ago

Do you have a document for this?

upvoted 1 times

Question #72

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. When you are ready to answer a question, click the Question button to return to the question.

Background -

Munson's Pickles and Preserves Farm is an agricultural cooperative corporation based in Washington, US, with farms located across the United States. The company supports agricultural production resources by distributing seeds, fertilizers, chemicals, fuel, and farm machinery to the farms.

Current Environment -

The company is migrating all applications from an on-premises datacenter to Microsoft Azure. Applications support distributors, farmers, and internal company staff.

Corporate website -

- The company hosts a public website located at <http://www.munsonspicklesandpreservesfarm.com>. The site supports farmers and distributors who request agricultural production resources.

Farms -

- The company created a new customer tenant in the Microsoft Entra admin center to support authentication and authorization for applications.

Distributors -

- Distributors integrate their applications with data that is accessible by using APIs hosted at <http://www.munsonspicklesandpreservesfarm.com/api> to receive and update resource data.

Requirements -

The application components must meet the following requirements:

Corporate website -

- The site must be migrated to Azure App Service.
- Costs must be minimized when hosting in Azure.

- Applications must automatically scale independent of the compute resources.
- All code changes must be validated by internal staff before release to production.
- File transfer speeds must improve, and webpage-load performance must increase.
- All site settings must be centrally stored, secured without using secrets, and encrypted at rest and in transit.
- A queue-based load leveling pattern must be implemented by using Azure Service Bus queues to support high volumes of website agricultural production resource requests.

Farms -

- Farmers must authenticate to applications by using Microsoft Entra ID.

Distributors -

- The company must track a custom telemetry value with each API call and monitor performance of all APIs.
- API telemetry values must be charted to evaluate variations and trends for resource data.

Internal staff -

- App and API updates must be validated before release to production.
- Staff must be able to select a link to direct them back to the production app when validating an app or API update.
- Staff profile photos and email must be displayed on the website once they authenticate to applications by using their Microsoft Entra ID.

Security -

- All web communications must be secured by using TLS/HTTPS.
- Web content must be restricted by country/region to support corporate compliance standards.
- The principle of least privilege must be applied when providing any user rights or process access rights.
- Managed identities for Azure resources must be used to authenticate services that support Microsoft Entra ID authentication.

Issues -

Corporate website -

- Farmers report HTTP 503 errors at the same time as internal staff report that CPU and memory usage are high.
- Distributors report HTTP 502 errors at the same time as internal staff report that average response times and networking traffic are high.
- Internal staff report webpage load sizes are large and take a long time to load.
- Developers receive authentication errors to Service Bus when they debug locally.

Distributors -

- Many API telemetry values are sent in a short period of time. Telemetry traffic, data costs, and storage costs must be reduced while preserving a statistically correct analysis of the data points sent by the APIs.

You need to configure all site configuration settings for the corporate website.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Create a managed identity.
- B. Update the role assignments for the Azure Key Vault.
- C. Create an Azure App Configuration store.
- D. Update the role assignments for the Azure App Configuration store.
- E. Create an Azure Key Vault.

Correct Answer: ABE*Community vote distribution*

ACD (80%)

10%

10%

✉️  **AzDeveloper**  3 months, 1 week ago**Selected Answer: ACD**

Why App Config?

Because is required to be: "All site settings must be centrally stored, secured without using secrets, and encrypted at rest and in transit."

Reference:

Does App Configuration encrypt my data?

<https://learn.microsoft.com/en-us/azure/azure-app-configuration/faq>

upvoted 5 times

✉️  **exeem** 3 months agoYes: <https://learn.microsoft.com/en-us/azure/azure-app-configuration/faq#does-app-configuration-encrypt-my-data>

upvoted 1 times

✉️  **FeriAZ**  2 months, 3 weeks ago**Selected Answer: ACD**

Create a managed identity:

Managed identities in Azure are used for Azure service resources, providing an identity for applications to use when connecting to other Azure services. This helps in securely accessing other Azure resources without having to manage credentials in your code.

Create an Azure App Configuration store:

Azure App Configuration provides a service to centrally manage application settings and feature flags. It's a great way to handle configurations that need to be shared across several components or environments.

Update the role assignments for the Azure App Configuration store:

Similar to the Key Vault, setting up appropriate role assignments for the Azure App Configuration store is about controlling access based on the principle of least privilege.

This action is crucial for ensuring that only authorized personnel or processes have access to the configuration data.

upvoted 3 times

✉️  **exeem** 3 months ago**Selected Answer: AC**

"All site settings must be centrally stored, secured without using secrets" - this excludes Azure Key Vault so we should focus on Azure App Config.

upvoted 1 times

✉️  **emysa341** 3 months ago

The correct options are ACE

"All site settings must be centrally stored, secured without using secrets"

This means the configs stored in the app config are not secrets.

upvoted 1 times

✉️  **my_nickname2** 3 months, 1 week ago**Selected Answer: BCE**

store app settings => App Config Store

secure Settings encrypted => KeyVault

access KeyVault => Update access assignments

upvoted 1 times

Question #73

Topic 4

You are developing an application that uses keys stored in Azure Key Vault.

You need to enforce a specific cryptographic algorithm and key size for keys stored in the vault.

What should you use?

- A. Secret versioning
- B. Azure Policy
- C. Key Vault Firewall
- D. Access policies

Correct Answer: B

Community vote distribution

B (100%)

✉️  **FeriAZ** 2 months, 3 weeks ago

Selected Answer: B

Azure Policy is a service in Azure that helps you enforce organizational standards and assess compliance at scale. With Azure Policy, you can apply policies on various resources, including Azure Key Vault, to ensure they comply with specific rules and requirements, like enforcing specific cryptographic algorithms and key sizes.

upvoted 1 times

✉️  **Tralalaaz204** 3 months, 1 week ago

Selected Answer: B

I agree.

See <https://learn.microsoft.com/en-us/azure/key-vault/policy-reference>

upvoted 3 times

✉️  **AzDeveloper** 3 months, 1 week ago

Selected Answer: B

B is correct

Reference: <https://learn.microsoft.com/en-us/azure/key-vault/general/azure-policy>

upvoted 1 times

✉️  **Ciupaz** 3 months, 1 week ago

Selected Answer: B

Azure Policy is the correct answer.

upvoted 1 times

Topic 5 - Question Set 5

Question #1

Topic 5

DRAG DROP -

You develop a web app that uses the tier D1 app service plan by using the Web Apps feature of Microsoft Azure App Service.

Spikes in traffic have caused increases in page load times.

You need to ensure that the web app automatically scales when CPU load is about 85 percent and minimize costs.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

NOTE: More than one order of answer choices is correct. You will receive credit for any of the correct orders you select.

Select and Place:

Actions

Answer Area

Configure the web app to the Premium App Service tier.

Configure the web app to the Standard App Service tier.

Enable autoscaling on the web app.

Add a Scale rule.

Switch to an Azure App Services consumption plan.

Configure a Scale condition.



Actions

Answer Area

Configure the web app to the Premium App Service tier.

Configure the web app to the Standard App Service tier.

Configure the web app to the Standard App Service tier.

Enable autoscaling on the web app.

Correct Answer:

Enable autoscaling on the web app.

Add a Scale rule.

Add a Scale rule.

Configure a Scale condition.



Step 1: Configure the web app to the Standard App Service Tier

The Standard tier supports auto-scaling, and we should minimize the cost.

Step 2: Enable autoscaling on the web app

First enable autoscale -

Step 3: Add a scale rule -

Step 4: Add a Scale condition -

Reference:

<https://docs.microsoft.com/en-us/azure/monitoring-and-diagnostics/monitoring-autoscale-get-started>

 **AmirCSUDH**  3 years, 5 months ago

The Provided Answer is Correct

upvoted 99 times

 **sgsvve** 2 years, 10 months ago

NO, NO, NO.... You can find it out yourself bij going to an appservice-> Scale Out -> Click Manual Scale (Now you have a default SCALE CONDITION) -> Click Add Scale Rule (Now you have a SCALE RULE)

upvoted 13 times

 **azurelearner666** 2 years, 10 months ago

No, you must create a condition (a default one is created for you) and then add scale rules.

upvoted 18 times

 **dhishkiyaau** 2 years, 7 months ago

the correct answer is :-

Configure the web app to the Standard App Service Tier

Enable autoscaling on the web app

Add a Scale condition

Add a scale rule

Please refer - <https://docs.microsoft.com/en-us/learn/modules/app-service-autoscale-rules/6-autoscale-a-web-app>

upvoted 46 times

 **altafpatel1984** 2 years, 4 months ago

No. This is for Monitor auto scale. Question is for Web App auto scale, and for that steps would be as follow:

Configure the web app to the Standard App Service Tier

Enable autoscaling on the web app

Add a Scale condition

Add a scale rule

Reference:

<https://docs.microsoft.com/en-us/learn/modules/app-service-autoscale-rules/6-autoscale-a-web-app>

upvoted 38 times

 **Mev4953** 2 years, 3 months ago

There is no "Add a scale condition". Do you mean "Configure a scale condition" ?

upvoted 3 times

 **Chiboy** 2 years, 2 months ago

It actually says " Add Scale Conditions" on the web app

upvoted 2 times

 **rkuifje** 2 years, 2 months ago

Seems you are right Altafpatel based on your link. Thanks.

upvoted 2 times

 **rkuifje** 2 years, 2 months ago

BTW, may be both answer are correct:

See question note: More than one order of answer choices is correct

upvoted 2 times

 **AbidooKing** 9 months, 1 week ago

Got this 2023July13. Score 917 and went with

Configure the web app to the Standard App Service Tier

Enable autoscaling on the web app

Add a Scale condition

Add a scale rule

upvoted 6 times

 **TEMPKAKAM**  3 years, 5 months ago

Configure the web app to the Standard App Service Tier

Enable autoscaling on the web app

Add a Scale condition

Add a scale rule

upvoted 40 times

profesorklaus 3 years, 2 months ago

There is nothing about adding scale condition. You have default scale condition so first configure it then add a rule (85%)
upvoted 5 times

Ciupaz **Most Recent** 4 months, 1 week ago

"Minimize costs" is the hint to choose Standard over Premium tier.
upvoted 1 times

kakajos 7 months ago

Upgrade to Standard App Service tier (it supports autoscaling)
Enable autoscaling
Add a scale condition
Create scale rule
Monitor autoscaling activity
upvoted 1 times

NombreFalso 1 year, 1 month ago

"NOTE: More than one order of answer choices is correct." Makes me think the order in which you add the rule or scale condition does not matter, nothing else could be in different order IMO.
upvoted 1 times

macobuzi 7 months, 3 weeks ago

Nice observation!
upvoted 1 times

OPT_001122 1 year, 5 months ago

Configure the web app to the Standard App Service Tier
Enable autoscaling on the web app
Add a Scale condition
Add a scale rule
upvoted 2 times

Eltooth 1 year, 9 months ago

Configure the web app to the Standard App Service Tier
Enable autoscaling on the web app
Add a Scale condition
Add a scale rule

<https://docs.microsoft.com/en-gb/learn/modules/scale-apps-app-service/4-autoscale-app-service>

upvoted 2 times

adilkhan 1 year, 2 months ago

100% correct!
upvoted 1 times

Evo_Morales 2 years ago

Sloppy answer set - this needs cleanup
upvoted 3 times

iamstudying 2 years, 1 month ago

answer is correct. les gettit, buddies. 6 figure here i come
upvoted 3 times

massnonn 2 years, 1 month ago

Once you enable autoscaling, you can edit the automatically created default scale condition, after you add scale rule
upvoted 1 times

sozturk88 2 years, 1 month ago

Configure the web app to the Standard App Service Tier
Enable autoscaling on the web app
Add a Scale condition
Add a scale rule

Reference:

<https://docs.microsoft.com/en-us/learn/modules/app-service-autoscale-rules/6-autoscale-a-web-app>

@altafpatel1984 is right
upvoted 1 times

oescm 2 years, 2 months ago

Got this one 02/2022. Went with most voted
upvoted 5 times

altafpatel1984 2 years, 4 months ago

Given answer is for Monitor auto scale. But question is for Web App auto scale and for this sequence will be as follow:

Configure the web app to the Standard App Service Tier

Enable autoscaling on the web app
Add a Scale condition
Add a scale rule

Reference: <https://docs.microsoft.com/en-us/learn/modules/app-service-autoscale-rules/6-autoscale-a-web-app>
upvoted 6 times

✉ **Pankaj78** 2 years, 5 months ago
Step 1: Configure the web app to the Standard App Service Tier
The Standard tier supports auto-scaling, and we should minimize the cost.

Step 2: Enable autoscaling on the web app

Step 4: Configure a Scale condition

Step 3: Add a scale rule
Scale rule is within the Scale condition.
upvoted 4 times

✉ **Jan91** 2 years, 5 months ago
Had this question in my exam this morning. It stated that there are 2 possible orders to answer. Most likely the scale condition and scale rule part is meant with this.
upvoted 3 times

✉ **catalene** 2 years, 6 months ago
I think, you only create "autoscaling rules app" with condition like ">85%" with Premium Plan. IT's Correct?
upvoted 1 times

✉ **zeugo** 4 months, 2 weeks ago
I think you need to read the documentation again ;)
upvoted 1 times

✉ **AOE** 2 years, 8 months ago
why not use consumption plan ?
upvoted 1 times

✉ **MontyKL** 2 years, 7 months ago
I have the same question. I think there is no such a thing as consumption plan for App service. Consumption plan is for Azure Functions. I might be wrong.
upvoted 4 times

✉ **the_odd_one** 6 months ago
Consumption plans are only for Azure functions
upvoted 1 times

Question #2

Topic 5

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution.

Determine whether the solution meets the stated goals.

You are developing and deploying several ASP.NET web applications to Azure App Service. You plan to save session state information and HTML output.

You must use a storage mechanism with the following requirements:

- ☞ Share session state across all ASP.NET web applications.
- ☞ Support controlled, concurrent access to the same session state data for multiple readers and a single writer.
- ☞ Save full HTTP responses for concurrent requests.

You need to store the information.

Proposed Solution: Enable Application Request Routing (ARR).

Does the solution meet the goal?

A. Yes

B. No

Correct Answer: B

Instead deploy and configure Azure Cache for Redis. Update the web applications.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/best-practices/caching#managing-concurrency-in-a-cache>

Community vote distribution

 B (100%)

 **gematsaljoa**  3 years, 4 months ago

B. NoNoNoNo

upvoted 31 times

 **Archimedes**  3 years, 2 months ago

The correct answer is NO. Application Request Routing is for controlling internet traffic in IIS using a proxy server. It doesn't deal with managing state information.

upvoted 18 times

 **clarionprogrammer** 3 years ago

As stated, "No" is the correct answer. Application Request Routing is for scaling and load balancing.

<https://www.iis.net/downloads/microsoft/application-request-routing>

upvoted 6 times

 **OPT_001122**  1 year, 5 months ago

Selected Answer: B

No - redis cache is the answer

upvoted 1 times

 **Eltooth** 1 year, 10 months ago

Selected Answer: B

B is correct answer. NO

upvoted 2 times

 **Eche_Ir73** 1 year, 11 months ago

Selected Answer: B

Es Correcto, es NO

upvoted 3 times

 **AzureDJ** 2 years, 1 month ago

B. No. ARR is not for session state. We typically use ARR module to host a reverse proxy on IIS.

upvoted 2 times

 **Freidrich** 2 years, 1 month ago

Selected Answer: B

The correct answer is B: No.

upvoted 2 times

 **Baskman** 2 years, 2 months ago

Correct answer: NO

Explanation: "Application Request Routing (ARR) [...] is a proxy-based routing module that forwards HTTP requests to content servers based on HTTP headers"

SOURCE: <https://docs.microsoft.com/en-us/iis/extensions/planning-for-arr/using-the-application-request-routing-module>

upvoted 1 times

 **glam** 2 years, 11 months ago

B. No...

upvoted 4 times

 **AzureAz204Fan** 2 years, 11 months ago

In Azure, you can store ASP.NET Session State in a SQL database, Azure Table Storage or an in-memory distributed cache such as REDIS.

<https://docs.microsoft.com/en-us/archive/msdn-magazine/2014/august/microsoft-azure-use-distributed-cache-in-microsoft-azure>

upvoted 6 times

 **pac1311** 3 years, 2 months ago

correct!

upvoted 3 times

Question #3

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution.

Determine whether the solution meets the stated goals.

You are developing and deploying several ASP.NET web applications to Azure App Service. You plan to save session state information and HTML output.

You must use a storage mechanism with the following requirements:

- ☞ Share session state across all ASP.NET web applications.
- ☞ Support controlled, concurrent access to the same session state data for multiple readers and a single writer.
- ☞ Save full HTTP responses for concurrent requests.

You need to store the information.

Proposed Solution: Deploy and configure an Azure Database for PostgreSQL. Update the web applications.

Does the solution meet the goal?

A. Yes

B. No

Correct Answer: B

Instead deploy and configure Azure Cache for Redis. Update the web applications.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/best-practices/caching#managing-concurrency-in-a-cache>

Community vote distribution

B (69%)

A (31%)

✉  **profesorklaus**  3 years, 2 months ago

The worst solution from a performance and scalability standpoint is to use a database backed session state provider.

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-aspnet-session-state-provider>

upvoted 26 times

✉  **ferut** 2 years, 11 months ago

Although Postgres will lack the performance compared to Redis, the solution will work. Any thought?

upvoted 8 times

✉  **lugospod** 2 years, 3 months ago

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-aspnet-session-state-provider>

Postgres is not listed as viable option on MS site. So I guess this is why it is not acceptable.

upvoted 8 times

✉  **warchoon** 1 year, 1 month ago

maybe worst but exists

[https://learn.microsoft.com/en-us/previous-versions/aspnet/ms178587\(v=vs.100\)](https://learn.microsoft.com/en-us/previous-versions/aspnet/ms178587(v=vs.100))

upvoted 2 times

✉  **venki21**  3 years, 4 months ago

But the question says, does it meet the goal? So using a postgres db to save the session does meet the goal. Any thoughts on this

upvoted 22 times

✉  **Pooochie** 3 years, 1 month ago

Got a similar question on udemy, where both solutions were marked as correct. Redis and Postgres

upvoted 8 times

✉  **ML333** 1 year, 10 months ago

Another brilliantly worded and completely unambiguous question, LOL, thanks MS! Questions should be clear (like is this a best practice or a recommended approach), because the question "does it meet the goal" is it does... but should you do this? ever? Never!

upvoted 3 times

✉  **SaneSage123** 8 months, 2 weeks ago

10000%

upvoted 1 times

✉  **Cornholioz** 3 years, 4 months ago

It will meet the goal, but my take is that it is a huge Overkill to introduce PostgreSQL.

upvoted 3 times

✉ **matejka** 3 years, 2 months ago

I think that it would be better to provide four options and you should pick the best one rather than four questions with YES/NO answers as there is no chance to review the further options (further questions).

upvoted 4 times

✉ **AshT1** 3 years, 2 months ago

"Save full HTTP responses" so PostgreSQL may not be a good choice.

upvoted 5 times

✉ **FeriAZ** **Most Recent** 2 months, 3 weeks ago

Selected Answer: B

Share Session State Across Applications: While PostgreSQL can store session state, sharing session state across multiple ASP.NET applications is not its primary use case. It requires additional configuration and programming to handle session state management.

Controlled, Concurrent Access for Multiple Readers and a Single Writer: PostgreSQL supports concurrent access, but managing this for session state data would require additional programming effort.

Save Full HTTP Responses for Concurrent Requests: PostgreSQL is not designed for caching full HTTP responses. It is primarily a relational database for structured data, not a cache for HTTP responses.

upvoted 1 times

✉ **dom271219** 6 months, 2 weeks ago

B. No

why they would choose this DB while Microsoft sells SQL DB?

upvoted 1 times

✉ **75daltonic** 8 months, 4 weeks ago

Selected Answer: B

I think No.

The question says: Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution.

There are three questions in ExamTopics about this:

1-Enable Application Request Routing (ARR).

2-Deploy and configure an Azure Database for PostgreSQL. Update the web applications.

3-Deploy and configure Azure Cache for Redis. Update the web applications.

And only one is the solution in order to the note at the beginning, so the best answer is Redis.

upvoted 7 times

✉ **ay_m** 8 months ago

Thanks for summing them up!

upvoted 2 times

✉ **JH81** 10 months ago

Selected Answer: B

MS is looking for Redis as the answer.

upvoted 1 times

✉ **thebarber87** 11 months, 3 weeks ago

Selected Answer: B

No - not best for the goal

upvoted 2 times

✉ **JamieS** 1 year ago

I'd do A, but B would work. Potentially poorly worded question if they want you to pick the best option.

upvoted 1 times

✉ **warchoon** 1 year, 1 month ago

ASP.NET session state is designed to enable you to store user session data in different sources. By default, session state values and information are stored in memory ... Another alternative is to store session data in a SQL Server database, where it can be shared by multiple Web servers.

Reference:

[https://learn.microsoft.com/en-us/previous-versions/aspnet/ms178587\(v=vs.100\)](https://learn.microsoft.com/en-us/previous-versions/aspnet/ms178587(v=vs.100))

upvoted 1 times

✉ **warchoon** 1 year, 1 month ago

=> <https://www.codeproject.com/Articles/18456/Sample-Session-Providers-with-PostgreSQL>

upvoted 1 times

✉ **alexine74** 1 year, 2 months ago

Selected Answer: B

B. No

An Azure Database for PostgreSQL is a fully-managed relational database service based on the PostgreSQL open-source database engine, it can be used to store data, but it's not designed to store session state information or HTML output, it will not meet the requirements of sharing

session state across all ASP.NET web applications, support controlled, concurrent access to the same session state data for multiple readers and a single writer, or save full HTTP responses for concurrent requests.

upvoted 3 times

✉ **rotimislaw** 1 year, 3 months ago

Selected Answer: B

The question description says "unique solution" so I'd say Redis is better and here the answer is B. NO

upvoted 2 times

✉ **mitu7** 1 year, 3 months ago

Selected Answer: B

Azure Database for PostgreSQL is a managed database service that can be used to store and retrieve data, but it is not designed for storing and managing session state information or HTML output. It does not provide the necessary features and capabilities to share session state across multiple web applications or support concurrent access to session state data.

upvoted 2 times

✉ **hubekpeter** 1 year, 4 months ago

Selected Answer: A

Yes, we used to do it that way when redis/memcached key value in memory store wasn't available.

upvoted 1 times

✉ **OPT_001122** 1 year, 5 months ago

Selected Answer: B

No - redis cache is the correct ans

upvoted 4 times

✉ **gmishra88** 1 year, 6 months ago

Providing both session state and output-caching is Redis. It is fine to keep session state in a DB or in CosmosDB, but the output caching is not a good idea. So, wait for Redis as the answer. But yes, all of this is a guess on what Microsoft thinks

upvoted 1 times

✉ **Knightie** 1 year, 7 months ago

Selected Answer: B

Deploy and configure an Azure Database for PostgreSQL. Update the web applications.

It didn't say storing the state or session or whatever else...

upvoted 3 times

✉ **Knightie** 1 year, 7 months ago

Instead deploy and configure Azure Cache for Redis. Update the web applications.

Compare to the Redis Default answer, it also didn't say saving state.

upvoted 1 times

✉ **Ami_Nou** 1 year, 10 months ago

Selected Answer: B

It says "Each question in the series contains a unique solution", so while this may work, deploying and configuring Azure Cache for Redis and updating the web applications is a better answer

upvoted 3 times

✉ **Amy_ant** 1 year, 7 months ago

"Each question in the series contains a unique solution" not "Each series of questions contains a unique (valid) solution" all they are stating is that that each question in the series is unique you are not answering the same question. A postgres db can be written to and read concurrently it can store session data.

Redis would be a faster solution, postgres would be more persistent.

upvoted 1 times

Question #4

Topic 5

HOTSPOT -

A company is developing a gaming platform. Users can join teams to play online and see leaderboards that include player statistics. The solution includes an entity named Team.

You plan to implement an Azure Redis Cache instance to improve the efficiency of data operations for entities that rarely change.

You need to invalidate the cache when team data is changed.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
void ClearCachedTeams()

{
    IDatabase cache = Connection.GetDatabase();
    ICache cache = Connection.GetDatabase();

    cache.KeyDelete("Team");
    cache.StringSet("Team", "");
    cache.ValueDelete("Team");
    cache.StringGet("Team", "");

    ViewBag.msg += "Team data removed from cache.";
}
```

Answer Area

```
void ClearCachedTeams()
```

```
{
```

```
    IDatabase cache = Connection.GetDatabase();
    ICache cache = Connection.GetDatabase();

    cache.KeyDelete("Team");
    cache.StringSet("Team", "");
    cache.ValueDelete("Team");
    cache.StringGet("Team", "");

    ViewBag.msg += "Team data removed from cache.";
}
```

Correct Answer:

Box 1: IDatabase cache = connection.GetDatabase();

Connection refers to a previously configured ConnectionMultiplexer.

Box 2: cache.StringSet("teams", "")

To specify the expiration of an item in the cache, use the TimeSpan parameter of StringSet. cache.StringSet("key1", "value1", TimeSpan.FromMinutes(90));

Reference:

<https://azure.microsoft.com/sv-se/blog/lap-around-azure-redis-cache-preview/> <https://docs.microsoft.com/en-us/cli/azure/webapp/config/container>

 **agueda**  3 years, 1 month ago

Same question on AZ203

Answer is:

```
IDatabase cache = Connection.GetDatabase();
cache.KeyDelete("teams")
```

<https://www.examtopics.com/discussions/microsoft/view/12706-exam-az-203-topic-5-question-9-discussion/>

upvoted 116 times

✉ **BrettusMaximus** 2 years, 11 months ago

It is KeyDelete. If you use StringSet you will set Teams to the value of Empty string and a future request to Teams will return an empty string and not null

upvoted 18 times

✉ **coffecold** 1 year, 6 months ago

Please see the GetFromList() in the tutorial.

if (String.IsNullOrEmpty(serializedTeams)) then the DB will be read instead of the cache.

Implementation with an empty string returned from cache if invalidated is no problem.

upvoted 2 times

✉ **MonkeyKing1024** 1 year, 1 month ago

```
https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-web-app-cache-aside-leaderboard#:~:text=void%20ClearCachedTeams()%0A%7B%0A%20%20%20%20IDatabase%20cache%20%3D%20Connection.GetDatabase()%3B%0A%20%20%20%20cache.KeyDelete(%22teamsList%22)%3B%0A%20%20%20%20cache.KeyDelete(%22teamsSortedSet%22)%3B%0A%20%20%20%20ViewBag.msg%20%2B%3D%20%22Team%20data%20removed%20from%20cache.%20%22%3B%0A%7D
```

upvoted 2 times

✉ **hobob** Highly Voted 3 years, 1 month ago

Answer is incorrect (for the 2nd part).

cache.KeyDelete() is the correct method for removing a key from the redis cache.

upvoted 32 times

✉ **Sukon_Desknot** 2 years, 7 months ago

Kindly provide references if you believe your knowledge is better than exam topics', we are not in the faith business 😊

upvoted 6 times

✉ **sas12321** 2 years, 7 months ago

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-web-app-cache-aside-leaderboard#update-the-teamscontroller-to-read-from-the-cache-or-the-database>

upvoted 22 times

✉ **Firo** Most Recent 8 months, 3 weeks ago

Answer is:

```
IDatabase cache = Connection.GetDatabase();
cache.KeyDelete("teams")
```

<https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-web-app-cache-aside-leaderboard>

upvoted 4 times

✉ **Kluk_Kluk** 1 year, 1 month ago

<https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-web-app-cache-aside-leaderboard>

void ClearCachedTeams()

{

```
IDatabase cache = Connection.GetDatabase();
cache.KeyDelete("teamsList");
cache.KeyDelete("teamsSortedSet");
ViewBag.msg += "Team data removed from cache. ";
```

}

upvoted 2 times

✉ **NeelParghi** 1 year, 2 months ago

KeyDelete Delete the key/value.

Ref: <https://learn.microsoft.com/en-us/training/modules/develop-for-azure-cache-for-redis/4-interact-redis-api>

upvoted 2 times

✉ **rolling_potato_** 1 year, 3 months ago

Was on my exam today (03-01-2023) I went with
IDatabase cache = Connection.GetDatabase();

cache.KeyDelete("teams")

upvoted 3 times

✉ **st0rmtrooperx** 1 year, 4 months ago

Got this on Dec 16th, 2022. I scored 921 and used given answer.

upvoted 2 times

✉ **coffecold** 1 year, 6 months ago

It is stringset, you want to have the key with the team name kept in your session.

Please see the GetFromList() code in the tutorial (<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-web-app-cache-aside-leaderboard>)

if (String.IsNullOrEmpty(serializedTeams)) then the DB will be read instead of the cache and the cache is set again with string set. Implementation with an empty string returned from cache if invalidated is no problem.

upvoted 2 times

✉ **serpevi** 1 year, 7 months ago

Got it on 09/2022, went with: IDatabase cache = connection.GetDatabase(); cache.KeyDelete("teams", ""); Score 927

upvoted 3 times

✉ **Eltooth** 1 year, 9 months ago

```
IDatabase cache = Connection.GetDatabase();
cache.KeyDelete("Team");
```

upvoted 2 times

✉ **Illeroiv** 1 year, 10 months ago

```
IDatabase cache = Connection.GetDatabase();
cache.KeyDelete("Team");
```

upvoted 1 times

✉ **liuliangzhou** 1 year, 10 months ago

Answer is:

```
IDatabase cache = Connection.GetDatabase();
cache.StringSet("Team")
```

Per the subject mentioned: You need to invalidate the cache when team data is changed.

It is changing value, not delete the object. we can update value for next time with no need re-create "Team" object.

<https://azure.microsoft.com/sv-se/blog/lat-around-azure-redis-cache-preview/>

upvoted 6 times

✉ **meoukg** 2 years, 1 month ago

Got it on 03/2022, I chose as below:

1. IDatabase cache = connection.GetDatabase();
2. cache.KeyDelete("teams", "")

upvoted 3 times

✉ **leonidn** 2 years, 3 months ago

```
IDatabase cache = Connection.GetDatabase();
cache.KeyDelete("teamsList");
```

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-web-app-cache-aside-leaderboard#update-the-teamscontroller-to-read-from-the-cache-or-the-database>

upvoted 3 times

✉ **BeshoyRomany** 2 years, 6 months ago

The Correct answer is

```
IDatabase cache = Connection.GetDatabase();
cache.KeyDelete("teams")
```

We don't need to use setTeams with "TimeSpan.FromMinutes(90)" becuse the requirement is:

--> You need to invalidate the cache when "team data is changed"

When data is changed only

So we have to cache.KeyDelete("teams") to clear the cach once the "team data is changed"

upvoted 3 times

✉ **myuv1131** 2 years, 7 months ago

The ClearCachedTeams method removes any cached team statistics from the cache.

```
void ClearCachedTeams()
{
    IDatabase cache = Connection.GetDatabase();
    cache.KeyDelete("teamsList");
    cache.KeyDelete("teamsSortedSet");
    ViewBag.msg += "Team data removed from cache. ";
}
```

upvoted 3 times

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-web-app-cache-aside-leaderboard#:~:text=The%20ClearCachedTeams%20method%20removes%20any%20cached%20team%20statistics%20from%20the%20cache.>

upvoted 1 times

✉ **kondapaturi** 2 years, 9 months ago

IDatabase cache = Connection.GetDatabase(); - we need to use the IDatabase interface

cache.KeyDelete("teams"); - To invalidate or remove the cache item, we can simply delete the key

upvoted 3 times

Question #5

DRAG DROP -

A company has multiple warehouses. Each warehouse contains IoT temperature devices which deliver temperature data to an Azure Service Bus queue.

You need to send email alerts to facility supervisors immediately if the temperature at a warehouse goes above or below specified threshold temperatures.

Which five actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

Add a logic app trigger that fires when one or more messages arrive in the queue.

Add a Recurrence trigger that schedules the app to run every 15 minutes.

Add an action that sends an email to specified personnel if the temperature is outside of those thresholds.

Add a trigger that reads IoT temperature data from a Service Bus queue.

Add a logic app action that fires when one or more messages arrive in the queue.

Add a condition that compares the temperature against the upper and lower thresholds.

Create a blank Logic app.

Add an action that reads IoT temperature data from the Service Bus queue.

Answer Area

Correct Answer:

Actions

Add a logic app trigger that fires when one or more messages arrive in the queue.

Add a Recurrence trigger that schedules the app to run every 15 minutes.

Add an action that sends an email to specified personnel if the temperature is outside of those thresholds.

Add a trigger that reads IoT temperature data from a Service Bus queue.

Add a logic app action that fires when one or more messages arrive in the queue.

Add a condition that compares the temperature against the upper and lower thresholds.

Create a blank Logic app.

Add an action that reads IoT temperature data from the Service Bus queue.

Answer Area

Create a blank Logic app.

Add a logic app action that fires when one or more messages arrive in the queue.

Add an action that reads IoT temperature data from the Service Bus queue.

Add a condition that compares the temperature against the upper and lower thresholds.

Add an action that sends an email to specified personnel if the temperature is outside of those thresholds.

Step 1: Create a blank Logic app.

Create and configure a Logic App.

Step 2: Add a logical app trigger that fires when one or more messages arrive in the queue.

Configure the logic app trigger.

Under Triggers, select When one or more messages arrive in a queue (auto-complete).

Step 3: Add an action that reads IoT temperature data from the Service Bus queue

Step 4: Add a condition that compares the temperature against the upper and lower thresholds.

Step 5: Add an action that sends an email to specified personnel if the temperature is outside of those thresholds

Reference:

<https://docs.microsoft.com/en-us/azure/iot-hub/iot-hub-monitoring-notifications-with-azure-logic-apps>

 **agueda** Highly Voted 3 years, 1 month ago

The 2nd step sholud be "logic app trigger". You need a trigger to fire the logic app

Ref: <https://docs.microsoft.com/en-us/azure/iot-hub/iot-hub-monitoring-notifications-with-azure-logic-apps>

upvoted 83 times

 **debanjan10** Highly Voted 2 years, 6 months ago

Order:

1. Create blank logic app
2. Add logic app trigger when one or more messages arrive in the queue
3. Action to read IoT Temp data
4. Compare upper and lower temp
5. Send emails

upvoted 57 times

 **Knightie** 1 year, 7 months ago

trigger instead of action, trigger is something kick starts asking someone to do the job..

upvoted 1 times

 **naivecoder786** Most Recent 1 year, 1 month ago

Logic Apps are out of Scope from AZ-204 now !

upvoted 10 times

 **warchoon** 1 year, 1 month ago

I think the trigger should bee after the action it uses. So 1 after 6.

Agree with the rest

upvoted 1 times

 **warchoon** 1 year, 1 month ago

Sorry

action should be created after trigger

<https://learn.microsoft.com/en-us/azure/logic-apps/quickstart-create-first-logic-app-workflow>

So the explanation is correct

upvoted 2 times

 **OPT_001122** 1 year, 4 months ago

this is a freaking question.

i try to remember it like a sentence

Blank Logic Arrives in the Q Acts Compares Sends

Order:

1. Blank Logic
2. Arrive in the Queue
3. Action
4. Compare
5. Send

upvoted 5 times

 **Lucky_me** 1 year, 10 months ago

The explanation is correct, the picture is incorrect.

upvoted 2 times

 **OPT_001122** 1 year, 4 months ago

Blank Logic Arrives in the Q Acts to Read Compares Sends

upvoted 1 times

 **liuliangzhou** 1 year, 10 months ago

Order: create app -> trigger -> action -> condition -> do

1. Create blank logic app
2. Add logic app trigger when one or more messages arrive in the queue
3. Action to read IoT Temp data
4. Compare upper and lower temp
5. Send emails

upvoted 5 times

 **ivan0590** 2 years ago

Admin, please correct the given answer.

The second step is wrong, is a "logic app trigger", not an action.

Even the explanation says is a trigger!

upvoted 1 times

✉ **kondapaturi** 2 years, 9 months ago

First , we need to create a blank logic app.
Then we need to add a trigger to the logic app which will fire when one or more messages arrives in the queue.
Then we need an action to read the IoT temperature data.
Then we need to have a condition that compares the temperature against the upper and lower thresholds.
And finally, we have to have an action that sends an email to the specified personnel.

upvoted 14 times

✉ **insanewriters** 2 years, 10 months ago

Logic apps are supposed to be out of scope as of March 2021. Can anybody confirm they have gotten this question recently?

upvoted 15 times

✉ **glam** 2 years, 11 months ago

Only Step 2 should be - Logic app triggers When one or more messages arrive in a queue .

Rest correct.

upvoted 5 times

✉ **Gajendran** 3 years ago

Steps and Answers are wrong. Steps given below the Answers are correct. Setp 2 should be - Logic app triggers When one or more messages arrive in a queue . Refer - <https://docs.microsoft.com/en-us/azure/iot-hub/iot-hub-monitoring-notifications-with-azure-logic-apps>

upvoted 2 times

✉ **Deputy_Cartman** 3 years ago

What is with this website and posting incorrect answers to questions and leaving them here?

upvoted 4 times

✉ **catsforthewin** 3 years ago

What is the correct answer for this question?

upvoted 2 times

✉ **jvyas** 3 years ago

Only step 2 needs to be replaced by " Add a trigger that fires when item arrives in queue" the rest of the steps are correct and in the same order as the provided answer.

upvoted 4 times

✉ **hobob** 3 years, 1 month ago

Answer is incorrect for step 2. It should be to add a Service Bus Trigger when items arrive in a queue (ref: <https://docs.microsoft.com/en-us/azure/connectors/connectors-create-api-servicebus#add-service-bus-trigger>)

upvoted 10 times

✉ **pmsiva** 3 years, 1 month ago

You are correct. Logic app must start with a Trigger.

upvoted 3 times

✉ **pmsiva** 3 years, 1 month ago

Add the RSS trigger

Every logic app must start with a trigger, which fires when a specific event happens, or when a specific condition is met. Each time the trigger finds new items, it fires and the Logic Apps engine creates a logic app instance that starts and runs your workflow. If the trigger doesn't find new items, the trigger doesn't fire and doesn't create an instance or run the workflow at this check.

upvoted 1 times

Question #6

DRAG DROP -

You develop an ASP.NET Core MVC application. You configure the application to track webpages and custom events.

You need to identify trends in application usage.

Which Azure Application Insights Usage Analysis features should you use? To answer, drag the appropriate features to the correct requirements.

Each feature may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Features	Answer Area	Requirement	Feature
Users		Which pages visited by users most often correlate to a product purchase?	Feature
Funnels		How does load time of the product display page affect a user's decision to purchase a product?	Feature
Impact		Which events most influence a user's decision to continue to use the application?	Feature
Retention		Are there places in the application that users often perform repetitive actions?	Feature
User Flows			

Correct Answer:

Features	Answer Area	Requirement	Feature
Users		Which pages visited by users most often correlate to a product purchase?	Users
Funnels		How does load time of the product display page affect a user's decision to purchase a product?	Impact
Impact		Which events most influence a user's decision to continue to use the application?	Retention
Retention		Are there places in the application that users often perform repetitive actions?	User Flows
User Flows			

Box 1: Users -

Box 2: Impact -

One way to think of Impact is as the ultimate tool for settling arguments with someone on your team about how slowness in some aspect of your site is affecting whether users stick around. While users may tolerate a certain amount of slowness, Impact gives you insight into how best to balance optimization and performance to maximize user conversion.

Box 3: Retention -

The retention feature in Azure Application Insights helps you analyze how many users return to your app, and how often they perform particular tasks or achieve goals. For example, if you run a game site, you could compare the numbers of users who return to the site after losing a game with the number who return after winning. This knowledge can help you improve both your user experience and your business strategy.

Box 4: User flows -

The User Flows tool visualizes how users navigate between the pages and features of your site. It's great for answering questions like:

- ⇒ How do users navigate away from a page on your site?
- ⇒ What do users click on a page on your site?
- ⇒ Where are the places that users churn most from your site?
- ⇒ Are there places where users repeat the same action over and over?

Incorrect Answers:

Funnel: If your application involves multiple stages, you need to know if most customers are progressing through the entire process, or if they are ending the process at some point. The progression through a series of steps in a web application is known as a funnel. You can use Azure Application Insights Funnels to gain insights into your users, and monitor step-by-step conversion rates.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/usage-impact>

✉  **Dinima**  3 years, 1 month ago

1st one Funnels
upvoted 57 times

✉  **clarionprogrammer** 3 years ago

The end-goal of a 'Funnel' is a product purchase. 'Users' is just about how users use your app.
upvoted 9 times

✉  **renzoku** 1 year, 9 months ago

but you need to know which page is visited, with funnel you have to know it previously (what specific page) and get how many people are using this page, I think that probably could be "Users" because you have to verify N° users that visited your pages
upvoted 1 times

✉  **KingChuang** 1 year, 3 months ago

On my exam 2022-12-26.
Chose:
Funnels
Impact
Retention
User Flows
upvoted 11 times

✉  **SachinV** 5 months ago

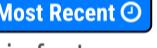
Agreed it's correct, in Nov182023, scored 962
upvoted 4 times

✉  **TonyMel** 1 year ago

correct, in 2023Mar24, score: 904/1000.
upvoted 4 times

✉  **kondapaturi**  2 years, 9 months ago

1.Funnels
2.Impact
3.Retention
4.User flow
upvoted 39 times

✉  **FeriAZ**  2 months, 3 weeks ago

Funnels: This feature will allow you to see how users progress through a series of steps (like viewing pages) leading up to a purchase.
Impact: This will help you understand the correlation between the performance metric (load time) and the user action (purchasing).
Retention: By analyzing when users return to the app and what they do, you can infer what keeps them coming back.
User Flows: This will show you the paths users take through your application and can highlight repetitive patterns in their journey.
upvoted 2 times

✉  **Ciupaz** 2 months, 3 weeks ago

FIRU to remember.
upvoted 2 times

✉  **JH81** 9 months, 3 weeks ago

Got this on 6/28/2023 and passed with 850. Went with highly voted answer.
upvoted 5 times

✉  **aragones** 11 months, 1 week ago

Got this 2023-05-12.

Had cases:

case: You need to configure authorization.
case: You need to ensure the app does not time out and processes the blob data
case study: VanArsdel Inc Canada
upvoted 3 times

✉  **OPT_001122** 1 year, 5 months ago

1.Funnels
2.Impact
3.Retention
4.User flow
upvoted 5 times

✉  **gmishra88** 1 year, 6 months ago

The answers given are fully correct. Absolutely sure.
upvoted 1 times

✉  **gmishra88** 1 year, 6 months ago

Except the first one, which I'm not sure:
Users it says : "Find out when people use your web app, what pages they're most interested in"
Funnels it says: "You can use Application Insights Funnels to gain insights into your users, and monitor step-by-step conversion rates."

Pages most interested: Users
Conversion/Buying product: Funnel

After deep meditation I will bet on Funnel is what microsoft guy thought as the right answer. Or it might be mean to increase mental health of people so that they meditate

upvoted 5 times

✉ **serpevi** 1 year, 7 months ago

Got it on 09/2022, went with 1.Funnels 2.Impact 3.Retention 4.User flow. Score 927

upvoted 2 times

✉ **Knightie** 1 year, 7 months ago

the answer is correct.. they all match the keywords of the individual paragraph.

1. no of visit.
2. loading time.
3. user coming back.
4. repeated page flow. (Are there places where users repeat the same action over and over?)

upvoted 3 times

✉ **ivan0590** 2 years ago

After some investigation, I think the given answer is correct.

The question is "Which pages visited by users most often correlate to a product purchase?".

So, if I'm right, they want to know how often visiting a page leads to a purchase and in which of those pages this happens more often. Let's say if I have a web app with a catalogue of 500 products and every product has a detail page from which the user has the option to purchase the product. Now, I want to know the top 10 product detail pages that ended up leading to a purchase.

Maybe I'm wrong, but I don't think that Funnels would be appropriate here. When you create a funnel, you have to specify the steps of the flow. I think that the main goal of a funnel is to analyse the transitions between steps in order to check if users get stuck in a step. Or in other words, how many people completed the flow.

Funnels are to be used in known and specific flows, like when a user signup. That's not what the question is about. The question is about knowing what pages tend to lead more to a purchase. This is not a specific flow, you don't know the exact steps.

upvoted 6 times

✉ **petitbilly** 2 years, 1 month ago

Got it in exam 03/22

upvoted 2 times

✉ **AzureDJ** 2 years, 1 month ago

Did you pass and did you go with?

- 1.Funnels (Funnels correlate to purchases, not general user events)
- 2.Impact
- 3.Retention
- 4.User flow

upvoted 2 times

✉ **Manivannan19** 2 years, 1 month ago

1. Users
2. Impact
3. Retention
4. User Flows

upvoted 2 times

✉ **barretowski** 2 years, 1 month ago

ta errado

upvoted 2 times

✉ **ytingyeu** 2 years, 1 month ago

Doesn't what Funnels care about is "if users stop somewhere during the whole process", not the correlation between pages and products?

For example, creating a customer ticket takes 5 steps and Funnel shows lots of visitors give up creating one at the 4th step. Then we know the UX in the 4th step might need some enhancement.

upvoted 1 times

✉ **leonidn** 2 years, 3 months ago

The Users feature does not allow to analyze flows. Funnels allows analyzing conversion rate through the specified flow. It does not allow us to find out what led to the final step. Users flow offers a means to identify user behavior. That is correlate the most to the first point.

1. UserFlow
2. Impact
3. Retention
4. UserFlow

upvoted 2 times

✉ **phvogel** 2 years, 5 months ago

First one funnels: <https://docs.microsoft.com/en-us/azure/azure-monitor/app/usage-funnels>

To quote: "Before you create your funnel, decide on the question you want to answer. For example, you might want to know how many users are viewing the home page, viewing a customer profile, and creating a ticket." Just replace "creating a ticket" with "buying a product"

upvoted 7 times

 **nonoss** 2 years, 6 months ago

- 1st one can't be "users", because users always returns a number of users verifying a certain condition, see the screenshot here
<https://docs.microsoft.com/en-us/azure/azure-monitor/app/usage-segmentation>

upvoted 1 times

Question #7

You develop a gateway solution for a public facing news API. The news API back end is implemented as a RESTful service and uses an OpenAPI specification.

You need to ensure that you can access the news API by using an Azure API Management service instance.

Which Azure PowerShell command should you run?

- A. Import-AzureRmApiManagementApi -Context \$ApiMgmtContext -SpecificationFormat "Swagger" -SpecificationPath \$SwaggerPath -Path \$Path
- B. New-AzureRmApiManagementBackend -Context \$ApiMgmtContext -Url \$Url -Protocol http
- C. New-AzureRmApiManagement -ResourceGroupName \$ResourceGroup -Name \$Name -Location \$Location -Organization \$Org -AdminEmail \$AdminEmail
- D. New-AzureRmApiManagementBackendProxy -Url \$ApiUrl

Correct Answer: D

New-AzureRmApiManagementBackendProxy creates a new Backend Proxy Object which can be piped when creating a new Backend entity.

Example: Create a Backend Proxy In-Memory Object

```
PS C:\>$secpassword = ConvertTo-SecureString "PlainTextPassword" -AsPlainText -Force
PS C:\>$proxyCreds = New-Object System.Management.Automation.PSCredential ("foo", $secpassword)
PS C:\>$credential = New-AzureRmApiManagementBackendProxy -Url "http://12.168.1.1:8080" -ProxyCredential $proxyCreds
PS C:\>$apimContext = New-AzureRmApiManagementContext -ResourceGroupName "Api-Default-WestUS" -ServiceName "contoso"
PS C:\>$backend = New-AzureRmApiManagementBackend -Context $apimContext -BackendId 123 -Url 'https://contoso.com/awesomeapi' -Protocol http -Title "first backend" -SkipCertificateChainValidation $true -Proxy $credential -Description "backend with proxy server"
```

Creates a Backend Proxy Object and sets up Backend

Incorrect Answers:

- A: The Import-AzureRmApiManagementApi cmdlet imports an Azure API Management API from a file or a URL in Web Application Description Language (WADL), Web Services Description Language (WSDL), or Swagger format.
- B: New-AzureRmApiManagementBackend creates a new backend entity in API Management.
- C: The New-AzureRmApiManagement cmdlet creates an API Management deployment in Azure API Management.

Reference:

<https://docs.microsoft.com/en-us/powershell/module/azurerm.apimanagement/new-azurermapimanagementbackendproxy?view=azurerm-powershell-6.13.0>

Community vote distribution

A (71%) C (25%) 4%

 **Vano6k**  3 years, 2 months ago

I think the correct answer is A because it is the one that takes advantage of the swagger definition of the API?

upvoted 83 times

 **pac1311** 3 years, 2 months ago

Think so too, the accepted answer has to be wrong.

upvoted 1 times

 **Eduary2015** 3 years, 1 month ago

It does not say that the API has swagger, D is correct.

upvoted 3 times

 **fadikh** 3 years, 1 month ago

Swagger and OpenAPI specs are the same

upvoted 14 times

 **trance13** 3 years ago

They are not the same.

upvoted 2 times

 **dhishkiyaau** 2 years, 7 months ago

OpenAPI = The specification itself, formerly known as Swagger specification. Swagger = Tools used in the implementation of OpenAPI.

upvoted 12 times

✉ **zero0** 2 years, 10 months ago

Totally agree.

find the 2 line code here.

<https://docs.microsoft.com/en-us/powershell/module/azurerm.apimanagement/import-azurermapiManagementapi?view=azurermmps-6.13.0>

upvoted 4 times

✉ **MiraA** 2 years, 6 months ago

There is an example of importing an API from an Open API Link using Import-AzApiManagementApi cmdlet:

<https://docs.microsoft.com/en-us/powershell/module/az.apimanagement/import-azapimanagementapi#example-4--import-an-api-from-a-open-api-link>

And as mentioned here:

<https://docs.microsoft.com/en-us/powershell/module/azurerm.apimanagement/import-azurermapiManagementapi>

the Import-AzureRmApiManagementApi cmdlet is obsoleted by Import-AzApiManagementApi cmdlet.

So I would select answer A.

upvoted 7 times

✉ **10x Highly Voted** 3 years, 2 months ago

D should be the answer - it is about accessing API - not creating it (eg based on swagger)

upvoted 15 times

✉ **XYZ2** 3 years ago

but you should have one first, that's why you need to create it before able to use

upvoted 1 times

✉ **borfavor** 2 years, 11 months ago

"The news API back end is implemented as a RESTful service and uses an OpenAPI specification." This clearly tells us the API already exists

upvoted 4 times

✉ **asdasdasg2** 2 years, 3 months ago

It exists in reality, but not on API management. API management API != backend API

upvoted 6 times

✉ **katrang Most Recent** 7 months, 1 week ago

A

When importing certain APIs, API Management configures the API backend automatically. For example, API Management configures the backend web service when importing:

An OpenAPI specification.

A SOAP API.

Azure resources, such as an HTTP-triggered Azure Function App or Logic App.

upvoted 1 times

✉ **uffuchs1** 1 year, 1 month ago**Selected Answer: A**

Since we already have an API back end, we need to import to Azure API Management.

<https://docs.microsoft.com/en-us/azure/api-management/scripts/powershell-import-api-and-add-to-product?toc=/powershell/module/toc.json>

upvoted 2 times

✉ **Sriniv** 1 year, 2 months ago

Answer is B reason

To configure an Azure API Management instance for a RESTful service that uses an OpenAPI specification, you can use the New-AzApiManagementBackend cmdlet to define the backend service and the Set-AzApiManagementApi cmdlet to add an API to the API Management instance. Here's an example Azure PowerShell command:

`New-AzApiManagementBackend -ResourceGroupName "myResourceGroup" -Name "myBackend" -Url "https://news-api.example.com" -Protocol "https" -Title "News API"``Set-AzApiManagementApi -Context "myApiManagementInstance" -SpecificationFormat "OpenApi" -SpecificationUrl "https://news-api.example.com/openapi.json" -Path "/news" -DisplayName "News API" -BackendUrl "https://news-api.example.com" -BackendProtocol "https" -BackendId "myBackend"`

upvoted 1 times

✉ **Sriniv** 1 year, 2 months ago

I think answer is b

`New-AzApiManagementBackend -ResourceGroupName "myResourceGroup" -Name "myBackend" -Url "https://news-api.example.com" -Protocol "https" -Title "News API"``Set-AzApiManagementApi -Context "myApiManagementInstance" -SpecificationFormat "OpenApi" -SpecificationUrl "https://news-api.example.com/openapi.json" -Path "/news" -DisplayName "News API" -BackendUrl "https://news-api.example.com" -BackendProtocol "https" -BackendId "myBackend"`

This command creates a new backend service named "myBackend" that points to the URL of the news API, and then creates a new API in the specified Azure API Management instance that references the OpenAPI specification for the news API. The API is mapped to the "/news" path and is associated with the backend service. The BackendId parameter specifies the ID of the backend service that was created earlier.

upvoted 1 times

✉ **OPT_001122** 1 year, 4 months ago

Selected Answer: A

A. Import-AzureRmApiManagementApi -Context \$ApiMgmtContext -SpecificationFormat "Swagger" -SpecificationPath \$SwaggerPath -Path \$Path

upvoted 2 times

✉ **P000JAAAAAA** 1 year, 3 months ago

not sure A or C ..

New-AzApiManagement -Name "myapim" -ResourceGroupName "myResourceGroup" -Location "West US" -Organization "Contoso" -AdminEmail "admin@contoso.com" Reference: <https://learn.microsoft.com/en-us/azure/api-management/powershell-create-service-instance>

upvoted 1 times

✉ **elequiel** 1 year, 5 months ago

Selected Answer: A

Letter A because you need to create using OpenAPI

upvoted 1 times

✉ **OPT_001122** 1 year, 5 months ago

Selected Answer: A

correct answer is A

upvoted 1 times

✉ **coffecold** 1 year, 6 months ago

-AzureRmApiManagement can be *-AzApiManagement in future exams
AzureRM will be replaced by AZ Powershell.

upvoted 2 times

✉ **gmishra88** 1 year, 6 months ago

Seems like Microsoft has no review mechanism. Look at how the vote is so divided. Microsoft should make this an open book exam. They can be assured people still will not get good scores. Isn't that their intention?

upvoted 2 times

✉ **macobuzi** 7 months, 3 weeks ago

More people will have to retake the exam, more money for Microsoft!

upvoted 2 times

✉ **Satish_Babu** 1 year, 8 months ago

Selected Answer: D

D is correct

upvoted 1 times

✉ **HumbleYolo** 1 year, 9 months ago

Selected Answer: C

whizlab has c

upvoted 4 times

✉ **vavra** 2 years ago

Selected Answer: C

It is a creation, and docs <https://docs.microsoft.com/en-us/powershell/module/azurerm.apimanagement/new-azurermapimanagement?view=azurermps-6.13.0> says there are obligatory params

-ResourceGroupName <String>
-Name <String>
-Location <String>
-Organization <String>
-AdminEmail <String>

It must be C.

upvoted 2 times

✉ **SivajiTheBoss** 2 years, 1 month ago

Selected Answer: A

A seems more promising

upvoted 1 times

✉ **ReniRechner** 2 years, 1 month ago

Selected Answer: A

- A: Import also creates. Importing also creates an already configured instance (opposed to C)
- B: We don't need another backend, we need a frontend for an existing backend
- C: creates a new APIM, but you would still need to manually configure it.
- D: is meant to proxy a public API (<https://petri.com/how-to-use-the-azure-api-management-to-proxy-a-public-api>)

=> A looks most promising

upvoted 7 times

 **fearoffree** 2 years, 2 months ago

It seems the trick here is to have the knowledge of API Proxy that could be used for verifying the API availability without a need to call the backend APIs. But the question is very vague!

upvoted 1 times

Question #8

You are creating a hazard notification system that has a single signaling server which triggers audio and visual alarms to start and stop. You implement Azure Service Bus to publish alarms. Each alarm controller uses Azure Service Bus to receive alarm signals as part of a transaction. Alarm events must be recorded for audit purposes. Each transaction record must include information about the alarm type that was activated.

You need to implement a reply trail auditing solution.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Assign the value of the hazard message SessionID property to the ReplyToSessionId property.
- B. Assign the value of the hazard message MessageId property to the DeliveryCount property.
- C. Assign the value of the hazard message SessionID property to the SequenceNumber property.
- D. Assign the value of the hazard message MessageId property to the CorrelationId property.
- E. Assign the value of the hazard message SequenceNumber property to the DeliveryCount property.
- F. Assign the value of the hazard message MessageId property to the SequenceNumber property.

Correct Answer: AD

D: CorrelationId: Enables an application to specify a context for the message for the purposes of correlation; for example, reflecting the MessageId of a message that is being replied to.

A: ReplyToSessionId: This value augments the ReplyTo information and specifies which SessionId should be set for the reply when sent to the reply entity.

Incorrect Answers:

B, E: DeliveryCount -

Number of deliveries that have been attempted for this message. The count is incremented when a message lock expires, or the message is explicitly abandoned by the receiver. This property is read-only.

C, F: SequenceNumber -

The sequence number is a unique 64-bit integer assigned to a message as it is accepted and stored by the broker and functions as its true identifier. For partitioned entities, the topmost 16 bits reflect the partition identifier. Sequence numbers monotonically increase and are gapless. They roll over to 0 when the 48-

64 bit range is exhausted. This property is read-only.

Reference:

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messages-payloads>

Community vote distribution

AD (100%)

 **Kitkit**  3 years, 2 months ago

Answer is correct.

ReplyToSessionId (reply-to-group-id) This value augments the ReplyTo information and specifies which SessionId should be set for the reply when sent to the reply entity.

CorrelationId (correlation-id) Enables an application to specify a context for the message for the purposes of correlation; for example, reflecting the MessageId of a message that is being replied to.

upvoted 42 times

 **rustycables**  2 years, 8 months ago

I have no clue what this is about, but I get the feeling we are supposed to arrive at the correct answer by elimination.

upvoted 29 times

 **coffecold** 1 year, 6 months ago

You're right : DeliveryCount, SequenceNumber are read only, two remaining.

upvoted 6 times

 **uffuchs**  1 year, 1 month ago

I believe A + D is correct

upvoted 1 times

 **Sriniv** 1 year, 2 months ago

Assign the value of the hazard message `MessageId` property to the `CorrelationId` property: When an alarm event is published to the Service Bus, set the `MessageId` property to a unique value that identifies the message. Then, when the alarm controller receives the message, set the `CorrelationId` property to the value of the `MessageId` property. This enables the controller to associate the message with the corresponding transaction record for auditing purposes.

upvoted 2 times

 **Sriniv** 1 year, 2 months ago

Assign the value of the hazard message `SessionId` property to the `ReplyToSessionId` property: When an alarm event is published to the Service Bus, set the `SessionId` property to a unique value that identifies the session. Then, when the alarm controller receives the message, set the `ReplyToSessionId` property to the value of the `SessionId` property. This enables the controller to reply to the message and maintain the context of the session.

upvoted 1 times

 **Sriniv** 1 year, 2 months ago

Therefore, the two actions that you should perform are A and D:

- A. Assign the value of the hazard message `SessionId` property to the `ReplyToSessionId` property.
- D. Assign the value of the hazard message `MessageId` property to the `CorrelationId` property.

upvoted 1 times

 **alexein74** 1 year, 2 months ago

Selected Answer: AD

Answer is correct.

upvoted 2 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: AD

Answer is correct.

upvoted 3 times

 **Praks13** 1 year, 7 months ago

AD

<https://docs.microsoft.com/en-us/training/modules/discover-azure-message-queue/5-messages-payloads-serialization>

upvoted 1 times

 **meoukg** 2 years, 1 month ago

<https://www.examtopics.com/discussions/microsoft/view/44385-exam-az-204-topic-4-question-11-discussion/>

Got it on 03/2022, I chose as below:

- A. Assign the value of the hazard message `SessionId` property to the `ReplyToSessionId` property.
- D. Assign the value of the hazard message `MessageId` property to the `CorrelationId` property.

upvoted 2 times

 **ScubaDiver123456** 2 years, 2 months ago

Selected Answer: AD

I arrived at A+D using the explanations given here

<https://docs.microsoft.com/en-us/learn/modules/discover-azure-message-queue/5-messages-payloads-serialization>

upvoted 3 times

 **ehurfheiz** 2 years, 3 months ago

Selected Answer: AD

Answer seems correct : AD

upvoted 2 times

 **glam** 2 years, 11 months ago

correct.

upvoted 3 times

 **SnakePlissken** 2 years, 11 months ago

Answer is correct. Nice example to show the purpose of `CorrelationId` and `ReplyToSessionId`.

upvoted 4 times

 **kwaazaar** 3 years ago

what reply entity? this question makes no sense.
a simple picture of the setup would greatly clarify things

upvoted 6 times

 **MrZoom** 3 years ago

Also, all the other answers are really kind of obviously wrong...

upvoted 2 times

Question #9

Topic 5

You are developing an Azure function that connects to an Azure SQL Database instance. The function is triggered by an Azure Storage queue.

You receive reports of numerous `System.InvalidOperationExceptions` with the following message:

‘Timeout expired. The timeout period elapsed prior to obtaining a connection from the pool. This may have occurred because all pooled connections were in use and max pool size was reached.’

You need to prevent the exception.

What should you do?

- A. In the `host.json` file, decrease the value of the `batchSize` option
- B. Convert the trigger to Azure Event Hub
- C. Convert the Azure Function to the Premium plan
- D. In the `function.json` file, change the value of the `type` option to `queueScaling`

Correct Answer: C

With the Premium plan the max outbound connections per instance is unbounded compared to the 600 active (1200 total) in a Consumption plan.

Note: The number of available connections is limited partly because a function app runs in a sandbox environment. One of the restrictions that the sandbox imposes on your code is a limit on the number of outbound connections, which is currently 600 active (1,200 total) connections per instance. When you reach this limit, the functions runtime writes the following message to the logs: Host thresholds exceeded:

Connections.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-functions/manage-connections> <https://docs.microsoft.com/en-us/azure/azure-functions/functions-scale#service-limits>

Community vote distribution

A (90%) 10%

✉  **VK7Az204**  3 years, 1 month ago

A is the correct answer
upvoted 72 times

✉  **djffcnhhyiyaigvvuy**  3 years ago

The error shown is a SQL Server error not an Azure functions error.
[https://forums.asp.net/t/2004198.aspx?](https://forums.asp.net/t/2004198.aspx)
`Timeout+expired+The+timeout+period+elapsed+prior+to+obtaining+a+connection+from+the+pool+This+may+have+occurred+because+all+poo`
`led+connections+were+in+use+and+max+pool+size+was+reached+`
The actual issue appears to be that you have too many simultaneous functions running, the solution is to limit the batch size (# functions that can run in parallel)
<https://social.msdn.microsoft.com/Forums/azure/en-US/a2955297-1c14-45f2-b799-6346b340519a/how-does-batchsize-works-in-hostjson?forum=AzureFunctions>
upvoted 29 times

✉  **rdemontis** 3 years ago

Correct! We have to consider we are speaking about Azure Function with Azure Storage Queue Trigger. This is a case where execution could happens in parallel based on the queue `batchSize` property
(see <https://docs.microsoft.com/en-us/azure/azure-functions/functions-bindings-storage-queue>).

I think this is the `batchSize` property referred in the question (and not that in the functions aggregator. That's another thing).

Besides, we have to consider the origin of the error message: Azure Sql Database!! So upgrading the plan the problem could only get worse thing.

upvoted 8 times

✉  **coffecold** 1 year, 6 months ago

This article explains well and easy : <https://medium.com/@hammadarif/throttling-the-scalability-of-azure-functions-v2-ddb8625eedc>
It gives also the answer why upgrading to premium won't help.
upvoted 1 times

✉  **coffecold** 1 year, 6 months ago

B.t.w usage in `host.json` :

```
{
  "queues": {
    "batchSize": 1,
    "newBatchThreshold": 7
  }
}
```

}

It has nothing to do with the aggregator batchsize

upvoted 2 times

 **NPE** Most Recent 5 months, 3 weeks ago

The answer should be A.

The error message shows that there is not enough connections, which means that the concurrency is too high. Too many instances are running parallel. So we have to reduce the concurrency of the app.

<https://github.com/Azure/azure-functions-host/wiki/host.json>

If you search for "batchSize" on the above official microsoft github wiki page for host.json, you will find the explanation for the batchSize:

```
// The number of queue messages to retrieve and process in
// parallel (per job function). The default is 16 and the maximum is 32.
"batchSize": 16,
```

upvoted 1 times

 **Lola2023** 7 months ago

Selected Answer: C

Explanation:

With the Premium plan the max outbound connections per instance is unbounded compared to the 600 active (1200 total) in a Consumption plan.

Note: The number of available connections is limited partly because a function app runs in a sandbox environment. One of the restrictions that the sandbox imposes on your code is a limit on the number of outbound connections, which is currently 600 active (1,200 total) connections per instance. When you reach this limit, the functions runtime writes the following message to the logs: Host thresholds exceeded: Connections.

upvoted 1 times

 **BaoNguyen2411** 8 months, 3 weeks ago

got this question on 29/06/2023

upvoted 3 times

 **macobuzi** 7 months, 3 weeks ago

What option did you choose?

upvoted 2 times

 **JOSEEVILLASAMIL** 1 year ago

Answer A, the error is at SQL Instance no a function error, and the questions does not say the plan you have... so you cannot assume that you are using consumption plan.

upvoted 2 times

 **Sriniv** 1 year, 2 months ago

None of the options provided are correct I think

upvoted 1 times

 **Sriniv** 1 year, 2 months ago

the most appropriate solution in this case is to optimize the function code and increase the Max Pool Size property of the connection string used by the function.

upvoted 1 times

 **alexine74** 1 year, 2 months ago

Selected Answer: A

A is the correct answer

upvoted 1 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: A

A is the correct answer

upvoted 3 times

 **gmishra88** 1 year, 6 months ago

This is a question for a machine. I have to remember all options in host.json. How many of these random questions can be asked by Microsoft.

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messages-payloads>

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messages-payloads>

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messages-payloads>

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messages-payloads>

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messages-payloads>

<https://learn.microsoft.com/en-us/azure/azure-functions/functions-host-json#aggregator>

According to the above link batchSize is about application insights. Another marvel from Microsoft to make it without application insights anywhere and cherry on top to ask that mistake in an exam. Especially when batching will make people think it has something to do with batching the calls. Microsoft wins again

<https://learn.microsoft.com/en-us/azure/azure-functions/functions-host-json#aggregator>

upvoted 3 times

 **gmishra88** 1 year, 6 months ago

This is the way to limit scale out, but guess that is not there in the options

<https://learn.microsoft.com/en-us/azure/azure-functions/event-driven-scaling#limit-scale-out>

upvoted 1 times

✉ **gmishra88** 1 year, 6 months ago

And yes, there is another batchSize that could be for the storage queue. But then we have to imagine that could be the unwritten scenario.
upvoted 1 times

✉ **Knightie** 1 year, 7 months ago

it can not be C when Pool connection is running out, it's able the DB connections pool is out, changing the function plan only allow more callers to access the DB but it does not increase the connections pool size, so it will still be out of connections in DB.
upvoted 1 times

✉ **Knightie** 1 year, 7 months ago

Selected Answer: A

It's the DB connection Pool..... you scale up the app also no use.. the connectors from DB Pool not release fast enough for the next app to operate.
upvoted 1 times

✉ **Dani_ac7** 1 year, 9 months ago

Selected Answer: A

A is the correct answer
upvoted 2 times

✉ **meoukg** 2 years, 1 month ago

Got it on 03/2022, I chose A. In the host.json file, decrease the value of the batchSize option
upvoted 4 times

✉ **ReniRechner** 2 years, 1 month ago

Selected Answer: C

A: seems not correct for me.
see: <https://docs.microsoft.com/en-us/azure/azure-functions/functions-host-json#aggregator>
it just seems to affect metrics. Besides an increase would help => more elements per function => less functions => less connections

B: might help (if scaling can be limited), but for itself is not a solution

C: does not help since source of problem (to many incoming connections to SQL Server) unchanged. There is no problem with number of outgoing connections in function

D: searching for "queueScaling" only resulted in AZ-204 question results ;-)

So what should I choose?

While doing research for point B I git to this site

<https://medium.com/microsoftazure/azure-functions-limiting-throughput-and-scalability-of-a-serverless-app-5b1c381491e3>

I found the hint that a premium plan can limit instances, which is verified here:

<https://docs.microsoft.com/en-us/azure/azure-functions/functions-premium-plan?tabs=portal#maximum-function-app-instances>

So I would go with C but for the option to limit scaling and not for the explanation in the result.

upvoted 1 times

✉ **iamstudying** 2 years, 1 month ago

I agree with your reason for not A.. batching is supposed to reduce number of IO requests and thus connections to SQL DB

The Medium post you referenced is outdated. You can actually limit the scaling on both Consumption and Premium plans:
<https://docs.microsoft.com/en-us/azure/azure-functions/event-driven-scaling#limit-scale-out>

I don't think there is a right answer here... pick A and prAy, buddies

upvoted 2 times

✉ **massnonn** 2 years, 1 month ago

Selected Answer: A

here: <https://stackoverflow.com/questions/56489553/host-json-meaning-of-batchsize>

upvoted 2 times

Question #10

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution.

Determine whether the solution meets the stated goals.

You are developing and deploying several ASP.NET web applications to Azure App Service. You plan to save session state information and HTML output.

You must use a storage mechanism with the following requirements:

- ☞ Share session state across all ASP.NET web applications.
- ☞ Support controlled, concurrent access to the same session state data for multiple readers and a single writer.
- ☞ Save full HTTP responses for concurrent requests.

You need to store the information.

Proposed Solution: Deploy and configure Azure Cache for Redis. Update the web applications.

Does the solution meet the goal?

A. Yes

B. No

Correct Answer: A

The session state provider for Azure Cache for Redis enables you to share session information between different instances of an ASP.NET web application.

The same connection can be used by multiple concurrent threads.

Redis supports both read and write operations.

The output cache provider for Azure Cache for Redis enables you to save the HTTP responses generated by an ASP.NET web application.

Note: Using the Azure portal, you can also configure the eviction policy of the cache, and control access to the cache by adding users to the roles provided. These roles, which define the operations that members can perform, include Owner, Contributor, and Reader. For example, members of the Owner role have complete control over the cache (including security) and its contents, members of the Contributor role can read and write information in the cache, and members of the Reader role can only retrieve data from the cache.

Reader role can only retrieve data from the cache.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/best-practices/caching>

Community vote distribution

A (100%)

 **Prakash4691**  3 years, 1 month ago

Correct answer

upvoted 33 times

 **THAMORAY**  2 years, 11 months ago

This question is part of a series of questions that present the same scenario. Each question in the series contains "" a unique solution""". Determine whether the solution meets the stated goals.

This is the unique solution

A: Yes.

upvoted 8 times

 **OPT_001122**  1 year, 5 months ago

Selected Answer: A

Correct answer

upvoted 3 times

 **Eltooth** 1 year, 9 months ago

Selected Answer: A

A is correct answer.

upvoted 3 times

 **mattvasc** 2 years, 1 month ago

Selected Answer: A

Redis will do the trick.

upvoted 1 times

 **glam** 2 years, 11 months ago

A. Yes

upvoted 5 times

 **roshansir** 2 years, 11 months ago

A is the correct answer

upvoted 3 times

Question #11

Topic 5

HOTSPOT -

You are debugging an application that is running on Azure Kubernetes cluster named cluster1. The cluster uses Azure Monitor for containers to monitor the cluster.

The application has sticky sessions enabled on the ingress controller.

Some customers report a large number of errors in the application over the last 24 hours.

You need to determine on which virtual machines (VMs) the errors are occurring.

How should you complete the Azure Monitor query? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
let startTimestamp = ;  
ago(1d)  
since(1d)  
totimespan(1d)  
date(now() - 1d)  
  
let ContainerIDs = KubePodInventory  
| where ClusterName == "Cluster1"  
| ;  
top ContainerID  
union ContainerID  
sample ContainerID  
distinct ContainerID  
  
ContainerLog  
| ;  
fork containerIDs  
where ContainerID in (ContainerIDs)  
restrict ContainerID in (ContainerIDs)  
join ContainerID == ContainerIDs.ContainerID  
  
| where TimeGenerated > startTimestamp  
| where LogEntrySource == "stderr"  
| ;  
project by Computer  
summarize by Computer  
partition count() by Computer  
summarize count() by Computer
```

Answer Area

```

let startTimestamp =
    ago(1d)
    since(1d)
    totimespan(1d)
    date(now() - 1d)

let ContainerIDs = KubePodInventory
| where ClusterName == "Cluster1"
    top ContainerID
    union ContainerID
    sample ContainerID
    distinct ContainerID

ContainerLog
| fork containerIDs
| where ContainerID in (ContainerIDs)
| restrict ContainerID in (ContainerIDs)
| join ContainerID == ContainerIDs.ContainerID
| where TimeGenerated > startTimestamp
| where LogEntrySource == "stderr"
    project by Computer
    summarize by Computer
    partition count() by Computer
    summarize count() by Computer

```

Box 1: ago(1d)

Box 2: distinct containerID -

Box 3: where ContainerID in (ContainerIDs)

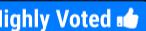
Box 4: summarize Count by Computer

Summarize: aggregate groups of rows

Use summarize to identify groups of records, according to one or more columns, and apply aggregations to them. The most common use of summarize is count, which returns the number of results in each group.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/log-query/get-started-queries> <https://docs.microsoft.com/en-us/azure/azure-monitor/log-query/query-optimization>

✉  **Marusyk**  3 years, 1 month ago

the answer is correct

upvoted 40 times

✉  **ReyPirata** 8 months ago

On my exam 2023-08-20. Scored 925

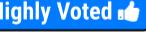
ago(1)

distinct ContainerID

where ContainerID in (ContainerIDs)

summarize count() by Computer

upvoted 5 times

✉  **pac1311**  3 years, 3 months ago

Isn't kubernetes out of scope for AZ-204?

upvoted 13 times

✉  **Chilred** 3 years, 3 months ago

No, i think there are still some question from az203

upvoted 3 times

✉  **phuria** 3 years, 2 months ago

You'll notice it really isn't about kubernetes per say, but how to query your logs from Log Analytics

upvoted 15 times

✉  **surprise0011** 1 year ago

yep, it is more about monitoring

upvoted 1 times

✉ **gmishra88** 1 year, 6 months ago

There are many questions from kubernetes

upvoted 2 times

✉ **hicloud** 7 months, 4 weeks ago

I guess if kubernetes is not the only product involved in the question then the question still appears

upvoted 1 times

✉ **WH16** Most Recent 8 months, 2 weeks ago

Kubernetes out of scope as of August 2023

upvoted 3 times

✉ **macobuzi** 7 months, 3 weeks ago

Better safe than sorry!

upvoted 2 times

✉ **rockstar89** 8 months, 3 weeks ago

Got it exam 31/07/2023

upvoted 5 times

✉ **juanckar** 9 months, 2 weeks ago

This was on the exam (July 2023). Went with proposed. Scored 917

upvoted 6 times

✉ **uffuchsi** 1 year, 1 month ago

The question asks for errors in last 24 hours, so startTimestamp should be 'ago (1d)'.

We need to distinct the logs for each container from set of containers. So, distinct based on ContainerID and where condition on container IDs. We need to determine the VM. So, summarize count() by computer.

<https://docs.microsoft.com/en-us/azure/azure-monitor/log-query/get-started-queries>

<https://docs.microsoft.com/en-us/azure/azure-monitor/log-query/query-optimization>

upvoted 1 times

✉ **Aterniss** 1 year, 4 months ago

Got this 15/12/2022

upvoted 5 times

✉ **surprise0011** 1 year ago

thanks for mentioning. Maybe Kubernetes is out of scope but this question is more about monitoring

upvoted 1 times

✉ **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 5 times

✉ **Knightie** 1 year, 7 months ago

the VMs here mean Containers.... sigh....

upvoted 1 times

✉ **petibilly** 2 years, 1 month ago

Got it in exam 03/22

upvoted 3 times

✉ **oescm** 2 years, 2 months ago

Got this one 02/2022, so this question is NOT out of scope. I went with the provided answer

upvoted 7 times

✉ **lugospod** 2 years, 3 months ago

Got this one 01/2022. Went with most voted (to avoid writing answers again)

upvoted 9 times

✉ **MasterQuestMaster** 2 years, 2 months ago

I got this as well.

upvoted 3 times

✉ **glam** 2 years, 11 months ago

correct

upvoted 4 times

✉ **diligent176** 3 years, 3 months ago

ago(1d) is correct.

date(now()-1d) is not correct syntax (query won't run or compile at all).

upvoted 6 times

✉ **HimanshuNankani** 2 years, 6 months ago

Should have told the correct syntax as well.
upvoted 1 times

✉ **m_siri** 3 years, 4 months ago

I think the answer to the 1st drop should be;
let start Timestamp = date(now() - 1d)

because we need to take the logs occurred within the last 24 hours, not for the previous day

upvoted 4 times

✉ **mitcheld** 3 years, 4 months ago

I think it should be ago.
Search for "Summarize: aggregate groups of rows" at <https://docs.microsoft.com/en-us/azure/azure-monitor/log-query/get-started-queries>
upvoted 1 times

✉ **ALittleBunny** 3 years, 4 months ago

Wrong. It should be ago(1d) as ago(a_timespan) subtracts the given timespan from the current UTC clock time.
<https://docs.microsoft.com/en-us/azure/azure-monitor/log-query/examples#search-application-level-events-described-as-cryptographic>
upvoted 6 times

✉ **stylebc** 3 years, 4 months ago

You are not right, correct is ago(1d), as from ago function description:
<https://docs.microsoft.com/en-us/azure/data-explorer/kusto/query/agofunction>
a_timespan - Interval to subtract from the current UTC clock time (now()).
upvoted 19 times

✉ **adilkhan** 1 year, 2 months ago

no ago is correct..!
upvoted 1 times

Question #12

HOTSPOT -

You plan to deploy a web app to App Service on Linux. You create an App Service plan. You create and push a custom Docker image that contains the web app to Azure Container Registry.

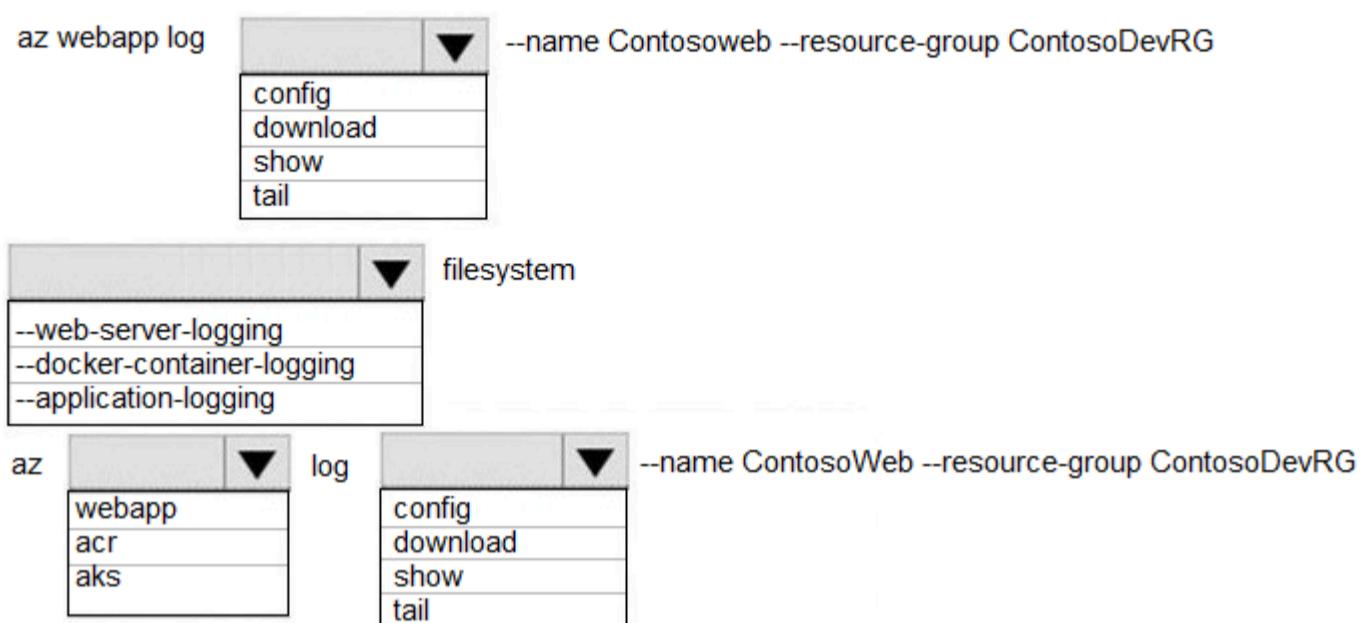
You need to access the console logs generated from inside the container in real-time.

How should you complete the Azure CLI command? To answer, select the appropriate options in the answer area.

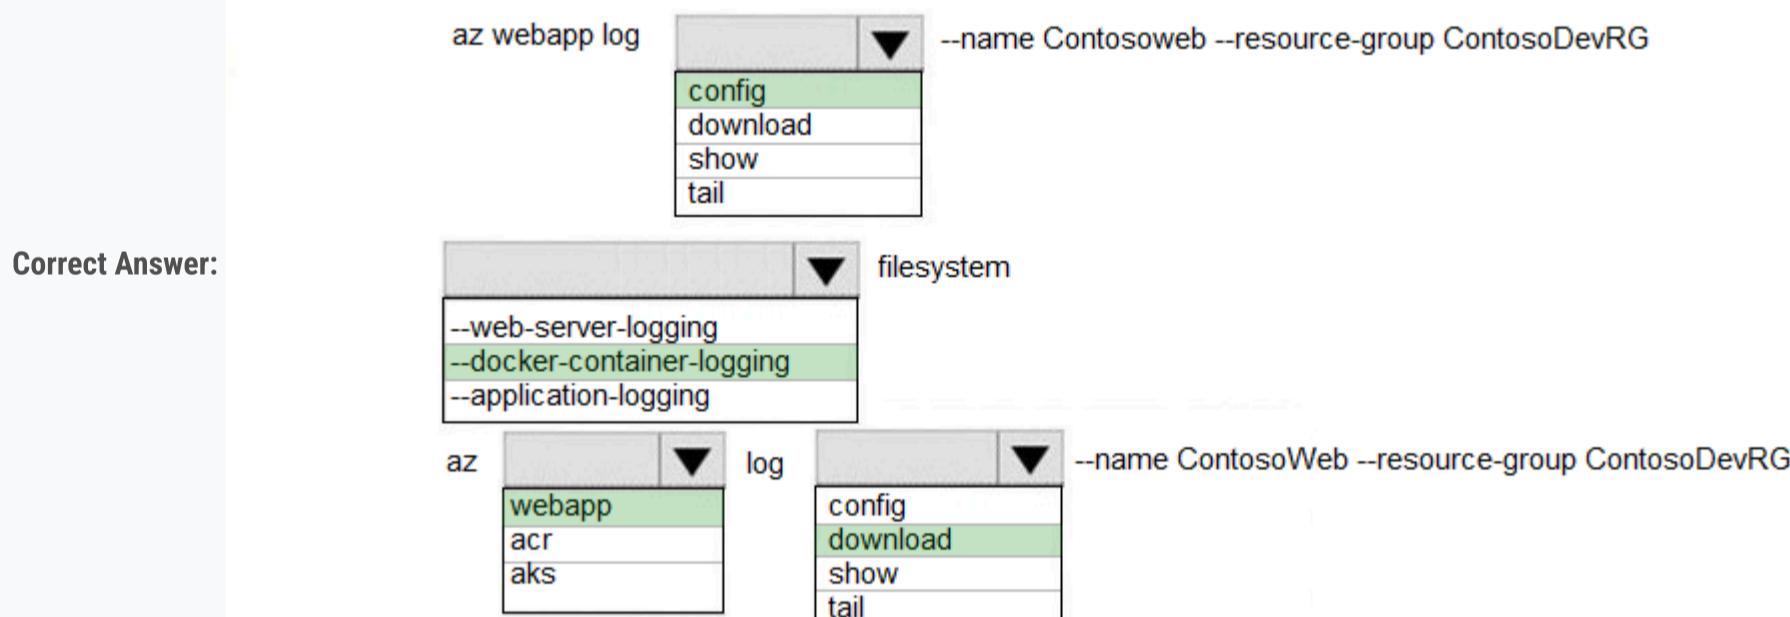
NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area



Answer Area



Box 1: config -

To Configure logging for a web app use the command:

az webapp log config

Box 2: --docker-container-logging

Syntax include:

az webapp log config [-docker-container-logging {filesystem, off}]

Box 3: webapp -

To download a web app's log history as a zip file use the command: az webapp log download

Box 4: download -

Reference:

<https://docs.microsoft.com/en-us/cli/azure/webapp/log>

granimtoula 3 years, 5 months ago

It never said download, it says "access real time"

Will go with TAIL:

https://docs.microsoft.com/en-us/cli/azure/webapp/log?view=azure-cli-latest#az_webapp_log_tail

upvoted 145 times

✉ **SachinV** 5 months ago

Agreed it's correct, in Nov182023, scored 962

upvoted 1 times

✉ **Kellyb85** 3 years, 5 months ago

I think download

in the question--> You plan to deploy a web app to App Service on Linux.

az-webapp-log-tail --> https://docs.microsoft.com/en-us/cli/azure/webapp/log?view=azure-cli-latest#az_webapp_log_tail --> This command may not work with web apps running on Linux.

upvoted 6 times

✉ **TakumaK** 2 years, 10 months ago

See the MSDN doc again. It obviously says the download is also "This command may not work with web apps running on Linux.". And the question clearly mentions "in real-time". Ok. let's think do you want to download the log zip file in real time? How?

upvoted 2 times

✉ **Ave** 3 years, 5 months ago

It says the same for download.

upvoted 10 times

✉ **iiiihhhh** 3 years, 3 months ago

There is a request for clarification on github:

WebApp:az webapp log tail and insufficient information about linux <https://github.com/Azure/azure-cli/issues/16203>

But there is no answer.

upvoted 1 times

✉ **Adol** 3 years, 3 months ago

You're right and it's tail according to a course in Udemy for the exact question

upvoted 9 times

✉ **AfroYeti** 3 years, 2 months ago

No see, that interesting because there is a notice on the port itself and it reads "tail might not be available for linux containers"

Now the question is, is that answer the "safe" answer based on the possibility that it might not be available?

upvoted 1 times

✉ **DSK** 3 years, 2 months ago

The suggested download command will not work because download needs a local file system location.

upvoted 1 times

✉ **monniq** 3 years, 2 months ago

It says "This command may not work with web apps running on Linux." on https://docs.microsoft.com/en-us/cli/azure/webapp/log?view=azure-cli-latest#az_webapp_log_tail

So, it can't be the answer.

upvoted 8 times

✉ **jay158** 2 years, 9 months ago

"This command may not work with web apps running on Linux." is removed https://docs.microsoft.com/en-us/cli/azure/webapp/log?view=azure-cli-latest#az_webapp_log_tail.

Warning is only for 'az webapp log download'

upvoted 6 times

✉ **Robert12345Robert** 2 years, 9 months ago

az webapp log download

Download a web app's log history as a zip file.

This command may not work with web apps running on Linux.

It is download that may not work, not tail.

upvoted 8 times

✉ **rdemontis** 3 years ago

I think TAIL is correct because it is the only way to access console log in real time. The sentence "This command may not work with web apps running on Linux" is present for every option, even for download, see https://docs.microsoft.com/en-us/cli/azure/webapp/log?view=azure-cli-latest#az_webapp_log_download

upvoted 4 times

✉ **RaviKS**  3 years, 3 months ago

Correct Answers are

config

docker-container-logging

webapp

tail

upvoted 71 times

✉  **FeriAZ** Most Recent 2 months, 3 weeks ago

TAIL

az webapp log tail -n ContosoWeb -g ContosoDevRG --docker-container-logging

The tail command is used to view streaming log data in real-time. This is likely what you need for accessing logs as they are generated.

upvoted 1 times

✉  **CarlosTheBoldest** 4 months, 1 week ago

I got this question on my exam, 2023Dec, go with what I remember was the most voted answer. Score 902, most of the questions were here, slightly different on wording because the Azure Ad <-> Entra Id change. Case was City Power & Light. Good luck!

Important tip, you have access to microsoft learn during the exam!

upvoted 1 times

✉  **juanckar** 9 months, 2 weeks ago

This was on the exam (July 2023). Went with highly voted. Scored 917

upvoted 5 times

✉  **JH81** 9 months, 3 weeks ago

Got this on 6/28/2023 and passed with 850. Went with highly voted answer.

upvoted 1 times

✉  **ExamDev** 10 months ago

Ok we got Linux and documentation says:

- az webapp log download

Download a web app's log history as a zip file.

!!! This command may not work with web apps running on Linux. !!!

<https://learn.microsoft.com/en-us/cli/azure/webapp/log?view=azure-cli-latest#az-webapp-log-download>

upvoted 1 times

✉  **Vmwarevirtual** 10 months, 3 weeks ago

Box 1: config

Box 2: --docker-container-logging

Box 3: webapp

Box 4 : Tail

Question appeared in the exam I took at 27/May 2023

upvoted 3 times

✉  **AzBeginner** 11 months ago

On Exam : 22 May 2023

upvoted 2 times

✉  **aragones** 11 months, 1 week ago

Got this 2023-05-12.

Had cases:

case: You need to configure authorization.

case: You need to ensure the app does not time out and processes the blob data

case study: VanArsdel Inc Canada

upvoted 2 times

✉  **uffuchi** 1 year, 1 month ago

To access the console logs generated from inside the container, first, turn on container logging by running the following command:

az webapp log config --name <app-name> --resource-group <resource-group-name> --docker-container-logging filesystem

Once container logging is turned on, run the following command to see the log stream:

az webapp log tail --name <app-name> --resource-group <resource-group-name>

If you don't see console logs immediately, check again in 30 seconds.

To stop log streaming at any time, type Ctrl+C.

<https://docs.microsoft.com/en-us/azure/app-service/configure-custom-container?pivots=container-linux#access-diagnostic-logs-1>

upvoted 4 times

✉  **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 4 times

✉  **Knightie** 1 year, 7 months ago

No Download, but Tail

<https://docs.microsoft.com/en-us/cli/azure/webapp/log?view=azure-cli-latest#az-webapp-log-tail>

upvoted 1 times

✉  **mahidharKollipara** 1 year, 8 months ago

it should be tail. Download does not work in linux.

upvoted 1 times

✉ **Eltooth** 1 year, 9 months ago

config
docker-container-logging
webapp
tail

upvoted 2 times

✉ **Eltooth** 1 year, 9 months ago

<https://docs.microsoft.com/en-us/cli/azure/webapp/log?view=azure-cli-latest#az-webapp-log-config>

upvoted 1 times

✉ **Sandeep12093** 1 year, 10 months ago

We don't want to download.also it said real time so definitely answer is "tail"

upvoted 1 times

✉ **petitbilly** 2 years, 1 month ago

Got it in exam 03/22

upvoted 2 times

Question #13

Topic 5

You develop and deploy an ASP.NET web app to Azure App Service. You use Application Insights telemetry to monitor the app. You must test the app to ensure that the app is available and responsive from various points around the world and at regular intervals. If the app is not responding, you must send an alert to support staff. You need to configure a test for the web app. Which two test types can you use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. integration
- B. multi-step web
- C. URL ping
- D. unit
- E. load

Correct Answer: BC

There are three types of availability tests:

- ⇒ URL ping test: a simple test that you can create in the Azure portal.
- ⇒ Multi-step web test: A recording of a sequence of web requests, which can be played back to test more complex scenarios. Multi-step web tests are created in Visual Studio Enterprise and uploaded to the portal for execution.
- ⇒ Custom Track Availability Tests: If you decide to create a custom application to run availability tests, the `TrackAvailability()` method can be used to send the results to Application Insights.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability>

Community vote distribution

BC (100%)

✉  **Ash111**  3 years, 5 months ago

Given answer is correct

upvoted 43 times

✉  **Secure01** 3 years, 4 months ago

True:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability>

There are three types of availability tests:

URL ping test
Multi-step web test
Custom Track Availability Tests

upvoted 11 times

✉  **pac1311**  3 years, 2 months ago

Correctemundo!

upvoted 10 times

✉  **BaoNguyen2411**  8 months, 3 weeks ago

got this question on 29/07/2023

upvoted 4 times

✉  **uffuchsi** 1 year, 1 month ago

B and C

upvoted 1 times

✉  **Sriniv** 1 year, 2 months ago

Note that multi-step web tests, load tests, and unit tests are not suitable for this scenario as they do not test the availability and responsiveness of the app from various points around the world at regular intervals. Multi-step web tests can simulate user scenarios, but they do not test the app from various points around the world. Load tests can test the performance of the app, but they do not test the availability and responsiveness of the app. Unit tests are used to test individual code components and are not suitable for testing the app as a whole.

upvoted 1 times

✉  **Sriniv** 1 year, 2 months ago

Answer is A and C reason is below

Integration tests: These tests simulate the user experience by using a browser to interact with the web app, and they can check for specific page

content, errors, and response times. You can configure integration tests using the Azure Portal or Visual Studio.

URL ping tests: These tests send a GET request to a specific URL of the web app to check for a valid HTTP response and response time. URL ping tests are the simplest and most common type of web app test. You can configure URL ping tests using the Azure Portal or the Application Insights API.

upvoted 1 times

 **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 6 times

 **Eltooth** 1 year, 9 months ago

Selected Answer: BC

B and C are correct answers.

upvoted 2 times

 **Eltooth** 1 year, 9 months ago

Multi-step web tests are deprecated:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/availability-multistep>

upvoted 1 times

 **oescm** 2 years, 2 months ago

Got this one 02/2022. Went with the given answers

upvoted 4 times

 **lugospod** 2 years, 3 months ago

Got similar one 01/2022. But tests needed to also test for SSL and custom headers.. went with multistep and Standard test (not ping or custom track).

upvoted 7 times

 **Henksc** 2 years, 3 months ago

Correct, however Multi-step web test is deprecated.

Source:

Multi-step web tests have been deprecated. We recommend using TrackAvailability() to submit custom availability tests instead of multi-step web tests. With TrackAvailability() and custom availability tests, you can run tests on any compute you want and use C# to easily author new tests.

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/availability-multistep>

upvoted 3 times

 **ehurfheiz** 2 years, 3 months ago

Selected Answer: BC

Answer seems to be correct : BC

upvoted 1 times

 **glam** 2 years, 11 months ago

correct

upvoted 4 times

 **roshansir** 2 years, 11 months ago

correct answer

upvoted 3 times

 **paru123456789** 3 years, 1 month ago

Answer:

Multi-step

Url ping

upvoted 5 times

 **10x** 3 years, 2 months ago

b. and c.

upvoted 3 times

Question #14

DRAG DROP -

A web service provides customer summary information for e-commerce partners. The web service is implemented as an Azure Function app with an HTTP trigger.

Access to the API is provided by an Azure API Management instance. The API Management instance is configured in consumption plan mode. All API calls are authenticated by using OAuth.

API calls must be cached. Customers must not be able to view cached data for other customers.

You need to configure API Management policies for caching.

How should you complete the policy statement?

Select and Place:

Targets	Answer Area				
Expect	<policies>				
Public	<inbound>				
Private	<base />				
Internal	<cache-lookup caching-type="	Target	" downstream-caching-type = "	Target	">
External	<vary-by-header>	Target			
Authorization	</vary-by-header>				
	</cache-lookup>				
	</inbound>				
	</policies>				

Correct Answer:

Targets	Answer Area				
Expect	<policies>				
Public	<inbound>				
Private	<base />				
Internal	<cache-lookup caching-type="	Internal	" downstream-caching-type = "	Private	">
External	<vary-by-header>	Authorization			
Authorization	</vary-by-header>				
	</cache-lookup>				
	</inbound>				
	</policies>				

Box 1: internal -

caching-type

Choose between the following values of the attribute:

- internal to use the built-in API Management cache,
- external to use the external cache as Azure Cache for Redis prefer-external to use external cache if configured or internal cache otherwise.

▪

Box 2: private -

downstream-caching-type

This attribute must be set to one of the following values.

- none - downstream caching is not allowed.
- private - downstream private caching is allowed.
- public - private and shared downstream caching is allowed.

Box 3: Authorization -

<vary-by-header>Authorization</vary-by-header>

<!-- should be present when allow-private-response-caching is "true"-->

Note: Start caching responses per value of specified header, such as Accept, Accept-Charset, Accept-Encoding, Accept-Language, Authorization, Expect, From,

Host, If-Match -

Reference:

<https://docs.microsoft.com/en-us/azure/api-management/api-management-caching-policies>

✉  **cgreen**  3 years, 5 months ago

Internal Cache isn't available for Consumption plan. So shouldn't it be External (using Redis)?

upvoted 91 times

✉  **Egger1992** 3 years, 5 months ago

Yes, you are right. In the feature-based comparison of the tiers, it says that built-in cache is not available. <https://docs.microsoft.com/en-us/azure/api-management/api-management-features>

And here it says that "internal" would be using the built-in cache. <https://docs.microsoft.com/en-us/azure/api-management/api-management-caching-policies>

upvoted 19 times

✉  **Ananth2020**  3 years, 5 months ago

Consumption Pricing Tier of API Management does not support Built-in Cache. Therefore caching-type should be set to "external".

Reference: <https://docs.microsoft.com/en-us/azure/api-management/api-management-features>

upvoted 36 times

✉  **BaoNguyen2411**  8 months, 3 weeks ago

got this question on 29/07/2023

upvoted 5 times

✉  **uffuchsi** 1 year, 1 month ago

Caching type must be external since consumption plan does not support internal caching.

Downstream caching type must be private as it is not shared.

Vary by header must be authorization as it is cached per customer.

<https://docs.microsoft.com/en-us/azure/api-management/api-management-caching-policies>

upvoted 7 times

✉  **adilkhan** 1 year, 2 months ago

<https://learn.microsoft.com/en-us/azure/api-management/cache-lookup-policy>

upvoted 1 times

✉  **OPT_001122** 1 year, 5 months ago

External

Private

Authorization

upvoted 11 times

✉  **Eltooth** 1 year, 9 months ago

External

Private

Authorization

<https://docs.microsoft.com/en-us/azure/api-management/api-management-caching-policies#attributes>

<https://docs.microsoft.com/en-us/azure/api-management/api-management-sample-cache-by-key#architecture>

upvoted 10 times

✉  **Sandeep12093** 1 year, 10 months ago

Yes, I agree consumption plan doesn't support internal caching..

Check microsoft pricing page for api management

upvoted 3 times

✉  **adone** 2 years, 1 month ago

I don't find any example with private downstream and a vary by header. For me the vary by header only make sense for public downstream. I opt for:

External

Public

Authorization

In normal situation, I would have put private downstream without a vary by header.

upvoted 2 times

✉  **oescm** 2 years, 2 months ago

Got this one 02/2022. Went with:

External

Private
Authorization
upvoted 16 times

✉ **leonidn** 2 years, 3 months ago
Private is because cached data is isolated between customers
External is because internal is not a feature of the Consumption plan
Authorization to enable caching responses per value of specified header
upvoted 9 times

✉ **Gautam47** 2 years, 7 months ago
Cache should be external
upvoted 3 times

✉ **chandru1dev** 2 years, 8 months ago
internal
caching-type Choose between the following values of the attribute:
- internal to use the built-in API Management cache,
- external to use the external cache as described in Use an external Azure Cache for Redis in Azure API Management,
- prefer-external to use external cache if configured or internal cache otherwise.
<https://docs.microsoft.com/en-us/azure/api-management/api-management-caching-policies>
upvoted 2 times

✉ **kondapaturi** 2 years, 9 months ago
External - When you look at the pricing plans for the Azure API management, when it comes to the Consumption plan, it only supports external cache.
Private - we need to ensure that the data is not available to other customers, we need to use the cache type of Private caching.
Authorization - the calls are being authenticated by OAuth, the caching can vary by the Authorization header.
upvoted 13 times

✉ **glam** 2 years, 11 months ago
Answers are:
External
Private
Authorization
upvoted 17 times

✉ **anandhprakash** 2 years, 11 months ago
Internal cache is not supported in consumption plan, so the cache type should be External.
Answers are:
External
Private
Authorization
upvoted 6 times

✉ **roshansir** 2 years, 11 months ago
External
Private
Authorization
upvoted 1 times

Question #15

You are developing applications for a company. You plan to host the applications on Azure App Services.

The company has the following requirements:

- ☞ Every five minutes verify that the websites are responsive.
- ☞ Verify that the websites respond within a specified time threshold. Dependent requests such as images and JavaScript files must load properly.
- ☞ Generate alerts if a website is experiencing issues.
- ☞ If a website fails to load, the system must attempt to reload the site three more times.

You need to implement this process with the least amount of effort.

What should you do?

- A. Create a Selenium web test and configure it to run from your workstation as a scheduled task.
- B. Set up a URL ping test to query the home page.
- C. Create an Azure function to query the home page.
- D. Create a multi-step web test to query the home page.
- E. Create a Custom Track Availability Test to query the home page.

Correct Answer: D

You can monitor a recorded sequence of URLs and interactions with a website via multi-step web tests.

Incorrect Answers:

A: Selenium is an umbrella project for a range of tools and libraries that enable and support the automation of web browsers.

It provides extensions to emulate user interaction with browsers, a distribution server for scaling browser allocation, and the infrastructure for implementations of the W3C WebDriver specification that lets you write interchangeable code for all major web browsers.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/availability-multistep>

Community vote distribution



✉ **kishe** 2 years, 10 months ago

IMO, B is correct, a URL ping test is enough.

upvoted 67 times

✉ **mattvasc** 2 years, 1 month ago

And how a URL ping will meet the requirement:

"☞ Verify that the websites respond within a specified time threshold. Dependent requests such as images and JavaScript files must load properly." ?

Definitely not B.

upvoted 9 times

✉ **ezelans** 2 years, 1 month ago

please, read this article and you will check it. There is a setting called "Test timeout"

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability>

upvoted 10 times

✉ **uffuchsi** 1 year, 1 month ago

URL ping test validate whether an endpoint is responding. It also measures the performance associated with that response. You can also set custom success criteria, coupled with more advanced features like parsing dependent requests and allowing for retries.

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability>

upvoted 3 times

✉ **KingChuang** 1 year, 3 months ago

On my exam 2022-12-26.

Chose: B

upvoted 3 times

✉ **AymanAkk** 1 year, 3 months ago

Thanks for mentioning the date. were there a lot of questions from exam topics ?

upvoted 3 times

✉ **TonyMel** 1 year ago

correct, in 2023Mar24, score: 904/1000.

upvoted 6 times

✉ **jay158** Highly Voted 2 years, 10 months ago

The Given Answer D is correct.

Ping URL, does not satisfies : Verify that the websites respond within a specified time threshold. Dependent requests such as images and JavaScript files must load properly

upvoted 26 times

✉ **anirbanzeus** 2 years, 10 months ago

Not really, all the objectives can be achieved through Ping-test. Multi-step test in a complex solution requiring more effort and dives into much deeper details that requested in the question.

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability#create-a-test> : All the options listed here map 1-1 with the question requirements.

upvoted 11 times

✉ **jay158** 2 years, 10 months ago

Second thought

Correct answer is B. ie URL Ping Test

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability>

upvoted 20 times

✉ **130nk3r5** Most Recent 4 months ago

Selected Answer: D

Because of "Dependent requests such as images and JavaScript files must load properly."

upvoted 1 times

✉ **Ciupaz** 5 months ago

Selected Answer: B

Every solution could be fine, but URL Ping Test requires the least amount of effort.

upvoted 1 times

✉ **katrang** 7 months, 1 week ago

Selected Answer: B

From the documentation:

The name URL ping test is a bit of a misnomer. These tests don't use the Internet Control Message Protocol (ICMP) to check your site's availability. Instead, they use more advanced HTTP request functionality to validate whether an endpoint is responding. They measure the performance associated with that response. They also add the ability to set custom success criteria, coupled with more advanced features like parsing dependent requests and allowing for retries.

upvoted 2 times

✉ **Paolo12** 7 months, 3 weeks ago

Selected Answer: B

In URL-Ping test it's possible to set dependent requests, retries, etc...

upvoted 1 times

✉ **aragones** 11 months, 1 week ago

Got this 2023-05-12.

FYI:

case: You need to configure authorization.

case: You need to ensure the app does not time out and processes the blob data

case study: VanArsdel Inc Canada

upvoted 4 times

✉ **motekim** 1 year ago

It looks like its B though: <https://learn.microsoft.com/en-us/previous-versions/azure/azure-monitor/app/monitor-web-app-availability>

upvoted 1 times

✉ **motekim** 1 year ago

if the answer is D. Create a multi-step web test to query the home page., then it's actually E. from Msoft site: We recommend using TrackAvailability to submit custom availability tests instead of multistep web tests. This option is the long-term supported solution for multi-request or authentication test scenarios. With TrackAvailability() and custom availability tests, you can run tests on any compute you want and use C# to easily author new tests. hmm

upvoted 1 times

✉ **NombreFalso** 1 year, 1 month ago

Selected Answer: D

Ping test seems to be the least effort solution

upvoted 1 times

✉ **Sriniv** 1 year, 2 months ago

Answer B as it satisfies all requirements. Below is the reason for URL ping test.

The Azure App Service has built-in URL ping testing that can be configured to check whether a website is responsive every five minutes, and

generate alerts if there are issues. The URL ping test can also check that the website responds within a specified time threshold and that dependent requests such as images and JavaScript files load properly. Additionally, the URL ping test will automatically retry up to three times if the website fails to load.

upvoted 2 times

 **adilkhan** 1 year, 2 months ago

B is correct because:

1. The test requests images, scripts, style files, and other files that are part of the webpage under test. The recorded response time includes the time taken to get these files. The test fails if any of these resources can't be successfully downloaded within the timeout for the whole test. If the option isn't enabled, the test only requests the file at the URL that you specified. Enabling this option results in a stricter check. The test might fail for cases that aren't noticeable from manually browsing through the site.

2. When the test fails, it's retried after a short interval. A failure is reported only if three successive attempts fail. Subsequent tests are then performed at the usual test frequency. Retry is temporarily suspended until the next success. This rule is applied independently at each test location. We recommend this option. On average, about 80 percent of failures disappear on retry.

<https://learn.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability>

upvoted 1 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: B

B. Set up a URL ping test to query the home page.

upvoted 2 times

 **Dhanishetty** 1 year, 5 months ago

Selected Answer: B

<https://learn.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability#:~:text=Setting,to%2016%20locations.>

upvoted 1 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: B

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability>

Test timeout - Decrease this value to be alerted about slow responses. The test is counted as a failure if the responses from your site haven't been received within this period. If you selected Parse dependent requests, all the images, style files, scripts, and other dependent resources must have been received within this period.

upvoted 1 times

 **nvtienanh** 1 year, 6 months ago

Selected Answer: E

Multi-step web tests have been deprecated. We recommend using TrackAvailability() to submit custom availability tests instead of multi-step web tests.

I think E is correct.

upvoted 2 times

 **azurepaul** 1 year, 7 months ago

I vote e based on the dependent files bit - and as called out in other comments:

Multi-step web tests have been deprecated. We recommend using TrackAvailability() to submit custom availability tests instead of multi-step web tests. With TrackAvailability() and custom availability tests, you can run tests on any compute you want and use C# to easily author new tests.

upvoted 3 times

Question #16

You develop and add several functions to an Azure Function app that uses the latest runtime host. The functions contain several REST API endpoints secured by using SSL. The Azure Function app runs in a Consumption plan.

You must send an alert when any of the function endpoints are unavailable or responding too slowly.

You need to monitor the availability and responsiveness of the functions.

What should you do?

- A. Create a URL ping test.
- B. Create a timer triggered function that calls `TrackAvailability()` and send the results to Application Insights.
- C. Create a timer triggered function that calls `GetMetric("Request Size")` and send the results to Application Insights.
- D. Add a new diagnostic setting to the Azure Function app. Enable the `FunctionAppLogs` and `Send to Log Analytics` options.

Correct Answer: B

You can create an Azure Function with `TrackAvailability()` that will run periodically according to the configuration given in `TimerTrigger` function with your own business logic. The results of this test will be sent to your Application Insights resource, where you will be able to query for and alert on the availability results data.

This allows you to create customized tests similar to what you can do via Availability Monitoring in the portal. Customized tests will allow you to write more complex availability tests than is possible using the portal UI, monitor an app inside of your Azure VNET, change the endpoint address, or create an availability test even if this feature is not available in your region.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/availability-azure-functions>

Community vote distribution

B (73%) A (27%)

✉  **mlantonis** Highly Voted 2 years, 10 months ago

Based on the doc, I believe B is correct.

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/availability-azure-functions>
upvoted 26 times

✉  **qzrkfrcglkzzvfrcro** Highly Voted 2 years, 9 months ago

I think the answer should be URL Ping Test.

It can send an alert when the website is unavailable or unresponsive. See test timeouts and alerts in the link.

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability>

upvoted 21 times

✉  **troy89** 2 years, 1 month ago

But you need a URL test per endpoint, so one URL test wouldn't be enough.

upvoted 2 times

✉  **Knightie** 1 year, 7 months ago

Ping Test only check for response code 200, it will not give you any info after the initial 200 ok code, contents and information will keep loading after the 200 code received, so if the 200 there but the page not responding or the background not processing or not completing, then the ping test will not be able to know...it will still be 200 but the response might be incomplete.

upvoted 2 times

✉  **Knightie** 1 year, 7 months ago

"responsiveness of the functions" <== this one will need more than the ping test

upvoted 3 times

✉  **ks321** 9 months ago

Note that the API is protected by SSL. Does URL ping test support auth?

upvoted 1 times

✉  **Ciupaz** Most Recent 5 months ago

Selected Answer: B

Based on the information, creating a timer triggered function that calls `TrackAvailability()` and sends the results to Application Insights is the best option for monitoring the availability and responsiveness of the functions.

upvoted 1 times

✉  **uffuchs1** 1 year, 1 month ago

It has to be C - The URL ping test uses advanced HTTP request functionality to validate whether an endpoint is responding. They measure the performance associated with that response. They also add the ability to set custom success criteria, coupled with more advanced features like

parsing dependent requests and allowing for retries.

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability#actions-if-you-see-failures>

upvoted 2 times

 **uffuchsi** 1 year, 1 month ago

Apologies - I meant A, not C

upvoted 2 times

 **Sriniv** 1 year, 2 months ago

Adding a new diagnostic setting (Option D) can enable logging for the Azure Function app, but it will not provide the ability to monitor the availability and responsiveness of the REST API endpoints or generate an alert when they are unavailable or responding too slowly.

upvoted 1 times

 **Sriniv** 1 year, 2 months ago

Creating a timer triggered function (Option B or C) can also be used to call the REST API endpoints and report the results to Application Insights. However, this approach requires additional development effort and maintenance overhead, which is not necessary for this scenario. Moreover, it will not provide an immediate alert when the function endpoints are unavailable or responding too slowly.

upvoted 1 times

 **Sriniv** 1 year, 2 months ago

the best option is to create a URL ping test to monitor the availability and responsiveness of the REST API endpoints. This can be easily configured in Azure Application Insights with minimal effort and can provide immediate alerts when the endpoints are unavailable or responding too slowly.

upvoted 2 times

 **Dhanishetty** 1 year, 5 months ago

Selected Answer: B

its B.....

upvoted 1 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: B

B is correct

upvoted 1 times

 **AZAdam22** 1 year, 8 months ago

Selected Answer: A

A - Because the URL ping test automatically measures availability and response times.

upvoted 2 times

 **Eltooth** 1 year, 9 months ago

Selected Answer: B

B is correct answer.

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/availability-azure-functions>

upvoted 3 times

 **yohannvf** 1 year, 11 months ago

Selected Answer: B

"You need to monitor the availability and responsiveness of the functions." You want it to be on application insights

upvoted 2 times

 **yohannvf** 1 year, 11 months ago

"You need to monitor the availability and responsiveness of the functions." You want it to be on application insights

upvoted 1 times

 **iamstudying** 2 years, 1 month ago

Selected Answer: A

I would go with A - URL Ping satisfy all the requirements. It does not check SSL validity BUT there is no requirement to do so, buddies

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/availability-standard-tests>

upvoted 2 times

 **adone** 2 years, 2 months ago

This scenario does not make any sense.

Consumption plans are not intended for high availability because the app may scale to zero when idle which may will leads to cold start of some minute(s) so it makes no sense to monitor it for availability. With URL Ping, you will pay for the compute resources generated by your monitoring request to your URL Ping which makes this option ridiculous with consumption plan. But monitoring your application with an URL Ping might keep your function app up and running if this is what you expect :D

The option B is the less worse but keep in mind that consumption plan are not intended for high availability (does not need to be monitored). They are intended for asynchronous executions, even high workloads and parallel executions, which does not require to be highly available. For high availability, choose Premium/Dedicated plans.

upvoted 9 times

 **leonidn** 2 years, 3 months ago

Selected Answer: B

I guess an SSL validity check is required. URL ping does not provide this feature. There is no Standard test option. So then, option B is the only one that is viable.

upvoted 3 times

 **ScubaDiver123456** 2 years, 2 months ago

I'm confused. I dont see a requirement to test SSL validity. It only says the API is protected by SSL.

upvoted 2 times

 **HiteshRamnani** 2 years, 3 months ago

B is correct

upvoted 1 times

Question #17

DRAG DROP -

You are developing an application to retrieve user profile information. The application will use the Microsoft Graph SDK.

The app must retrieve user profile information by using a Microsoft Graph API call.

You need to call the Microsoft Graph API from the application.

In which order should you perform the actions? To answer, move all actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

- Create an authentication provider.
- Create a new instance of the GraphServiceClient.
- Invoke the request to the Microsoft Graph API.
- Register the application with the Microsoft identity platform.
- Build a client by using the client app ID.

Answer Area



Correct Answer:

Actions

Answer Area

- Register the application with the Microsoft identity platform.
- Build a client by using the client app ID.
- Create an authentication provider.
- Create a new instance of the GraphServiceClient.
- Invoke the request to the Microsoft Graph API.

Step 1: Register the application with the Microsoft identity platform.

To authenticate with the Microsoft identity platform endpoint, you must first register your app at the Azure app registration portal

Step 2: Build a client by using the client app ID

Step 3: Create an authentication provider

Create an authentication provider by passing in a client application and graph scopes.

Code example:

```
DeviceCodeProvider authProvider = new DeviceCodeProvider(publicClientApplication, graphScopes);
// Create a new instance of GraphServiceClient with the authentication provider.
```

```
GraphServiceClient graphClient = new GraphServiceClient(authProvider);
```

Step 4: Create a new instance of the GraphServiceClient

Step 5: Invoke the request to the Microsoft Graph API

Reference:

<https://docs.microsoft.com/en-us/graph/auth-v2-service>

<https://docs.microsoft.com/en-us/graph/sdks/create-client>

kishe Highly Voted 2 years, 10 months ago

Answer is correct

upvoted 46 times

warchoon 1 year, 1 month ago

Answer is correct, but the question is not.

3,4 are the parts of 2 not actually steps

upvoted 2 times

 **BishopeL**  2 years, 7 months ago

I got this in my exam this morning. The provided answer is correct.
upvoted 34 times

 **BaoNguyen2411**  8 months, 3 weeks ago

got this question on 29/06/2023
upvoted 5 times

 **warchoon** 1 year, 1 month ago

Check out the JavaScript example <https://learn.microsoft.com/en-us/graph/sdks/create-client?tabs=Javascript>
The C# sample does not clearly describe the authentication provider.
upvoted 1 times

 **OPT_001122** 1 year, 4 months ago

The answer is correct - RBCCI
upvoted 10 times

 **AAcademy** 1 year, 4 months ago

I lov that, thanks
upvoted 3 times

 **NombreFalso** 1 year, 1 month ago

Hahaha, that's so classic, creating acronyms for memorizing stuff like this.
upvoted 4 times

 **TheExamMaster2020** 1 year, 5 months ago

Did my exam on 15th November 2022. This question was on it.
upvoted 7 times

 **OPT_001122** 1 year, 4 months ago

Thanks for mentioning the date
upvoted 2 times

Question #18

DRAG DROP -

You develop and deploy an Azure Logic App that calls an Azure Function app. The Azure Function App includes an OpenAPI (Swagger) definition and uses an Azure Blob storage account. All resources are secured by using Azure Active Directory (Azure AD).

The Logic App must use Azure Monitor logs to record and store information about runtime data and events. The logs must be stored in the Azure Blob storage account.

You need to set up Azure Monitor logs and collect diagnostics data for the Azure Logic App.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

- Create action groups and alert rules.
- Create a Log Analytics workspace.
- Install the Logic Apps Management solution.
- Add a diagnostic setting to the Azure Function App.
- Create an Azure storage account.
- Add a diagnostic setting to the Azure Logic App.

Answer Area



Correct Answer:

Actions

- Create action groups and alert rules.
- Add a diagnostic setting to the Azure Function App.
- Create an Azure storage account.

Answer Area

- Create a Log Analytics workspace.
- Install the Logic Apps Management solution.
- Add a diagnostic setting to the Azure Logic App.



Step 1: Create a Log Analytics workspace

Before you start, you need a Log Analytics workspace.

Step 2: Install the Logic Apps Management solution

To set up logging for your logic app, you can enable Log Analytics when you create your logic app, or you can install the Logic Apps Management solution in your

Log Analytics workspace for existing logic apps.

Step 3: Add a diagnostic setting to the Azure Logic App

Set up Azure Monitor logs -

1. In the Azure portal, find and select your logic app.
2. On your logic app menu, under Monitoring, select Diagnostic settings > Add diagnostic setting.

Reference:

<https://docs.microsoft.com/en-us/azure/logic-apps/monitor-logic-apps-log-analytics>

jadda Highly Voted 2 years, 10 months ago

After the March 26 changes, Logic apps are out of scope for the AZ-204 exam

upvoted 33 times

- ✉ **dhishkiyaau** 2 years, 7 months ago
Can you refer the source?
upvoted 2 times
- ✉ **mandusya** 2 years, 3 months ago
confirmed
<https://query.prod.cms.rt.microsoft.com/cms/api/am/binary/RE4oZ7B>
upvoted 4 times
- ✉ **HiteshRammnani** 2 years, 3 months ago
thank you
upvoted 1 times
- ✉ **hawgie** 7 months ago
Got this one today. So much for being out of scope. 2023-09-19
upvoted 9 times
- ✉ **jay158** Highly Voted 2 years, 9 months ago
Answer is correct
<https://docs.microsoft.com/en-us/azure/logic-apps/monitor-logic-apps-log-analytics>
upvoted 12 times
- ✉ **raymond_abcd** Most Recent 2 months, 2 weeks ago
correct, answer is given here:
<https://learn.microsoft.com/en-us/azure/logic-apps/monitor-workflows-collect-diagnostic-data?tabs=consumption#install-management-solution>
upvoted 1 times
- ✉ **JH81** 9 months, 3 weeks ago
I got this on 6/29/2023 and passed with 850. Answer is correct.
upvoted 10 times
- ✉ **OPT_001122** 1 year, 4 months ago
The given answer is correct.
CIA
upvoted 5 times
- ✉ **Knightie** 1 year, 7 months ago
eliminate action and alert.
eliminate storage account creation because the question already has an account.
eliminate diagnostics on function app, you need to put it on the logic app instead.
Then arrange the rest to get the answer.
upvoted 4 times
- ✉ **learner06** 2 years, 7 months ago
don't we need to create Azure storage account in order to store logs in blob?
upvoted 2 times
- ✉ **xahah2222** 2 years, 5 months ago
Such an account is already there
upvoted 2 times
- ✉ **Jurgen1234** 2 years, 9 months ago
First one should Azure storage account as you need to store them in a blob storage
upvoted 1 times
- ✉ **Jurgen1234** 2 years, 9 months ago
Nevermind you already have a storage account
upvoted 1 times
- ✉ **markra** 2 years, 10 months ago
The given answer is correct. See the reference provided in a solution.
upvoted 4 times

Question #19

DRAG DROP -

You develop an application. You plan to host the application on a set of virtual machines (VMs) in Azure.

You need to configure Azure Monitor to collect logs from the application.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

Create a Log Analytics workspace.

Install agents on the VM and VM scale set to be monitored.

Send console logs.

Add a VMInsights solution.

Create an Application Insights resource.

Answer Area**Correct Answer:****Actions**

Send console logs.

Answer Area

Create a Log Analytics workspace.

Add a VMInsights solution.

Install agents on the VM and VM scale set to be monitored.

Create an Application Insights resource.

Step 1: Create a Log Analytics workspace.

First create the workspace.

Step 2: Add a VMInsights solution.

Before a Log Analytics workspace can be used with VM insights, it must have the VMInsights solution installed.

Step 3: Install agents on the VM and VM scale set to be monitored.

Prior to onboarding agents, you must create and configure a workspace. Install or update the Application Insights Agent as an extension for Azure virtual machines and VM scale sets.

Step 4: Create an Application Insights resource

Sign in to the Azure portal, and create an Application Insights resource.

Home > New > Application Insights >

Application Insights

Monitor web app performance and usage

[Basics](#) [Tags](#) [Review + create](#)

Create an Application Insights resource to monitor your live web application. With Application Insights, you have full observability into your application across all components and dependencies of your complex distributed architecture. It includes powerful analytics tools to help you diagnose issues and to understand what users actually do with your app. It's designed to help you continuously improve performance and usability. It works for apps on a wide variety of platforms including .NET, Node.js and Java EE, hosted on-premises, hybrid, or any public cloud. [Learn More](#)

PROJECT DETAILS

Select a subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ

Visual Studio Enterprise

Resource Group * ⓘ

My_Resource_Group

[Create new](#)

INSTANCE DETAILS

Name * ⓘ

My_AppInsights_Resource

Region * ⓘ

(US) West US 2

Resource Mode * ⓘ

Classic **Workspace-based**

WORKSPACE DETAILS

Subscription * ⓘ

Visual Studio Enterprise

Log Analytics Workspace * ⓘ

my-workspace-name [westus2]

[Review + create](#)[« Previous](#)[Next : Tags >](#)

Once a workspace-based Application Insights resource has been created, configuring monitoring is relatively straightforward.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/vm/vminsights-configure-workspace> <https://docs.microsoft.com/en-us/azure/azure-monitor/app/create-workspace-resource>

✉ **aradice** Highly Voted 2 years, 9 months ago

correct <https://docs.microsoft.com/en-us/azure/azure-monitor/vm/vminsights-configure-workspace?tabs=CLI>
<https://docs.microsoft.com/en-us/azure/azure-monitor/app/create-workspace-resource>

upvoted 23 times

✉ **vilainchien** 2 years, 1 month ago

correct <https://docs.microsoft.com/en-us/azure/azure-monitor/vm/vminsights-overview#configuring-vm-insights>
upvoted 3 times

✉ **CarlosTheBoldest** 4 months, 1 week ago

The list is correct but it is on the wrong order.
1. Create Log Analytics workspace (in Azure Portal)
2. Create Application Insights resource (in Monitor, Application Insights with workspace attached)
3. Add VMInsights solution (== activate VMInsights in Monitor, choose VM's, attach workspace).
4. Install agents on VM
upvoted 4 times

✉ **rolling_potato_** Highly Voted 1 year, 3 months ago

Was on my exam today (3rd of Jan 2023)

upvoted 8 times

✉ **adilkhan** 1 year, 2 months ago

what was your answer and your score please?
upvoted 1 times

✉ **Ciupaz** Most Recent 2 months, 3 weeks ago

CAIC to remember

upvoted 3 times

✉ **WH16** 8 months, 1 week ago

This was on the exam (9th August 2023). Went with proposed. Scored 910.

upvoted 3 times

✉ **coffecold** 1 year, 6 months ago

You must make the workspace in the portal

In Monitor you must add Application insights with the workspace

In Monitor you must activate VMInsights and add your VM's and assign a workspace to it

Then install and run the agents in each VM

1. Create Log Analytics workspace (in Azure Portal)

2. Create Application Insights resource (in Monitor, Application Insights with workspace attached)

3. Add VMInsights solution (== activate VMInsights in Monitor, choose VM's, attach workspace).

4. Install agents on VM

upvoted 8 times

✉ **wsguruprasath** 2 years ago

Create Log Analytics workspace.

Add VMInsights solution to workspace.

Install agents on virtual machine and virtual machine scale set to be monitored.

<https://docs.microsoft.com/en-us/azure/azure-monitor/vm/vminsights-overview>

upvoted 1 times

✉ **RajMasilamani** 2 years, 7 months ago

Configuring VM insights

The steps to configure VM insights are as follows. Follow each link for detailed guidance on each step:

1. Create Log Analytics workspace.

2. Add VMInsights solution to workspace.

3. Install agents on virtual machine and virtual machine scale set to be monitored.

upvoted 5 times

✉ **ning** 2 years, 8 months ago

Why application insights???

upvoted 1 times

✉ **nyanko** 2 years, 2 months ago

Is VMInsights required to monitor application log ?

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/azure-vm-vmss-apps>

upvoted 1 times

✉ **ning** 2 years, 8 months ago

Add more stuff ...

If you want to monitor VM, CPU, memory, disk, you should use VM insights, if you want to monitor an app running inside VM, you need application insights ...

Not sure what is the question asking ...

upvoted 3 times

✉ **jvyas** 2 years, 7 months ago

It does say "You need to configure Azure Monitor to collect logs from the application." Hope this helps to discern.

upvoted 1 times

✉ **ning** 2 years, 8 months ago

I think this is the answer

1. Create log name space

2. Create app insights

3. Install app insights agents on VM

4. Send console log for the app

upvoted 4 times

✉ **Maharaja** 2 years, 6 months ago

Thank you. Can you provide some ms doc link supporting your answer?

upvoted 2 times

✉ **adilkhan** 1 year, 2 months ago

why 2 is application insights?

upvoted 1 times

✉ **ning** 2 years, 7 months ago

After lots of research ...

1. Log analytics name space
 2. Application insights
 3. VM solutions
 4. Install agents
- upvoted 5 times

 **adilkhan** 1 year, 2 months ago
why 2 is application insights?
upvoted 1 times

 **finnishr** 2 years, 8 months ago
100% correct <https://docs.microsoft.com/en-us/azure/azure-monitor/vm/vminsights-overview>
upvoted 7 times

Question #20

Topic 5

You develop and deploy an Azure App Service web app. The app is deployed to multiple regions and uses Azure Traffic Manager. Application Insights is enabled for the app.

You need to analyze app uptime for each month.

Which two solutions will achieve the goal? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. Azure Monitor logs
- B. Application Insights alerts
- C. Azure Monitor metrics
- D. Application Insights web tests

Correct Answer: *BD*

Reference:

<https://azure.microsoft.com/en-us/blog/creating-a-web-test-alert-programmatically-with-application-insights/>

Community vote distribution

AC (52%) BC (26%) CD (17%) 4%

✉  **coffecold** Highly Voted 1 year, 6 months ago

A and C. Metrics will give you the uptime. Logs will give you the causes of the downtime.

wrong B: Alerts are not required

wrong D: Web tests has nothing to do with uptime.

upvoted 21 times

✉  **130nk3r5** 4 months ago

Application Insights web tests Application Insights web tests can be used to monitor the availability of your app. It sends requests to your application at regular intervals from points around the world. It alerts you when your app is slow or failing, and provides diagnostic information to help you find and fix the issue. This can be used to analyze the uptime of your app for each month.

upvoted 1 times

✉  **Ciupaz** 5 months, 2 weeks ago

Not correct for D:

These web tests can help you track the uptime and availability of your application across different regions and provide valuable insights into its performance.

upvoted 1 times

✉  **qwerty112233** Highly Voted 1 year, 7 months ago

"You need to analyze app uptime for each month." how web tests and alerts will help you to analyze previous data?

upvoted 9 times

✉  **edmond** 1 year, 4 months ago

Certainly web tests don't help, but for alerts, you can get use of the availability alerts, refer to <https://learn.microsoft.com/en-us/azure/azure-monitor/app/availability-alerts>.

I don't think the Metrics logs work here, hence BC makes more sense.

upvoted 2 times

✉  **edmond** 1 year, 4 months ago

Sorry I meant the Monitor logs don't work in this case, the Monitor metrics and Application Insight alerts make more sense here.

upvoted 2 times

✉  **micro9000** 1 year, 5 months ago

We are analyzing the app uptime not the resources or errors that our azure resources or the application generates, if you go to your App Insight instance -> Availability then click the SLA Report you will see the uptime report of your application there. That is the SLA report of your application.

upvoted 2 times

✉  **FeriAZ** Most Recent 2 months, 3 weeks ago

Selected Answer: AC

Azure Monitor logs (A) - because they allow for the collection and complex querying of detailed operational data, including uptime.

Azure Monitor metrics (C) - as they provide real-time and historical data analysis capabilities, suitable for assessing uptime over a specified period like a month.

upvoted 1 times

✉  **130nk3r5** 4 months ago

Selected Answer: AD

The two solutions that will achieve the goal are:

Azure Monitor logs Azure Monitor logs can be used to analyze the uptime of the app. It collects telemetry and other data from your Azure and non-Azure resources, and provides a consolidated view for you to analyze. You can create queries to analyze the data and derive insights about the uptime of your app.

Application Insights web tests Application Insights web tests can be used to monitor the availability of your app. It sends requests to your application at regular intervals from points around the world. It alerts you when your app is slow or failing, and provides diagnostic information to help you find and fix the issue. This can be used to analyze the uptime of your app for each month.

Please note that while Azure Monitor metrics (C) and Application Insights alerts (B) are useful for monitoring and alerting, they do not provide the necessary data for analyzing app uptime over a period of a month.

upvoted 1 times

 **Ciupaz** 5 months ago

Selected Answer: CD

For me: C and D.

D. Application Insights web tests

Application Insights web tests allow you to simulate user traffic to your web app and monitor its performance. You can use Application Insights web tests to measure the uptime of your web app, and you can view this data for each month.

C. Azure Monitor metrics

Azure Monitor metrics provide a way to collect and analyze data about the performance and health of your Azure resources. You can use Azure Monitor metrics to track the uptime of your Azure App Service web app, and you can view this data for each month.

upvoted 1 times

 **EliteAllen** 8 months, 3 weeks ago

Selected Answer: AC

A. Azure Monitor logs: By querying logs, you can analyze detailed information about app uptime and other metrics over a specified time period, such as a month.

C. Azure Monitor metrics: Metrics provide numerical data about your system's operation, including uptime. You can use Azure Monitor metrics to track and analyze the availability and responsiveness of your app over time, including on a monthly basis.

upvoted 3 times

 **motekim** 1 year ago

I like A, D why?

A: <https://www.azure.cn/en-us/support/sla/application-insights/index.html>

d: <https://learn.microsoft.com/en-us/azure/azure-monitor/app/availability-overview>

upvoted 1 times

 **Sriniv** 1 year, 2 months ago

Definitely not A &B due to following reason

Option A, Azure Monitor logs, is not the best solution for analyzing app uptime. While Azure Monitor logs can provide information on app usage and performance, they may not provide the necessary information to calculate the app uptime percentage for each month.

Option B, Application Insights alerts, can be useful for notifying you of any issues with the app, but they may not provide the necessary information to calculate the app uptime percentage for each month.

upvoted 1 times

 **Osip21** 1 year, 2 months ago

Selected Answer: CD

ChatGPT:

C. Azure Monitor metrics

D. Application Insights web tests

Azure Monitor metrics can be used to track the availability of the app by analyzing the uptime percentage. Application Insights web tests can be configured to periodically check the availability of the app and provide uptime data.

upvoted 3 times

 **alexein74** 1 year, 2 months ago

Selected Answer: AC

A. Azure Monitor logs

C. Azure Monitor metrics

Azure Monitor logs: Azure Monitor logs provides detailed diagnostic and operational information from Azure resources and allows you to store and analyze the logs. By querying Azure Monitor logs, you can analyze app uptime for each month by searching for uptime and downtime events.

Azure Monitor metrics: Azure Monitor metrics allow you to monitor the performance and health of your Azure resources. With Azure Monitor metrics, you can analyze app uptime by looking at the percentage of availability for the app, and you can set up alerts if uptime falls below a certain threshold.

upvoted 2 times

 **rotimislaw** 1 year, 3 months ago

Selected Answer: BC

B. Azure Monitor metrics - can be set to count uptime

C. Application Insights alerts - flag when servers go down

upvoted 3 times

 **KingChuang** 1 year, 3 months ago

On my exam 2022-12-26.

Chose: A, C

upvoted 4 times

 **bertwin** 1 year, 4 months ago

Selected Answer: AC

A & C i would say

<https://learn.microsoft.com/en-us/azure/azure-monitor/essentials/metrics-supported>

upvoted 2 times

 **TheExamMaster2020** 1 year, 5 months ago

Did my exam on 15th November 2022. This question was on it.

upvoted 3 times

 **somename20221106** 1 year, 5 months ago

What is the correct answer ?

upvoted 3 times

 **OPT_001122** 1 year, 4 months ago

correct ans?

upvoted 1 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: AC

A and C.

upvoted 3 times

 **Yumi21** 1 year, 5 months ago

Selected Answer: AC

Shouldn't it be Monitor logs instead of alerts?

upvoted 1 times

 **g2000** 1 year, 3 months ago

<https://learn.microsoft.com/en-us/azure/azure-monitor/app/monitor-web-app-availability#alerts>

upvoted 1 times

 **sof123** 1 year, 5 months ago

BC - Azure monitor metrics and Application insights alerts

upvoted 1 times

Question #21

DRAG DROP -

You develop and deploy an Azure App Service web app. The web app accesses data in an Azure SQL database.

You must update the web app to store frequently used data in a new Azure Cache for Redis Premium instance.

You need to implement the Azure Cache for Redis features.

Which feature should you implement? To answer, drag the appropriate feature to the correct requirements. Each feature may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Answer Area

Features	Requirement	Feature
horizontal partitioning	Create a data structure for storing collections of related items	Feature
channel	Create a data structure for the most recently accessed cache items	Feature
list	Send messages through a high-performance publisher/subscriber mechanism	Feature
set		

Correct Answer:

Answer Area

Features	Requirement	Feature
horizontal partitioning	Create a data structure for storing collections of related items	set
channel	Create a data structure for the most recently accessed cache items	list
list	Send messages through a high-performance publisher/subscriber mechanism	channel
set		

Reference:

<https://www.red-gate.com/simple-talk/development/dotnet-development/overview-of-azure-cache-for-redis/> <https://docs.microsoft.com/en-us/azure/architecture/best-practices/caching>

baroo1 Highly Voted 1 year, 1 month ago

This is correct!

<https://docs.microsoft.com/en-us/azure/architecture/best-practices/caching>

If you look for the right words (related > set | high-performance > channel | recent > list) in documentation, it helps you validate this answer. Ignore the other comments as it only leads to confusion.

upvoted 12 times

juancar Highly Voted 9 months, 2 weeks ago

This was on the exam (July 2023). Went with proposed. Scored 917

upvoted 10 times

130nk3r5 Most Recent 4 months ago

Here are the appropriate features for each requirement:

Create a data structure for storing collections of related items: Set A set in Redis is an unordered collection of unique items. It is suitable for storing collections of related items.

Create a data structure for the most recently accessed cache items: List A list in Redis is an ordered collection of items. It can be used to store the most recently accessed cache items, as you can easily add new items to the front of the list and remove the least recently accessed items from the end of the list.

Send messages through high-performance publisher/subscriber mechanism: Channel A channel in Redis is used for the publish/subscribe messaging pattern. It allows you to send messages through a high-performance publisher/subscriber mechanism.

upvoted 2 times

Sriniv 1 year, 2 months ago

represents a collection of unique items, where each item can only appear once in the set. While a set can be used to store a collection of related items, it may not be the most suitable data structure for this requirement. This is because a set does not maintain the order of the items and does not provide an efficient way to search for a specific item.

upvoted 1 times

✉ **Sriniv** 1 year, 2 months ago

Here are the features that match the given requirements:

Feature: horizontal partitioning

Requirement: Create a data structure for storing collections of related items

Feature: list

Requirement: Create a data structure for the most recently accessed cache items

Feature: channel

Requirement: Send messages through a high-performance publisher/subscriber mechanism

upvoted 2 times

✉ **aruni_mishra** 1 year, 4 months ago

Set: for related items

<https://learn.microsoft.com/en-us/azure/architecture/best-practices/caching#use-tags-to-cross-correlate-cached-items>

List: to find the most recently accessed items

<https://learn.microsoft.com/en-us/azure/architecture/best-practices/caching#find-recently-accessed-items>

upvoted 5 times

✉ **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 8 times

✉ **OPT_001122** 1 year, 5 months ago

thnx for the date

upvoted 1 times

✉ **Knightie** 1 year, 7 months ago

<https://redis.io/docs/data-types/sets/>

upvoted 1 times

✉ **qwerty112233** 1 year, 7 months ago

Correct

<https://docs.microsoft.com/en-us/azure/architecture/best-practices/caching>

upvoted 3 times

✉ **sghaha** 1 year, 11 months ago

caching (in Korean)

<https://docs.microsoft.com/ko-kr/azure/architecture/best-practices/caching>

upvoted 2 times

Question #22

Topic 5

You are developing an ASP.NET Core Web API web service. The web service uses Azure Application Insights for all telemetry and dependency tracking. The web service reads and writes data to a database other than Microsoft SQL Server. You need to ensure that dependency tracking works for calls to the third-party database. Which two dependency telemetry properties should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Telemetry.Context.Cloud.RoleInstance
- B. Telemetry.Id
- C. Telemetry.Name
- D. Telemetry.Context.Operation.Id
- E. Telemetry.Context.Session.Id

Correct Answer: BD

Example:

```
public async Task Enqueue(string payload)
{
    // StartOperation is a helper method that initializes the telemetry item
    // and allows correlation of this operation with its parent and children. var operation = telemetryClient.StartOperation<DependencyTelemetry>
    ("enqueue " + queueName);

    operation.Telemetry.Type = "Azure Service Bus";
    operation.Telemetry.Data = "Enqueue " + queueName;
    var message = new BrokeredMessage(payload);
    // Service Bus queue allows the property bag to pass along with the message.
    // We will use them to pass our correlation identifiers (and other context)
    // to the consumer.
    message.Properties.Add("ParentId", operation.Telemetry.Id);
    message.Properties.Add("RootId", operation.Telemetry.Context.Operation.Id);
```

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/custom-operations-tracking>

Community vote distribution

BD (100%)

✉  **cyberbull**  3 years, 6 months ago

the Given answer is correct IMHO

upvoted 45 times

✉  **10x** 3 years, 2 months ago

Yes. Makes sense to me too.

upvoted 8 times

✉  **mlantonis**  2 years, 10 months ago

Correct Answer: B and D

```
message.Properties.Add("ParentId", operation.Telemetry.Id);
message.Properties.Add("RootId", operation.Telemetry.Context.Operation.Id);
```

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/custom-operations-tracking#enqueue>

upvoted 18 times

✉  **alexein74**  1 year, 2 months ago

Selected Answer: BD

Answer is correct

upvoted 1 times

✉  **OPT_001122** 1 year, 5 months ago

Selected Answer: BD

correct ans

upvoted 1 times

 **AzureDJ** 2 years, 1 month ago

Answer is correct

upvoted 2 times

 **Ravindu** 2 years, 6 months ago

The given answer is correct.

upvoted 1 times

 **glam** 2 years, 11 months ago

Correct

upvoted 2 times

 **Immy_08** 2 years, 11 months ago

The given answer is correct, also available on udemy course

upvoted 3 times

 **vb3d** 3 years ago

Options form the answer mentioned in the code on <https://docs.microsoft.com/en-us/azure/azure-monitor/app/custom-operations-tracking#enqueue>

upvoted 2 times

 **paru123456789** 3 years, 1 month ago

Answer:

Telemetry.Id

Telemetry.Context.Operation.Id

upvoted 3 times

Question #23

HOTSPOT -

You are using Azure Front Door Service.

You are expecting inbound files to be compressed by using Brotli compression. You discover that inbound XML files are not compressed. The files are 9 megabytes (MB) in size.

You need to determine the root cause for the issue.

To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statement	Yes	No
-----------	-----	----

The file MIME type is supported by the service.

Edge nodes must be purged of all cache assets.

The compression type is supported.

Answer Area

Statement	Yes	No
-----------	-----	----

Correct Answer: The file MIME type is supported by the service.

Edge nodes must be purged of all cache assets.

The compression type is supported.

Box 1: No -

Front Door can dynamically compress content on the edge, resulting in a smaller and faster response to your clients. All files are eligible for compression.

However, a file must be of a MIME type that is eligible for compression list.

Box 2: No -

Sometimes you may wish to purge cached content from all edge nodes and force them all to retrieve new updated assets. This might be due to updates to your web application, or to quickly update assets that contain incorrect information.

Box 3: Yes -

These profiles support the following compression encodings: Gzip (GNU zip), Brotli

Reference:

<https://docs.microsoft.com/en-us/azure/frontdoor/front-door-caching>

 **bensagar10**  3 years, 5 months ago

1. Yes
 2. No Refer - <https://docs.microsoft.com/en-us/azure/frontdoor/front-door-caching>, You don't need to purge the cache assets. Here the issue is that the file size needs to be less than 8MB
 3. Yes
- upvoted 184 times

✉ **Tealon** 3 years, 4 months ago

I agree with this answer.
upvoted 1 times

✉ **profesorklaus** 3 years, 3 months ago

I agree also.
upvoted 1 times

✉ **AfroYeti** 3 years, 2 months ago

You could argue that compression isn't supported due to the file size

"Additionally, the file must also be between 1 KB and 8 MB in size, inclusive."

upvoted 2 times

✉ **jkes80** 2 years, 9 months ago

Compression isn't but the answer states "compression TYPE". And Brotli is supported.
upvoted 4 times

✉ **Juanlu** 3 years, 1 month ago

I Agree !
upvoted 3 times

✉ **AakashNeedsEmAll** Highly Voted 3 years, 5 months ago

application/xml is a MIME type and therefore, is supported. Answer 1 should be yes.
upvoted 26 times

✉ **FeriAZ** Most Recent 2 months ago

1. Yes
2. No
3. Yes
upvoted 1 times

✉ **WH16** 8 months, 1 week ago

This was on the exam (9th August 2023). Went with proposed. Scored 910.
upvoted 2 times

✉ **sundeeprch1** 9 months, 3 weeks ago

1. No
2. Yes
3. Yes
upvoted 1 times

✉ **wrak** 1 year, 1 month ago

IMHO, Yes/No answers should be based on the details provided in the question, hence:
1. Yes - application/xml is supported
2. No - if we will purge cache it will still not compress the file
3. Yes - statement is about compression type and Brotli is supported
upvoted 3 times

✉ **rolling_potato_** 1 year, 3 months ago

Was on my exam today (3rd of Jan 2023), went with given answer
upvoted 3 times

✉ **adilkhan** 1 year, 2 months ago

you mean No, Yes and Yes? your score please
upvoted 1 times

✉ **OPT_001122** 1 year, 5 months ago

Why are the admins not correcting the answers having gone thru the highly voted discussion comments!!!
upvoted 6 times

✉ **OPT_001122** 1 year, 5 months ago

Yes
No
Yes
upvoted 2 times

✉ **Eltooth** 1 year, 9 months ago

Yes
No
Yes
upvoted 5 times

✉ **minaritochuck** 2 years ago

<https://docs.microsoft.com/en-us/azure/frontdoor/standard-premium/how-to-compression>

In Azure Front Door, only eligible files are compressed. To be eligible for compression, a file must:

- Be of a MIME type
- Be larger than 1 KB
- Be smaller than 8 MB

These profiles support the following compression encodings:

- gzip (GNU zip)
- brotli

If the request supports more than one compression type, brotli compression takes precedence.

I guess the root cause is because xml is 9MB in size.

1. YES : xml = MIME type
2. NO : I don't think purging cache is necessary..
3. YES - brotli is supported

upvoted 6 times

 **Azprep** 2 years ago

1. Yes
2. No
3. Yes

upvoted 2 times

 **Prasu69** 2 years ago

1. NO - the file must be of a MIME type to be eligible for compression. application/xml or text/xml are not supported.
2. Yes - size should be in 1- 8MB chunks. (9 MB divides into 8MB chunk and 1 MB chunk)
3. Yes - Brotli Compression method is supported.

upvoted 1 times

 **Prasu69** 2 years ago

1. NO - the file must be of a MIME type to be eligible for compression. application/xml or text/xml are not supported.
2. No - size should be in 1- 8MB chunks.
3. Yes - Brotli Compression method is supported.

upvoted 1 times

 **petibilly** 2 years, 1 month ago

Got it in exam 03/22

upvoted 3 times

 **SivajiTheBoss** 2 years, 1 month ago

Question: What is the root cause for the issue, So

Answer:

1. No
2. No
3. Yes (because the file size greater than 8 mb)

upvoted 1 times

 **kozchris** 2 years, 1 month ago

no - the files are getting stored, there is no blocking of the XML files

yes - maybe files were cached prior to some setting getting configured

yes - you need to make sure the compression is supported.

upvoted 1 times

Question #24

HOTSPOT -

You are developing an Azure App Service hosted ASP.NET Core web app to deliver video-on-demand streaming media. You enable an Azure Content Delivery

Network (CDN) Standard for the web endpoint. Customer videos are downloaded from the web app by using the following example URL:

<http://www.contoso.com/content.mp4?quality=1>.

All media content must expire from the cache after one hour. Customer videos with varying quality must be delivered to the closest regional point of presence

(POP) node.

You need to configure Azure CDN caching rules.

Which options should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Setting	Action
Caching behavior	<div style="border: 1px solid black; padding: 5px; display: inline-block;"> Bypass cache Override Set if missing </div>
Cache expiration duration	<div style="border: 1px solid black; padding: 5px; display: inline-block;"> 1 second 1 minute 1 hour 1 day </div>
Query string caching behavior	<div style="border: 1px solid black; padding: 5px; display: inline-block;"> Ignore query strings Bypass caching for query strings Cache every unique URL </div>

Answer Area

Setting	Action
Caching behavior	<div style="border: 1px solid black; padding: 5px; display: inline-block;"> Bypass cache Override Set if missing </div>
Correct Answer: Cache expiration duration	<div style="border: 1px solid black; padding: 5px; display: inline-block;"> 1 second 1 minute 1 hour 1 day </div>
Query string caching behavior	<div style="border: 1px solid black; padding: 5px; display: inline-block;"> Ignore query strings Bypass caching for query strings Cache every unique URL </div>

Box 1: Override -

Override: Ignore origin-provided cache duration; use the provided cache duration instead. This will not override cache-control: no-cache.

Set if missing: Honor origin-provided cache-directive headers, if they exist; otherwise, use the provided cache duration.

Incorrect:

Bypass cache: Do not cache and ignore origin-provided cache-directive headers.

Box 2: 1 hour -

All media content must expire from the cache after one hour.

Box 3: Cache every unique URL -

Cache every unique URL: In this mode, each request with a unique URL, including the query string, is treated as a unique asset with its own cache. For example, the response from the origin server for a request for example.ashx?q=test1 is cached at the POP node and returned for subsequent caches with the same query string. A request for example.ashx?q=test2 is cached as a separate asset with its own time-to-live setting.

Incorrect Answers:

Bypass caching for query strings: In this mode, requests with query strings are not cached at the CDN POP node. The POP node retrieves the asset directly from the origin server and passes it to the requestor with each request.

Ignore query strings: Default mode. In this mode, the CDN point-of-presence (POP) node passes the query strings from the requestor to the origin server on the first request and caches the asset. All subsequent requests for the asset that are served from the POP ignore the query strings until the cached asset expires.

Reference:

<https://docs.microsoft.com/en-us/azure/cdn/cdn-query-string>

✉  **RaviKS**  3 years, 3 months ago

Given Answers are correct

upvoted 70 times

✉  **mlantonis**  2 years, 10 months ago

Box 1: Override

Override: Ignore origin-provided cache duration; use the provided cache duration instead. This will not override cache-control: no-cache.

Set if missing: Honor origin-provided cache-directive headers, if they exist; otherwise, use the provided cache duration.

Bypass cache: Do not cache and ignore origin-provided cache-directive headers.

Box 2: 1 hour

All media content must expire from the cache after one hour.

upvoted 29 times

✉  **mlantonis** 2 years, 10 months ago

Reference:

<https://docs.microsoft.com/en-us/azure/cdn/cdn-query-string>

<https://docs.microsoft.com/en-us/azure/cdn/cdn-caching-rules>

upvoted 9 times

✉  **mlantonis** 2 years, 10 months ago

Box 3: Cache every unique URL

Cache every unique URL: In this mode, each request with a unique URL, including the query string, is treated as a unique asset with its own cache. For example, the response from the origin server for a request for example.ashx?q=test1 is cached at the POP node and returned for subsequent caches with the same query string. A request for example.ashx?q=test2 is cached as a separate asset with its own time-to-live setting.

Bypass caching for query strings: In this mode, requests with query strings are not cached at the CDN POP node. The POP node retrieves the asset directly from the origin server and passes it to the requestor with each request.

Ignore query strings: Default mode. In this mode, the CDN point-of-presence (POP) node passes the query strings from the requestor to the origin server on the first request and caches the asset. All subsequent requests for the asset that are served from the POP ignore the query strings until the cached asset expires.

upvoted 20 times

✉  **AndySmith**  5 months, 2 weeks ago

On exam 3-Nov-2023. Went with proposed answer - 932/1000.

upvoted 1 times

✉  **Ciupaz** 5 months ago

How do you remember all 375 questions?

upvoted 1 times

✉  **jaf19f** 8 months, 1 week ago

I got this question (12-Aug-2023) and I chose the given answer - 932 passed

upvoted 3 times

✉ **rolling_potato_** 1 year, 3 months ago

Got it in my exam today (3rd of Jan)

upvoted 2 times

✉ **adilkhan** 1 year, 2 months ago

are these questions enough to pass the exam?

upvoted 1 times

✉ **petitbilly** 2 years, 1 month ago

Got it in exam 03/22

upvoted 3 times

✉ **AzureDJ** 2 years, 1 month ago

Hope your exam went well. This answer looks right. Did you choose it?

upvoted 1 times

✉ **HimanshuNankani** 2 years, 6 months ago

1. Override - because we are explicitly defining the duration to be 1 hour .

We DO NEED Cache (hence it cannot be bypass cache)

We need EXACT 1 HOUR (so can't use Set If Missing which could change the specific value)

2. 1 Hour as required

3. We do need to 'Cache every unique url' because we need to provide cache for different qualities that is specified in the query parameter.

upvoted 10 times

✉ **MK22** 2 years, 8 months ago

Provided answer is correct

upvoted 3 times

✉ **glam** 2 years, 11 months ago

correct.

upvoted 2 times

✉ **rakelcoelho** 3 years ago

Didn't understand. On the first box it can be «override» or «set if missing»?

upvoted 4 times

✉ **titombo** 3 years ago

Me too, but according MS page the concepts are:

Override: Ignore origin-provided cache duration; use the provided cache duration instead. This will not override cache-control: no-cache.
Set if missing: Honor origin-provided cache-directive headers, if they exist; otherwise, use the provided cache duration.

What I don't understand is:

origin-provided cache duration: is the cache provided in Azure, right?

So in other words it would something like:

Override cache means that it will always use the cache provided in Azure.

Set if missing, this means it can use the client cache request in the request header property.

Am I wrong?

upvoted 3 times

✉ **Cricketer** 2 years, 12 months ago

Not really, origin means your backend servers

upvoted 1 times

✉ **profesorklaus** 3 years, 2 months ago

What hasn't been mentioned here is when you try to do on azure portal you should use cdn from Akamai or Verizon but not from Microsoft.

Microsoft has default caching rules which can not be changed using portal

upvoted 1 times

✉ **yerisof** 3 years, 4 months ago

why is this not ignore query string

upvoted 1 times

✉ **fredrodlima** 3 years, 4 months ago

I guess is because each video with different qualities should be cached for upcoming requests to it from users

upvoted 19 times

✉ **Tealon** 3 years, 4 months ago

That is correct.

upvoted 3 times

✉ **TakumaK** 2 years, 10 months ago

As the item in the question, the query has the quality field. this is intended to maintain the cache based on the quality.
upvoted 2 times

 **Nyaku** 3 years, 4 months ago

For more reference to the answer .. <https://docs.microsoft.com/en-us/azure/cdn/cdn-caching-rules>

upvoted 11 times

Question #25

HOTSPOT -

You are developing an ASP.NET Core time sheet application that runs as an Azure Web App. Users of the application enter their time sheet information on the first day of every month.

The application uses a third-party web service to validate data.

The application encounters periodic server errors due to errors that result from calling a third-party web server. Each request to the third-party server has the same chance of failure.

You need to configure an Azure Monitor alert to detect server errors unrelated to the third-party service. You must minimize false-positive alerts.

How should you complete the Azure Resource Manager template? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

```

"type": "Microsoft.Insights/metricAlerts",
"properties": {
  "criteria": {
    "odata.type": ". . .",
    "allOf": [
      {
        "criterionType": ",  


```

```

        "metricName": ",  


```

```

        "alertSensitivity": ",  


```

]] }

Correct Answer:

```

"type": "Microsoft.Insights/metricAlerts",
"properties": {
  "criteria": {
    "odata.type": ". . .",
    "allOf": [
      {
        "criterionType": ",  


```

```

        "metricName": ",  


```

```

        "alertSensitivity": ",  


```

]] }

Box 1: DynamicThresholdCriterion

Box 2: Http5xx -

Server errors are in the 5xx range.

Client errors are in the 4xx range

Box 3: Low -

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/alerts/alerts-dynamic-thresholds>

DynamicThresholdCriterion

Http5XX

Low

upvoted 13 times

✉ **coffecold** Highly Voted 1 year, 6 months ago

DynamicThresholdCriterion: Dynamic thresholds in metric alerts use advanced machine learning to learn metric Http5XX, just a name, we are monitoring server errors, so 500 range error responses

AlertSensitivity: would set is to high. "You must minimize false-positive alerts" The thresholds will be tight and close to the metric series pattern. An alert rule will be triggered on the smallest deviation, resulting in more alerts.

upvoted 7 times

✉ **Yumi21** 1 year, 5 months ago

I think sensitivity should be low due to "false positive" requirements, if we set to high we report on every smallest deviation, and it means there is a high chance of false positive

upvoted 11 times

✉ **b1ade** 1 year, 4 months ago

Agree. Definition according to MS:

High: The thresholds will be tight and close to the metric series pattern. An alert rule will be triggered on the smallest deviation, resulting in more alerts.

Medium: The thresholds will be less tight and more balanced. There will be fewer alerts than with high sensitivity (default).

Low: The thresholds will be loose with more distance from metric series pattern. An alert rule will only trigger on large deviations, resulting in fewer alerts.

upvoted 6 times

✉ **warchoon** Most Recent 1 year, 1 month ago

"SingleResourceMultipleMetricCriteria" is used for "odata.type", not "criterionType".

upvoted 3 times

✉ **ozzy123** 1 year, 2 months ago

It says "errors unrelated to the third-party service". So, I think it should be Http4xx. Not 5xx.

upvoted 1 times

✉ **surprise0011** 1 year ago

but you need to detect server errors:

"You need to configure an Azure Monitor alert to detect server errors..."

upvoted 1 times

✉ **sghaha** 1 year, 11 months ago

about alerts-dynamic-thresholds (in Korean)

<https://docs.microsoft.com/ko-kr/azure/azure-monitor/alerts/alerts-dynamic-thresholds>

upvoted 2 times

✉ **ArturKon** 1 year, 7 months ago

You can provide us links in other language. Just change ko-kr in URL to, (for ex.) en-us

upvoted 5 times

✉ **kerpse** 1 year, 5 months ago

Or in preferred language by just removing the language code, e.g.

<https://docs.microsoft.com/azure/azure-monitor/alerts/alerts-dynamic-thresholds>

upvoted 4 times

Question #26

Topic 5

You are developing a web application that uses Azure Cache for Redis. You anticipate that the cache will frequently fill and that you will need to evict keys.

You must configure Azure Cache for Redis based on the following predicted usage pattern: A small subset of elements will be accessed much more often than the rest.

You need to configure the Azure Cache for Redis to optimize performance for the predicted usage pattern.

Which two eviction policies will achieve the goal?

NOTE: Each correct selection is worth one point.

- A. noevasion
- B. allkeys-lru
- C. volatile-lru
- D. allkeys-random
- E. volatile-ttl
- F. volatile-random

Correct Answer: BC

B: The allkeys-lru policy evict keys by trying to remove the less recently used (LRU) keys first, in order to make space for the new data added.

Use the allkeys-lru policy when you expect a power-law distribution in the popularity of your requests, that is, you expect that a subset of elements will be accessed far more often than the rest.

C: volatile-lru: evict keys by trying to remove the less recently used (LRU) keys first, but only among keys that have an expire set, in order to make space for the new data added.

Note: The allkeys-lru policy is more memory efficient since there is no need to set an expire for the key to be evicted under memory pressure.

Reference:

<https://redis.io/topics/lru-cache>

Community vote distribution

BC (100%)

 **sghaha**  1 year, 11 months ago

my choices are BC

upvoted 9 times

 **qwert112233** 1 year, 7 months ago

agree!

upvoted 3 times

 **alexein74**  1 year, 2 months ago

Selected Answer: BC

B. allkeys-lru

C. volatile-lru

allkeys-lru : This eviction policy evicts the least recently used (LRU) key among all keys, regardless of the key's time-to-live (TTL) value. This policy is suitable when you have a small subset of elements that will be accessed much more often than the rest, as the most frequently used elements will be kept in cache while the least frequently used elements will be evicted.

volatile-lru: This eviction policy evicts the least recently used (LRU) key among keys with an expiring TTL. This policy is suitable when you have a small subset of elements that will be accessed much more often than the rest and also when you want to expire keys that are not used frequently.

upvoted 7 times

 **alinuxguru70**  7 months ago

allkeys-lru would include volatile-lru, it is redundant. Therefore B allkeys-lru and E volatile-ttl.

upvoted 1 times

 **Ciupaz** 7 months ago

"least recently used", not "less recently used". Take care.

upvoted 1 times

 **AymanAkk** 1 year, 3 months ago

Selected Answer: BC

correct answers.

upvoted 1 times

 **KingChuang** 1 year, 3 months ago

On my exam 2022-12-26.

Chose: B, C

upvoted 2 times

 **OPT_001122** 1 year, 4 months ago

Selected Answer: BC

B. allkeys-lru

C. volatile-lru

upvoted 2 times

 **ExamTopicsAnkur** 1 year, 4 months ago

Selected Answer: BC

<https://redis.io/docs/reference/eviction/>

upvoted 2 times

 **TheExamMaster2020** 1 year, 5 months ago

Did my exam on 15th November 2022. This question was on it.

upvoted 4 times

Question #27

DRAG DROP -

An organization has web apps hosted in Azure.

The organization wants to track events and telemetry data in the web apps by using Application Insights.

You need to configure the web apps for Application Insights.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions	Answer Area
Configure the Azure App Service SDK for the app	
Configure the Application Insights SDK in the app	
Copy the connection string	
Create an Azure Machine Learning workspace	
Create an Application Insights resource	

Correct Answer:

Actions	Answer Area
Configure the Azure App Service SDK for the app	
Create an Azure Machine Learning workspace	
Create an Application Insights resource	

Step 1: Create an Application Insights resource

Creating an Application Insights workspace-based resource is a prerequisite.

Step 2: Copy the connection string

A connection string identifies the resource that you want to associate with your telemetry data. It also allows you to modify the endpoints that your resource will use as a destination for your telemetry. You'll need to copy the connection string and add it to your application's code or to an environment variable.

Step 3: Configure the Application Insights SDK in the app

The Application Insights SDK for ASP.NET Core can monitor your applications no matter where or how they run.

Install the Application Insights SDK NuGet package for ASP.NET Core.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/asp-net-core>

 **osagiestar** Highly Voted 1 year, 6 months ago

The given answer is correct.

1. Create an Application Insights resource
 2. Copy the instrumentation key
 3. Install the SDK in your app
- <https://learn.microsoft.com/en-us/azure/azure-monitor/app/create-new-resource>

upvoted 17 times

 **NombreFalso** Highly Voted 1 year, 1 month ago

Why would you make a Machine Learning Workspace LOL

upvoted 9 times

✉  **Ciupaz** Most Recent 5 months ago

"Connection string" is different from "instrumentation key". This question is bad formulated.

Anyway the steps are:

- 1) Create an Application Insights resource
- 2) Copy the instrumentation key
- 3) Install the SDK in your app

upvoted 1 times

✉  **threeslice91** 9 months, 2 weeks ago

The order of 2nd and 3rd does it matter?

upvoted 2 times

✉  **Heighte** 1 year ago

bad question, connection string != instrumentation key

upvoted 4 times

✉  **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 6 times

✉  **willchenxa** 1 year, 7 months ago

1. Create resource
 2. install Application insight SDK
 3. setup connection string
- according to provided link: <https://docs.microsoft.com/en-us/azure/azure-monitor/app/asp-net-core>

upvoted 5 times

✉  **AymanAkk** 1 year, 3 months ago

yes for a setup , but here we have the " copy " word

upvoted 2 times

✉  **finnishr** 1 year, 7 months ago

Correctmundo.

upvoted 2 times

Question #28

Topic 5

An organization hosts web apps in Azure. The organization uses Azure Monitor.

You discover that configuration changes were made to some of the web apps.

You need to identify the configuration changes.

Which Azure Monitor log should you review?

- A. AppServiceAppLogs
- B. AppServiceEnvironmentPlatformLogs
- C. AppServiceConsoleLogs
- D. AppServiceAuditLogs

Correct Answer: B

The log type AppServiceEnvironmentPlatformLogs handles the App Service Environment: scaling, configuration changes, and status logs.

Incorrect:

AppServiceAppLogs contains logs generated through your application.

AppServiceAuditLogs logs generated when publishing users successfully log on via one of the App Service publishing protocols.

Reference:

<https://docs.microsoft.com/en-us/azure/app-service/troubleshoot-diagnostic-logs>

Community vote distribution

 B (100%)

✉  **serpевi**  1 year, 7 months ago

Got it 10/09/2022. Went with B. Score 927

upvoted 15 times

✉  **coffecold**  1 year, 6 months ago

AppServiceEnvironmentPlatformLogs only logs when app services are hosted in an App Service Environment which is a deployment of Azure App Service into a subnet in your Azure Virtual Network (VNet). It is a Premium service option

The problem is, all other answers are wrong :

AppServiceAppLogs : logs from inside your app

AppServiceConsoleLogs : logs from the console

AppServiceAuditLogs : logs who are logged in or out.

but...

AppServiceFileAuditLogs will do the job for premium tiers.

<https://ruslany.net/2020/01/how-to-monitor-azure-app-service-content-changes-and-publishing-activity/>

upvoted 5 times

✉  **Dilmurod** 1 year ago

AppServiceEnvironmentPlatformLogs - App Service Environment: scaling, configuration changes, and status logs

AppServiceFileAuditLogs - File changes made to the site content; only available for Premium tier and above

So, the correct answer is AppServiceEnvironmentPlatformLogs

upvoted 2 times

✉  **BaoNguyen2411**  8 months, 3 weeks ago

got this question on 29/06/2023

upvoted 3 times

✉  **OPT_001122** 1 year, 4 months ago

Selected Answer: B

You discover that configuration changes were made to some of the web apps.

hence the ans is correct B

upvoted 2 times

✉  **seckinfo** 1 year, 5 months ago

Ref: <https://learn.microsoft.com/en-us/azure/app-service/troubleshoot-diagnostic-logs#supported-log-types>

upvoted 3 times

✉  **qwerty112233** 1 year, 7 months ago

Correct

upvoted 2 times

Question #29

Topic 5

You develop and deploy an Azure App Service web app to a production environment. You enable the Always On setting and the Application Insights site extensions. You deploy a code update and receive multiple failed requests and exceptions in the web app. You need to validate the performance and failure counts of the web app in near real time. Which Application Insights tool should you use?

- A. Profiler
- B. Smart Detection
- C. Live Metrics Stream
- D. Application Map
- E. Snapshot Debugger

Correct Answer: C

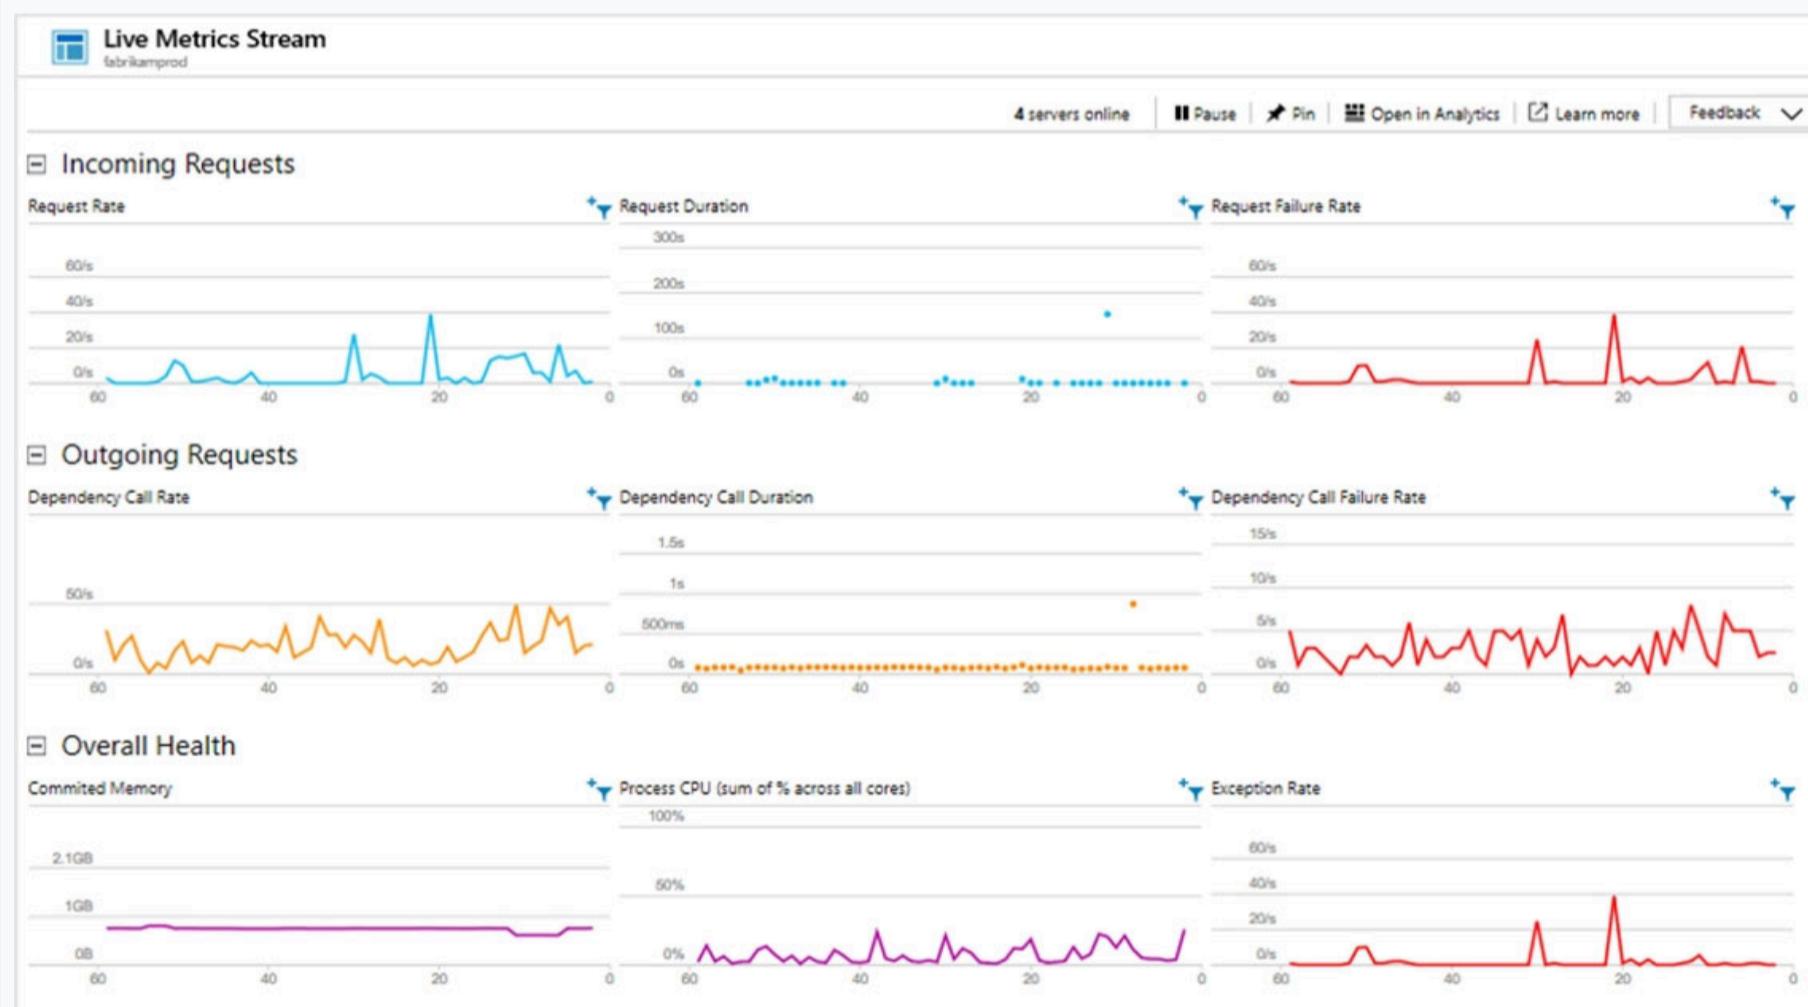
Live Metrics Stream -

Deploying the latest build can be an anxious experience. If there are any problems, you want to know about them right away, so that you can back out if necessary. Live Metrics Stream gives you key metrics with a latency of about one second.

With Live Metrics Stream, you can:

* Validate a fix while it's released, by watching performance and failure counts.

* Etc.



Incorrect:

* Profiler

Azure Application Insights Profiler provides performance traces for applications running in production in Azure. Profiler:

Captures the data automatically at scale without negatively affecting your users.

Helps you identify the `hot` code path spending the most time handling a particular web request.

* Snapshot debugger

When an exception occurs, you can automatically collect a debug snapshot from your live web application. The snapshot shows the state of source code and variables at the moment the exception was thrown. The Snapshot Debugger in Azure Application Insights monitors exception telemetry from your web app. It collects snapshots on your top-throwing exceptions so that you have the information you need to diagnose issues in production.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/live-stream>

Community vote distribution

C (75%)

B (25%)

✉  **adilkhan**  1 year, 2 months ago

C is correct

upvoted 6 times

✉  **g2000**  1 year, 3 months ago

the key is near real time.

<https://learn.microsoft.com/en-us/azure/azure-monitor/app/live-stream>

upvoted 2 times

✉  **Tornik** 1 year, 3 months ago

Selected Answer: B

Isnt it B ?

"Smart detection automatically warns you of potential performance problems and failure anomalies in your web application. It performs proactive analysis of the telemetry that your app sends to Application Insights. "

<https://learn.microsoft.com/en-us/azure/azure-monitor/alerts/proactive-diagnostics>

upvoted 2 times

✉  **Tornik** 1 year, 3 months ago

This is the same question just Smart Detection as answer, so which is it ? I think it might be Live Metrics Stream after all as it says "Near Realtime" while Smart Detection send out notification which is not realtime ?

<https://www.pass4success.com/microsoft/discussions/exam-az-204-topic-11-question-59-discussion>

upvoted 1 times

✉  **OPT_001122** 1 year, 4 months ago

Selected Answer: C

C. Live Metrics Stream

upvoted 3 times

✉  **Perplex** 1 year, 5 months ago

Selected Answer: C

correct

upvoted 3 times

✉  **gmishra88** 1 year, 6 months ago

Another beautiful question with a red-herring: "always on". People who know that is needed for setting Insights profiler will want to answer that. A common thing in many exams, but Microsoft is taking this personally and checking how to trick

upvoted 1 times

✉  **Mousavi** 1 year, 6 months ago

why not A ?

upvoted 1 times

✉  **Ciupaz** 5 months ago

Profiler is not suitable for monitoring overall performance or identifying failures.

upvoted 1 times

✉  **kampatra** 1 year, 7 months ago

Ans Correct

upvoted 1 times

Question #30

HOTSPOT -

You deploy an ASP.NET web app to Azure App Service.

You must monitor the web app by using Application Insights.

You need to configure Application Insights to meet the requirements.

Which feature should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area**Requirement**

Automatically warn you of potential performance problems and failure anomalies in the web app.

Feature

- Smart Detection
- Snapshot Debugger
- Profiler
- Multi-step test

Automatically collect the state of the source code and variables when an exception is thrown in the web app.

- Smart Detection
- Snapshot Debugger
- Profiler
- Multi-step test

Capture performance traces of the web app without negatively affecting users of the web app.

- Smart Detection
- Snapshot Debugger
- Profiler
- Multi-step test

Correct Answer:

Answer Area**Requirement**

Automatically warn you of potential performance problems and failure anomalies in the web app.

Feature

- Smart Detection
- Snapshot Debugger
- Profiler
- Multi-step test

Automatically collect the state of the source code and variables when an exception is thrown in the web app.

- Smart Detection
- Snapshot Debugger
- Profiler
- Multi-step test

Capture performance traces of the web app without negatively affecting users of the web app.

- Smart Detection
- Snapshot Debugger
- Profiler
- Multi-step test

Box 1: Smart Detection -

Smart detection automatically warns you of potential performance problems and failure anomalies in your web application. It performs proactive analysis of the telemetry that your app sends to Application Insights. If there is a sudden rise in failure rates, or abnormal patterns in client or server performance, you get an alert. This feature needs no configuration. It operates if your application sends enough telemetry.

Box 2: Snapshot Debugger -

When an exception occurs, you can automatically collect a debug snapshot from your live web application. The snapshot shows the state of source code and variables at the moment the exception was thrown. The Snapshot Debugger in Azure Application Insights monitors exception telemetry from your web app. It collects snapshots on your top-throwing exceptions so that you have the information you need to diagnose issues in production.

Box 3: Profiler -

Azure Application Insights Profiler provides performance traces for applications running in production in Azure. Profiler:

Captures the data automatically at scale without negatively affecting your users.

Helps you identify the `hot` code path spending the most time handling a particular web request.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/proactive-diagnostics> <https://docs.microsoft.com/en-us/azure/azure-monitor/snapshot-debugger/snapshot-debugger> <https://docs.microsoft.com/en-us/azure/azure-monitor/profiler/profiler-overview>

✉  **OPT_001122**  1 year, 4 months ago

Answer is correct

upvoted 10 times

✉  **KingChuang** 1 year, 3 months ago

On my exam 2022-12-26.

Chose: Given Answer (1, 2, 3)

upvoted 5 times

✉  **ExamNow**  1 year, 7 months ago

1. Smart detection
2. Snapshot
3. user profile

upvoted 7 times

✉  **TheExamMaster2020**  1 year, 5 months ago

Did my exam on 15th November 2022. This question was on it.

upvoted 4 times

✉  **devsb** 1 year, 7 months ago

Answer is correct

upvoted 4 times

Question #31

Topic 5

You are building a web application that performs image analysis on user photos and returns metadata containing objects identified. The image analysis is very costly in terms of time and compute resources. You are planning to use Azure Redis Cache so duplicate uploads do not need to be reprocessed.

In case of an Azure data center outage, metadata loss must be kept to a minimum.

You need to configure the Azure Redis cache instance.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Configure Azure Redis with AOF persistence.
- B. Configure Azure Redis with RDB persistence.
- C. Configure second storage account for persistence.
- D. Set backup frequency to the minimum value.

Correct Answer: BD

RDB persistence - When you use RDB persistence, Azure Cache for Redis persists a snapshot of your cache in a binary format. The snapshot is saved in an Azure Storage account. The configurable backup frequency determines how often to persist the snapshot. If a catastrophic event occurs that disables both the primary and replica cache, the cache is reconstructed using the most recent snapshot.

Note: Azure Cache for Redis supports zone redundant configurations in the Premium and Enterprise tiers. A zone redundant cache can place its nodes across different Azure Availability Zones in the same region. It eliminates data center or AZ outage as a single point of failure and increases the overall availability of your cache.

Incorrect:

Not A: Zone redundancy doesn't support AOF persistence or work with geo-replication currently.

Not C: No need for a second storage account.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-how-to-premium-persistence>

Community vote distribution

AC (69%)

BD (31%)

✉  **finnishr**  1 year, 7 months ago

The key here is "In case of an Azure data center outage, metadata loss must be kept to a minimum."

So the correct answer is AC.

upvoted 20 times

✉  **ShoaibAnwar** 1 month, 2 weeks ago

Azure Redis Cache does not allow for direct configuration to a second storage account for persistence. Persistence and backup configurations are managed within the service's settings, particularly in the premium tier where you can configure data persistence and backup frequency.

upvoted 1 times

✉  **Bear_Polar** 1 year, 4 months ago

I think this is correct.

"RDB is NOT good if you need to minimize the chance of data loss in case Redis stops working". source:
<https://redis.io/docs/management/persistence/>

since the question doesn't mention ZRS for Redis cache, we can use AOF with a 2nd storage to store replica cache data. check the config with AOF: <https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-how-to-premium-persistence#set-up-data-persistence>

upvoted 5 times

✉  **serpevi**  1 year, 7 months ago

got it 10/09/2022. went with AC. Score 927

upvoted 14 times

✉  **dzits** 1 year, 5 months ago

Very confused why it's AC if it states on MS website that Zone redundancy doesn't support AOF persistence or work with geo-replication currently. <https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-how-to-zone-redundancy> (as mentioned by coffecold). I would go with provided answer B,D unless someone can prove why this wouldn't be the case otherwise?

upvoted 2 times

OrfeasTsk **Most Recent** 1 month, 4 weeks ago

Selected Answer: AC

In my opinion this question refers to the durability of the data and not to the availability. On the one hand AOF persistence does not support ZRS (availability) but on the other hand AOF persistence ensures the minimum data loss (durability)

upvoted 1 times

AndySmith 5 months, 3 weeks ago

Selected Answer: AC

Technically here could be just 2 valid answer:

AC and BD.

But, we need to select the option mostly matching the requirements.

Since, they ask to minimize data loss (rather than speed of recovery) then AC is the right answer, because:

- 1) With AOF persistence every write transaction is logged. So, in theory we can recover everything up to the moment of failure.
- 2) Additionally, we'll enable Secondary Storage for AOF, which will help to restore Redis image if primary zone was fully destroyed.

Why RDB is worse? Even if we minimize the period of snapshotting to minimum (15 mins), then we can lose up to 15 mins of data in case of data center outage. Though advantage is that we can restore the snapshot much faster than from the AOF logs.

You can find full info about AOF and RDB settings in Azure Redis Cache on this link:

<https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-how-to-premium-persistence?tabs=premium>

upvoted 5 times

Dianahu 8 months, 3 weeks ago

<https://redis.io/docs/management/persistence/> Ok, so what should I use? If you care a lot about your data, but still can live with a few minutes of data loss in case of disasters, you can simply use RDB alone.

upvoted 1 times

nekkilodeon 12 months ago

Selected Answer: BD

B,D is correct

Zone redundancy doesn't support AOF persistence or work with geo-replication currently", meaning it shouldn't be A because "In case of an Azure data center outage, metadata loss must be kept to a minimum.

<https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-how-to-zone-redundancy>

upvoted 2 times

lh2022 1 year ago

<https://learn.microsoft.com/en-us/training/modules/develop-for-azure-cache-for-redis/3-configure-azure-cache-redis>

The Premium tier allows you to persist data in two ways to provide disaster recovery:

RDB persistence takes a periodic snapshot and can rebuild the cache using the snapshot in case of failure.

AOF persistence saves every write operation to a log that is saved at least once per second. This creates bigger files than RDB but has less data loss.

So AOF has less data loss which satisfies the requirement.

So AOF with data written to a second storage account with ZRS should do the job.

I would also go with AC

upvoted 1 times

damianadalid 11 months ago

Can we just use that second Storage account based in a different region WITHOUT ZRS? That is the main breaking point about not choosing AOF.

upvoted 1 times

damianadalid 11 months ago

Even more, the restriction is: "Currently, you can't use Append-only File (AOF) persistence or geo-replication with multiple replicas (more than one replica)". So, as long as we have just one replica set, can we use AOF without issues?

<https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-how-to-multi-replicas>

upvoted 1 times

adilkhan 1 year, 1 month ago

A,C chatGPT!

upvoted 1 times

wrak 1 year, 1 month ago

AOF leverage on the second storage account: With AOF persistence enabled, write operations to the cache are saved to the named storage account (or accounts if you've configured a second storage account). If there's a catastrophic failure that takes down both the primary and replica cache, the stored AOF log is used to rebuild the cache.

upvoted 1 times

NombreFalso 1 year, 1 month ago

Goodluck!

upvoted 1 times

Sriniv 1 year, 2 months ago

Options C and D are not relevant for this scenario. Configuring a second storage account is not necessary since Redis cache already has persistence options available. Setting backup frequency to minimum value is also not necessary since AOF and RDB persistence options already ensure that data is regularly saved to disk.

upvoted 1 times

 **alexein74** 1 year, 2 months ago

Selected Answer: AC

finnishr is right

upvoted 1 times

 **toysky731** 1 year, 3 months ago

Selected Answer: AC

finnishr is correct

upvoted 1 times

 **edmond** 1 year, 4 months ago

Selected Answer: BD

So many misleading guys, don't answer if you are not sure! Why do you think it's AC?

"Zone redundancy doesn't support AOF persistence or work with geo-replication currently", meaning it shouldn't be A because "In case of an Azure data center outage, metadata loss must be kept to a minimum".

And it's meaningless to use another storage account so it should not be C.

upvoted 3 times

 **aruni_mishra** 1 year, 4 months ago

B & D:

"In case of an Azure data center outage, metadata loss must be kept to a minimum": this can only be achieved by "RDB persistence". period. <https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-how-to-zone-redundancy>.

Backup Frequency: This interval starts counting down after the previous backup operation successfully completes. When it elapses, a new backup starts.

<https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-how-to-premium-persistence#which-persistence-model-should-i-choose>

upvoted 3 times

 **OPT_001122** 1 year, 4 months ago

Selected Answer: AC

A. Configure Azure Redis with AOF persistence.

C. Configure second storage account for persistence.

upvoted 2 times

 **OPT_001122** 1 year, 5 months ago

the admin must correct the answers based on the highly voted discussion

upvoted 3 times

 **Ciupaz** 10 months, 1 week ago

I agree.

upvoted 1 times

Question #32

You are developing an Azure-based web application. The application goes offline periodically to perform offline data processing. While the application is offline, numerous Azure Monitor alerts fire which result in the on-call developer being paged. The application must always log when the application is offline for any reason. You need to ensure that the on-call developer is not paged during offline processing. What should you do?

- A. Add Azure Monitor alert processing rules to suppress notifications.
- B. Disable Azure Monitor Service Health Alerts during offline processing.
- C. Create an Azure Monitor Metric Alert.
- D. Build an Azure Monitor action group that suppresses the alerts.

Correct Answer: D

You can use alert processing rules to add action groups or remove (suppress) action groups from your fired alerts.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/alerts/alerts-action-rules>

Community vote distribution

 A (89%) 11%

✉  **finnishr**  1 year, 7 months ago

Selected Answer: A

I'd choose A.

upvoted 19 times

✉  **scrafy**  1 year, 7 months ago

It's A 100% sure

<https://docs.microsoft.com/en-us/azure/azure-monitor/alerts/alerts-processing-rules?tabs=portal>

upvoted 9 times

✉  **130nk3r5**  4 months ago

A and D are technically both correct....

upvoted 1 times

✉  **dddddd111** 6 months ago

Selected Answer: D

D. as of 18Oct23

upvoted 2 times

✉  **katrang** 7 months, 1 week ago

Selected Answer: A

In all these cases, an alert processing rule provides an easy way to suppress notifications.

upvoted 1 times

✉  **warchoon** 1 year, 1 month ago

Selected Answer: A

The explanation points to A

upvoted 1 times

✉  **adilkhan** 1 year, 1 month ago

D. Build an Azure Monitor action group that suppresses the alerts. (Chat GPT)

To ensure that the on-call developer is not paged during offline processing while still logging when the application is offline, you can create an Azure Monitor action group that suppresses the alerts during that time. This action group can be configured to temporarily suspend alerts or send them to a different channel, such as an email or a different paging system.

upvoted 1 times

✉  **warchoon** 1 year, 1 month ago

ChatGPT is incorrect as usual

upvoted 2 times

✉  **NombreFalso** 1 year, 1 month ago

The WebApp is based

upvoted 1 times

 **SATHTECH** 1 year, 4 months ago

Selected Answer: D

ANSWER D

Suppression: This action removes all the action groups from the affected fired alerts. So, the fired alerts won't invoke any of their action groups, not even at the end of the maintenance window. Those fired alerts will still be visible when you list your alerts in the portal, Azure Resource Graph, API, or PowerShell. The suppression action has a higher priority over the Apply action groups action. If a single fired alert is affected by different alert processing rules of both types, the action groups of that alert will be suppressed.

URL : <https://learn.microsoft.com/en-us/azure/azure-monitor/alerts/alerts-processing-rules?tabs=portal#suppress-notifications-during-planned-maintenance>

upvoted 1 times

 **Bear_Polar** 1 year, 4 months ago

does your source refer to ans A?

upvoted 2 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: A

It's A

upvoted 3 times

 **gmishra88** 1 year, 6 months ago

You can use alert processing rules to add action groups or remove (suppress) action groups from your fired alerts.

upvoted 2 times

Question #33

Topic 5

You are developing an online game that includes a feature that allows players to interact with other players on the same team within a certain distance. The calculation to determine the players in range occurs when players move and are cached in an Azure Cache for Redis instance.

The system should prioritize players based on how recently they have moved and should not prioritize players who have logged out of the game.

You need to select an eviction policy.

Which eviction policy should you use?

- A. allkeys-lru
- B. volatile-lru
- C. allkeys-lfu
- D. volatile-ttl

Correct Answer: A

Community vote distribution



✉ **baroo1** Highly Voted 1 year, 1 month ago

Selected Answer: B

Changed my mind;

There must be a way to tell our redis that logged off users must not be prioritized.

Sample: User A moves and then automatically logs-off. With allkeys-lru we can't distinguish this particularity. With volatile-lru we can tell our redis what are good candidates to be removed using different TTL values.

upvoted 12 times

✉ **exeem** 3 months ago

A lot of assumptions here but I think it make sense

upvoted 1 times

✉ **TonyMel** 1 year ago

correct, in 2023Mar24, score: 904/1000.

upvoted 2 times

✉ **surprise0011** 1 year ago

totally agree with you my friend

upvoted 2 times

✉ **raymond_abcd** Most Recent 2 months, 2 weeks ago

Selected Answer: A

The default TTL for Redis is infinite, so for volatile-lru to be the correct answer, the config must have the TTL set to some expiration value. Nothing is mentioned in the question about TTL. So it doesn't seem likely that this is the correct answer.

upvoted 1 times

✉ **FeriAZ** 2 months, 3 weeks ago

Selected Answer: B

volatile-lru:

This policy also evicts the least recently used keys but only among those with a set TTL(time to live). Keys without a TTL are not considered for eviction.

This policy could be appropriate if you set a TTL for players' data and reset or update this TTL when they move. Players who log out can have their TTL removed or expired, making them ineligible for eviction under this policy.

upvoted 2 times

✉ **alirasouli** 3 months ago

Selected Answer: A

I vote for allkeys-lru. volatile-lru is helpful to deal with or correlate a concept with expiration time. We have the cache of recent players, and what they do then (for example, logging off) does not affect the eviction policy.

upvoted 1 times

✉ **alirasouli** 3 months ago

The question is confusing, and I carefully re-read it again. If the question is about having a way to de-prioritize the players who are logged off, then a second parameter, like expiration time, is helpful, and volatile-lru would be the right choice.

upvoted 1 times

 **AlexChacon99** 6 months ago

Hi! I just passed my AZ 204 exam, I have a question, I don't want my employer to know my score, only that I passed, is there any way they can figure out my score if I only show them the certificate PDF ? Thank you!

upvoted 1 times

 **Ciupaz** 5 months ago

No, they can't.

upvoted 1 times

 **dom271219** 6 months ago

Selected Answer: A

There's no link with ttl and players who have logged out of the game.

In Redis, you can set the TTL (Time To Live) for keys using commands such as EXPIRE, PEXPIRE, or similar commands that allow you to specify a lifespan for a key. The TTL specifies the period for which the key will be retained in the Redis database before being automatically removed.

upvoted 2 times

 **mic_** 6 months, 3 weeks ago

I think the answer is correct. We can remove the related entry from the cache after logoff or simply grab the next entry from the cache if the user is inactive. Volatile policies could even prevent caching, see

<https://redis.io/docs/reference eviction/>

The policies volatile-lru, volatile-lfu, volatile-random, and volatile-ttl behave like noeviction if there are no keys to evict matching the prerequisites.

upvoted 1 times

 **unraval** 7 months ago

Selected Answer: B

Think it should be B : volatile-lru

volatile-lru : Removes least recently used keys with the expire field set to true.

The expire field for the logged out players can be marked true indicating Redis to evict those entries

upvoted 3 times

 **sundeeprch1** 9 months, 3 weeks ago

ChatGPT answer option is volatile-ttl

upvoted 1 times

 **aragones** 11 months, 1 week ago

Got this 2023-05-12.

Had cases:

case: You need to configure authorization.

case: You need to ensure the app does not time out and processes the blob data

case study: VanArsdel Inc Canada

upvoted 2 times

 **adilkhan** 1 year, 1 month ago

The "volatile-ttl" eviction policy evicts keys based on their expiration time, which is set when the key is created or updated. This policy is the most suitable for this scenario since it allows the system to prioritize recently moved players by setting a shorter expiration time for their cached location data. It also ensures that players who have logged out of the game are automatically removed from the cache when their data expires.

Chat GPT

upvoted 3 times

 **baroo1** 1 year, 1 month ago

Selected Answer: A

allkeys-lru: Keeps most recently used keys; removes least recently used (LRU) keys

volatile-lru: Removes least recently used keys with the expire field set to true.

Source: <https://redis.io/docs/reference eviction/>

Not prioritizing logged off players does not mean to remove them! volatile-lru will not be the best solution at all.

upvoted 2 times

 **NombreFalso** 1 year, 2 months ago

Selected Answer: A

I think it is A because I don't understand the need to be restricted to only entries with TTL.

upvoted 2 times

 **adilkhan** 1 year, 2 months ago

B looks correct b/c of log out condition.

upvoted 1 times

 **alexein74** 1 year, 2 months ago

Selected Answer: B

B. volatile-lru

The eviction policy "volatile-lru" evicts the least recently used (LRU) key among keys with an expiring time-to-live (TTL). This policy is suitable for

your scenario because it prioritizes players based on how recently they have moved and does not prioritize players who have logged out of the game.

upvoted 1 times

 **Osip21** 1 year, 2 months ago

Selected Answer: D

"volatile-ttl" is a type of eviction policy that uses a time-to-live (TTL) value for each cache entry. This policy automatically evicts entries from the cache after their TTL expires. This can be useful for scenarios where you want to ensure that stale or outdated data is not kept in the cache for too long. In your case, you can set a shorter TTL for players who have moved recently and a longer TTL for players who have not moved recently or logged out of the game, which allow the system to prioritize players based on their activity.

upvoted 2 times

 **JustHereToLearn** 1 year, 2 months ago

Selected Answer: B

I think the answer should be B based on the facts. "The system should prioritize players based on how recently they have moved" this means we will use LRU but not we have to decide between allkeys-lru or valatile-lru.

The statement "should not prioritize the players who logged out of game" provides a hint that there needs to be some to evict logged out users like using the expire to true.

Reference: <https://redis.io/docs/reference/eviction/>

upvoted 4 times

Question #34

Topic 5

You develop an Azure App Service web app and deploy to a production environment. You enable Application Insights for the web app.

The web app is throwing multiple exceptions in the environment.

You need to examine the state of the source code and variables when the exceptions are thrown.

Which Application Insights feature should you configure?

- A. Smart detection
- B. Profiler
- C. Snapshot Debugger
- D. Standard test

Correct Answer: C

Community vote distribution

C (100%)

✉  **alexein74**  1 year, 2 months ago

Selected Answer: C

C. Snapshot Debugger

The Snapshot Debugger is an Application Insights feature that allows you to examine the state of the source code and variables when an exception is thrown. It captures a snapshot of the call stack, local variables, and object state at the time of the exception. With this feature, you can easily identify the cause of the exception and fix the issue.

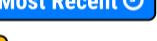
upvoted 7 times

✉  **Jhilphis**  1 year, 3 months ago

Selected Answer: C

Correct answer

upvoted 5 times

✉  **james2033**  1 month ago

Selected Answer: C

'The Snapshot Debugger in Application Insights: ...' at <https://learn.microsoft.com/en-us/azure/azure-monitor/snapshot-debugger/snapshot-debugger>

upvoted 1 times

✉  **Ciupaz** 6 months, 3 weeks ago

The Snapshot Debugger in Azure is a debugging tool and feature that helps developers diagnose and troubleshoot issues in their applications running on Azure App Service without impacting the application's availability or performance. It allows you to capture snapshots of the application's state at specific points in time when an issue or exception occurs, allowing you to inspect the state of variables and the call stack to identify the root cause of problems.

upvoted 1 times

Question #35

DRAG DROP

You develop and deploy a Java application to Azure. The application has been instrumented by using the Application Insights SDK.

The telemetry data must be enriched and processed before it is sent to the Application Insights service.

You need to modify the telemetry data.

Which Application Insights SDK features should you use? To answer, drag the appropriate features to the correct requirements. Each feature may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Features	Answer Area	Requirement	Feature
Sampling		Reduce the volume of telemetry without affecting statistics.	
Telemetry initializer		Enrich telemetry with additional properties or override an existing one.	
Telemetry processor		Completely replace or discard a telemetry item.	
Telemetry channel			

Correct Answer:	Answer Area	Requirement	Feature
		Reduce the volume of telemetry without affecting statistics.	Sampling
		Enrich telemetry with additional properties or override an existing one.	Telemetry initializer
		Completely replace or discard a telemetry item.	Telemetry processor

✉️  **whiteblack**  1 year, 3 months ago

Correct answers

Link: <https://learn.microsoft.com/en-us/azure/azure-monitor/app/api-filtering-sampling>

upvoted 22 times

✉️  **Rodikito**  8 months ago

what is 'Case study: city and Lights'??

upvoted 1 times

✉️  **Rodikito** 7 months, 3 weeks ago

sorry, I don't read all exam topics (I didn't have contributor yet). Case studies are in the last pages of this dump

upvoted 1 times

✉️  **kvtborad** 8 months, 2 weeks ago

I got this question on 6th August 2023. chose given answers. passed with 904. I got Case study: city and Lights. All questions are from ExamTopics.

upvoted 4 times

✉️  **WH16** 8 months, 2 weeks ago

Hi bro, how many questions in case study?

upvoted 2 times

✉️  **Rodikito** 7 months, 3 weeks ago

Normally, 1 case study per exam, with 5 questions

upvoted 3 times

✉️  **macobuzi** 7 months, 2 weeks ago

5 - 12 question

upvoted 2 times

Question #36

HOTSPOT

You develop new functionality in a web application for a company that provides access to seismic data from around the world. The seismic data is stored in Redis Streams within an Azure Cache for Redis instance.

The new functionality includes a real-time display of seismic events as they occur.

You need to implement the Azure Cache for Redis command to receive seismic data.

How should you complete the command? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

XLEN
XREAD
XRANGE

BLOCK 0
COUNT 0
BLOCK -1
COUNT -1

STREAMS seismicData

\$
(&
(0-0 +
(0-0 -

Answer Area

Correct Answer:

XLEN
XREAD
XRANGE

BLOCK 0
COUNT 0
BLOCK -1
COUNT -1

STREAMS seismicData

\$
(&
(0-0 +
(0-0 -

✉  **Michael2023**  1 year, 2 months ago

puuh... knowing all this little things/ remembering all this little things for the exam sounds impossible :(

upvoted 55 times

✉  **macobuzi** 7 months, 3 weeks ago

Exactly! That's why we study the exam in question bank :)

upvoted 11 times

✉  **Ciupaz** 6 months, 3 weeks ago

Agree!

upvoted 3 times

✉  **2fa**  1 year, 3 months ago

Correct.

BLOCK 0 means that we will wait infinitely (timeout 0 means it will never expire) for new items in a stream.

\$ ID means that we will receive only new messages, starting from the time we started listening

<https://redis.io/docs/data-types/streams-tutorial/#listening-for-new-items-with-xread>

upvoted 18 times

✉  **piotrbugla**  1 month, 3 weeks ago

This

<https://learn.microsoft.com/en-us/training/modules/azure-redis-publish-subscribe-streams/>

Should be part of AZ-204 learning path

upvoted 2 times

✉ **130nk3r5** 4 months ago

Answer is correct!

To receive seismic data in real-time from Redis Streams, you should use the XREAD command.

The XREAD command is a blocking command that allows a client to wait for new data to be added to a stream. This is perfect for real-time applications as it allows the client to receive new data as soon as it is available.

The command should be completed as follows:

```
XREAD BLOCK 0 STREAMS seismic_data $
```

Here, "BLOCK 0" means the command will block indefinitely until new data is available. "STREAMS" is followed by the name of the stream(s) to read from, in this case "seismic_data". The "\$" symbol is a special ID which represents the last message in the stream, so the command will only return new messages that arrive after the command is run.

upvoted 3 times

✉ **subarna530** 4 months ago

XREAD, BLOCK 0, and \$ is the answer that I memorized! I have no clue why!

upvoted 2 times

✉ **dddddd111** 6 months ago

I failed on my first attempt. I only got 646/1000. I chose C#. MS has recently updated the AZ-204 (Oct 2023), so please be guided that some highly voted answers here might not be correct.

upvoted 2 times

✉ **AndySmith** 5 months, 3 weeks ago

Did they add any new Qs outside of this "tutorial"?

Or they just updated existing Qs - e.g. answer options, etc.

upvoted 1 times

✉ **BaoNguyen2411** 8 months, 3 weeks ago

got this question on 29/06/2023

upvoted 6 times

✉ **brwilli1** 9 months ago

This was on the test 07/14/2023

upvoted 6 times

✉ **KarthikCS** 12 months ago

These AZ exam questions are some of the most ridiculous ones ever. Looks like MS wants you to fail unless you senselessly memorize stupid things. I wonder how dumb the team that sets these questions must be.

upvoted 12 times

✉ **surprise0011** 1 year ago

wtf is that haha

upvoted 9 times

✉ **mmdex** 1 year, 1 month ago

This question is so ridiculous that I actually remember its answer better than any other.

upvoted 14 times

✉ **uffuchsi** 1 year, 1 month ago

Agree with the others here - seems quite silly to be expected to remember detail like this.

Would be interested to know if anyone received this question in their exam

upvoted 4 times

✉ **NombreFalso** 1 year, 1 month ago

Feels like anything goes and they can ask you about anything that's ever existed

upvoted 10 times

Question #37

You develop an ASP.NET Core app that uses Azure App Configuration. You also create an App Configuration containing 100 settings.

The app must meet the following requirements:

- Ensure the consistency of all configuration data when changes to individual settings occur.
- Handle configuration data changes dynamically without causing the application to restart.
- Reduce the overall number of requests made to App Configuration APIs.

You must implement dynamic configuration updates in the app.

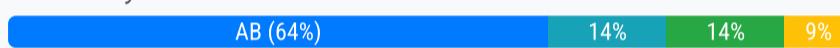
What are two ways to achieve this goal? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- Create and register a sentinel key in the App Configuration store. Set the refreshAll parameter of the Register method to true.
- Increase the App Configuration cache expiration from the default value.
- Decrease the App Configuration cache expiration from the default value.
- Create and configure Azure Key Vault. Implement the Azure Key Vault configuration provider.
- Register all keys in the App Configuration store. Set the refreshAll parameter of the Register method to false.
- Create and implement environment variables for each App Configuration store setting.

Correct Answer: AB

Community vote distribution



✉ **whiteblack** Highly Voted 1 year, 3 months ago

Correct Answers

Link: <https://learn.microsoft.com/en-us/azure/azure-app-configuration/enable-dynamic-configuration-aspnet-netfx>

upvoted 9 times

✉ **halfway** Highly Voted 1 year ago

Selected Answer: AB

Quotes from the reference below:

"A sentinel key is a key that you update after you complete the change of all other keys. Your app monitors the sentinel key. When a change is detected, your app refreshes all configuration values. This approach helps to ensure the consistency of configuration in your app and reduces the overall number of requests made to your App Configuration store, compared to monitoring all keys for changes."

"You can add a call to the refreshOptions.SetCacheExpiration method to specify the minimum time between configuration refreshes. In this example, you use the default value of 30 seconds. Adjust to a higher value if you need to reduce the number of requests made to your App Configuration store."

Reference: <https://learn.microsoft.com/en-us/azure/azure-app-configuration/enable-dynamic-configuration-aspnet-core?tabs=core6x>
upvoted 8 times

✉ **130nk3r5** Most Recent 4 months ago

Selected Answer: AB

A / B

Creating and registering a sentinel key in the App Configuration store with the refreshAll parameter set to true will ensure the consistency of all configuration data when changes to individual settings occur. This is because when the sentinel key changes, all configuration data is refreshed.

Increasing the App Configuration cache expiration from the default value will reduce the overall number of requests made to App Configuration APIs. This is because the application will cache the configuration data for a longer period of time before it needs to request updated data from the App Configuration APIs.

upvoted 1 times

✉ **athli** 5 months, 1 week ago

Selected Answer: AB

AB

- Increase the refresh timeout, especially if your configuration values do not change frequently.

- Watch a single sentinel key, rather than watching individual keys. Refresh all configuration only if the sentinel key changes
 Ref: <https://learn.microsoft.com/en-us/azure/azure-app-configuration/howto-best-practices#reduce-requests-made-to-app-configuration>
 upvoted 1 times

✉ **AndySmith** 5 months, 3 weeks ago

It's AB for sure, since it fully follows the best practices of AzureAppConfig service:
<https://learn.microsoft.com/en-us/azure/azure-app-configuration/howto-best-practices#reduce-requests-made-to-app-configuration>

It can't be D since there is no any word about a secret in the Q.

upvoted 1 times

✉ **dddddd111** 6 months ago

Selected Answer: AD

Option B. Increase the App Configuration cache expiration from the default value is a valid consideration for reducing the number of requests made to App Configuration APIs, but it does not directly address the requirement of dynamic configuration updates without causing the application to restart.

In fact, the default behavior of Azure App Configuration already includes caching to minimize the number of requests to the configuration store. Adjusting the cache expiration might help reduce the frequency of requests, but it won't handle dynamic updates without application restarts.

For achieving dynamic configuration updates without restarting the application, options A (sentinel key) and D (Azure Key Vault configuration provider) are more relevant. These options allow you to implement dynamic configuration changes and ensure consistency without requiring application restarts.

upvoted 1 times

✉ **drkpnsn** 7 months, 3 weeks ago

Selected Answer: AD

I had to go for A & D.

For B "You can add a call to the refreshOptions.SetCacheExpiration method to specify the minimum time between configuration refreshes. In this example, you use the default value of 30 seconds. Adjust to a higher value if you need to reduce the number of requests made to your App Configuration store." explanation makes no sense for me.

upvoted 1 times

✉ **Arr14487** 8 months, 2 weeks ago

Selected Answer: AF

Option B, Increase the App Configuration cache expiration from the default value, is not a good solution for dynamic configuration updates.

upvoted 2 times

✉ **adilkhan** 11 months, 4 weeks ago

A,B 100% correct answer!

upvoted 2 times

✉ **adilkhan** 12 months ago

B makes no sense A and D are correct!

upvoted 1 times

✉ **Dats1987** 1 year ago

Selected Answer: AB

AB , looks correct.

upvoted 2 times

✉ **adilkhan** 1 year, 1 month ago

A. Create and register a sentinel key in the App Configuration store. Set the refreshAll parameter of the Register method to true.

D. Create and configure Azure Key Vault. Implement the Azure Key Vault configuration provider. Chat GPT

B n C are related to Cache and has nothing to do with the question. Chat GPT

upvoted 1 times

✉ **CODE_STS** 1 year, 1 month ago

Selected Answer: AD

A. Create and register a sentinel key in the App Configuration store. Set the refreshAll parameter of the Register method to true: By creating a sentinel key in the App Configuration store, you can ensure that all configuration data remains consistent when changes to individual settings occur. You can register the sentinel key and set the refreshAll parameter of the Register method to true to refresh all settings in the App Configuration store.

D. Create and configure Azure Key Vault. Implement the Azure Key Vault configuration provider: The Azure Key Vault configuration provider enables you to store configuration data in Azure Key Vault. By using the Azure Key Vault configuration provider, you can implement dynamic configuration updates in your app without causing the app to restart. This approach also reduces the overall number of requests made to the App Configuration APIs.

upvoted 1 times

✉ **Sriniv** 1 year, 2 months ago

The two ways to achieve dynamic configuration updates in an ASP.NET Core app that uses Azure App Configuration are:

A. Create and register a sentinel key in the App Configuration store. Set the refreshAll parameter of the Register method to true.

D. Create and configure Azure Key Vault. Implement the Azure Key Vault configuration provider.

upvoted 1 times

 **SEEKER31** 1 year, 2 months ago

Selected Answer: AB

correct answer

upvoted 1 times

 **OlavG** 1 year, 2 months ago

Selected Answer: AB

Correct

upvoted 1 times

 **adilkhan** 1 year, 2 months ago

AB is correct!

<https://learn.microsoft.com/en-us/azure/azure-app-configuration/enable-dynamic-configuration-aspnet-core?tabs=core6x>

upvoted 3 times

Question #38

HOTSPOT

You develop and deploy an Azure App Service web app that connects to Azure Cache for Redis as a content cache. All resources have been deployed to the East US 2 region.

The security team requires the following audit information from Azure Cache for Redis:

- The number of Redis client connections from an associated IP address.
- Redis operations completed on the content cache.
- The location (region) in which the Azure Cache for Redis instance was accessed.

The audit information must be captured and analyzed by a security team application deployed to the Central US region.

You need to log information on all client connections to the cache.

Which configuration values should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area**Requirement**

Store log information.

Configuration value

- Log Analytics workspace
- Blob Storage account
- Data Lake Storage Gen2 Storage account
- Event hub

Requirement

Enable client connection logging.

- Diagnostic setting
- Managed identity
- App registration
- Environment variable

Answer Area**Requirement**

Store log information.

Configuration value

- Log Analytics workspace
- Blob Storage account
- Data Lake Storage Gen2 Storage account
- Event hub

Correct Answer:**Requirement**

Enable client connection logging.

- Diagnostic setting
- Managed identity
- App registration
- Environment variable

👤 **halfway** Highly Voted 1 year ago

Correct.

Only Log Analytics supports cross-region logging

Reference: <https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-monitor-diagnostic-settings?tabs=basic-standard-premium#enable-connection-logging-using-the-azure-portal>

upvoted 17 times

👤 **kakajos** Highly Voted 6 months, 4 weeks ago

Log Analytics Workspace - used to store log information.

Diagnostic setting - diagnostic logs are streamed to that workspace as soon as new event data is generated

upvoted 6 times

Question #39

Topic 5

You develop and deploy a web app to Azure App Service. The Azure App Service uses a Basic plan in a single region.

Users report that the web app is responding slow. You must capture the complete call stack to help identify performance issues in the code. Call stack data must be correlated across app instances. You must minimize cost and impact to users on the web app.

You need to capture the telemetry.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Restart all apps in the App Service plan.
- B. Enable Application Insights site extensions.
- C. Upgrade the Azure App Service plan to Premium.
- D. Enable Profiler.
- E. Enable the Always On setting for the app service.
- F. Enable Snapshot debugger.
- G. Enable remote debugging.

Correct Answer: DEF

Community vote distribution

BDF (52%)

BDE (48%)

✉  **paunski7**  12 months ago

B E F

B. Enable Application Insights site extensions: Application Insights provides end-to-end insights into the web app's performance and usage. Enabling the Application Insights site extension allows you to collect telemetry data, including call stack traces, to help identify performance issues in the code.

E. Enable the Always On setting for the app service: Enabling the Always On setting ensures that the web app is always running, even when there is no traffic. This reduces the impact of cold starts and improves the web app's response time.

F. Enable Snapshot debugger: Enabling Snapshot debugger allows you to capture the complete call stack and local variables for an exception that occurred in the web app. This helps to identify performance issues and diagnose issues in the code.

upvoted 15 times

✉  **paunski7** 12 months ago

Option A, restarting all apps in the App Service plan, and option G, enabling remote debugging, are not necessary for capturing telemetry and identifying performance issues in the web app.

Option C, upgrading the Azure App Service plan to Premium, is not necessary for capturing telemetry and identifying performance issues in the web app, and it would significantly increase the cost of the App Service plan.

Option D, enabling Profiler, is not necessary as it is an outdated feature and has been replaced by the more powerful and comprehensive Snapshot debugger.

upvoted 5 times

✉  **damianadalid** 11 months ago

Hello, could you please provide the source where you found that Profiler is being replaced by Snapshot debugger?

upvoted 3 times

✉  **JohnStivenson** 6 months, 3 weeks ago

E - Enable the Always On, how this will fulfill "You need to capture the telemetry."?

upvoted 2 times

✉  **JohnStivenson** 6 months, 3 weeks ago

I think, it's BDF:

- Enable Application Insights site extensions
- Enable Profiler
- Enable Snapshot debugger

upvoted 3 times

✉ **JH81** Highly Voted 10 months ago

Selected Answer: BDE

B, D, E

Gotta enable App Insights.

Always On can solve the performance issues.

Profiler can help find performance issues.

You can eliminate Premium upgrade because it says minimize costs, snapshot debugger is for debugging exceptions and has nothing to do with performance issues, restarting app impacts users and it says minimize impact.

upvoted 7 times

✉ **the_odd_one** 5 months, 4 weeks ago

Dosent say anything about needing to solve the performance issue, for me it is B, D, F

upvoted 1 times

✉ **Kotsume** Most Recent 2 months, 1 week ago

Selected Answer: BDE

B: There is no need to discuss.

D: from msLearn,

> With Application Insights Profiler, you can capture, identify, and view performance traces for your application running in Azure, regardless of the scenario. The Profiler trace process occurs automatically, at scale, and doesn't negatively affect your users.

<https://learn.microsoft.com/en-US/azure/azure-monitor/snapshot-debugger/snapshot-debugger>

E: If you want to use Application Insights Profiler, you have to set Always on for App Services.

<https://learn.microsoft.com/en-US/azure/azure-monitor/profiler/profiler#verify-the-always-on-setting-is-enabled>

F: This is not the answer. from mslearn,

> with Snapshot Debugger, you can automatically collect a debug snapshot when an EXCEPTION OCCURS in your live .NET application.

<https://learn.microsoft.com/en-US/azure/azure-monitor/snapshot-debugger/snapshot-debugger>

upvoted 1 times

✉ **FeriAZ** 2 months, 2 weeks ago

Selected Answer: BDF

Enable Application Insights site extensions (B): To collect comprehensive telemetry and performance data.

Enable Profiler (D): To capture detailed performance traces and call stacks for slow requests.

Enable Snapshot debugger (F): To take snapshots of the app in production, capturing the state and call stack at specific points without heavily impacting users.

upvoted 2 times

✉ **yosbelsd** 4 months, 3 weeks ago

Selected Answer: BDE

You need to activate Profiler (D) and to do this, as prerequisites, Application Insights (B) and AlwaysOn (E) must be enabled

upvoted 3 times

✉ **CarlosTheBoldest** 5 months ago

Selected Answer: BDF

The alwayson option does not help us gathering the needed data, so should be BDF

upvoted 3 times

✉ **mathunit1** 5 months, 1 week ago

Oh goody, it's one of those questions where no one can agree on a definite answer. Greaaaat.

upvoted 2 times

✉ **Ciupaz** 6 months, 3 weeks ago

Selected Answer: BDF

For me:

- Enable Application Insights site extensions
- Enable Profiler
- Enable Snapshot debugger

upvoted 2 times

✉ **kakajos** 6 months, 4 weeks ago

BDE is the correct answer. Profiler capture data automatically without negatively affecting users.

upvoted 1 times

✉ **Paolo12** 7 months, 3 weeks ago

Always On is available from Standard tier

upvoted 2 times

✉ **NPE_** 5 months, 2 weeks ago

No it is available from Basic tier

upvoted 1 times

 **EliteAllen** 8 months ago

Selected Answer: BDF

B. Enable Application Insights site extensions.

This will allow you to collect telemetry from your web app, including performance metrics, failed requests, and dependencies.

D. Enable Profiler.

The Application Insights Profiler will help you identify performance issues in your app's code by capturing telemetry about the app's slowest operations.

F. Enable Snapshot debugger.

The Snapshot Debugger takes snapshots of your in-production apps when exceptions occur, allowing you to see exactly what went wrong without affecting traffic of your production application.

upvoted 5 times

 **Dianahu** 8 months, 3 weeks ago

D) Enable profiler

With Application Insights Profiler, you can capture and view performance traces for your application in all these dynamic situations. The process occurs automatically at scale and doesn't negatively affect your users. Profiler captures the following information so that you can easily identify performance issues while your app is running in Azure (<https://learn.microsoft.com/en-us/azure/azure-monitor/profiler/profiler-overview>)

E) enable always on (to enable profiler -> <https://learn.microsoft.com/en-us/azure/azure-monitor/profiler/profiler>)

F) Snapshot debugger - to collect call stacks for your application - <https://learn.microsoft.com/en-us/azure/azure-monitor/snapshot-debugger/media/snapshot-debugger/enablement-ui.png>

upvoted 3 times

 **Dreamer999** 10 months, 1 week ago

I think "Always On" is an option to fix the issue. But the question is about collecting metrics. So I think "Always On" should not be selected. I would choose Profiler, Snapshot debugger, Enable application Insights site extensions

upvoted 2 times

 **damianadalid** 11 months ago

I think it should be B,E,D :

-B and E make the system able to use Application Insights via setting the App Service to always On and enabling Application Insights extensions.

-D: Profiler is the most suitable one regarding finding performance issues instead of Snapshot debugger, which is more related to exceptions and source code state when errors.

upvoted 3 times

 **dy0917** 11 months, 1 week ago

Should be B D E

F. Enable Snapshot debugger for debugging purposes.

D. Profiler for monitoring performance.

upvoted 4 times

 **sachinrikhe** 12 months ago

This is new question. It was not here yesterday

upvoted 2 times

 **adilkhan** 12 months ago

answers look correct!

upvoted 1 times

Question #40

Topic 5

You are building an application to track cell towers that are available to phones in near real time. A phone will send information to the application by using the Azure Web PubSub service. The data will be processed by using an Azure Functions app. Traffic will be transmitted by using a content delivery network (CDN).

The Azure function must be protected against misconfigured or unauthorized invocations.

You need to ensure that the CDN allows for the Azure function protection.

Which HTTP header should be on the allowed list?

- A. Authorization
- B. WebHook-Request-Callback
- C. Resource
- D. WebHook-Request-Origin

Correct Answer: D

Community vote distribution

D (67%) A (33%)

 **halfway**  1 year ago

Selected Answer: D

<https://learn.microsoft.com/en-us/azure/azure-web-pubsub/howto-develop-eventhandler#upstream-and-validation>
upvoted 12 times

 **AndySmith**  5 months, 2 weeks ago

On exam 3-Nov-2023. Went with most-voted answer - 932/1000.
upvoted 2 times

 **AndySmith** 5 months, 3 weeks ago

Most likely, it's A, since they ask about "Unauthorized protection" instead of "Abuse Protection", which are both supported by underlying "CloudEvents HTTP protocols" (web-hooks).
<https://github.com/cloudevents/spec/blob/v1.0/http-webhook.md#4-abuse-protection>

So, if we read attentively about Abuse Protection / webhook-request-origin, we can find this excerpt:

"It is important to understand is that the handshake does not aim to establish an authentication or authorization context. It only serves to protect the sender from being told to a push to a destination that is not expecting the traffic. While this specification mandates use of an authorization model, this mandate is not sufficient to protect any arbitrary website from unwanted traffic if that website doesn't implement access control and therefore ignores the Authorization header."

And this is Authorization header description in this protocol:

<https://github.com/cloudevents/spec/blob/v1.0/http-webhook.md#3-authorization>

upvoted 3 times

 **dddddd111** 6 months ago

Selected Answer: A

The Authorization header (Option A) is widely used for authentication in HTTP requests. It is a standard header for including credentials, tokens, or other authentication information. By allowing the Authorization header on the CDN, you can ensure that only requests with valid authorization tokens or credentials can invoke your Azure function, providing a secure method of protection against unauthorized invocations.

upvoted 3 times

 **Ciupaz** 6 months, 4 weeks ago

Selected Answer: A

Why not A?

To ensure that the Azure Function is protected against misconfigured or unauthorized invocations when using a CDN, you should allow the "Authorization" HTTP header.

upvoted 3 times

 **JH81** 10 months ago

This cannot possibly be a real question the way it is written.

upvoted 3 times

 **adilkhan** 11 months, 4 weeks ago

- B. Enable Application Insights site extensions.
- E. Enable the Always On setting for the app service.
- D. Enable Profiler. chat gpt

upvoted 1 times

 **damianadalid** 11 months ago
wrong questions, this relates to #39
upvoted 3 times

 **adilkhan** 12 months ago
D is correct
upvoted 3 times

Question #41

Topic 5

You are developing an Azure App Service web app.

The web app must securely store session information in Azure Redis Cache.

You need to connect the web app to Azure Redis Cache.

Which three Azure Redis Cache properties should you use? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Access key
- B. SSL port
- C. Subscription name
- D. Location
- E. Host name
- F. Subscription id

Correct Answer: ABE

✉  **AbidooKing**  9 months, 1 week ago

Answer is correct.

<https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-web-app-howto#retrieve-host-name-ports-and-access-keys-from-the-azure-portal>

upvoted 6 times

✉  **anupmaurya**  7 months ago

When connecting a web app to Azure Redis Cache, you need to specify certain properties. Here are three Azure Redis Cache properties you should use:

Hostname: This property represents the fully qualified domain name (FQDN) or the endpoint URL of the Azure Redis Cache instance. It is used to establish a connection between the web app and the cache. The hostname typically follows the format `yourcache.redis.cache.windows.net`.

Port: This property specifies the port number used for communication with the Azure Redis Cache instance. The default port for Redis is 6379, but you can configure it differently if needed.

Access key: This property is the primary or secondary access key for the Azure Redis Cache instance.

upvoted 5 times

✉  **laurorucker**  9 months, 2 weeks ago

When connecting a web app to Azure Redis Cache, you need to specify certain properties. Here are three Azure Redis Cache properties you should use:

Hostname: This property represents the fully qualified domain name (FQDN) or the endpoint URL of the Azure Redis Cache instance. It is used to establish a connection between the web app and the cache. The hostname typically follows the format `yourcache.redis.cache.windows.net`.

Port: This property specifies the port number used for communication with the Azure Redis Cache instance. The default port for Redis is 6379, but you can configure it differently if needed.

Access key: This property is the primary or secondary access key for the Azure Redis Cache instance. It acts as a password for authentication purposes. The access key is used by the web app to authorize access to the cache and securely store session information.

upvoted 2 times

✉  **AbidooKing** 9 months, 1 week ago

At least let us know when you are posting chatGPT-shit here. This time is seems to be correct.

<https://learn.microsoft.com/en-us/azure/azure-cache-for-redis/cache-web-app-howto#retrieve-host-name-ports-and-access-keys-from-the-azure-portal>

upvoted 5 times

Question #42

HOTSPOT

You are developing several microservices to run on Azure Container Apps.

You need to monitor and diagnose the microservices.

Which features should you use? To answer, select the appropriate feature in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Requirement

View console logs from a container in near real-time.

Feature

Log streaming
Container console
Azure Monitor metrics
Azure Monitor Log Analytics

Debug the microservice from inside the container.

Container console
Azure Monitor metrics
Azure Container Registry
Azure Monitor Log Analytics

Answer Area

Requirement

View console logs from a container in near real-time.

Feature

Log streaming
Container console
Azure Monitor metrics
Azure Monitor Log Analytics

Correct Answer:

Debug the microservice from inside the container.

Container console
Azure Monitor metrics
Azure Container Registry
Azure Monitor Log Analytics

 **jaf19f**  8 months, 1 week ago

I got this question (12-Aug-2023) and I chose the given answers - 932 passed
upvoted 10 times

 **AbidooKing**  9 months, 1 week ago

Answer is correct!

<https://learn.microsoft.com/en-us/azure/container-apps/observability>
upvoted 7 times

 **AndySmith** 6 months ago

From this page:
Log Streaming - View streaming system and console logs from a container in near real-time.
Container console - Connect to the Linux console in your containers to debug your application from inside the container.
upvoted 2 times

 **75daltonic** 8 months, 4 weeks ago

Correct.
It is not deployed yet. So:

Development and test

During the development and test phase, real-time access to your containers' application logs and console is critical for debugging issues. Container Apps provides:

Log streaming: View real-time log streams from your containers.

Container console: Access the container console to debug your application.

<https://learn.microsoft.com/en-us/azure/container-apps/observability#development-and-test>

upvoted 3 times

 **Maskoo** Most Recent 2 months, 3 weeks ago

I think it's correct but first box could be container console too

If we connect to the container console we would see console logs as well

upvoted 1 times

 **CarlosTheBoldest** 4 months, 1 week ago

I got this question on my exam, 2023Dec, go with what I remember was the most voted answer. Score 902, most of the questions were here, slightly different on wording because the Azure Ad <-> Entra Id change. Case was City Power & Light. Good luck!

Important tip, you have access to microsoft learn during the exam!

upvoted 2 times

 **Ciupaz** 4 months, 1 week ago

In near real-time -> log streaming

Inside the container -> container console

upvoted 1 times

 **laurorucker** 9 months, 1 week ago

Azure Monitor Logs: Azure Monitor Logs, also known as Azure Log Analytics, allows you to collect, analyze, and view log data from various sources, including containerized applications. By configuring Azure Monitor Logs for your microservices, you can stream and view console logs in near real-time. You can create log queries and use the Live Stream feature to monitor logs as they are generated, enabling you to observe the output from the console logs in real-time.

upvoted 1 times

 **laurorucker** 9 months, 1 week ago

For me both would be azure monitor logs

upvoted 1 times

Question #43

Topic 5

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. When you are ready to answer a question, click the Question button to return to the question.

Background -

VanArsdel, Ltd. is a global office supply company. The company is based in Canada and has retail store locations across the world. The company is developing several cloud-based solutions to support their stores, distributors, suppliers, and delivery services.

Current environment -

Corporate website -

The company provides a public website located at <http://www.vanarsdelltd.com>. The website consists of a React JavaScript user interface, HTML, CSS, image assets, and several APIs hosted in Azure Functions.

Retail Store Locations -

The company supports thousands of store locations globally. Store locations send data every hour to an Azure Blob storage account to support inventory, purchasing and delivery services. Each record includes a location identifier and sales transaction information.

Requirements -

The application components must meet the following requirements:

Corporate website -

- Secure the website by using SSL.
- Minimize costs for data storage and hosting.
- Implement native GitHub workflows for continuous integration and continuous deployment (CI/CD).
- Distribute the website content globally for local use.
- Implement monitoring by using Application Insights and availability web tests including SSL certificate validity and custom header value verification.
- The website must have 99.95 percent uptime.

Retail store locations -

- Azure Functions must process data immediately when data is uploaded to Blob storage. Azure Functions must update Azure Cosmos DB by using native SQL language queries.
- Audit store sale transaction information nightly to validate data, process sales financials, and reconcile inventory.

Delivery services -

- Store service telemetry data in Azure Cosmos DB by using an Azure Function. Data must include an item id, the delivery vehicle license plate, vehicle package capacity, and current vehicle location coordinates.
- Store delivery driver profile information in Azure Active Directory (Azure AD) by using an Azure Function called from the corporate website.

Inventory services -

The company has contracted a third-party to develop an API for inventory processing that requires access to a specific blob within the retail store storage account for three months to include read-only access to the data.

Security -

- All Azure Functions must centralize management and distribution of configuration data for different environments and geographies, encrypted by using a company-provided RSA-HSM key.
- Authentication and authorization must use Azure AD and services must use managed identities where possible.

Issues -

Retail Store Locations -

- You must perform a point-in-time restoration of the retail store location data due to an unexpected and accidental deletion of data.
- Azure Cosmos DB queries from the Azure Function exhibit high Request Unit (RU) usage and contain multiple, complex queries that exhibit high point read latency for large items as the function app is scaling.

You need to test the availability of the corporate website.

Which two test types can you use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. Standard
- B. URL ping
- C. Custom testing using the TrackAvailability API method
- D. Multi-step

Correct Answer: AC

Community vote distribution

AC (100%)

 **arunkuml**  4 months ago

Got it in the exam 14/12/23. Scored 912/1000. All questions are from ExamTopics. Case study - VanArsdel, Ltd (11 questions)

upvoted 5 times

 **Firo**  8 months, 3 weeks ago

Selected Answer: AC

The answer is correct

A-Standard

C-Custom Testing - TrackAvailability

<https://learn.microsoft.com/en-us/azure/azure-monitor/app/availability-overview#types-of-tests>

upvoted 5 times

 **AndySmith**  5 months, 2 weeks ago

On exam 3-Nov-2023. Went with proposed answer - 932/1000.

upvoted 3 times

 **AndySmith** 5 months, 3 weeks ago

A and C per this requirement: "Implement monitoring by using Application Insights and availability web tests including SSL certificate validity and custom header value verification."

A - Standard test: This single request test is similar to the URL ping test. It includes TLS/SSL certificate validity, proactive lifetime check, HTTP request verb.

C - Custom TrackAvailability test: If you decide to create a custom application to run availability tests, you can use the TrackAvailability() method to send the results to Application Insights.

<https://learn.microsoft.com/en-us/azure/azure-monitor/app/availability-overview#types-of-tests>

upvoted 5 times

 **ENGs** 6 months, 3 weeks ago

Selected Answer: AC

A and C

upvoted 2 times

 **drkpns** 8 months ago

Because it is for the local use, my understanding is it is in the intranet, not internet. So url ping & multi step cannot be used.

<https://learn.microsoft.com/en-us/azure/azure-monitor/app/availability-overview#types-of-tests>

upvoted 2 times

 **lorenaizzo** 8 months ago

The answer must meet the following rule "• Implement monitoring by using Application Insights and availability web tests including SSL certificate validity and custom header value verification."

So Standard Test can check if the certificate is valid. And creating a custom application and sending the results using a Custom TrackAvailabilityTest.

About the local site vs internet...

"The older classic tests, URL ping test and multi-step web test, rely on the DNS infrastructure of the public internet to resolve the domain names of the tested endpoints. If you're using private DNS, you must ensure that the public domain name servers can resolve every domain name of your test. When that's not possible, you can use custom TrackAvailability tests instead."

I understand that you can use any of them, just need to pay attention to the DNS entries (and routing)

upvoted 1 times

 **lyggwtwtxnxhznebw** 8 months, 1 week ago

Selected Answer: AC

Based on <https://learn.microsoft.com/en-us/azure/azure-monitor/app/availability-overview#types-of-tests> the answer is correct.

upvoted 4 times

 **Arr14487** 8 months, 2 weeks ago

given answer , i think its wrong.
it should be urlping and multi step

upvoted 2 times

Question #44

Topic 5

You have an Azure API Management (APIM) Standard tier instance named APIM1 that uses a managed gateway.

You plan to use APIM1 to publish an API named API1 that uses a backend database that supports only a limited volume of requests per minute. You also need a policy for API1 that will minimize the possibility that the number of requests to the backend database from an individual IP address you specify exceeds the supported limit.

You need to identify a policy for API1 that will meet the requirements.

Which policy should you use?

- A. ip-filter
- B. quota-by-key
- C. rate-limit-by-key
- D. rate-limit

Correct Answer: C

Community vote distribution

C (100%)

 **dddddd111** 6 months ago

Selected Answer: C

C is correct.

upvoted 1 times

 **Ciupaz** 6 months, 4 weeks ago

Selected Answer: C

The "rate-limit-by-key" policy in Azure API Management allows you to apply rate limiting based on a specific key, which can be an IP address or any other attribute that you want to use for rate limiting.

upvoted 2 times

 **macobuzi** 7 months, 3 weeks ago

Selected Answer: C

C is correct:

Rate limits are usually used to protect against short and intense volume bursts. For example, if you know your backend service has a bottleneck at its database with a high call volume, you could set a rate-limit-by-key policy to not allow high call volume by using this setting.

<https://learn.microsoft.com/en-us/azure/api-management/api-management-sample-flexible-throttling#rate-limits-and-quotas>

upvoted 3 times

 **drogbix** 8 months, 2 weeks ago

Selected Answer: C

Correct Answer.

<https://learn.microsoft.com/en-us/azure/api-management/api-management-policies>

upvoted 1 times

Question #45

Topic 5

You develop a web application that sells access to last-minute openings for child camps that run on the weekends. The application uses Azure Application Insights for all alerting and monitoring.

The application must alert operators when a technical issue is preventing sales to camps.

You need to build an alert to detect technical issues.

Which alert type should you use?

- A. Metric alert using multiple time series
- B. Metric alert using dynamic thresholds
- C. Log alert using multiple time series
- D. Log alert using dynamic thresholds

Correct Answer: B

Community vote distribution

B (100%)

✉️  **Ciupaz**  5 months, 2 weeks ago

Selected Answer: B

I think is correct.

<https://learn.microsoft.com/en-us/azure/azure-monitor/alerts/alerts-dynamic-thresholds>

upvoted 5 times

✉️  **AndySmith**  5 months, 2 weeks ago

On exam 3-Nov-2023. Went with proposed answer - 932/1000.

upvoted 5 times

✉️  **vovap0vovap**  3 weeks, 5 days ago

I think that requirements here is quote opposite that "Considerations when using dynamic thresholds" in <https://learn.microsoft.com/en-us/azure/azure-monitor/alerts/alerts-dynamic-thresholds> - which basically describe long time stable series

upvoted 1 times

Question #46

Topic 5

Case study -

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study -

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. When you are ready to answer a question, click the Question button to return to the question.

Background -

Munson's Pickles and Preserves Farm is an agricultural cooperative corporation based in Washington, US, with farms located across the United States. The company supports agricultural production resources by distributing seeds, fertilizers, chemicals, fuel, and farm machinery to the farms.

Current Environment -

The company is migrating all applications from an on-premises datacenter to Microsoft Azure. Applications support distributors, farmers, and internal company staff.

Corporate website -

- The company hosts a public website located at <http://www.munsonspicklesandpreservesfarm.com>. The site supports farmers and distributors who request agricultural production resources.

Farms -

- The company created a new customer tenant in the Microsoft Entra admin center to support authentication and authorization for applications.

Distributors -

- Distributors integrate their applications with data that is accessible by using APIs hosted at <http://www.munsonspicklesandpreservesfarm.com/api> to receive and update resource data.

Requirements -

The application components must meet the following requirements:

Corporate website -

- The site must be migrated to Azure App Service.
- Costs must be minimized when hosting in Azure.

- Applications must automatically scale independent of the compute resources.
- All code changes must be validated by internal staff before release to production.
- File transfer speeds must improve, and webpage-load performance must increase.
- All site settings must be centrally stored, secured without using secrets, and encrypted at rest and in transit.
- A queue-based load leveling pattern must be implemented by using Azure Service Bus queues to support high volumes of website agricultural production resource requests.

Farms -

- Farmers must authenticate to applications by using Microsoft Entra ID.

Distributors -

- The company must track a custom telemetry value with each API call and monitor performance of all APIs.
- API telemetry values must be charted to evaluate variations and trends for resource data.

Internal staff -

- App and API updates must be validated before release to production.
- Staff must be able to select a link to direct them back to the production app when validating an app or API update.
- Staff profile photos and email must be displayed on the website once they authenticate to applications by using their Microsoft Entra ID.

Security -

- All web communications must be secured by using TLS/HTTPS.
- Web content must be restricted by country/region to support corporate compliance standards.
- The principle of least privilege must be applied when providing any user rights or process access rights.
- Managed identities for Azure resources must be used to authenticate services that support Microsoft Entra ID authentication.

Issues -

Corporate website -

- Farmers report HTTP 503 errors at the same time as internal staff report that CPU and memory usage are high.
- Distributors report HTTP 502 errors at the same time as internal staff report that average response times and networking traffic are high.
- Internal staff report webpage load sizes are large and take a long time to load.
- Developers receive authentication errors to Service Bus when they debug locally.

Distributors -

- Many API telemetry values are sent in a short period of time. Telemetry traffic, data costs, and storage costs must be reduced while preserving a statistically correct analysis of the data points sent by the APIs.

You need to implement an aggregate of telemetry values for distributor API calls.

Which Application Insights API method should you use?

- A. TrackEvent
- B. TrackDependency
- C. TrackMetric
- D. TrackException
- E. TrackTrace

Correct Answer: C

Community vote distribution C (100%)

✉  **AzDeveloper** 3 months, 1 week ago

It seems C is correct answer

<https://learn.microsoft.com/en-us/azure/azure-monitor/app/data-model-complete>

upvoted 1 times

✉  **Ciupaz** 3 months, 1 week ago

Selected Answer: C

The correct answer is C. TrackMetric.

upvoted 2 times

Question #47

Topic 5

DRAG DROP

-

Case study

-

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

-

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. When you are ready to answer a question, click the Question button to return to the question.

Background

-

Munson's Pickles and Preserves Farm is an agricultural cooperative corporation based in Washington, US, with farms located across the United States. The company supports agricultural production resources by distributing seeds, fertilizers, chemicals, fuel, and farm machinery to the farms.

Current Environment

-

The company is migrating all applications from an on-premises datacenter to Microsoft Azure. Applications support distributors, farmers, and internal company staff.

Corporate website

-

- The company hosts a public website located at <http://www.munsonspicklesandpreservesfarm.com>. The site supports farmers and distributors who request agricultural production resources.

Farms

-

- The company created a new customer tenant in the Microsoft Entra admin center to support authentication and authorization for applications.

Distributors

-

- Distributors integrate their applications with data that is accessible by using APIs hosted at

<http://www.munsonspicklesandpreservesfarm.com/api> to receive and update resource data.

Requirements

-

The application components must meet the following requirements:

Corporate website

- The site must be migrated to Azure App Service.
- Costs must be minimized when hosting in Azure.
- Applications must automatically scale independent of the compute resources.
- All code changes must be validated by internal staff before release to production.
- File transfer speeds must improve, and webpage-load performance must increase.
- All site settings must be centrally stored, secured without using secrets, and encrypted at rest and in transit.
- A queue-based load leveling pattern must be implemented by using Azure Service Bus queues to support high volumes of website agricultural production resource requests.

Farms

-

- Farmers must authenticate to applications by using Microsoft Entra ID.

Distributors

-

- The company must track a custom telemetry value with each API call and monitor performance of all APIs.
- API telemetry values must be charted to evaluate variations and trends for resource data.

Internal staff

-

- App and API updates must be validated before release to production.
- Staff must be able to select a link to direct them back to the production app when validating an app or API update.
- Staff profile photos and email must be displayed on the website once they authenticate to applications by using their Microsoft Entra ID.

Security

-

- All web communications must be secured by using TLS/HTTPS.
- Web content must be restricted by country/region to support corporate compliance standards.
- The principle of least privilege must be applied when providing any user rights or process access rights.
- Managed identities for Azure resources must be used to authenticate services that support Microsoft Entra ID authentication.

Issues

-

Corporate website

-

- Farmers report HTTP 503 errors at the same time as internal staff report that CPU and memory usage are high.
- Distributors report HTTP 502 errors at the same time as internal staff report that average response times and networking traffic are high.
- Internal staff report webpage load sizes are large and take a long time to load.
- Developers receive authentication errors to Service Bus when they debug locally.

Distributors

- Many API telemetry values are sent in a short period of time. Telemetry traffic, data costs, and storage costs must be reduced while preserving a statistically correct analysis of the data points sent by the APIs.

You need to correct the internal staff issue with webpages.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

Create an Azure Content Delivery Network profile.

Create an Azure Content Delivery Network origin group.

Configure Azure Content Delivery Network compression.

Create an Azure Content Delivery Network endpoint.

Configure a new Azure Content Delivery Network origin.

Add the Azure Content Delivery Network origin to the origin group.

Answer area**Answer area**

Create an Azure Content Delivery Network profile.

Correct Answer: Create an Azure Content Delivery Network endpoint.

Configure Azure Content Delivery Network compression.

✉ **AzDeveloper** 3 months, 1 week ago

it seems correct.

<https://learn.microsoft.com/en-us/azure/cdn/cdn-improve-performance>

upvoted 3 times

✉ **Jedi** 3 months, 1 week ago

Answer looks correct. Profile and Endpoint are covered here as the two steps: <https://learn.microsoft.com/en-us/azure/cdn/cdn-create-new-endpoint>

The text mentions "webpage load sizes are large" so compression matches that.

upvoted 2 times

Question #48

Topic 5

HOTSPOT

-

Case study

-

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

-

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. When you are ready to answer a question, click the Question button to return to the question.

Background

-

Munson's Pickles and Preserves Farm is an agricultural cooperative corporation based in Washington, US, with farms located across the United States. The company supports agricultural production resources by distributing seeds, fertilizers, chemicals, fuel, and farm machinery to the farms.

Current Environment

-

The company is migrating all applications from an on-premises datacenter to Microsoft Azure. Applications support distributors, farmers, and internal company staff.

Corporate website

-

- The company hosts a public website located at <http://www.munsonspicklesandpreservesfarm.com>. The site supports farmers and distributors who request agricultural production resources.

Farms

-

- The company created a new customer tenant in the Microsoft Entra admin center to support authentication and authorization for applications.

Distributors

-

- Distributors integrate their applications with data that is accessible by using APIs hosted at

<http://www.munsonspicklesandpreservesfarm.com/api> to receive and update resource data.

Requirements

-

The application components must meet the following requirements:

Corporate website

- The site must be migrated to Azure App Service.
- Costs must be minimized when hosting in Azure.
- Applications must automatically scale independent of the compute resources.
- All code changes must be validated by internal staff before release to production.
- File transfer speeds must improve, and webpage-load performance must increase.
- All site settings must be centrally stored, secured without using secrets, and encrypted at rest and in transit.
- A queue-based load leveling pattern must be implemented by using Azure Service Bus queues to support high volumes of website agricultural production resource requests.

Farms

-

- Farmers must authenticate to applications by using Microsoft Entra ID.

Distributors

-

- The company must track a custom telemetry value with each API call and monitor performance of all APIs.
- API telemetry values must be charted to evaluate variations and trends for resource data.

Internal staff

-

- App and API updates must be validated before release to production.
- Staff must be able to select a link to direct them back to the production app when validating an app or API update.
- Staff profile photos and email must be displayed on the website once they authenticate to applications by using their Microsoft Entra ID.

Security

-

- All web communications must be secured by using TLS/HTTPS.
- Web content must be restricted by country/region to support corporate compliance standards.
- The principle of least privilege must be applied when providing any user rights or process access rights.
- Managed identities for Azure resources must be used to authenticate services that support Microsoft Entra ID authentication.

Issues

-

Corporate website

-

- Farmers report HTTP 503 errors at the same time as internal staff report that CPU and memory usage are high.
- Distributors report HTTP 502 errors at the same time as internal staff report that average response times and networking traffic are high.
- Internal staff report webpage load sizes are large and take a long time to load.
- Developers receive authentication errors to Service Bus when they debug locally.

Distributors

- Many API telemetry values are sent in a short period of time. Telemetry traffic, data costs, and storage costs must be reduced while preserving a statistically correct analysis of the data points sent by the APIs.

You need to correct the errors for farmers and distributors.

Which solution should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Farmer and distributor errors

Issue	Solution
Farmers' errors	<ul style="list-style-type: none"> Scale up the App Service plan to Premium. Add an App Service staging deployment slot. Configure the App Service Local Cache feature. Create an Azure Content Delivery Network profile and endpoint.
Distributors' errors	<ul style="list-style-type: none"> Scale up the App Service plan to Premium. Configure the App Service Local Cache feature. Restart the application from the App Service portal. Create a custom autoscale rule to increase the instance count.

Farmer and distributor errors		
	Issue	Solution
Correct Answer:	Farmers' errors	<ul style="list-style-type: none"> Scale up the App Service plan to Premium. Add an App Service staging deployment slot. Configure the App Service Local Cache feature. Create an Azure Content Delivery Network profile and endpoint.
	Distributors' errors	<ul style="list-style-type: none"> Scale up the App Service plan to Premium. Configure the App Service Local Cache feature. Restart the application from the App Service portal. Create a custom autoscale rule to increase the instance count.

 **D4Byte** 3 months ago

Seems to me this is a scaling up and scaling out solution. Which makes sense in MS's business objectives.
upvoted 1 times

 **AzDeveloper** 3 months, 1 week ago

it seems correct.
Reference: <https://learn.microsoft.com/en-us/azure/app-service/troubleshoot-http-502-http-503>
upvoted 2 times

 **AzDeveloper** 3 months ago

sorry for distributors' errors: Creating a custom autoscale rule to increase the instance count may help to solve the error
upvoted 3 times

 **Ciupaz** 3 months, 1 week ago

Correct.

Farmers' errors:

HTTP 503 errors: These errors are most likely due to insufficient compute resources on the App Service plan. Scaling up the App Service plan to Premium will provide more CPU and memory resources, which will help to handle the increased demand and reduce the likelihood of HTTP 503 errors.

Distributors' errors:

HTTP 502 errors: These errors are typically caused by communication problems between the application and the underlying infrastructure. Restarting the application from the App Service portal may resolve the issue.

upvoted 1 times

Question #49

HOTSPOT

-

Case study

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

- To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. When you are ready to answer a question, click the Question button to return to the question.

Background

-

Munson's Pickles and Preserves Farm is an agricultural cooperative corporation based in Washington, US, with farms located across the United States. The company supports agricultural production resources by distributing seeds, fertilizers, chemicals, fuel, and farm machinery to the farms.

Current Environment

-

The company is migrating all applications from an on-premises datacenter to Microsoft Azure. Applications support distributors, farmers, and internal company staff.

Corporate website

-

- The company hosts a public website located at <http://www.munsonspicklesandpreservesfarm.com>. The site supports farmers and distributors who request agricultural production resources.

Farms

-

- The company created a new customer tenant in the Microsoft Entra admin center to support authentication and authorization for applications.

Distributors

-

- Distributors integrate their applications with data that is accessible by using APIs hosted at

<http://www.munsonspicklesandpreservesfarm.com/api> to receive and update resource data.

Requirements

-

The application components must meet the following requirements:

Corporate website

-

- The site must be migrated to Azure App Service.
- Costs must be minimized when hosting in Azure.
- Applications must automatically scale independent of the compute resources.
- All code changes must be validated by internal staff before release to production.
- File transfer speeds must improve, and webpage-load performance must increase.
- All site settings must be centrally stored, secured without using secrets, and encrypted at rest and in transit.
- A queue-based load leveling pattern must be implemented by using Azure Service Bus queues to support high volumes of website agricultural production resource requests.

Farms

-

- Farmers must authenticate to applications by using Microsoft Entra ID.

Distributors

-

- The company must track a custom telemetry value with each API call and monitor performance of all APIs.
- API telemetry values must be charted to evaluate variations and trends for resource data.

Internal staff

-

- App and API updates must be validated before release to production.
- Staff must be able to select a link to direct them back to the production app when validating an app or API update.
- Staff profile photos and email must be displayed on the website once they authenticate to applications by using their Microsoft Entra ID.

Security

-

- All web communications must be secured by using TLS/HTTPS.
- Web content must be restricted by country/region to support corporate compliance standards.
- The principle of least privilege must be applied when providing any user rights or process access rights.
- Managed identities for Azure resources must be used to authenticate services that support Microsoft Entra ID authentication.

Issues

-

Corporate website

-

- Farmers report HTTP 503 errors at the same time as internal staff report that CPU and memory usage are high.
- Distributors report HTTP 502 errors at the same time as internal staff report that average response times and networking traffic are high.
- Internal staff report webpage load sizes are large and take a long time to load.
- Developers receive authentication errors to Service Bus when they debug locally.

Distributors

- Many API telemetry values are sent in a short period of time. Telemetry traffic, data costs, and storage costs must be reduced while preserving a statistically correct analysis of the data points sent by the APIs.

You need to resolve the authentication errors for developers.

Which Service Bus security configuration should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Azure Service Bus security configuration

Security configuration setting

Azure role-based access control (RBAC) role

Service Bus scope

Security configuration value

Owner
Contributor
Service Bus Data Owner
Service Bus Data Sender

Queue
Namespace
Subscription
Resource group

Azure Service Bus security configuration

Security configuration setting

Azure role-based access control (RBAC) role

Correct Answer:

Service Bus scope

Security configuration value

Owner
Contributor
Service Bus Data Owner
Service Bus Data Sender

Queue
Namespace
Subscription
Resource group

oskx2 4 weeks ago

The answer is correct. I have found the exact instruction from microsoft:

<https://learn.microsoft.com/en-us/azure/service-bus-messaging/service-bus-dotnet-get-started-with-queues?tabs=passwordless#add-microsoft-entra-user-to-azure-service-bus-owner-role>

"Add your Microsoft Entra user name to the Azure Service Bus Data Owner role at the Service Bus namespace level."
upvoted 2 times

raymond_abcd 2 months, 2 weeks ago

Think the answer is correct,

reference: <https://learn.microsoft.com/en-us/azure/service-bus-messaging/service-bus-managed-service-identity>

upvoted 2 times

FeriAZ 2 months, 2 weeks ago

correct.

Azure RBAC Role: Service Bus Data Owner. This role is appropriate for developers who need to debug Service Bus as it provides comprehensive access to the data plane, including both sending and receiving messages, which are essential capabilities for debugging.

Service Bus Scope: Namespace. The Namespace scope is recommended if the developers need to work across multiple queues or topics within

the same Service Bus namespace. If the debugging is limited to a single queue, then "Queue" could be more appropriate. However, "Namespace" offers more flexibility without being overly broad.

upvoted 3 times

 **emysa341** 3 months ago

"using Azure Service Bus queues to support high volumes of website agricultural production resource requests."
Service bus data sender --> principle of least privilege
queue

upvoted 1 times

 **AzDeveloper** 3 months ago

Correct answer is Service Bus Data Owner, Namespace

upvoted 1 times

 **AzDeveloper** 2 months, 4 weeks ago

Sorry the correct answer is : Service Bus Data Sender, Namespace
upvoted 1 times

 **azure1282** 2 months, 3 weeks ago

IMO - Scope should be Queue.
Any reason why you are suggesting Namespace?
upvoted 1 times

 **emysa341** 3 months, 1 week ago

Contributor (The principle of least privilege must be applied when providing any user rights or process access rights.)
Queue (A queue-based load leveling pattern must be implemented by using Azure Service Bus queues to support high volumes of website agricultural production resource requests.)

upvoted 3 times

Topic 6 - Question Set 6

Question #1

Topic 6

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You are developing an Azure solution to collect point-of-sale (POS) device data from 2,000 stores located throughout the world. A single device can produce 2 megabytes (MB) of data every 24 hours. Each store location has one to five devices that send data.

You must store the device data in Azure Blob storage. Device data must be correlated based on a device identifier. Additional stores are expected to open in the future.

You need to implement a solution to receive the device data.

Solution: Provision an Azure Service Bus. Configure a topic to receive the device data by using a correlation filter.

Does the solution meet the goal?

A. Yes

B. No

Correct Answer: A

A message is raw data produced by a service to be consumed or stored elsewhere. The Service Bus is for high-value enterprise messaging, and is used for order processing and financial transactions.

Reference:

<https://docs.microsoft.com/en-us/azure/event-grid/compare-messaging-services>

Community vote distribution

B (83%)

A (17%)

✉ **Spooky7** Highly Voted 2 years, 10 months ago

First of all - that question is not about THE BEST solution but about THE VALID solution. So responses like "Event Hub is the best choice" doesn't bring anything to the table.

That being said - is the provided solution valid? I would say NO, because:

- topics allows multiple subscribers, and here we need to process each event once
- correlation filter is for subscriptions, not topics
- even when assuming there is typo in the question and correlation filter is defined on the subscription level - it still is not a valid solution, because new stores can be opened in the future with many new device identifiers which you can't know in advance. Besides that filter make no sense in this scenario whatsoever, you just need to save data in storage account and basically partition it by device identifier.

upvoted 46 times

✉ **alexmalex02** 3 months, 3 weeks ago

But it is about the best solution: "Each question in the series contains a unique solution that might meet the stated goals."

Of course, as we cannot see all possible answers in exam before picking the best one, it is all about hope and pray.

upvoted 2 times

✉ **RDTAus** 3 months, 2 weeks ago

I think you might have missed the sentence after that "Some question sets might have more than one correct solution, while others might not have a correct solution." It's NOT just about the best solution.

upvoted 1 times

✉ **kabs** Highly Voted 3 years, 9 months ago

Answer should be event hub

upvoted 28 times

✉ **Tmm58** 3 years, 4 months ago

Why Event Hub? This is clearly about delivering messages (with the actual data) and not events.

upvoted 24 times

✉ **xofowi5140** 3 years, 9 months ago

The event hub message size limit is 1MB.

upvoted 26 times

✉ **MaheshBeeravelli** 3 years, 9 months ago

I agree with xofowi5140. Given answer is correct.

upvoted 13 times

✉ **Adol** 3 years, 3 months ago

Message size of 1MB has no conflict with 2MB of messages in 24 hours! So still Event hub when it comes to IoT or when many devices sending data with device id as identifier

upvoted 5 times

✉ **profesorklaus** 3 years, 2 months ago

Of course it has a conflict. One store can send message as a batch with size 2MB.

upvoted 8 times

✉ **TalesFromEarth** 3 years, 9 months ago

So is in service hub. Maximum message size: 256 KB for Standard tier, 1 MB for Premium tier.

upvoted 3 times

✉ **Dirk** 3 years, 8 months ago

256 KB is the max size of the message queue, But this is not about the messages delivered per 24 hours.

<https://stackoverflow.com/questions/22384193/azure-service-bus-message-size-technology-limit-and-pricing/22384741#:~:text=1%20Answer&text=Service%20Bus%20queues%20support%20a,messages%20held%20by%20a%20queue.>

In the question it's about the size in 24 hours.

upvoted 4 times

✉ **oskx2** **Most Recent** 4 weeks ago

Answer is incorrect. It is easier to implement azure event hub and enable capture to save it to azure storage account. Also, if the question does not have a typo, you can't configure filters on the topic (should be subscription).

upvoted 1 times

✉ **raymond_abcd** 2 months, 2 weeks ago

Selected Answer: B

I think it is B because the of the message size of 2 MB, It is not mentioned that the service bus depends on the Premium tier which is max 100MB message size for a topic. While standard only can hold 256kb message

upvoted 1 times

✉ **CarlosTheBoldest** 4 months, 1 week ago

I got this question on my exam, 2023Dec, go with what I remember was the most voted answer. Score 902, most of the questions were here, slightly different on wording because the Azure Ad <-> Entra Id change. Case was City Power & Light. Good luck!
Important tip, you have access to microsoft learn during the exam!

upvoted 1 times

✉ **fefbb52** 4 months, 1 week ago

whar is the correct answer?

upvoted 1 times

✉ **dddddd111** 6 months ago

Selected Answer: B

Answer is No.

upvoted 1 times

✉ **lugia4000** 6 months, 3 weeks ago

These question series came out on my exam 2023-09-29

upvoted 3 times

✉ **adilkhan** 1 year, 1 month ago

One approach could be to use Azure IoT Hub to receive the data from the POS devices, and then use Azure Functions or Azure Stream Analytics to process and store the data in Azure Blob storage based on the device identifier. ChatGPT

upvoted 1 times

✉ **Michael2023** 1 year, 2 months ago

Events vs Message

This is clearly a message (2MB) from multiple devices which needs to be stored in Azure Blob Storage... that means we need to decide between Azure Service Bus or Azure Queue Storage queues (if available option) because these are message-based solutions...

for me Azure Service Bus sounds like a good fit... take a look at the comparison page of these Azure Services

<https://learn.microsoft.com/en-us/azure/event-grid/compare-messaging-services>

upvoted 1 times

✉ **adilkhan** 1 year, 2 months ago

Here we are using Azure Blob storage. We store the data in the blob storage and name of the blob we store in the Azure service bus to process it but we don't need topic here.

upvoted 2 times

✉ **rolling_potato_** 1 year, 3 months ago

This series of questions I got on my exam today (03/01/2023) I chose no for Service Bus and Event Grid and Yes for Event HUB (score 921)

upvoted 9 times

OPT_001122 1 year, 4 months ago

Selected Answer: B

Correct ans - event hub (There are multiple questions with different answers)

Event Hub for this scenario - 2000 POS devices

upvoted 1 times

hubekpeter 1 year, 4 months ago

Selected Answer: A

I vote for A. It's POS terminal. We are talking about financial transactions that need to be reliably delivered and processed at least once. All the limitations mentioned below can be easily overcome. Azure service - bus service is intended for enterprise applications that require transactions, ordering, duplicate detection, and instantaneous consistency.

upvoted 2 times

somename20221106 1 year, 5 months ago

Answer will be Yes - As Service Bus is the recommended for order processing and Financials
POS - is order process/finance.

<https://learn.microsoft.com/en-us/azure/event-grid/compare-messaging-services>

upvoted 1 times

Knightie 1 year, 5 months ago

remember service bus is an always correct answer in all of the event and messaging scenario.

upvoted 2 times

Knightie 1 year, 5 months ago

Correlation Filters - A CorrelationFilter holds a set of conditions that are matched against one or more of an arriving message's user and system properties. A common use is to match against the CorrelationId property, but the application can also choose to match against the following properties:

ContentType

Label

MessageId

ReplyTo

ReplyToSessionId

SessionId

To

any user-defined properties.

A match exists when an arriving message's value for a property is equal to the value specified in the correlation filter. For string expressions, the comparison is case-sensitive. If you specify multiple match properties, the filter combines them as a logical AND condition, meaning for the filter to match, all conditions must match.

Answer A, nothing to stop us from grouping up the topics and you can have whosoever interested to subscribe to the message of the grouped devices.

upvoted 1 times

Knightie 1 year, 5 months ago

Selected Answer: A

Yes, it can be used so... even if you use the correlation id as topic, the point is. it is still workable, right? who knows how the pub/sub model they wanna be? may be the finance, boss and account wanna be the 3 subscribers for all different topics? unless there is a limitation of topics... otherwise, nothing is stopping the correlation_id to be the topic.

upvoted 2 times

Question #2

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen. You are developing an Azure solution to collect point-of-sale (POS) device data from 2,000 stores located throughout the world. A single device can produce 2 megabytes (MB) of data every 24 hours. Each store location has one to five devices that send data.

You must store the device data in Azure Blob storage. Device data must be correlated based on a device identifier. Additional stores are expected to open in the future.

You need to implement a solution to receive the device data.

Solution: Provision an Azure Event Grid. Configure event filtering to evaluate the device identifier.

Does the solution meet the goal?

A. Yes

B. No

Correct Answer: B

Instead use an Azure Service Bus, which is used order processing and financial transactions.

Note: An event is a lightweight notification of a condition or a state change. Event hubs is usually used reacting to status changes.

Reference:

<https://docs.microsoft.com/en-us/azure/event-grid/compare-messaging-services>

Community vote distribution

B (86%)

14%

 **xRiot007**  3 years, 9 months ago

For the ones above me. Sensors do not send events, they send messages containing specific data that has been gathered. This makes automatically the solution incorrect, because you need a Service Bus to collect them. Event Grids and Event Hubs won't do the job here.
upvoted 44 times

 **muggins** 1 year, 10 months ago

Except these are Point of Sale devices, not sensors.

upvoted 2 times

 **xRiot007** 1 year, 10 months ago

It's irrelevant, you have a device that sends *messages* containing valuable business information. The service bus implementation is even used in Microsoft's own documentation as an example for POS data gathering implementation
upvoted 2 times

 **Nabeelcp**  3 years, 7 months ago

Event Grid and Event Hubs are basically for event based communication. here the scenario is more suited to Message based communication. so the answer is correct

upvoted 15 times

 **ferut** 2 years, 11 months ago

The question doesn't mention 'message', it's just plain 'data'. Events can also be data.

upvoted 3 times

 **raymond_abcd**  2 months, 2 weeks ago

Selected Answer: B

It is B because of the message size of 1 MB.

Max message size: 512 KB (for a topic)

Max event size within an array: 1 MB

upvoted 1 times

 **Weam** 4 months, 3 weeks ago

Selected Answer: B

Event sources send events to Azure Event Grid in an array, which can have several event objects. When posting events to an Event Grid topic, the array can have a total size of up to 1 MB. Each event in the array is limited to 1 MB. If an event or the array is greater than the size limits, you receive the response 413 Payload Too Large. Operations are charged in 64 KB increments though. So, events over 64 KB incur operations charges as though they were multiple events. For example, an event that is 130 KB would incur operations as though it were three separate events.

<https://learn.microsoft.com/en-us/azure/event-grid/event-schema>

upvoted 1 times

 **vovap0vovap** 3 weeks, 5 days ago

That POS terminals. not even close to 64K
upvoted 1 times

✉ **ddd111** 6 months ago

Selected Answer: A
I think the answer is Yes.
upvoted 1 times

✉ **Eltooth** 1 year, 9 months ago

Selected Answer: B
B is correct answer.
Event Hub would be a solution here NOT Event Grid.
upvoted 3 times

✉ **meoukg** 2 years, 1 month ago

Got it on 03/2022, chose B. No
upvoted 1 times

✉ **Netspud** 2 years, 1 month ago

Selected Answer: B
Answer is no.
upvoted 1 times

✉ **Aratilra** 2 years, 3 months ago

What is the right ans ??????????
upvoted 3 times

✉ **chingdm** 2 years, 1 month ago

I would go for Service Bus because it supports transactional data. POS usually connects to inventory at real time, so make sense for a transaction.
upvoted 1 times

✉ **ning** 2 years, 8 months ago

No, for any subs, you can have max 100 event grid topics, but there are 2000 stores and 5 POS for each store, for each topic, you can have max 500 subs. The only possible way, is to multiple those, to generate 50,000 subs for 100 topics. But, this would be a horrible situation
upvoted 2 times

✉ **Knightie** 1 year, 7 months ago

This is the reason for b, else Event Grid makes sense too. It works with a pub/sub topic and event handler based operation. It can work in the scenario, but too little topics for 2000 stores with 5 devices...100 topics is not efficient at all.
upvoted 1 times

✉ **Ram0202** 2 years, 8 months ago

max size allowed in service bus is 1 mb ,ques is for 2 mb so go for event hub
<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted>
upvoted 1 times

✉ **altafpatel1984** 2 years, 5 months ago

max size allowed for event hub is 1 MB
upvoted 1 times

✉ **renzoku** 1 year, 9 months ago

2MB for day, each message could be a lot less
upvoted 1 times

✉ **kondapaturi** 2 years, 9 months ago

Answer – No, The Azure Event Grid service is used to receiving events and would not fit the purpose of the requirement.
upvoted 1 times

✉ **Hamg** 2 years, 11 months ago

I think many of clarionprogrammer's answers can be very confusing
upvoted 1 times

✉ **eobo** 3 years ago

Correct Answer is : B

Reference: <https://azure.microsoft.com/en-us/blog/events-data-points-and-messages-choosing-the-right-azure-messaging-service-for-your-data/>

A sale recorded in a point-of-sale solution is both a financial record and an inventory tracking record, and not a mere event. It's recorded in a ledger, which will eventually be merged into a centralized accounting system, often via several integration bridges, and the information must not be lost on the way. The sales information, possibly expressed as separate messages to keep track of the stock levels at the point of sale, and across the sales region, may be used to initiate automated resupply orders with order status flowing back to the point of sale. Service Bus queues or topic subscriptions are ideal for this use-case

upvoted 8 times

 **paru123456789** 3 years, 1 month ago

Answer: No

upvoted 3 times

 **agcertif** 3 years, 3 months ago

An Event is a lightweight notification that indicates that something happened.

A Message contains the data itself (not just a reference to that data)

=> Correct answer

upvoted 2 times

 **Cornholioz** 3 years, 4 months ago

Just curious: what if IoT Hub or Event Hub was provided as a possible option.

Agree that the scenario is messages based and not event based.

But wouldn't the use case still hold good to have maybe the IoT Hub be the right pick too?

upvoted 1 times

Question #3

DRAG DROP -

You manage several existing Logic Apps.

You need to change definitions, add new logic, and optimize these apps on a regular basis.

What should you use? To answer, drag the appropriate tools to the correct functionalities. Each tool may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Answer Area

Tools	Functionality	Tool
Logic Apps Designer	Edit B2B workflows	
Code View Editor	Edit definitions in JSON	
Enterprise Integration Pack	Visually add functionality	

Answer Area

Tools	Functionality	Tool
Logic Apps Designer	Edit B2B workflows	Enterprise Integration Pack
Code View Editor	Edit definitions in JSON	Code View Editor
Enterprise Integration Pack	Visually add functionality	Logic Apps Designer

Box 1: Enterprise Integration Pack

For business-to-business (B2B) solutions and seamless communication between organizations, you can build automated scalable enterprise integration workflows by using the Enterprise Integration Pack (EIP) with Azure Logic Apps.

Box 2: Code View Editor -

Edit JSON - Azure portal -

1. Sign in to the Azure portal.
2. From the left menu, choose All services. In the search box, find "logic apps", and then from the results, select your logic app.
3. On your logic app's menu, under Development Tools, select Logic App Code View.
4. The Code View editor opens and shows your logic app definition in JSON format.

Box 3: Logic Apps Designer -

Reference:

<https://docs.microsoft.com/en-us/azure/logic-apps/logic-apps-enterprise-integration-overview> <https://docs.microsoft.com/en-us/azure/logic-apps/logic-apps-author-definitions>

✉  **kimalto452**  2 years, 11 months ago
start from march 2021 no logic apps in the exam!
upvoted 42 times

✉  **Marusyk**  3 years, 1 month ago
The answer is correct
upvoted 39 times

✉  **Ciupaz**  4 months, 1 week ago
Hope to get this question in my exam.
upvoted 2 times

✉  **olowyinka** 2 years, 3 months ago

Correct answer 🔥

upvoted 5 times

✉  **JSrikant** 2 years, 6 months ago

That's correct Answer

upvoted 4 times

Question #4

Topic 6

A company is developing a solution that allows smart refrigerators to send temperature information to a central location. The solution must receive and store messages until they can be processed. You create an Azure Service Bus instance by providing a name, pricing tier, subscription, resource group, and location.

You need to complete the configuration.

Which Azure CLI or PowerShell command should you run?

A.

```
az servicebus queue create
--resource-group fridge-rg
--namespace-name fridge-ns
--name fridge-q
```

B.

```
New-AzureRmResourceGroup
-Name fridge-rg
-Location fridge-loc
```

C.

```
az servicebus namespace create
--resource-group fridge-rg
--name fridge-ns
--location fridge-loc
```

D.

```
connectionString=$(az servicebus namespace authorization-rule keys list
--resource-group fridge-rg
--namespace-name fridge-ns
--query primaryConnectionString --output tsv)
```

Correct Answer: A

A service bus instance has already been created (Step 2 below). Next is step 3, Create a Service Bus queue.

Note:

Steps:

Step 1: # Create a resource group

```
resourceGroupName="myResourceGroup"
```

```
az group create --name $resourceGroupName --location eastus
```

Step 2: # Create a Service Bus messaging namespace with a unique name namespaceName=myNameSpace\$RANDOM az servicebus namespace create --resource-group \$resourceGroupName --name \$namespaceName --location eastus

Step 3: # Create a Service Bus queue

```
az servicebus queue create --resource-group $resourceGroupName --namespace-name $namespaceName --name BasicQueue
```

Step 4: # Get the connection string for the namespace

```
connectionString=$(az servicebus namespace authorization-rule keys list --resource-group $resourceGroupName --namespace-name $namespaceName --name
$RootManageSharedAccessKey --query primaryConnectionString --output tsv)
```

Reference:

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-quickstart-cli>

✉  **SnakePlissken**  2 years, 11 months ago

I think the Service Bus has already been created and now they ask you to complete the configuration. The next step is creating the queue. In fact, all the steps are shown:

- B. Create group.
- C. Create Service Bus.
- A. Create Queue. <-- Correct answer.
- D. Get connectionstring.

upvoted 61 times

✉  **xRiot007** 1 year, 10 months ago

Your steps are correct.

You need a resource group before creating any resource.

The service bus that you create in Step 2 will be a resource in that group.

Inside an Azure Service Bus you can create queues and topics

Inside the same Azure Service Bus, you will have the primary and secondary connection strings. You can get the connection before or after you created your queue.

In code, you will connect to the service bus and then specify the name of the topic or queue to the service client.

upvoted 2 times

✉  **SnakePlissken** 2 years, 11 months ago

Got this question at the exam and scored 100% on Azure Storage, so I'm sure this is correct.
upvoted 21 times

✉  **SnakePlissken** 2 years, 10 months ago

I'm really sorry, Service Bus falls in another section where I scored 85%, so I can't be sure if this is the correct answer.
upvoted 17 times

✉  **RavindraDevkhile**  3 years ago

It Ask for create an Azure Service Bus instance by providing not the Service bus Queue
So correct Answes should be C
upvoted 22 times

✉  **Bear_Polar** 1 year, 4 months ago

As C location part is incorrect (fridge-loc is not a variable neither a valid location). So, A must be the ans (assume that the queue has been created)
upvoted 2 times

✉  **damianadalid** 11 months ago

In Azure Portal explation to set up Azure Service Bus, it forces you to create a namespace as you start creating the instance:
<https://learn.microsoft.com/en-us/azure/service-bus-messaging/service-bus-quickstart-portal> So I would say the namespace is already there and therefore the answer is A (create the queue).
upvoted 1 times

✉  **Michael2023**  1 year, 2 months ago

Got this on my exam 30.1.2023 went with Highly Voted answer score 876
upvoted 2 times

✉  **OPT_001122** 1 year, 5 months ago

A or C? looking at question C looks correct but here A is highly voted
upvoted 1 times

✉  **student85** 1 year, 5 months ago

I think A is correct <https://learn.microsoft.com/en-us/cli/azure/servicebus/queue?view=azure-cli-latest>
upvoted 1 times

✉  **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022
upvoted 3 times

✉  **rcamara32** 2 years, 3 months ago

C is the correct answer
upvoted 5 times

✉  **tiggertag2000** 2 years, 4 months ago

C is the correct answer
<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-tutorial-topics-subscriptions-cli>
upvoted 2 times

✉  **tiggertag2000** 2 years, 4 months ago

I take that back, the service bus is already created, so A is the right answer
upvoted 3 times

✉  **hinoue** 2 years, 5 months ago

question says You create an Azure Service Bus instance by providing a name, pricing tier, subscription, resource group, and location
pricing tier(sku) can be set at the creation of namespace, not queue

So, I think answer is C.

https://docs.microsoft.com/en-us/cli/azure/servicebus/namespace?view=azure-cli-latest#az_servicebus_namespace_create
https://docs.microsoft.com/en-us/cli/azure/servicebus/queue?view=azure-cli-latest#az_servicebus_queue_create
upvoted 4 times

✉  **noro5** 2 years, 3 months ago

It says also 'You need to complete the configuration.' so I assume the Service Bus instance is in place
upvoted 3 times

✉  **Franz22** 2 years, 6 months ago

Answer A is correct, because if we look at the documentation of Microsoft where are listed all the steps to create a Service Bus:
<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-quickstart-portal>

We can see that when we create a Service Bus, we MUST specify a namespace so the ServiceBus instance is the namespace itself and the question says "You create a service bus instance... You need to complete the configuration" --> So, in order to complete the configuration, we have to create a Queue.

upvoted 8 times

✉  **RajMasilamani** 2 years, 7 months ago

The answer should be A.
connection string and the queue name. You use them to send and receive messages.
<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-quickstart-cli>

upvoted 2 times

 **ning** 2 years, 8 months ago

No idea what this question is asking ...

upvoted 11 times

 **warchoon** 1 year, 1 month ago

It's asking which command is correct. The only one has no errors and provides requirements

upvoted 1 times

 **mc0re** 2 years, 8 months ago

According to Microsoft, the namespace must be created before the queue. So the correct answer must be C.

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-quickstart-cli>

upvoted 2 times

 **Mr_X** 2 years, 8 months ago

Commands B&C have set wrong location this will throw an error.

D does not make sense to me at all. Thats nonsense.

upvoted 6 times

 **Helpnosense** 2 years, 9 months ago

--location fridge-loc is definitely wrong for az servicebus namespace create command. Whoever thinks C is right, just try the command to see if CLI will succeed or not.

A is correct.

upvoted 7 times

 **Laaptu** 2 years, 9 months ago

Option A is correct, az servicebus queue create --resource-group myresourcegroup --namespace-name mynamespace --name myqueue

https://docs.microsoft.com/en-us/cli/azure/servicebus/queue?view=azure-cli-latest#az_servicebus_queue_create

upvoted 2 times

 **manojchavan** 2 years, 10 months ago

Question is poorly worded. I think what is asked here is that service bus instance is already created and now you need to complete the configuration to start using the bus. In this case, you will need to create Queue and hence correct answer is A.

upvoted 3 times

Question #5

HOTSPOT -

You are developing an application that uses Azure Storage Queues.

You have the following code:

```
CloudStorageAccount storageAccount = CloudStorageAccount.Parse
(CloudConfigurationManager.GetSetting("StorageConnectionString"));
CloudQueueClient queueClient = storageAccount.CreateCloudQueueClient();

CloudQueue queue = queueClient.GetQueueReference("appqueue");
await queue.CreateIfNotExistsAsync();

CloudQueueMessage peekedMessage = await queue.PeekMessageAsync();
if (peekedMessage != null)
{
    Console.WriteLine("The peeked message is: {0}", peekedMessageAsString);
}
CloudQueueMessage message = await queue.GetMessageAsync();
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statement	Yes	No
The code configures the lock duration for the queue.	<input type="radio"/>	<input type="radio"/>
The last message read remains in the queue after the code runs.	<input type="radio"/>	<input type="radio"/>
The storage queue remains in the storage account after the code runs.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statement	Yes	No
The code configures the lock duration for the queue.	<input type="radio"/>	<input checked="" type="radio"/>
The last message read remains in the queue after the code runs.	<input checked="" type="radio"/>	<input type="radio"/>
The storage queue remains in the storage account after the code runs.	<input checked="" type="radio"/>	<input type="radio"/>

Box 1: No -

The QueueDescription.LockDuration property gets or sets the duration of a peek lock; that is, the amount of time that the message is locked for other receivers.

The maximum value for LockDuration is 5 minutes; the default value is 1 minute.

Box 2: Yes -

You can peek at the message in the front of a queue without removing it from the queue by calling the PeekMessage method.

Box 3: Yes -

Reference:

<https://docs.microsoft.com/en-us/azure/storage/queues/storage-dotnet-how-to-use-queues> <https://docs.microsoft.com/en-us/dotnet/api/microsoft.servicebus.messaging.queuedescription.lockduration>

 **rajwit** Highly Voted 3 years, 3 months ago

Given answer is correct

upvoted 46 times

 **igorole** Highly Voted 3 years, 4 months ago

GetMessageAsync:

Gets a message from the queue using the default request options. This operation marks the retrieved message as invisible in the queue for the default visibility timeout period.

Only marks the message is invisible but does not delete.

upvoted 34 times

 **muggins** 1 year, 10 months ago

Thank you. Provided answer's reasoning for 'Yes' is incorrect.

upvoted 3 times

 **raymond_abcd** Most Recent 2 months, 2 weeks ago

Given answer is correct.

The GetMessage method is part of the peek-lock pattern. It retrieves a message without actually removing them from the queue until you explicitly complete, abandon, or dead-letter them.

upvoted 1 times

 **TaoLu** 1 year, 2 months ago

when you use the GetMessageAsync method to retrieve a message from an Azure Storage queue, the message is removed from the queue. If you want to retrieve a message without removing it from the queue, you can use the PeekMessageAsync method instead.

upvoted 4 times

 **warchoon** 1 year, 1 month ago

This is not a simple Queue. This is another class that does not delete the message when it is read.

upvoted 2 times

 **OPT_001122** 1 year, 4 months ago

Given answer is correct

upvoted 1 times

 **xRiot007** 1 year, 10 months ago

No - we are not configuring anything here

Yes - a message that is retrieved from a queue and not marked as completely processed will return to the same queue after a given amount of time, no matter if you use Peek or Get.

Yes - the queue will stay there until we decide to delete it ourselves.

upvoted 4 times

 **Miroshi** 2 years, 2 months ago

2nd answer is No

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.azure.storage.queue.cloudqueue.getmessageasync?view=azure-dotnet-legacy>

You Take the msg after you peeked.

upvoted 2 times

 **Netspud** 2 years, 1 month ago

Your link appears to imply that the message is not deleted unless additional params are included. Which would suggest the answer is as quoted, yes. (Not no as you suggest)

upvoted 3 times

 **ReniRechner** 2 years, 1 month ago

peek: get the message, don't lock, don't delete
get: get the message, lock the message (make it invisible for some seconds)
delete: delete the message

intention is: if get would also delete the message from the queue, an error in the function would render the message unhandled. Thus first get, process than delete.

"at least handled once"

upvoted 9 times

 **Yazhu** 2 years, 3 months ago

I dont understand here.. its a repetitive of 9th question.. its a copy paste question
why diff answers? can someone clarify me?

upvoted 1 times

✉ **Yazhu** 2 years, 3 months ago
commented in wrong window..sorry
upvoted 1 times

✉ **nonoss** 2 years, 5 months ago
Box 3 : Correct answer
<https://docs.microsoft.com/en-us/azure/storage/queues/storage-dotnet-how-to-use-queues?tabs=dotnet#delete-a-queue>
upvoted 1 times

✉ **paru123456789** 3 years, 1 month ago
Answer:
No
Yes
Yes
upvoted 12 times

✉ **bhushan_786** 3 years, 3 months ago
Can someone confirm if the given answers are correct or not??
upvoted 2 times

✉ **diligent176** 3 years, 3 months ago
They are correct, N, Y, Y.
The code does not configure lock duration.
The message will remain because GetMessageAsync does not remove it.
The queue will also remain after execution.
upvoted 14 times

✉ **igorole** 3 years, 4 months ago
Previous comment can be ignored, sorry, the right interface is here:
<https://docs.microsoft.com/en-us/dotnet/api/microsoft.azure.storage.queue.cloudqueue.getmessageasync?view=azure-dotnet-legacy>
upvoted 2 times

✉ **igorole** 3 years, 4 months ago
LastLine, GetMessageAsync() is not even part of the interface, this code just won't compile.
<https://docs.microsoft.com/en-us/dotnet/api/azure.storage.queues.queueclient?view=azure-dotnet>
upvoted 1 times

✉ **Kobee** 3 years, 4 months ago
You're wrong. It's CloudQueue, not QueueClient
upvoted 4 times

✉ **gematsaljoa** 3 years, 4 months ago
1. X
2. O
3. O
upvoted 5 times

✉ **JVTM** 3 years, 5 months ago
2) explanation comments a peekMessage (which leave message in queue). But the last command is getMessageAsync() which reads the message out of the queue. So, it is not in queue anymore.
upvoted 3 times

✉ **thomas204** 3 years, 5 months ago
According to me you need to call DeleteMessageAsync() to remove the message from the queue.
upvoted 21 times

✉ **ashuyop** 3 years, 5 months ago
thats right
<https://docs.microsoft.com/en-us/azure/storage/queues/storage-dotnet-how-to-use-queues?tabs=dotnet#de-queue-the-next-message>
upvoted 8 times

✉ **nonoss** 2 years, 5 months ago
Dequeue a message from a queue in two steps. When you call ReceiveMessages, you get the next message in a queue. A message returned from ReceiveMessages becomes invisible to any other code reading messages from this queue. By default, this message stays invisible for 30 seconds. To finish removing the message from the queue, you must also call DeleteMessage. This two-step process of removing a message assures that if your code fails to process a message due to hardware or software failure, another instance of your code can get the same message and try again. Your code calls DeleteMessage right after the message has been processed.
upvoted 1 times

✉ **Magneton** 3 years, 5 months ago
GetMessageAsync() does not get out the message from the queue
upvoted 31 times

Question #6

Topic 6

A company is developing a solution that allows smart refrigerators to send temperature information to a central location. The solution must receive and store messages until they can be processed. You create an Azure Service Bus instance by providing a name, pricing tier, subscription, resource group, and location.

You need to complete the configuration.

Which Azure CLI or PowerShell command should you run?

A.

```
az group create
--name fridge-rg
--location fridge-loc
```

B.

```
New-AzureRmServiceBusNamespace
-ResourceGroupName fridge-rg
-NamespaceName fridge-ns
-Location fridge-loc
```

C.

```
New-AzureRmServiceBusQueue
-ResourceGroupName fridge-rg
-NamespaceName fridge-ns
-Name fridge-q
-EnablePartitioning $False
```

D.

```
az servicebus namespace create
--resource-group fridge-rg
--name fridge-rg
--location fridge-loc
```

Correct Answer: C

A service bus instance has already been created (Step 2 below). Next is step 3, Create a Service Bus queue.

Note:

Steps:

Step 1: # Create a resource group

```
resourceGroupName="myResourceGroup"
az group create --name $resourceGroupName --location eastus
```

Step 2: # Create a Service Bus messaging namespace with a unique name namespaceName=myNameSpace\$RANDOM az servicebus namespace create --resource-group \$resourceGroupName --name \$namespaceName --location eastus

Step 3: # Create a Service Bus queue

```
az servicebus queue create --resource-group $resourceGroupName --namespace-name $namespaceName --name BasicQueue
```

Step 4: # Get the connection string for the namespace

```
connectionString=$(az servicebus namespace authorization-rule keys list --resource-group $resourceGroupName --namespace-name $namespaceName --name
```

```
RootManageSharedAccessKey --query primaryConnectionString --output tsv)
```

Reference:

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-quickstart-cli>

  **manojchavan** Highly Voted 2 years, 10 months ago

Question is poorly worded. I think what is asked here is that service bus instance is already created and now you need to complete the configuration to start using the bus. In this case, you will need to create Queue and hence correct answer is C

upvoted 31 times

  **Franz22** Highly Voted 2 years, 6 months ago

This question is equal to the question #9, the ONLY difference here is that we have to pick the Powershell answer, because it's the only that has the creation of the service bus queue. Go to the question #9 to see my answer why we should create the queue.

upvoted 13 times

  **AbidooKing** Most Recent 9 months, 1 week ago

Got this 2023July13. Went with given answer. score 917

upvoted 4 times

  **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 3 times

 **Eltooth** 1 year, 9 months ago

C is correct answer.

upvoted 6 times

 **[Removed]** 2 years ago

Agree with the accepted answer but given question 9, this seems less like it is testing your knowledge and more an attempt to trick you.

upvoted 4 times

 **SivajiTheBoss** 2 years, 1 month ago

This is the command to create Service bus namespace.

New-AzServiceBusNamespace -ResourceGroupName ContosoRG -Name ContosoSBusNS -Location eastus

Option C is wrong New-AzureRM....

upvoted 1 times

 **SivajiTheBoss** 2 years, 1 month ago

To create ServiceBusQueue this is the command in Powershell

New-AzServiceBusQueue -ResourceGroupName ContosoRG -NamespaceName ContosoSBusNS -Name ContosoOrdersQueue

upvoted 1 times

 **SivajiTheBoss** 2 years, 1 month ago

Correct Answer C: As needs to create a instance of Queue

upvoted 1 times

 **edengoforit** 2 years, 2 months ago

The description is almost the same as the question #9. Hence, after creating a namespace, we should create a Service Bus Queue thus provided answer is correct. Answer should be C

upvoted 3 times

 **Richard2021** 2 years, 6 months ago

the correct answer is B.

upvoted 2 times

 **Illumielle** 2 years, 7 months ago

This is a repeat of question 9. But it seems to have different answers.

upvoted 2 times

 **ReniRechner** 2 years, 1 month ago

this is because there is more than 1 way to do it...

upvoted 2 times

 **ning** 2 years, 8 months ago

No idea what question is this ...

Steps should be 1. create a resource group 2. create a service bus namespace 3. create a service bus queue 4. get connection string for the queue

upvoted 2 times

 **ning** 2 years, 8 months ago

There is no such thing called service bus instance ...

Only Service Bus Namespace, Service Bus Queue Or Service Bus Topic

upvoted 1 times

 **ReniRechner** 2 years, 1 month ago

namespace = instance

upvoted 1 times

 **warchoon** 1 year, 1 month ago

Everything created is an instance. It's a trick

upvoted 1 times

 **mc0re** 2 years, 8 months ago

A namespace must be created before creating a queue. So why not B or D?

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-quickstart-cli>

upvoted 1 times

 **asdasdasg2** 2 years, 4 months ago

The question states that you have already created the service bus, which is what "creating a namespace" does. So therefore the next step is to create the queue

upvoted 2 times

 **if54uran** 2 years, 10 months ago

This question does not seem to make much sense

upvoted 4 times

✉  **jokergester** 3 years ago

There is a requirement to provide an instance and yet the correct answer does not have the prior creation of the namespace. "C" would be ideal if the namespace is created prior to the script.

upvoted 2 times

✉  **Zidimirite** 3 years ago

"You create an Azure Service Bus instance by providing a name, pricing tier, subscription, resource group, and location."

They seem to ask you to do this, but in this case I think they mean that you've done that and next "You need to complete the configuration.", which is adding a queue to the just created namespace. So C is right.

It has to be right, A makes no sense and B and D both do the same thing.

upvoted 14 times

✉  **rustycables** 2 years, 8 months ago

I don't know but guessed D. D seems more like a CLI command, and it also has location - which is specified in the question. C does not have location and does not look like a typical CLI command that begins AZ.

upvoted 2 times

✉  **satyadharma** 2 years, 6 months ago

It is a powershell command which is mentioned in question as well. C is correct answer

upvoted 2 times

Question #7

Topic 6

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen. You are developing an Azure Service application that processes queue data when it receives a message from a mobile application. Messages may not be sent to the service consistently.

You have the following requirements:

- ⇒ Queue size must not grow larger than 80 gigabytes (GB).
- ⇒ Use first-in-first-out (FIFO) ordering of messages.
- ⇒ Minimize Azure costs.

You need to implement the messaging solution.

Solution: Use the .Net API to add a message to an Azure Storage Queue from the mobile application. Create an Azure Function App that uses an Azure Storage

Queue trigger.

Does the solution meet the goal?

A. Yes

B. No

Correct Answer: B

Create an Azure Function App that uses an Azure Service Bus Queue trigger.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-functions/functions-create-storage-queue-triggered-function>

Community vote distribution

B (100%)

✉  **Laaptu**  2 years, 9 months ago

Answer is B (NO),

Becz FIFO is supported by Service bus queue and Service bus queue should be use for data less than 80GB

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted>

upvoted 25 times

✉  **xRiot007** 1 year, 10 months ago

The most important thing is not capacity, but the FIFO guarantee that is provided only by the service bus and not the service queue. The link you provided is specifying that as well.

upvoted 5 times

✉  **ay_m** 8 months ago

great link, thanks a lot!

upvoted 1 times

✉  **elequier**  1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 5 times

✉  **cb98160**  1 month ago

Selected Answer: B

No, Azure Queue doesn't support FIFO!

upvoted 1 times

✉  **raymond_abcd** 2 months, 2 weeks ago

Selected Answer: B

Storage Queue is not guaranteed FIFO, and also should be used for size larger than 80 gigabytes. For this scenario a service bus queue is a better option.

<https://learn.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted>

upvoted 1 times

✉  **adilkhan** 1 year, 1 month ago

The solution uses Azure Storage Queue, which provides FIFO ordering of messages, and can handle a large volume of messages. By using the .Net API to add a message to the Azure Storage Queue, the mobile application can send messages to the queue reliably.

The Azure Function App with a Storage Queue trigger can process the messages in the queue as soon as they are added, ensuring that messages are processed in the order in which they are received. This helps meet the FIFO requirement.

Additionally, by using Azure Storage Queue, the solution can ensure that the queue size does not grow larger than 80 GB, as specified in the requirement. This can help minimize Azure costs as well, as storage usage is a key factor in determining Azure costs.

Overall, the proposed solution meets all the requirements and can provide a reliable, scalable, and cost-effective messaging solution for the given scenario. YES ChatGPT

upvoted 2 times

 **adilkhan** 1 year, 1 month ago

ChatGPT is wrong. B is correct as Azure Service Bus Queue can ensure that the queue size does not grow larger than 80 GB

upvoted 3 times

 **paunski7** 12 months ago

ChatGPT answered NO - 3.5V paid

upvoted 2 times

 **Bear_Polar** 1 year, 4 months ago

Selected Answer: B

Correct

upvoted 1 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: B

B is correct answer.

upvoted 2 times

 **Eltooth** 1 year, 9 months ago

Selected Answer: B

B is correct answer.

upvoted 3 times

 **heisenberg33** 2 years, 2 months ago

Selected Answer: B

Answer is No. Ref: <https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted>

upvoted 1 times

 **lugospod** 2 years, 3 months ago

Got this one 01/2022. Went with NO

upvoted 4 times

 **kondapaturi** 2 years, 9 months ago

Answer is No

upvoted 4 times

 **idrisfl** 3 years ago

Shouldn't this be Yes? I'm not sure where "Service Bus" came into the picture in their answer.

upvoted 3 times

 **idrisfl** 3 years ago

I take back my previous comment. 'No' is the correct answer. It says 'Queue size should not grow larger than 80 GB'

upvoted 9 times

 **Didi123** 3 years ago

agreed, also FiFo is only guaranteed in Service Bus Queues and not in Azure Storage Queues

upvoted 11 times

 **Franz22** 2 years, 6 months ago

I think that here, the size is not the thing that should let you think "oh, it's not the right answer", because you can use storage queue even with less than 80gb of data. The answer "No" should be given because the storage queue doesn't support FIFO.

upvoted 8 times

 **warchoon** 1 year, 1 month ago

And it has no Push-style API.

<https://learn.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted#:~:text=Push%2Dstyle%20API,-No,-Yes%0A%0AOur%20.NET>

upvoted 1 times

Question #8

DRAG DROP -

You develop software solutions for a mobile delivery service. You are developing a mobile app that users can use to order from a restaurant in their area. The app uses the following workflow:

1. A driver selects the restaurants for which they will deliver orders.
2. Orders are sent to all available drivers in an area.
3. Only orders for the selected restaurants will appear for the driver.
4. The first driver to accept an order removes it from the list of available orders.

You need to implement an Azure Service Bus solution.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

Create a single Service Bus topic.

Create a Service Bus Namespace for each restaurant for which a driver can receive messages.

Create a single Service Bus subscription.

Create a Service Bus subscription for each restaurant for which a driver can receive orders.

Create a single Service Bus Namespace.

Create a Service Bus topic for each restaurant for which a driver can receive messages.

Answer Area

Correct Answer:

Actions

Create a single Service Bus topic.

Create a Service Bus Namespace for each restaurant for which a driver can receive messages.

Create a single Service Bus subscription.

Create a Service Bus subscription for each restaurant for which a driver can receive orders.

Create a single Service Bus Namespace.

Create a Service Bus topic for each restaurant for which a driver can receive messages.

Answer Area

Create a single Service Bus Namespace.

Create a Service Bus topic for each restaurant for which a driver can receive messages.

Create a Service Bus subscription for each restaurant for which a driver can receive orders.

Box 1: Create a single Service Bus Namespace

To begin using Service Bus messaging entities in Azure, you must first create a namespace with a name that is unique across Azure. A namespace provides a scoping container for addressing Service Bus resources within your application.

Box 2: Create a Service Bus Topic for each restaurant for which a driver can receive messages.

Create topics.

Box 3: Create a Service Bus subscription for each restaurant for which a driver can receive orders.

Topics can have multiple, independent subscriptions.

Reference:

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messaging-overview>

 **NH**  3 years, 1 month ago

Create a single Service Bus Namespace.

Create a single Service Bus Topic.

Create a Service Bus subscription for each restaurant for which a driver can receive orders.

upvoted 158 times

 **lugospod** 2 years, 3 months ago

One thing to note in this scenario if we have 20 drivers... and 19 of them are busy (talking to their spouse, ...) it will take 19 timeouts for that one order to finally get picked up by the 20th driver. So yes, this is less evil than the rest of the options, but still I would rather use queues so

that each driver SEES all of the orders... by using subscriptions the driver doesn't have a list..s/he only sees the current message that was delegated by the round robin algorithm.

upvoted 1 times

✉ **BrettusMaximus** 2 years, 11 months ago

If you create a single topic, why bother creating a topic at all as it has all the data

upvoted 5 times

✉ **SnakePlissken** 2 years, 11 months ago

Sorry BrettusMaximus, but I think you have to read the documentation about Service Bus first... There's a good reason that NH is highly voted!

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messaging-overview>

upvoted 1 times

✉ **kimalto452** 2 years, 11 months ago

SnakePlissken you need read better documentation) There no reason for have only one topic...

upvoted 3 times

✉ **coffecold** 1 year, 6 months ago

Please come with the phrase in the article that one topic is the way to go instead of multiple topics.

Still I would go for one topic per restaurant.

upvoted 1 times

✉ **ferut** 2 years, 11 months ago

I agree with BrettusMaximus.

Although one topic is workable but multiple topics are more efficient.

Consider the 1 topic scenario. On the driver's application, it will remember the restaurants it subscribes. Because there's only one topic, the app will receive all messages from all restaurants, and the app should filter out messages not in interest.

When using multiple topics, the driver's queue will be much cleaner, receiving only messages of interest.

All drivers subscribe to the same restaurant will receive the same message. Basically each driver will have their own queue and the message will be deleted from the queue once it's processed.

I think the 'removing' part should be a separate message sent to all subscribers (finally only the restaurant can remove it from the list upon accepting a driver's request).

upvoted 5 times

✉ **lugospod** 2 years, 3 months ago

in this case you get a concurrency problem.. because now you introduced an additional layer of deciding which DRIVER clicked first, and additional component that has to notify the rest of the drivers that the message they received is no longer valid or introduce a new API to check if the order is still valid.. all in all, smells...

upvoted 1 times

✉ **MiraA** 2 years, 6 months ago

Note there is a limit of "Number of subscriptions per topic" set to "2,000 per-topic for the Standard tier and Premium tier".

<https://docs.microsoft.com/en-us/azure/azure-resource-manager/management/azure-subscription-service-limits#service-bus-limits>

Does this mean 2.000 restaurants at most? :-))

upvoted 2 times

✉ **mattvasc** 2 years, 1 month ago

Did you mean 2.000 drivers at most?

upvoted 3 times

✉ **sashasasha** 2 years, 10 months ago

If you create a single Service Bus Topic, two different drivers who subscribed to the same restaurant will each receive a copy of the order that could process independently and it's in conflict with the condition: "4. The first driver to accept an order removes it from the list of available orders."

upvoted 5 times

✉ **jay158** 2 years, 10 months ago

In fact all the drivers, who have subscribed to restaurant will get notification.

The first driver, who accepts it, will click the order, and it shall be dequeued.

upvoted 11 times

✉ **hstml** 2 years, 7 months ago

That is completely correct. This question is also in Whizlabs where the explanation is the same as jay158s.

upvoted 2 times

✉ **lugospod** 2 years, 3 months ago

No it won't get dequeued because this is not a queue.

There is no option for SUBSCRIPTION PER DRIVER, only PER RESTAURANT or a SINGLE subscription. So only SINGLE listener will receive the notification and until s/he accepts it, rejects it or it times out no one else will see the message. If this were a QUEUE than it would be as you described. This question sucks on so many levels - too many open questions arise to be able to give a valid response, plus, the final one, I would rather use queues for this problem than subscriptions...

upvoted 4 times

✉ **coffecold** 1 year, 6 months ago

Why ? One order is placed solely with the restaurant topic applicable. No copy. Then one subscriber locks the message, and then deletes it. Once locked it cannot be accessed by other subscribers.

upvoted 2 times

✉ **Opimer** Highly Voted 3 years, 1 month ago

Correct. If you create only one topic there won't be any filtering of messages depending on the restaurant. So every driver will have access to all the the orders. With a topic by restaurant, message from one restaurant will go to only one topic, from which the driver can choose to subscribe.

upvoted 18 times

✉ **Phenr** 3 years ago

I don't think that's correct. You can filter the messages on subscription, so the subscription can get a message from a specific restaurant even with just one topic created.

upvoted 8 times

✉ **rdemontis** 3 years ago

Exactly and creating one topic for each restaurant doesn't cover two of the four requirements:

- "Orders are sent to all available drivers in an area"
- "The first driver to accept an order removes it from the list of available orders"

This means that one restaurant can be server by more drivers. So how can you think to manage order acceptance from one driver if you duplicate the same in more topics? It's a big deal that cause a lot of overwork not necessary.

It's all simpler by managing all orders with a topic and using the filters for each subscription to allow drivers to view only the orders that compete with them. I also inform you that on a paid test platform such as <https://www.testpreptraining.com/> there is the same question and the correct answer requires only one topic.

upvoted 7 times

✉ **kwaazaar** 3 years ago

so how would the driver remove the order if he has his own subscription? The order would still be visible in other subscriptions.

upvoted 1 times

✉ **anvimi** 2 years, 11 months ago

multiple drivers can share a single restaurant subscription thus only one driver will handle an order

upvoted 3 times

✉ **mandynotmandy** 2 years, 1 month ago

exactly, you can't have each driver on their own sub because that duplicates the orders into each sub queues, meaning one guy accepts the order, another guy still sees it in their own sub queue. the subs has to be on restaurant level, and you can do that by filters on the sub

upvoted 1 times

✉ **coffecold** 1 year, 6 months ago

I agree. A driver (subscriber) can subscribe to a topic once he is near the restaurant and un-subscribe if he is at a distance. I don't get the high vote of NH.

upvoted 2 times

✉ **FeriAZ** Most Recent 2 months, 2 weeks ago

Create a single Service Bus Namespace. Create a single Service Bus Topic. Create a Service Bus subscription for each restaurant for which a driver can receive orders.

Topics are used for one-to-many communication with multiple subscribing systems. By creating a topic for each restaurant, you can efficiently route messages (orders) to drivers who have selected that they will deliver orders for those restaurants. This setup allows for orders to be sent to a specific set of drivers (those interested in deliveries for a particular restaurant).

upvoted 1 times

✉ **arunkuml** 4 months ago

Got it in the exam 14/12/23. Went with provided answer. Scored 912/1000. All questions are from ExamTopics. Case study - VanArsdel, Ltd (11 questions)

upvoted 1 times

✉ **Weam** 4 months, 3 weeks ago

Create a single Service Bus Namespace.

Create a single Service Bus Topic.

Create a Service Bus subscription for each restaurant for which a driver can receive orders.

Explanation:

we create different topics if we are having different categories of messages such as Consider an e-commerce platform that handles order notifications, inventory updates, and customer support messages. Using separate topics, such as "order-notifications," "inventory-updates," and "customer-support-messages," allows for clear separation of message categories and efficient routing to respective consumers.

However in this case, we don't have different topics, we only have one topic which is order processing and we can enable partitioning on this topic so messages can be under the same partition with the restaurant id

upvoted 2 times

✉ **AndySmith** 5 months, 2 weeks ago

On exam 3-Nov-2023. Went with most-voted answer - 932/1000.

1) single namespace

- 2) single topic
- 3) subscription for each restaurant.

upvoted 2 times

 **jaf19f** 8 months, 1 week ago

I got this question (12-Aug-2023) and I chose the given answers - 932 passed

upvoted 4 times

 **juanckar** 9 months, 2 weeks ago

This was on the exam (July 2023). Went with single/single/subscription for each restaurant. Scored 917

upvoted 4 times

 **warchoon** 1 year, 1 month ago

Correct

- 1: Namespace is about the solution
- 2: Topics are about restaurant messages
- 3: Subscriptions are about driver's orders

upvoted 2 times

 **apparaog99** 1 year, 3 months ago

Got in exam 12/31

upvoted 2 times

 **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 5 times

 **bipinboops** 1 year, 7 months ago

The provided answer is the most correct. There is a 2000 subscription limit for topics, so you would only allow 2000 drivers TOTAL for the entire app if you chose this as a solution. You need one topic per restaurant or the application cannot scale. You can absolutely have 1 topic, that will still work as intended, it's just not the best solution. Source: <https://docs.microsoft.com/en-us/azure/azure-resource-manager/management/azure-subscription-service-limits#service-bus-limits>

upvoted 3 times

 **zhongzi** 1 year, 10 months ago

It's the same question in ESI. The given answer is the same as highly voted.

upvoted 1 times

 **AZ204Cert** 2 years ago

Got this on 04/05/22

upvoted 1 times

 **iamstudying** 2 years, 1 month ago

BUDDIES, listen listen..

1. Single Service Bus Namespace
2. Topic per restaurant (10k limit of restaurants).

(A single topic would work, but requires subscription filters or another component to decide which drivers receive which restaurant orders. But this isn't specified so imo not an option)

3. Single subscription per topic. Each drivers who wants to deliver for a restaurant will periodically peek the subscription for messages (no lock, message not removed). Once they decide they want to deliver the order, they dequeue the message and away they go, other drivers who attempt to do so will not be able to.

(At first, I also thought multiple subscriptions - but this introduces another complexity of syncing the orders across all the subscriptions if someone accepts the order delivery first... how can we notify other subscriptions? we can't.)

All in all, agree the solution sucks a\$\$ but this is the best bet.

upvoted 6 times

 **coffecold** 1 year, 6 months ago

Multi subscriptions (drivers) per topic. How can we notify other subscriptions?> the message is locked.

<https://learn.microsoft.com/en-us/rest/api/servicebus/peek-lock-message-non-destructive-read>

upvoted 1 times

 **petibilly** 2 years, 1 month ago

Got it in exam 03/22

upvoted 1 times

 **leonidn** 2 years, 3 months ago

Single topic is less complex and more flexible solution from the client applications which submit orders. A single topic is less complex from the perspective of infrastructure management effort.

Create a single Service Bus Namespace.

Create a single Service Bus Topic.

Create a Service Bus subscription for each restaurant for which a driver can receive orders.

upvoted 3 times

Question #9

HOTSPOT -

You develop a news and blog content app for Windows devices.

A notification must arrive on a user's device when there is a new article available for them to view.

You need to implement push notifications.

How should you complete the code segment? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
string notificationHubName = "contoso_hub";
string notificationHubConnection = "connection_string";
hub =
NotificationHubClient
NotificationHubClientSettings
NotificationHubJob
NotificationDetails

.
.

GetInstallation
CreateClientFromConnectionString
CreateOrUpdateInstallation
PatchInstallation

(notificationHubConnection, notificationHubName);
string windowsToastPayload =
@"<toast><visual><binding template=""ToastText01""><text id=""1"">" +
@"New item to view" + @"</text></binding></visual></toast>";
try
{
var result =
    await hub.
        v (windowsToastPayload);
SendWindowsNativeNotificationAsync
SubmitNotificationHubJobAsync
ScheduleNotificationAsync
SendAppleNativeNotificationAsync
...
}
catch (System.Exception ex)
{
...
}
...
```

Correct Answer:

Answer Area

```

string notificationHubName = "contoso_hub";
string notificationHubConnection = "connection_string";
hub =
NotificationHubClient
NotificationHubClientSettings
NotificationHubJob
NotificationDetails

.
.
.

NotificationHubClient
NotificationHubClientSettings
NotificationHubJob
NotificationDetails

GetInstallation
CreateClientFromConnectionString
CreateOrUpdateInstallation
PatchInstallation

(notificationHubConnection, notificationHubName);
string windowsToastPayload =
@"<toast><visual><binding template=""ToastText01""><text id=""1"">" +
@"New item to view" + @"/text></binding></visual></toast>";
try
{
var result =
    await hub.
SendWindowsNativeNotificationAsync
SubmitNotificationHubJobAsync
ScheduleNotificationAsync
SendAppleNativeNotificationAsync
}

.
.
.

}
catch (System.Exception ex)
{
.
.
.

}
.
.
.

Box 1: NotificationHubClient -

```

Box 2: NotificationHubClient -

Box 3: CreateClientFromConnectionString

// Initialize the Notification Hub

NotificationHubClient hub = NotificationHubClient.CreateClientFromConnectionString(listenConnString, hubName);

Box 4: SendWindowsNativeNotificationAsync

Send the push notification.

var result = await hub.SendWindowsNativeNotificationAsync(windowsToastPayload);

Reference:

<https://docs.microsoft.com/en-us/azure/notification-hubs/notification-hubs-push-notification-registration-management><https://github.com/MicrosoftDocs/azure-docs/blob/master/articles/app-service-mobile/app-service-mobile-windows-store-dotnet-get-started-push.md>

✉  **sumitg**  3 years ago

Answer is correct.

upvoted 42 times

✉  **Tom87**  3 years ago

The answer is correct, but the first reference is not very helpful and the second one doesn't even exists. Here are references of the methods used in the answer:

https://docs.microsoft.com/en-us/dotnet/api/microsoft.azure.notificationhubs.notificationhubclient.createclientfromconnectionstring?view=azure-dotnet#Microsoft_Azure_NotificationHubs_NotificationHubClient_CreateClientFromConnectionString_System_String_System_String_

https://docs.microsoft.com/en-us/dotnet/api/microsoft.azure.notificationhubs.notificationhubclient.sendwindowsnativenotificationasync?view=azure-dotnet#Microsoft_Azure_NotificationHubs_NotificationHubClient_SendWindowsNativeNotificationAsync_System_String_

upvoted 18 times

✉  **Razvan_C**  4 months, 3 weeks ago

Is this one still a valid question for AZ-204? Didn't see anything about Notification Hubs in study guide anymore.

upvoted 2 times

✉ **bp_a_user** 3 months, 3 weeks ago

....and it is only this question that seems not to appear in the study guide :)

upvoted 1 times

✉ **Sivaramakrishnan** 2 years, 1 month ago

This the GitHub link

<https://github.com/uglide/azure-content/blob/master/articles/app-service-mobile/app-service-mobile-windows-store-dotnet-get-started-push.md>

upvoted 4 times

✉ **florodevops** 2 years, 2 months ago

1- NotificationHubClient 2 - NotificationHubClient 3- CreateClientFromConnectionString 4- SendWindowsNativeNotificationAsync

upvoted 6 times

Question #10

You are developing an Azure messaging solution.

You need to ensure that the solution meets the following requirements:

- ⇒ Provide transactional support.
- ⇒ Provide duplicate detection.
- ⇒ Store the messages for an unlimited period of time.

Which two technologies will meet the requirements? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. Azure Service Bus Topic
- B. Azure Service Bus Queue
- C. Azure Storage Queue
- D. Azure Event Hub

Correct Answer: AB

The Azure Service Bus Queue and Topic has duplicate detection.

Enabling duplicate detection helps keep track of the application-controlled MessageId of all messages sent into a queue or topic during a specified time window.

Incorrect Answers:

C: There is just no mechanism that can query a Storage queue and find out if a message with the same contents is already there or was there before.

D: Azure Event Hub does not have duplicate detection

Reference:

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/duplicate-detection>

Community vote distribution

AB (80%) BC (20%)

✉  **matejka**  3 years, 2 months ago

Answer is correct. Queue Storage does not provide transactional support and Event Hub can't be configured to store events for infinite time.
upvoted 40 times

✉  **mvr** 3 years, 2 months ago

Event Hub can be configured to store events, with Event Hub Capture.
Makes the answer still correct though.

<https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-capture-overview>

upvoted 1 times

✉  **cbn** 3 years, 2 months ago

Event hubs (with capture) are not supposed to retain data forever.
<https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-faq#what-is-the-maximum-retention-period-for-events>
upvoted 3 times

✉  **ferut** 2 years, 11 months ago

The stress is on 'messaging' which Event Hub is not (it's eventing)
upvoted 6 times

✉  **coffecold** 1 year, 6 months ago

Agree, answer is correct, but I think duplicate detection is the key phrase.
upvoted 3 times

✉  **dddddd111**  6 months ago

Selected Answer: AB

Given answers are correct. ChatGPT apologized for A and D after I asked to justify the answer.

upvoted 2 times

✉  **mkahmann** 7 months, 1 week ago

Selected Answer: AB

Agreed

upvoted 1 times

OPT_001122 1 year, 4 months ago

Selected Answer: AB

- A. Azure Service Bus Topic
- B. Azure Service Bus Queue

upvoted 1 times

RochaG2 1 year, 6 months ago

Selected Answer: BC

correct

upvoted 2 times

RochaG2 1 year, 6 months ago

Sorry I meant AB!!

upvoted 2 times

[Removed] 2 years, 2 months ago

Does C not become an option because of unlimited storing?
Because < 80gb queue storage can only be used.

upvoted 3 times

[Removed] 2 years, 2 months ago

* > 80gb (more than 80gb)

upvoted 1 times

Knightie 1 year, 5 months ago

it's unlimited period of time, not storage size.

upvoted 1 times

Lucario95 2 years, 3 months ago

Selected Answer: AB

Answer is correct

upvoted 4 times

Question #11

DRAG DROP -

You develop a gateway solution for a public facing news API.

The news API back end is implemented as a RESTful service and hosted in an Azure App Service instance.

You need to configure back-end authentication for the API Management service instance.

Which target and gateway credential type should you use? To answer, drag the appropriate values to the correct parameters. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Values	Answer Area
Azure Resource	Configuration parameter
HTTP(s) endpoint	Target
Basic	Gateway credentials
Client cert	

Correct Answer:

Values	Answer Area
	Configuration parameter
HTTP(s) endpoint	Target
Basic	Gateway credentials

Box 1: Azure Resource -

Box 2: Client cert -

API Management allows to secure access to the back-end service of an API using client certificates.

Reference:

<https://docs.microsoft.com/en-us/rest/api/apimanagement/apimanagementrest/azure-api-management-rest-api-backend-entity>

✉  **Kuna_Lambo**  3 years, 1 month ago

#1 Http ?

upvoted 43 times

✉  **AzureAz204Fan** 2 years, 11 months ago

Yes. Target has only two options: 1) Azure Logic App 2) HTTP(s) endpoint. In the question it is given that backend is a REST API which means Target has to be HTTP

upvoted 9 times

✉  **ning** 2 years, 7 months ago

On the API level there are two targets, Logic App and HTTPS endpoint, on the operation level two targets are Azure Resource and HTTPS endpoint. So sure what level is this question asking, if API level then HTTPS, if operation level that will be azure resource, since it is hosted in Azure.

upvoted 6 times

✉  **tevivi8222** 3 years, 1 month ago

You might be right: <https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-mutual-certificates#configure-an-api-to-use-client-certificate-for-gateway-authentication>

upvoted 21 times

✉ **Molte** 2 years, 3 months ago

but why cert and not basic?
upvoted 4 times

✉ **Knightie** 1 year, 7 months ago

Public Facing.. so not all incoming request will have their known credentials. So just cert to ensure the caller is the rightful app not the user right.
upvoted 1 times

✉ **adilkhan** 1 year, 2 months ago

configure back-end authentication and to do that use cert
upvoted 2 times

✉ **Laaptu** 2 years, 9 months ago

agreed
1. Https(s) endpoint,
2. Client cert
upvoted 17 times

✉ **Gregoryhouse2020** Highly Voted 2 years, 10 months ago

Answer :

1. Https(s) endpoint,
2. Client cert

Reference : <https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-mutual-certificates#configure-an-api-to-use-client-certificate-for-gateway-authentication>

upvoted 33 times

✉ **unraval** Most Recent 7 months ago

Reference for Target has only two options: 1) Azure Logic App 2) HTTP(s) endpoint ??

upvoted 2 times

✉ **applepie** 8 months, 3 weeks ago

got this question today - 7/30/2023
upvoted 3 times

✉ **adilkhan** 1 year, 2 months ago

Https and client certificate. 100%

<https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-mutual-certificates#configure-an-api-to-use-client-certificate-for-gateway-authentication>

upvoted 3 times

✉ **OPT_001122** 1 year, 4 months ago

1. Https(s) endpoint,
2. Client cert
upvoted 6 times

✉ **POOOJAAAAAA** 1 year, 3 months ago

am always searching for this opt_001122 person's comment to know the correct answer.. hope he/she's answers are correct.
upvoted 4 times

✉ **Ciupaz** 10 months, 2 weeks ago

Me too.

upvoted 1 times

✉ **gmishra88** 1 year, 6 months ago

Both Azure resource and Http(s) endpoint can be used because it is an azure resource Rest api. Wonder what Microsoft expects.
I'm clueless seeing Target and Gateway. No idea what they are till I saw the portal. Do I have to remember the unfortunate name the Microsoft guys chose for these UI elements?

upvoted 2 times

✉ **Pize** 1 year, 9 months ago

1.Http
2. Basic or Client cert would fit
upvoted 1 times

✉ **Eltooth** 1 year, 9 months ago

HTTPs) endpoint
Client cert
upvoted 3 times

✉ **Sandeep12093** 1 year, 10 months ago

It's definitely since target has only options https and logic app
1. Https(s) endpoint,
2. Client cert
upvoted 3 times

✉ **Azprep** 2 years ago

1. Http End point
2. Client certificate

<https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-mutual-certificates#configure-an-api-to-use-client-certificate-for-gateway-authentication>

upvoted 1 times

✉ **[Removed]** 2 years, 1 month ago

1. Https(s) endpoint,
2. Client cert

<https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-mutual-certificates>

upvoted 1 times

✉ **lugospod** 2 years, 3 months ago

Got this one 01/2022. Went with http and certificate. but from my perspective basic should also fit.

upvoted 7 times

✉ **debanjan10** 2 years, 6 months ago

Target: HTTP endpoints
Gateway creds: Client cert

upvoted 3 times

✉ **ning** 2 years, 8 months ago

This question is valid? It is so poorly worded. Not sure what is this asking for?

I guess for the first on HTTP(S), for the second one, either basic or certificate is accepted by azure. Nothing is mentioned in the question which one is preferred.

upvoted 8 times

✉ **cool_tool** 2 years, 8 months ago

1. HTTPS
2. CLIENT CERT

<https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-mutual-certificates#configure-an-api-to-use-client-certificate-for-gateway-authentication>

upvoted 5 times

✉ **kondapaturi** 2 years, 9 months ago

1. Https(s) endpoint,
2. Client cert

upvoted 5 times

Question #12

HOTSPOT -

You are creating an app that uses Event Grid to connect with other services. Your app's event data will be sent to a serverless function that checks compliance.

This function is maintained by your company.

You write a new event subscription at the scope of your resource. The event must be invalidated after a specific period of time.

You need to configure Event Grid.

What should you do? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Option	Value
WebHook event delivery	<ul style="list-style-type: none"> SAS tokens Key authentication Management Access Control
Topic publishing	<ul style="list-style-type: none"> ValidationCode handshake ValidationURL handshake JWT token

Answer Area

Option	Value
WebHook event delivery	<ul style="list-style-type: none"> SAS tokens Key authentication Management Access Control
Topic publishing	<ul style="list-style-type: none"> ValidationCode handshake ValidationURL handshake JWT token

Box 1: SAS tokens -

Custom topics use either Shared Access Signature (SAS) or key authentication. Microsoft recommends SAS, but key authentication provides simple programming, and is compatible with many existing webhook publishers.

In this case we need the expiration time provided by SAS tokens.

Box 2: ValidationCode handshake -

Event Grid supports two ways of validating the subscription: ValidationCode handshake (programmatic) and ValidationURL handshake (manual).

If you control the source code for your endpoint, this method is recommended.

Incorrect Answers:

ValidationURL handshake (manual): In certain cases, you can't access the source code of the endpoint to implement the ValidationCode handshake. For example, if you use a third-party service (like Zapier or IFTTT), you can't programmatically respond with the validation code.

Reference:

<https://docs.microsoft.com/en-us/azure/event-grid/security-authentication>

 **jay158**  2 years, 10 months ago

Dropdowns placement in Answer area is incorrect
WebHook Event Delivery ---ValidationCode handsShake
Topic publishing --- SAS Tokens

<https://docs.microsoft.com/en-us/azure/event-grid/concepts>
<https://docs.microsoft.com/en-us/azure/event-grid/webhook-event-delivery>
upvoted 48 times

✉️ **Alex_Wackoo** 2 years, 10 months ago

Agreed and your last link is not working so here:
<https://docs.microsoft.com/en-us/azure/event-grid/webhook-event-delivery>
upvoted 7 times

✉️ **applepie** 8 months, 3 weeks ago

got this question today, go with this answer - 7/30/2023, score 895/1000
upvoted 7 times

✉️ **TonyMel** 1 year ago

correct, in 2023Mar24, score: 904/1000.
upvoted 4 times

✉️ **NightshadeRC** Highly Voted 8 months, 3 weeks ago

Had this question in today's exam: 2023-07-26.
I think the drop-downs were in the order they are here, not 100% on it though. I went with the recommended answer anyway. Passed 940.
upvoted 5 times

✉️ **warchoon** Most Recent 1 year, 1 month ago

"The event must be invalidated after a specific period of time." - Asynchronous handshake => validationUrl
<https://learn.microsoft.com/en-us/azure/event-grid/webhook-event-delivery#:~:text=Asynchronous%20handshake>
upvoted 2 times

✉️ **18Marks** 1 year, 4 months ago

Got this in exam. Use given answer
upvoted 3 times

✉️ **jvyas** 2 years, 6 months ago

Below link for Authenticating clients that publish events to Azure Event Grid topics using access key or Shared Access Signature (SAS) token.
<https://docs.microsoft.com/en-us/azure/event-grid/authenticate-with-access-keys-shared-access-signatures>
upvoted 2 times

✉️ **warchoon** 1 year, 1 month ago

So what's the answer? Keys or tokens?
Keys are from SAS too.
SAS token can be expired after "specific period of time".
upvoted 1 times

✉️ **Simbetie** 2 years, 8 months ago

Topic publishing to be matched with the 1st dropdown....the answer must be Management Access Control. Azure Event Grid allows you to control the level of access given to different users to do various management operations such as list event subscriptions, create new ones, and generate keys. Event Grid uses Azure's role-based access control (RBAC).
upvoted 2 times

✉️ **ning** 2 years, 8 months ago

This question is worded poorly, the handshake only happens when subscription is setup ...

If you are talking about how subscription should be authenticated when receiving messages, then webhook should use jwt
upvoted 2 times

✉️ **ning** 2 years, 7 months ago

For Topic Publishing <https://docs.microsoft.com/en-us/azure/event-grid/authentication-overview>
upvoted 3 times

✉️ **ning** 2 years, 7 months ago

For web hook delivery <https://docs.microsoft.com/en-us/azure/event-grid/secure-webhook-delivery>
upvoted 2 times

Question #13

HOTSPOT -

You are working for Contoso, Ltd.

You define an API Policy object by using the following XML markup:

```
<set-variable name="bodySize" value="@{context.Request.Headers["Content-Length"] [0]}"/>
<choose>
  <when condition="@{int.Parse(context.Variables.GetValueOrDefault<string> ("bodySize"))<512000}">
  </when>
  <otherwise>
    <rewrite-uri template="/put"/>
    <set-backend-service base-url="http://contoso.com/api/9.1/" />
  </otherwise>
</choose>
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statement	Yes	No
The XML segment belongs in the <inbound> section of the policy.	<input type="radio"/>	<input type="radio"/>
If the body size is >256k, an error will occur.	<input type="radio"/>	<input type="radio"/>
If the request is http://contoso.com/api/9.2/, the policy will retain the higher version.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statement	Yes	No
The XML segment belongs in the <inbound> section of the policy.	<input checked="" type="radio"/>	<input type="radio"/>
If the body size is >256k, an error will occur.	<input type="radio"/>	<input checked="" type="radio"/>
If the request is http://contoso.com/api/9.2/, the policy will retain the higher version.	<input type="radio"/>	<input checked="" type="radio"/>

Box 1: Yes -

Use the set-backend-service policy to redirect an incoming request to a different backend than the one specified in the API settings for that operation. Syntax:

```
<set-backend-service base-url="base URL of the backend service" />
```

Box 2: No -

The condition is on 512k, not on 256k.

Box 3: No -

The set-backend-service policy changes the backend service base URL of the incoming request to the one specified in the policy.

Reference:

<https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies>

  paru123456789  3 years, 1 month ago

Answer:

Yes

No

No
upvoted 72 times

✉ **clarionprogrammer** 3 years ago

This is correct.

upvoted 4 times

✉ **KingChuang** 1 year, 3 months ago

On my exam 2022-12-26.

Chose: Y/N/N

upvoted 5 times

✉ **TonyMel** 1 year ago

correct, in 2023Mar24, score: 904/1000.

upvoted 3 times

✉ **Cornholioz** **Highly Voted** 3 years, 3 months ago

It's doing nothing in the When condition. So if msg is <512, it does nothing. The otherwise gets executed only for >512 ??? So it will set backend service only for msgs >512? Question doesn't talk about this case.

Also, like someone said in the other discussion for this question:

technically >256k can be >512k

Poorly framed question!!!

upvoted 14 times

✉ **Knightie** 1 year, 7 months ago

not framed, just tricked.. if body > 512k wther you use 1024 base of 1000 base, nevermind, it just change the base URL and use /put instead.. it does not give an error, that's what otherwise for.

upvoted 3 times

✉ **MiraA** 2 years, 6 months ago

I think that no action in <when> means nothing needs to be changed (for requests with body < 512000) so the back-end service URL remains as specified in the policy.

And only large requests (> 512000) will be redirected to API 9.1 using <set-backend-service> in <otherwise> branch.

<https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies#SetBackendService>

upvoted 8 times

✉ **NightshadeRC** **Most Recent** 8 months, 3 weeks ago

Had this question in today's exam: 2023-07-26

upvoted 5 times

✉ **AlexeyG** 1 year, 2 months ago

Got this in 16/02/2023

upvoted 3 times

✉ **edengoforit** 2 years, 2 months ago

I think the first question should be backend instead of inbound?

Backend service: <!-- statements to be applied before the request is forwarded to the backend service go here -->

upvoted 3 times

✉ **gabavo** 2 years ago

rewrite-uri policy can appear only in inbound scope, so the answer is correct

<https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies#usage-8>

upvoted 4 times

✉ **mattvasc** 2 years, 1 month ago

Makes sense and if we check the set-var docs on Microsoft we can see that it can be used under inbound, outbound or backend scope:

<https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-mutual-certificates#configure-an-api-to-use-client-certificate-for-gateway-authentication>

upvoted 1 times

✉ **lugospod** 2 years, 3 months ago

Got this one 01/2022. Went with yes no no

upvoted 5 times

✉ **Pomphard** 3 years ago

For statements 2 & 3, I think they try to emphasize whether something will happen or may happen

>256K may be larger than 512K, so it may happen but there's no guarantee as the size might be in between these values

The base URL may be retained if the condition is met, but there's no guarantee as it might be rewritten

So I'll go with yes, no, no

upvoted 3 times

✉ **Araneus** 3 years, 4 months ago

I would say "it depends" for the third question. Since the set-backend-service policy is in the otherwise block, it would only rewrite the URL for requests with a body size of ≥ 512000 . Or am I missing a clue here?

upvoted 9 times

✉ **Cornholioz** 3 years, 4 months ago

I think Araneus is right. Easy to miss. Looks like a basic WHEN / OTHERWISE condition to me. If bodysize limit is not hit, it will never reach the OTHERWISE scope and hence will not retain a higher version. I think.

upvoted 4 times

✉ **Zidimirite** 3 years ago

You mean will retain a higher version, right? Since the OTHERWISE scope doesn't get executed hence the set-backend-service isn't set to 9.1? 9.1 is the lower version.

upvoted 6 times

✉ **iiiihhhh** 3 years, 4 months ago

There is another discussion for the question: <https://www.examtopics.com/discussions/microsoft/view/23300-exam-az-204-topic-5-question-8-discussion/>

upvoted 3 times

✉ **ahadjithoma** 3 years, 4 months ago

Use the set-backend-service policy to redirect an incoming request to a different backend than the one specified in the API settings for that operation. This policy changes the backend service base URL of the incoming request to the one specified in the policy.

Ref: <https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies>

So it looks like why don't even care if any of the WHEN / OTHERWISE conditions are met, since the 9.2 is not specified anywhere in the given code snippet

upvoted 4 times

✉ **deborahQ** 1 year, 5 months ago

to my understanding, if the body size is less than 512K, it could use higher version (if it has and requested, which is not mentioned here).

upvoted 1 times

Question #14

You are developing a solution that will use Azure messaging services.

You need to ensure that the solution uses a publish-subscribe model and eliminates the need for constant polling.

What are two possible ways to achieve the goal? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. Service Bus
- B. Event Hub
- C. Event Grid
- D. Queue

Correct Answer: AC

It is strongly recommended to use available messaging products and services that support a publish-subscribe model, rather than building your own. In Azure, consider using Service Bus or Event Grid. Other technologies that can be used for pub/sub messaging include Redis, RabbitMQ, and Apache Kafka.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/patterns/publisher-subscriber>

Community vote distribution

AC (100%)

 **Amankothari**  3 years, 9 months ago

Correct Ans is A & C

<https://docs.microsoft.com/en-us/azure/architecture/patterns/publisher-subscriber#issues-and-considerations>

upvoted 52 times

 **rongcheng** 3 years, 8 months ago

"In Azure, consider using Service Bus or Event Grid. "

upvoted 9 times

 **coffecold** 1 year, 6 months ago

Agree, but publish/subscribe is a feature of event-hub as well, but event-based receiving is not. That is why option B falls off.

upvoted 3 times

 **lion2k6**  3 years, 2 months ago

The article has been updated, the new text says "In Azure, consider using Service Bus, Event Hubs or Event Grid."

So A, B and C should be all correct.

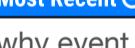
<https://docs.microsoft.com/en-us/azure/architecture/patterns/publisher-subscriber#issues-and-considerations>

upvoted 49 times

 **coffecold** 1 year, 6 months ago

Event Hubs : event-based receiving is not possible, you still need to poll. So B falls off

upvoted 1 times

 **raymond_abcd**  2 months, 2 weeks ago

Still no convinced why event hub is not a correct answer. Some say that the event hub consumer is polling? I read the following in the doc:

Any entity that reads event data from an event hub is an event consumer. All Event Hubs consumers connect via the AMQP 1.0 session and events are delivered through the session as they become available. The client doesn't need to poll for data availability.

upvoted 1 times

 **AndySmith** 5 months, 2 weeks ago

On exam 3-Nov-2023. Went with proposed answer - 932/1000.

upvoted 1 times

 **juanckar** 9 months, 2 weeks ago

This was on the exam (July 2023). Went with AC. Scored 917

upvoted 3 times

 **OPT_001122** 1 year, 5 months ago

Selected Answer: AC

A & C Correct

upvoted 3 times

✉  **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 5 times

✉  **AZAdam22** 1 year, 10 months ago

Selected Answer: AC

A and C as both service bus and event grid implement publish-subscribe and don't require polling.

upvoted 3 times

✉  **renzoku** 1 year, 9 months ago

exactly, Event HUB needs to poll and process an event

upvoted 1 times

✉  **liuliangzhou** 1 year, 10 months ago

A, B and C should be all correct.

It is strongly recommended to use available messaging products and services that support a publish-subscribe model, rather than building your own. In Azure, consider using Service Bus, Event Hubs or Event Grid. Other technologies that can be used for pub/sub messaging include Redis, RabbitMQ, and Apache Kafka.

<https://docs.microsoft.com/en-us/azure/architecture/patterns/publisher-subscriber>

upvoted 2 times

✉  **meoukg** 2 years, 1 month ago

Got it on 03/2022, I chose A. Service Bus & C. Event Grid

upvoted 3 times

✉  **heisenberg33** 2 years, 2 months ago

Selected Answer: AC

Correct answer is A and C Ref: <https://docs.microsoft.com/en-us/azure/architecture/patterns/publisher-subscriber#issues-and-considerations>.

Wrong answer B (eliminate the need for constant polling)Ref: <https://docs.microsoft.com/en-us/azure/connectors/connectors-create-api-azure-event-hubs#trigger-polling-behavior>

upvoted 8 times

✉  **wk__** 2 years, 2 months ago

It seems that the selected answers are correct: "Existing technologies. It is strongly recommended to use available messaging products and services that support a publish-subscribe model, rather than building your own. In Azure, consider using Service Bus, Event Hubs or Event Grid. Other technologies that can be used for pub/sub messaging include Redis, RabbitMQ, and Apache Kafka."

<https://docs.microsoft.com/en-us/azure/architecture/patterns/publisher-subscriber>

upvoted 2 times

✉  **lugospod** 2 years, 3 months ago

Got this one 01/2022. Went with most voted (to avoid writing answers again)

upvoted 3 times

✉  **kondapaturi** 2 years, 9 months ago

Service bus

Event Grid

upvoted 3 times

✉  **mprokopy** 2 years, 11 months ago

A & C

Triggers of Event hub are long-polling, so B is not correct.

<https://docs.microsoft.com/en-us/azure/connectors/connectors-create-api-azure-event-hubs#trigger-polling-behavior>

upvoted 6 times

✉  **Frakandel** 2 years, 11 months ago

I believe MS is rather picky about the distinction between events and messages, that why I go for Service Bus and Queues... However Event Hubs and Grid are also capable of delivering the requested functionality (as far as I understand them... but I'm still studying on them :-)

upvoted 5 times

✉  **ferut** 2 years, 11 months ago

Agree, see:

<https://docs.microsoft.com/en-us/azure/event-grid/compare-messaging-services?toc=https%3A%2F%2Fdocs.microsoft.com%2Fen-us%2Fazure%2Fservice-bus-messaging%2Ftoc.json&bc=https%3A%2F%2Fdocs.microsoft.com%2Fen-us%2Fazure%2Fbread%2Ftoc.json>

upvoted 2 times

✉  **iamstudying** 2 years, 1 month ago

Buddies, although there is a distinction between events and messages, all of these services (Service Bus Topics, Service Bus Queues, Event Grid, Event Hubs) are under the "Azure Messaging" umbrella.

Queues do not satisfy the pub-sub model and Event Hub is long polling <https://docs.microsoft.com/en-us/azure/connectors/connectors-create-api-azure-event-hubs#trigger-polling-behavior>

upvoted 2 times

✉  **wtkwsk** 3 years ago

A, B and C are all correct:

Existing technologies. It is strongly recommended to use available messaging products and services that support a publish-subscribe model, rather than building your own. In Azure, consider using Service Bus, Event Hubs or Event Grid. Other technologies that can be used for pub/sub messaging include Redis, RabbitMQ, and Apache Kafka.

(<https://docs.microsoft.com/en-us/azure/architecture/patterns/publisher-subscriber>)

upvoted 4 times

Question #15

A company is implementing a publish-subscribe (Pub/Sub) messaging component by using Azure Service Bus. You are developing the first subscription application.

In the Azure portal you see that messages are being sent to the subscription for each topic. You create and initialize a subscription client object by supplying the correct details, but the subscription application is still not consuming the messages.

You need to ensure that the subscription client processes all messages.

Which code segment should you use?

- A. await subscriptionClient.AddRuleAsync(new RuleDescription(RuleDescription.DefaultRuleName, new TrueFilter()));
- B. subscriptionClient = new SubscriptionClient(ServiceBusConnectionString, TopicName, SubscriptionName);
- C. await subscriptionClient.CloseAsync();
- D. subscriptionClient.RegisterMessageHandler(ProcessMessagesAsync, messageHandlerOptions);

Correct Answer: D

Using topic client, call RegisterMessageHandler which is used to receive messages continuously from the entity. It registers a message handler and begins a new thread to receive messages. This handler is waited on every time a new message is received by the receiver.

```
subscriptionClient.RegisterMessageHandler(ReceiveMessagesAsync, messageHandlerOptions);
```

Reference:

<https://www.c-sharpcorner.com/article/azure-service-bus-topic-and-subscription-pub-sub/>

Community vote distribution

D (100%)

✉  **Ritesh073**  3 years, 7 months ago

Agreed on D

Clientsubscriber object initialization is done, so B incorrect.

Here Nothing to do with Rule and closeclient so A and C incorrect

upvoted 63 times

✉  **KingChuang** 1 year, 3 months ago

On my exam 2022-12-26.

Chose: D

upvoted 3 times

✉  **paru123456789**  3 years, 1 month ago

Answer: D

upvoted 7 times

✉  **Vmwarevirtual**  10 months, 3 weeks ago

Appeared in my exam on 27 May 2023.

I chose 'D'

upvoted 4 times

✉  **AlexeyG** 1 year, 2 months ago

Got this in 16/02/2023

upvoted 2 times

✉  **OPT_001122** 1 year, 5 months ago

Selected Answer: D

D is correct

upvoted 4 times

✉  **coffecold** 1 year, 6 months ago

Old library : Microsoft.Azure.ServiceBus v5.2.0

I think nowadays (Azure.Messaging.ServiceBus v7.11.0) you simply would use :

```
processor = client.CreateProcessor(queueName, new ServiceBusProcessorOptions());
```

```
processor.ProcessMessageAsync += MessageHandler;
```

Still : the anwer is correct

upvoted 5 times

✉  **Perplex** 1 year, 9 months ago

Selected Answer: D

D is correct

upvoted 4 times

meoukg 2 years, 1 month ago

Got it on 03/2022, I chose D. `subscriptionClient.RegisterMessageHandler(ProcessMessagesAsync, messageHandlerOptions);`
upvoted 4 times

vulht 2 years, 1 month ago

Agreed on D when we use `Microsoft.Azure.ServiceBus` package.

Furthermore, with the new version (`Azure.Messaging.ServiceBus`), we use a processor to receive/handle messages.

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-dotnet-get-started-with-queues#add-the-code-to-receive-messages-from-the-queue>

upvoted 3 times

Question #16

Topic 6

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen. You are developing an Azure Service application that processes queue data when it receives a message from a mobile application. Messages may not be sent to the service consistently.

You have the following requirements:

- ⇒ Queue size must not grow larger than 80 gigabytes (GB).
- ⇒ Use first-in-first-out (FIFO) ordering of messages.
- ⇒ Minimize Azure costs.

You need to implement the messaging solution.

Solution: Use the .Net API to add a message to an Azure Storage Queue from the mobile application. Create an Azure VM that is triggered from Azure Storage

Queue events.

Does the solution meet the goal?

A. Yes

B. No

Correct Answer: B

Don't use a VM, instead create an Azure Function App that uses an Azure Service Bus Queue trigger.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-functions/functions-create-storage-queue-triggered-function>

Community vote distribution

B (100%)

✉  **Lkk51**  3 years, 10 months ago

the reason that the solution is service bus because of this message "Use first-in-first-out (FIFO) ordering of messages."

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted#foundational-capabilities>

Storage account does not have ordering guarantee

upvoted 54 times

✉  **edengoforit** 2 years, 2 months ago

Absolutely correct. There are so many questions related to this. FIFO and business transactions. If anything related, it seems to be the Service Bus

upvoted 2 times

✉  **xRiot007** 3 years, 9 months ago

I also think the explanation should be corrected. Size <80GB + FIFO guarantee is what makes Service Bus a good option, not the existence of a VM or not.

upvoted 19 times

✉  **lautaro_sw** 3 years, 8 months ago

I think the reason for not using a VM is the 'Minimize Azure costs' requirement

upvoted 17 times

✉  **rdemontis** 3 years ago

I totally agree with you all. Only a doubt for VM usage: In addition to save costs I think it's not possible to interact with a service bus queue. in than case i believe you have to use a storage queue.

upvoted 1 times

✉  **Nabeelcp**  3 years, 7 months ago

Answer is correct . but explanation is confusing ..

upvoted 7 times

✉  **Codenob0b** 3 years, 7 months ago

It is confusing indeed. But the reason the explanation pointed out about VMs is because of VMs are expensive and won't match with the "Minimize azure costs" requirement. Azure Functions is the best way to handle Service Bus Queue events by utilizing its trigger functionality and providing minimal costs to it.

upvoted 8 times

✉ **PoTermin** 2 years, 8 months ago

Your explanation sounds good.

upvoted 2 times

✉ **applepie** **Most Recent** 8 months, 3 weeks ago

got this question today, answer N - 7/30/2023, score 895/1000

upvoted 2 times

✉ **juanckar** 9 months, 2 weeks ago

This was on the exam (July 2023). Went with NO. Scored 917

upvoted 3 times

✉ **OPT_001122** 1 year, 4 months ago

Selected Answer: B

B. No - is the correct

upvoted 1 times

✉ **elequiel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 2 times

✉ **Eltooth** 1 year, 9 months ago

Selected Answer: B

B is correct answer.

NO

upvoted 2 times

✉ **Udoyen** 2 years, 8 months ago

This is the kind of question we need ...

upvoted 4 times

✉ **paru123456789** 3 years, 1 month ago

Answer: NO

upvoted 1 times

✉ **svaza** 3 years, 1 month ago

I think the given option of using Azure Storage queue is right, refer this document

<https://docs.microsoft.com/en-us/azure/storage/queues/scalability-targets#scale-targets-for-queue-storage>

- You can store upto 500TB of data in queue, queue works in FIFO.

The approach of using VM for processing the queue messages is overkill. So given question the answer is B.
Still its a stilly question

upvoted 1 times

✉ **Tom87** 3 years ago

Queues in general should be FIFO. But not Azure Storage Queue. Ordering is not guaranteed. See this:

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted#foundational-capabilities>

upvoted 2 times

✉ **_nma** 3 years, 4 months ago

Answer: A. Yes

1.Your solution requires the queue to provide a guaranteed first-in-first-out (FIFO) ordered delivery.

2.Your queue size won't grow larger than 80 GB.

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted#consider-using-service-bus-queues>

upvoted 1 times

✉ **_nma** 3 years, 4 months ago

Apologies, wrong discussion. This was meant for the net question.

upvoted 1 times

Question #17

Topic 6

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen. You are developing an Azure Service application that processes queue data when it receives a message from a mobile application. Messages may not be sent to the service consistently.

You have the following requirements:

- ⇒ Queue size must not grow larger than 80 gigabytes (GB).
- ⇒ Use first-in-first-out (FIFO) ordering of messages.
- ⇒ Minimize Azure costs.

You need to implement the messaging solution.

Solution: Use the .Net API to add a message to an Azure Service Bus Queue from the mobile application. Create an Azure Windows VM that is triggered from

Azure Service Bus Queue.

Does the solution meet the goal?

A. Yes

B. No

Correct Answer: B

Don't use a VM, instead create an Azure Function App that uses an Azure Service Bus Queue trigger.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-functions/functions-create-storage-queue-triggered-function>

Community vote distribution

B (100%)

✉  **cbn**  3 years, 2 months ago

Using Service bus is fine, however having a Windows VM does not address the cost requirement.

Answer is 'NO'

upvoted 44 times

✉  **Udoyen** 2 years, 8 months ago

So if I may ask, if the cost wasn't an issue would Windows VM be okay?

upvoted 6 times

✉  **dohnalr** 2 years, 5 months ago

Yes. It would be ok

upvoted 3 times

✉  **coffecold** 1 year, 6 months ago

I think every solution would be OK then.

upvoted 2 times

✉  **neosri**  3 years, 11 months ago

answer explanation is wrong. It should be azure service bus queues and not storage queues trigger for Azure functions

upvoted 25 times

✉  **applepie**  8 months, 3 weeks ago

got this question today, answer N - 7/30/2023, score 895/1000

upvoted 2 times

✉  **OPT_001122** 1 year, 4 months ago

Selected Answer: B

Azure Service Bus Queue trigger

upvoted 1 times

✉  **hubekpeter** 1 year, 4 months ago

Selected Answer: B

Queue size must not grow larger than 80 gigabytes (GB) - yes that's fine

Use first-in-first-out (FIFO) ordering of messages. - yes, you can get it with service bus

Minimize Azure costs & Create an Azure Windows VM that is triggered from Azure Service Bus Queue - Firstly there's nothing like an AzureVM

triggering, but Azure Functions triggering instead. Secondly - using vm with queue is more expensive, but of course it depends on multiple factors.

upvoted 3 times

✉ **Eltooth** 1 year, 9 months ago

Selected Answer: B

B is correct answer.

NO

upvoted 3 times

✉ **Dev666** 2 years ago

Selected Answer: B

I choose B as think A will be expensive

upvoted 1 times

✉ **meoukg** 2 years, 1 month ago

Got it on 03/2022, I chose B. No

upvoted 1 times

✉ **altafpatel1984** 2 years, 4 months ago

Here I guess following line meant to mention that VM will unnecessarily increase cost.

"Messages may not be sent to the service consistently."

upvoted 2 times

✉ **altafpatel1984** 2 years, 4 months ago

If B is answer then question is confusing because theriotically VM is not bad idea since there is no mentioned of cost saving etc.

upvoted 1 times

✉ **grappigekerkel** 2 years, 2 months ago

Requirement clearly says minimize Azure costs

upvoted 2 times

✉ **Idkhow** 2 years, 10 months ago

why do we need the Azure Windows VM for? Service bus is already fine

upvoted 3 times

✉ **silva_831** 2 years, 10 months ago

The given answer is incorrect. Service Bus queue meets the requirements

upvoted 1 times

✉ **Cholo981** 2 years, 10 months ago

StorageQueue doens't guarantees FIFO. Must be ServiceBus.

upvoted 1 times

✉ **paru123456789** 3 years, 1 month ago

Answer: NO

upvoted 2 times

✉ **_nma** 3 years, 4 months ago

Answer: A. Yes

1. Your solution requires the queue to provide a guaranteed first-in-first-out (FIFO) ordered delivery.

2. Your queue size won't grow larger than 80 GB.

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted#consider-using-service-bus-queues>

upvoted 2 times

✉ **Cornholioz** 3 years, 3 months ago

Wrong because of the VM use in the solution which is an overkill and increases cost.

upvoted 10 times

✉ **javayung** 3 years, 5 months ago

Answer is No:

As a solution architect/developer, you should consider using Storage queues when:

Your application must store over 80 GB of messages in a queue.

upvoted 2 times

✉ **Dirk** 3 years, 8 months ago

Yes - I think explanation is wrong, but answer is correct.

See this link: <https://www.serverless360.com/blog/azure-storage-queue-vs-service-bus-queue#:~:text=Message%20Size,incoming%20messages%20throwing%20an%20exception>.

upvoted 3 times

✉ **cloud_exam1** 3 years, 5 months ago

I think the answer is wrong. In the requirement, said 'Queue size must not grow larger than 80 gigabytes (GB)'. And FIFO. So should be Service Bus Queue.

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted>
upvoted 1 times

 **Cornholioz** 3 years, 4 months ago

In the famous words of Gandalf, "cloud_exam1, you shall not pass"!

upvoted 7 times

Question #18

DRAG DROP -

You are developing a REST web service. Customers will access the service by using an Azure API Management instance.

The web service does not correctly handle conflicts. Instead of returning an HTTP status code of 409, the service returns a status code of 500.

The body of the status message contains only the word conflict.

You need to ensure that conflicts produce the correct response.

How should you complete the policy? To answer, drag the appropriate code segments to the correct locations. Each code segment may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Policy segments

server
context
on-error
set-status
when-error
override-status

Answer Area

```
< Policy segment >
<base />
<choose>
  <when condition = " @ Policy segment .Response.StatusCode == 500
    && Policy segment .LastError.Message.Contains
      <return-response>
        < Policy segment >
      </return-response>
    </when>
    <otherwise />
  </choose>
< Policy segment >
```

Correct Answer:

Policy segments

server
context
on-error
set-status
when-error
override-status

Answer Area

```
< on-error >
<base />
<choose>
  <when condition = " @ context .Response.StatusCode == 500
    && context .LastError.Message.Contains
      <return-response>
        < set-status >
      </return-response>
    </when>
    <otherwise />
  </choose>
< on-error >
```

Box 1: on-error -

Policies in Azure API Management are divided into inbound, backend, outbound, and on-error.

If there is no on-error section, callers will receive 400 or 500 HTTP response messages if an error condition occurs.

Box 2: context -

Box 3: context -

Box 4: set-status -

The return-response policy aborts pipeline execution and returns either a default or custom response to the caller. Default response is 200 OK with no body.

Custom response can be specified via a context variable or policy statements.

Syntax:

```
<return-response response-variable-name="existing context variable">
<set-header/>
```

```
<set-body/>
<set-status/>
</return-response>
```

Box 5: on-error -

Reference:

<https://docs.microsoft.com/en-us/azure/api-management/api-management-error-handling-policies> <https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies>

✉  **jokergester**  3 years ago

Just want to add that "set-status" also has code and reason (e.g. set-status code="409" reason="Conflict").

<https://docs.microsoft.com/en-us/azure/api-management/api-management-advanced-policies#SetStatus>

upvoted 39 times

✉  **SnakePlissken** 2 years, 11 months ago

Thanks. set-status on its own seemed a bit strange.
Minor detail: Last on-error must be preceded by a slash.

upvoted 9 times

✉  **armoniyem**  2 years, 5 months ago

THE ANSWER IS CORRECT.

upvoted 15 times

✉  **Maskoo**  2 months, 3 weeks ago

Given answer is correct

But tbh as a developer in this situation I would fix the API to return a 409 in case of conflicts instead of tweaking it using a Azure Policy...

upvoted 2 times

✉  **jaf19f** 8 months, 1 week ago

I got this question (12-Aug-2023) and I chose the given answers - 932 passed

upvoted 3 times

✉  **BaoNguyen2411** 8 months, 3 weeks ago

got this question on 29/06/2023

upvoted 3 times

✉  **AbidooKing** 9 months, 1 week ago

Got this 2023July13. Went with given and scored 917.

upvoted 3 times

✉  **JH81** 9 months, 3 weeks ago

I got this on 6/29/2023 and passed with 850. Answer is correct.

upvoted 2 times

✉  **Indhumathi_Saravanan** 1 year ago

Question was on 19-MAR-2023 exam

upvoted 2 times

✉  **seila1331** 1 year, 4 months ago

Did my exam on 8/12/2022. This question was on it.

upvoted 4 times

✉  **TheExamMaster2020** 1 year, 5 months ago

Did my exam on 15th November 2022. This question was on it.

upvoted 3 times

✉  **at_sayali** 1 year, 5 months ago

Got this in 11/2022

Went with on-error, context, context, set-status, override-status, on-error

One more blank space after <otherwise />, I selected override-status

upvoted 5 times

✉  **Ciupaz** 3 months ago

Override-status where?

upvoted 1 times

✉  **elequel** 1 year, 6 months ago

Got it in exam 20/10/2022

upvoted 4 times

Question #19

DRAG DROP -

You are a developer for a Software as a Service (SaaS) company. You develop solutions that provide the ability to send notifications by using Azure Notification Hubs.

You need to create sample code that customers can use as a reference for how to send raw notifications to Windows Push Notification Services (WNS) devices.

The sample code must not use external packages.

How should you complete the code segment? To answer, drag the appropriate code segments to the correct locations. Each code segment may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Code segments

- raw
- windows
- windowsphone
- application/xml
- application/json
- application/octet-stream

Answer Area

```
var endpoint = "...";
var payload = "...";
var request = new HttpRequestMessage(HttpMethod.Post, endpoint);
request.Headers.Add("X-WNS-Type", "wns/raw");
request.Headers.Add("ServiceBusNotification-Format", "Code segment");
request.Content = new StringContent(payload, Encoding.UTF8, "Code segment");
var client = new HttpClient();
await client.SendAsync(request);
```

Correct Answer:

Code segments

- raw
- windows
- windowsphone
- application/xml
- application/json

Answer Area

```
var endpoint = "...";
var payload = "...";
var request = new HttpRequestMessage(HttpMethod.Post, endpoint);
request.Headers.Add("X-WNS-Type", "wns/raw");
request.Headers.Add("ServiceBusNotification-Format", "windows");
request.Content = new StringContent(payload, Encoding.UTF8, "application/octet-stream");
var client = new HttpClient();
await client.SendAsync(request);
```

Box 1: windows -

Example code:

```
var request = new HttpRequestMessage(method, $"{resourceUri}?api-version=2017-04"); request.Headers.Add("Authorization", createToken(resourceUri, KEY_NAME, KEY_VALUE));
request.Headers.Add("X-WNS-Type", "wns/raw");
request.Headers.Add("ServiceBusNotification-Format", "windows"); return request;
```

Box 2: application/octet-stream -

Example code capable of sending a raw notification:

```
string resourceUri = $"https://{{NH_NAMESPACE}}.servicebus.windows.net/{{HUB_NAME}}/messages/"; using (var request = CreateHttpRequest(HttpMethod.Post, resourceUri))
{
    request.Content = new StringContent(content, Encoding.UTF8, "application/octet-stream");
    request.Content.Headers.ContentType.CharSet = string.Empty;
    var httpClient = new HttpClient();
    var response = await httpClient.SendAsync(request);
```

```
Console.WriteLine(response.StatusCode);  
}  
  
Reference:  
https://stackoverflow.com/questions/31346714/how-to-send-raw-notification-to-azure-notification-hub/31347901
```

✉  **markra**  2 years, 10 months ago

Answer is correct.

<https://docs.microsoft.com/en-us/rest/api/notificationhubs/send-wns-native-notification>

upvoted 46 times

✉  **lugospod**  2 years, 3 months ago

Got this one 01/2022. Went with most voted (to avoid writing answers again)

upvoted 11 times

✉  **arunkuml**  4 months ago

Got it in the exam 14/12/23. Scored 912/1000. All questions are from ExamTopics. Case study - VanArsdel, Ltd (11 questions)

upvoted 2 times

✉  **AndySmith** 5 months, 2 weeks ago

On exam 3-Nov-2023. Went with proposed answer - 932/1000.

upvoted 2 times

✉  **BaoNguyen2411** 8 months, 3 weeks ago

got this question on 29/06/2023

upvoted 3 times

✉  **juanckar** 9 months, 2 weeks ago

This was on the exam (July 2023). Went with proposed. Scored 917

upvoted 3 times

✉  **Indhumathi_Saravanan** 1 year ago

Got this question - 19-Mar-2023

upvoted 1 times

✉  **at_sayali** 1 year, 5 months ago

Answer is correct.

Got this question in 10/2022

upvoted 2 times

✉  **gmishra88** 1 year, 6 months ago

Notification hub and logic apps are removed from scope

upvoted 4 times

✉  **Mev4953** 2 years, 3 months ago

about second box:

Set to application/json; charset=utf-8 or application/xml.

If the notification type (X-WNS-Type) is wns/raw, set to application/octet-stream

upvoted 8 times

Question #20

Topic 6

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen. You are developing an Azure solution to collect point-of-sale (POS) device data from 2,000 stores located throughout the world. A single device can produce

2 megabytes (MB) of data every 24 hours. Each store location has one to five devices that send data.

You must store the device data in Azure Blob storage. Device data must be correlated based on a device identifier. Additional stores are expected to open in the future.

You need to implement a solution to receive the device data.

Solution: Provision an Azure Event Hub. Configure the machine identifier as the partition key and enable capture.

Does the solution meet the goal?

A. Yes

B. No

Correct Answer: A

Reference:

<https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-programming-guide>

Community vote distribution

A (50%)

B (50%)

✉  **edengoforit**  2 years, 2 months ago

When the question comes in the exam, then we need to pray to God

upvoted 64 times

✉  **huislaw** 2 years, 1 month ago

Let's pray the question doesn't appear.

upvoted 11 times

✉  **jay158**  2 years, 10 months ago

Answer is NO

How many partitions event hub can have? NOT MORE than 1024

<https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-quotas>

upvoted 28 times

✉  **vovap0vovap** 3 weeks, 2 days ago

Why do you need partition for every store or terminal?

upvoted 1 times

✉  **DefaultName2** 2 years, 5 months ago

Answer is YES :)

Why do you want to have so many partitions? For me, it would be pretty strange to have so many concurrent groups.

In the Event HUB output group you can have whether you want a Service bus, account storage, stream analytics (to aggregate data) and more

...

upvoted 9 times

✉  **hromadkam** 1 year, 8 months ago

I have to agree. You don't need as many partitions as there are devices. Partition key is used in function to determine partition (like hashing function). So you could have 2000 devices and for example 10 partitions.

<https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-features#publishing-an-event>

upvoted 2 times

✉  **hromadkam** 1 year, 8 months ago

So the answer is YES

upvoted 1 times

✉  **troy89** 2 years, 1 month ago

But when you use Event Hubs cluster then you can purchase up to 20 Capacity Units and every Capacity Unit can have 2000 partitions.

upvoted 2 times

✉  **AndresCH** 2 years, 6 months ago

"does not reach goal". See topic5 / question 5

upvoted 2 times

✉ **Mev4953** 2 years, 3 months ago

Question 5 is not the same. Asking about "Provision an Azure Notification Hub. Register all devices with the hub."

upvoted 1 times

✉ **oskx2** **Most Recent** 3 weeks, 2 days ago

The recommendation from Microsoft: <https://learn.microsoft.com/en-us/azure/event-hubs/event-hubs-features#mapping-of-events-to-partitions>

"A per-device or user unique identity makes a good partition key, but other attributes such as geography can also be used to group related events into a single partition."

The correct answer is YES.

upvoted 1 times

✉ **FeriAZ** 2 months, 2 weeks ago

Selected Answer: A

The solution scales well to accommodate the current load and future expansions as more stores open.

Using the device identifier as the partition key ensures that data correlation is straightforward and efficient.

By enabling the Capture feature, the data is automatically saved to Azure Blob storage, meeting the storage requirement.

Event Hubs has a feature called Capture, which automatically captures the streaming data and saves it to a storage account (Azure Blob storage or Azure Data Lake Storage), without requiring any additional coding or services to process the data. This directly aligns with the requirement to store the device data in Azure Blob storage.

upvoted 1 times

✉ **zeugo** 4 months ago

The theoretical partition limit for a Eventhub is 1024, and it can hold up to 1000 event hubs per namespace. It's totally fine to scale out within a namespace/scope to add more partitions to handle incoming data. plz see the different tiers:

<https://learn.microsoft.com/en-us/azure/event-hubs/event-hubs-quotas>

and read about namespaces: <https://learn.microsoft.com/en-us/azure/event-hubs/event-hubs-features>

upvoted 1 times

✉ **zeugo** 4 months ago

So the answer is yes

upvoted 1 times

✉ **dom271219** 5 months, 3 weeks ago

Selected Answer: B

Azure Event Hub is primarily designed for handling high-throughput, real-time event streaming, such as telemetry data and event-driven architectures. It's not the best choice for storing large volumes of static data like 2 MB of POS device data per device per day

upvoted 2 times

✉ **nardk** 7 months, 1 week ago

Selected Answer: B

Jay158 is right

upvoted 1 times

✉ **Dianahu** 8 months, 3 weeks ago

<https://www.examtopics.com/discussions/microsoft/view/22503-exam-az-204-topic-6-question-2-discussion/> is not the solution

<https://www.examtopics.com/discussions/microsoft/view/25308-exam-az-204-topic-6-question-29-discussion/> also not the solution so it must be the hub

upvoted 1 times

✉ **EliteAllen** 8 months, 3 weeks ago

Selected Answer: A

Azure Event Hub is suitable for the given scenario because it is designed to handle large-scale data ingestion from numerous devices, allowing efficient correlation based on device identifiers, and can automatically capture and store the data in Azure Blob storage.

upvoted 1 times

✉ **adilkhan** 1 year, 1 month ago

Yes is correct

upvoted 1 times

✉ **tmuras** 1 year, 2 months ago

This question was on exam on 17/02/23.

Answered Yes. Score 860

upvoted 4 times

✉ **Sulzirsha** 1 year, 2 months ago

Partition key is consistently used to assign related events to a partition. Partition key doesn't create unique partitions. So answer is 'Yes'.

upvoted 3 times

✉ **Twiddling5623** 1 year, 2 months ago

Selected Answer: B

Answer is No. According to the solution each machine identifier is going to be a partition key and with 1 to 5 machines this will create at least 2000 partitions. The max is 1024 for the biggest price tier.

Ref: <https://learn.microsoft.com/en-us/azure/event-hubs/event-hubs-quotas#basic-vs-standard-vs-premium-vs-dedicated-tiers>

upvoted 1 times

 **Idenouter** 1 year, 3 months ago

Selected Answer: A

1000 Partition Key doesn't lead to 1000 partitions.

upvoted 3 times

 **charthouse** 1 year, 3 months ago

Is there a difference between device identifier (Solution) and machine identifier (Requirement)?

upvoted 1 times

 **edmond** 1 year, 4 months ago

Selected Answer: A

1000 Partition Key doesn't lead to 1000 partitions. As hromadkam mentioned, it's simply something like hashing.

upvoted 1 times

 **hromadkam** 1 year, 8 months ago

The answer is YES.

You don't need as many partitions as there are devices. Partition key is used in function to determine partition (like hashing function). So you could have 2000 devices and for example 10 partitions.

<https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-features#publishing-an-event>

upvoted 1 times

 **FatSeal** 1 year, 7 months ago

the solution says "Configure the machine identifier as the partition key and enable capture." so you can't really do it your way.

upvoted 1 times

 **Marek79** 1 year, 5 months ago

The answer is YES (A).

You can use a partition key to map incoming event data into specific partitions for the purpose of data organization. The partition key is a sender-supplied value passed into an event hub. It is processed through a static hashing function, which creates the partition assignment.

If you don't specify a partition key when publishing an event, a round-robin assignment is used.

<https://learn.microsoft.com/en-us/azure/event-hubs/event-hubs-features#mapping-of-events-to-partitions>

upvoted 1 times

Question #21

DRAG DROP -

You are developing an Azure solution to collect inventory data from thousands of stores located around the world. Each store location will send the inventory data hourly to an Azure Blob storage account for processing.

The solution must meet the following requirements:

- ⇒ Begin processing when data is saved to Azure Blob storage.
- ⇒ Filter data based on store location information.
- ⇒ Trigger an Azure Logic App to process the data for output to Azure Cosmos DB.
- ⇒ Enable high availability and geographic distribution.
- ⇒ Allow 24-hours for retries.
- ⇒ Implement an exponential back off data processing.

You need to configure the solution.

What should you implement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Select and Place:

Technologies

Azure Event Hub

Azure Event Grid

Azure Service Bus

Azure Blob Storage

Azure App Service

Azure Logic App

Answer Area**Object****Technology**

Event Source

Technology

Event Receiver

Technology

Event Handler

Technology

Correct Answer:**Technologies**

Azure Event Hub

Answer Area**Object****Technology**

Event Source

Azure Event Grid

Event Receiver

Azure Logic App

Event Handler

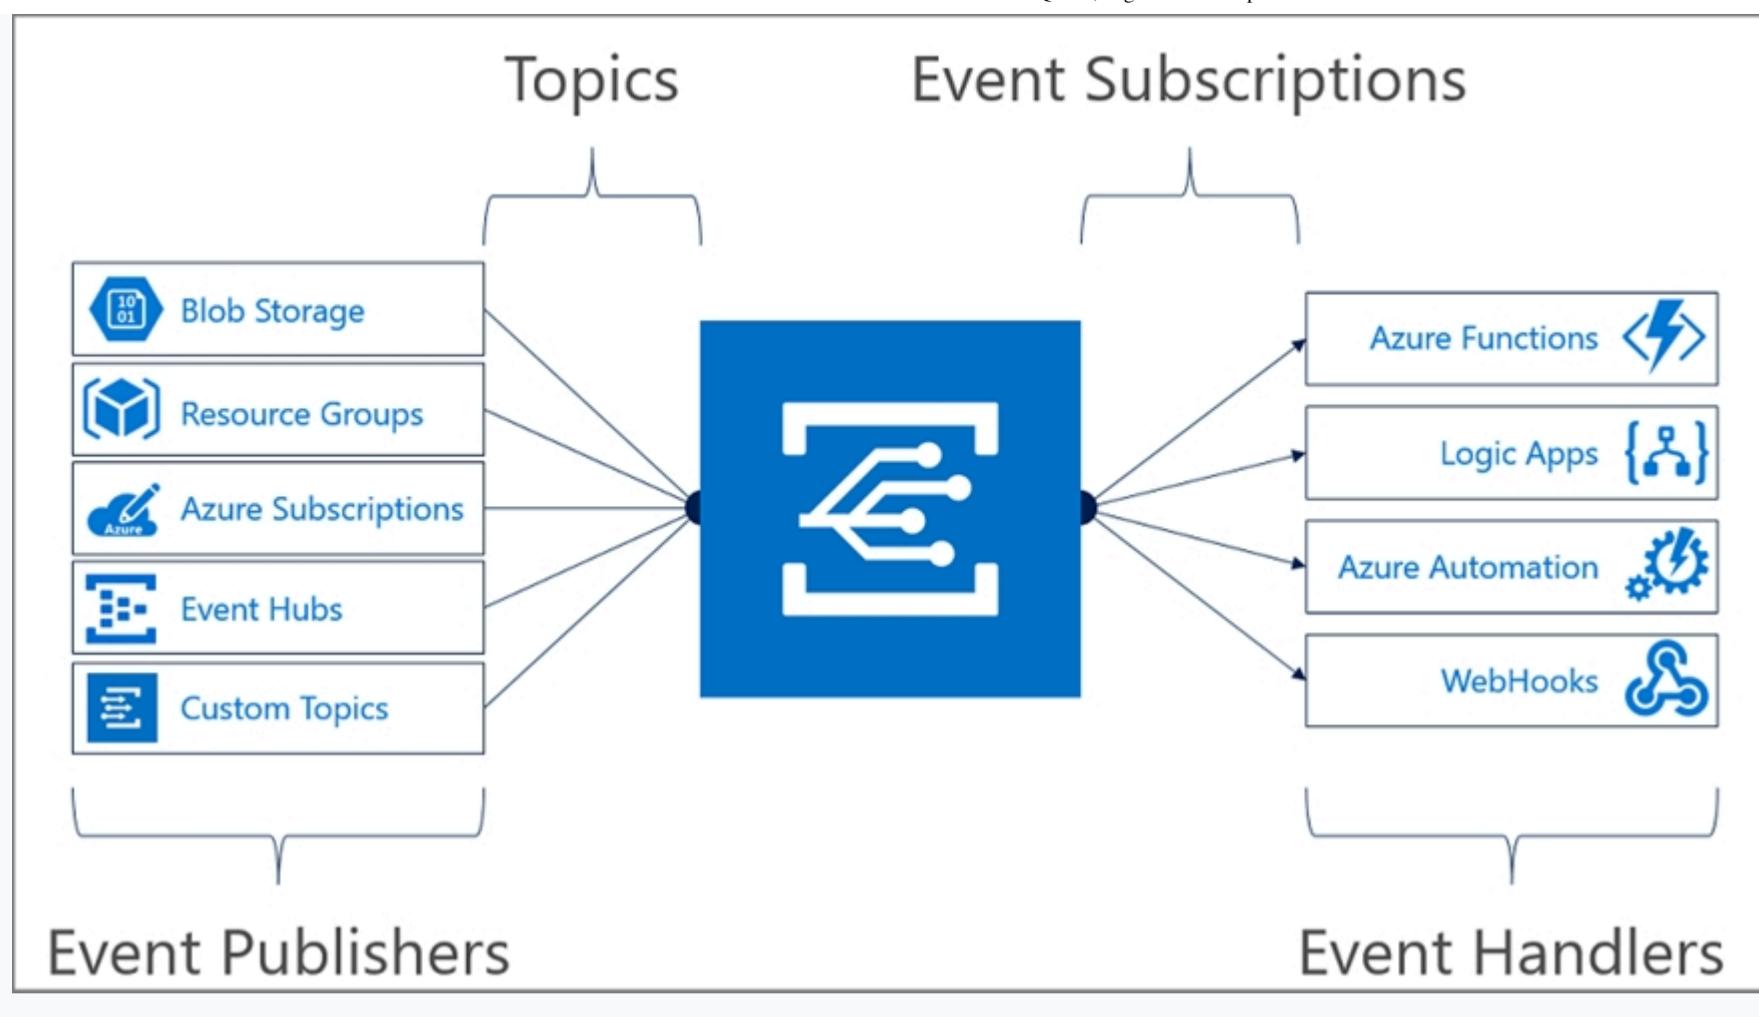
Azure Service Bus

Box 1: Azure Event Grid -

Blob storage events are pushed using Azure Event Grid to subscribers such as Azure Functions, Azure Logic Apps, or even to your own http listener. Event Grid provides reliable event delivery to your applications through rich retry policies and dead-lettering.

Box 2: Azure Logic App -

Event Grid uses event subscriptions to route event messages to subscribers. This image illustrates the relationship between event publishers, event subscriptions, and event handlers.



Event Publishers

Event Handlers

Box 3: Azure Service Bus -

The Event Grid service doesn't store events. Instead, events are stored in the Event Handlers, including ServiceBus, EventHubs, Storage Queue, WebHook endpoint, or many other supported Azure Services.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/blobs/storage-blob-event-overview> <https://docs.microsoft.com/en-us/java/api/overview/azure/messaging-eventgrid-readme>

✉️  **Jurgen1234**  2 years, 9 months ago

Source -> blob storage
Receiver -> event grid
Handler -> logic app
upvoted 121 times

✉️  **unilldreams** 2 years, 4 months ago

Shouldn't receiver be Azure service bus , we are not just sending event but also inventory data (Message type).
<https://docs.microsoft.com/en-us/azure/event-grid/compare-messaging-services#comparison-of-services>

upvoted 1 times

✉️  **Molte** 2 years, 3 months ago

Also "Filter data based on store location information." indicates its a message right? Furthermore ServiceBus has the ability to filter message: <https://docs.microsoft.com/en-us/azure/service-bus-messaging/topic-filters>

upvoted 1 times

✉️  **Mev4953** 2 years, 3 months ago

An event handler is the place where the event is sent. The handler takes some further action to process the event. Several Azure services are automatically configured to handle events and Azure Service Bus is one of them.

<https://docs.microsoft.com/en-us/azure/event-grid/handler-service-bus>

upvoted 2 times

✉️  **jay158** 2 years, 9 months ago

Receiver ??
<https://docs.microsoft.com/en-us/azure/event-grid/overview>
upvoted 6 times

✉️  **PhilLI** 2 years, 3 months ago

Indeed when you look at it from the perspective of the Event Grid the question seems answerable. Your link is good. Receiver is implicitly mentioned here:
"Event sources.
Currently, the following Azure services support sending events to Event Grid."

upvoted 1 times

✉️  **TonyMel** 1 year ago

correct, in 2023Mar24, score: 904/1000.

upvoted 2 times

✉️  **TaoLu** 1 year, 2 months ago

You are 100% correct

upvoted 1 times

 **argoth** Highly Voted 2 years, 9 months ago

I'm always wondering who is the big brain behind the examtopics question solutions...

upvoted 47 times

 **oskx2** Most Recent 3 weeks, 2 days ago

The answer is wrong.

From microsoft: <https://learn.microsoft.com/en-us/azure/event-grid/event-handlers>

The supported event handlers is logic app.

And <https://azure.microsoft.com/en-us/products/event-grid>

"Event Grid connects data sources and event handlers."

So, event grid is the receiver and logic app is the event handler. By elimination, az storage is the event source.

upvoted 1 times

 **arunkuml** 4 months ago

Got it in the exam 14/12/23. Scored 912/1000. All questions are from ExamTopics. Case study - VanArsdel, Ltd (11 questions)

upvoted 4 times

 **JH81** 9 months, 3 weeks ago

I got this on 6/29/2023 and passed with 850. Went with highly voted answer.

upvoted 4 times

 **aragones** 11 months, 1 week ago

Got this 2023-05-12.

FYI:

case: You need to configure authorization.

case: You need to ensure the app does not time out and processes the blob data

case study: VanArsdel Inc Canada

upvoted 1 times

 **Nhiendo** 1 year, 3 months ago

I thiinh blob storage,event grid, logic app

upvoted 1 times

 **meoukg** 2 years, 1 month ago

Got it on 03/2022, chose Source: Blob storage, Receiver: Event Grid, Handler: Logic App.

upvoted 7 times

 **vruizrob** 1 year, 5 months ago

And your response was ok?

upvoted 2 times

 **MikeAWS** 6 months ago

how can he knows that? when you pass/fail an exam, all you get is total number of points and some statistics (in %) with chapters!

upvoted 3 times

 **edengoforit** 2 years, 3 months ago

Event Grid is a fully-managed event routing service and the first of its kind. Azure Event Grid greatly simplifies the development of event-based applications and simplifies the creation of serverless workflows. Using a single service, Azure Event Grid manages all routing of events from any source, to any destination, for any application.

upvoted 1 times

 **d_programmer** 2 years, 4 months ago

Receiver should be event hub as the message needs to be retained for 24 hours

upvoted 2 times

 **koolexam** 2 years, 4 months ago

Event Grid -> Reliability - 24-hour retry with exponential backoff to make sure events are delivered.

<https://docs.microsoft.com/en-us/azure/event-grid/overview#capabilities>

upvoted 3 times

 **jvyas** 2 years, 5 months ago

I think the given answer tries to meet all the requirements solution must meet in your responses.

upvoted 1 times

 **Franz22** 2 years, 6 months ago

I think that the right solution should be:

Source -> blob storage

Receiver -> event grid

Handler -> logic app

Because:

- 1)the source of the event, which means WHO has created the event and who is publishing this event, is the Blob Storage
- 2)The receiver of this event, which means who should receive this event, is the Event Grid. In the event Grid we can create Topic. Topic is where all publishers write messages --> so the Blob Storage write into a topic which is inside an Event Grid.
- 3)Handler of this event, which means who should handle this event, is the Logic App by creating a subscription in the event grid and connecting the Logic App to that subscription.

upvoted 18 times

✉ **catalene** 2 years, 6 months ago

Why not "Event Hub" like Event Receiver?

upvoted 4 times

✉ **ning** 2 years, 8 months ago

No idea what is event receiver? No where in documentation mentioned that!

Source --> Blob Storage, Handler --> Logic App

upvoted 1 times

✉ **ranjitklive** 2 years, 8 months ago

<https://www.cognizantsoftvision.com/blog/azure-event-grid-vs-event-hubs/>

upvoted 3 times

✉ **7ack** 2 years, 9 months ago

According to the diagram, shouldn't the source be Blob storage, the receiver Event Grid, and the handler Logic app?

upvoted 16 times

Question #22

You are creating an app that will use CosmosDB for data storage. The app will process batches of relational data.

You need to select an API for the app.

Which API should you use?

- A. MongoDB API
- B. Table API
- C. SQL API
- D. Cassandra API

Correct Answer: *Incorrect Answer:*

For relational data you will need the SQL API

A: The MongoDB API is not used for relational data.

B: The Table API only supports data in the key/value format

D: The Cassandra API only supports OLTP (Online Transactional Processing) and not batch processing.

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/choose-api>

Community vote distribution

C (100%)

✉  **hicloud**  7 months, 4 weeks ago

Finally a question I can answer without any knowledge of Azure!

upvoted 8 times

✉  **DongKG**  1 year, 11 months ago

Selected Answer: C

Because the app will process batches of relational data

upvoted 6 times

✉  **Ciupaz**  6 months, 3 weeks ago

Selected Answer: C

Why not Table API?

The Table API in Azure Cosmos DB is specifically designed for storing semi-structured or NoSQL data in a schema-less fashion, not for storing traditional relational data. Here are some reasons why the Table API is not a good choice for handling relational data.

upvoted 1 times

✉  **escame** 1 year, 2 months ago

Selected Answer: C

C is correct

upvoted 3 times

✉  **mark_1234** 1 year, 2 months ago

Selected Answer: C

SQL is the only relational data storage type there.

upvoted 1 times

✉  **OPT_001122** 1 year, 5 months ago

Selected Answer: C

C - correct answer

upvoted 3 times

✉  **Chief2110** 1 year, 7 months ago

Selected Answer: C

C should be the answer

upvoted 3 times

✉  **Frefen** 1 year, 7 months ago

Selected Answer: C

Relational data

upvoted 4 times

Question #23

HOTSPOT -

You are developing a .NET application that communicates with Azure Storage.

A message must be stored when the application initializes.

You need to implement the message.

How should you complete the code segment? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
CloudStorageAccount storageAccount = CloudStorageAccount.Parse
(CloudConfigurationManager.GetSetting ("StorageConnectionString"));

pVar1 = storageAccount. ( ) ;

CloudQueueClient
CloudTableClient
CloudQueue
CloudTable

pVar2 = pVar1. ( "contoso-storage" );

CloudQueueClient
CloudTableClient
CloudQueue
CloudTable

try
{
    await pVar2.CreateIfNotExistsAsync();
}
catch (StorageException x)
{
    throw;
}
CloudQueueMessage cloudQueueMessage = new CloudQueueMessage("App Launch: <iUserID>");
await pVar2.AddMessageAsync(cloudQueueMessage);
```

Correct Answer:

Answer Area

```

CloudStorageAccount storageAccount = CloudStorageAccount.Parse
(CloudConfigurationManager.GetSetting ("StorageConnectionString"));

pVar1 = storageAccount. ▼ (); ▼
CloudQueueClient
CloudTableClient
CloudQueue
CloudTable

pVar2 = pVar1. ▼ (); ▼ ("contoso-storage");
CloudQueueClient
CloudTableClient
CloudQueue
CloudTable
CreateCloudQueueClient
CreateCloudTableClient
GetQueueReference
GetTableReference

try
{
    await pVar2.CreateIfNotExistsAsync();
}
catch (StorageException x)
{
    throw;
}
CloudQueueMessage cloudQueueMessage = new CloudQueueMessage ("App Launch: <iUserID>");
await pVar2.AddMessageAsync (cloudQueueMessage);

```

Reference:

<https://docs.microsoft.com/en-us/azure/storage/queues/storage-dotnet-how-to-use-queues?tabs=dotnetv11>

✉ **BogdanG** Highly Voted 1 year, 11 months ago

Correct, just that this is deprecated

For .NET V12 SDK is as follows :

```

string connectionString = ConfigurationManager.AppSettings["StorageConnectionString"];
QueueClient queueClient = new QueueClient(connectionString, queueName);
queueClient.CreateIfNotExists();

```

upvoted 16 times

✉ **juanckar** Highly Voted 9 months, 2 weeks ago

This was on the exam (July 2023). Went with proposed. Scored 917

upvoted 8 times

✉ **arunkuml** Most Recent 4 months ago

Got it in the exam 14/12/23. Scored 912/1000. All questions are from ExamTopics. Case study - VanArsdel, Ltd (11 questions)

upvoted 2 times

✉ **huislaw** 1 year, 11 months ago

Answer is correct:

1. Create the queue client
2. Retrieve a reference to a queue

upvoted 5 times

Question #24

Topic 6

HOTSPOT -

A software as a service (SaaS) company provides document management services. The company has a service that consists of several Azure web apps. All

Azure web apps run in an Azure App Service Plan named PrimaryASP.

You are developing a new web service by using a web app named ExcelParser. The web app contains a third-party library for processing Microsoft Excel files.

The license for the third-party library stipulates that you can only run a single instance of the library.

You need to configure the service.

How should you complete the script? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
Set-AzAppServicePlan ` 
  -ResourceGroupName $rg ` 
  -Name "PrimaryASP" `
```

NumberOfSites 1
PerSiteScaling \$true
TargetWorkerCount = 1
MaxNumberOfWorkers = 1
SiteConfig.NumberOfWorkers = 1

```
$app = Get-AzWebApp ` 
  -ResourceGroupName $rg ` 
  -Name "ExcelParser"
```

```
$app. 

|                                |
|--------------------------------|
| NumberOfSites 1                |
| PerSiteScaling \$true          |
| TargetWorkerCount = 1          |
| MaxNumberOfWorkers = 1         |
| SiteConfig.NumberOfWorkers = 1 |


```

```
Set-AzWebApp $app
```

Answer Area

```
Set-AzAppServicePlan ` 
  -ResourceGroupName $rg ` 
  -Name "PrimaryASP" `
```

NumberOfSites 1	▼
PerSiteScaling \$true	▼
TargetWorkerCount = 1	▼
MaxNumberOfWorkers = 1	▼
SiteConfig.NumberOfWorkers = 1	▼

Correct Answer:

```
$app = Get-AzWebApp ` 
  -ResourceGroupName $rg ` 
  -Name "ExcelParser"
```

\$app.	▼
NumberOfSites 1	▼
PerSiteScaling \$true	▼
TargetWorkerCount = 1	▼
MaxNumberOfWorkers = 1	▼
SiteConfig.NumberOfWorkers = 1	▼

```
Set-AzWebApp $app
```

Reference:

<https://docs.microsoft.com/en-us/azure/app-service/manage-scale-per-app>

✉  **huislaw**  1 year, 11 months ago

Answer is correct, -PerSiteScaling \$true; \$app.SiteConfig.NumberOfWorkers = 1
upvoted 13 times

✉  **WH16**  8 months, 1 week ago

This was on the exam (9th August 2023). Went with proposed.
upvoted 3 times

✉  **70PineApple** 1 year, 1 month ago

Got this in exam today..20/02/23
upvoted 1 times

✉  **rotimislaw** 1 year, 3 months ago

Why PerSiteScaling \$true is needed?
upvoted 3 times

✉  **Heighte** 1 year ago

i guess not to impact other apps on the service plan
upvoted 3 times

✉  **unraval** 7 months ago

Enable per-app scaling with an existing App Service Plan by passing in the -PerSiteScaling \$true parameter to the Set-AzAppServicePlan cmdlet.
<https://learn.microsoft.com/en-us/azure/app-service/manage-scale-per-app>
upvoted 1 times

✉  **OPT_001122** 1 year, 4 months ago

Correct
PerSiteScaling \$true
\$app.SiteConfig.NumberOfWorkers = 1
upvoted 2 times

 **LGWJ12** 1 year, 5 months ago

yes, the Answer is correct.

upvoted 3 times

Question #25

DRAG DROP -

You have an application that provides weather forecasting data to external partners. You use Azure API Management to publish APIs.

You must change the behavior of the API to meet the following requirements:

- ⇒ Support alternative input parameters
- ⇒ Remove formatting text from responses
- ⇒ Provide additional context to back-end services

Which types of policies should you implement? To answer, drag the policy types to the correct requirements. Each policy type may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Policy types		Answer Area	
		Requirement	Policy type
<input type="checkbox"/> Inbound		Support alternative input parameters.	<input type="checkbox"/> policy type
<input type="checkbox"/> Outbound		Remove formatting text from responses.	<input type="checkbox"/> policy type
<input type="checkbox"/> Backend		Provide additional context to back-end services.	<input type="checkbox"/> policy type

Correct Answer:

Policy types		Answer Area	
		Requirement	Policy type
<input type="checkbox"/> Inbound		Support alternative input parameters.	<input type="checkbox"/> Inbound
<input type="checkbox"/> Outbound		Remove formatting text from responses.	<input type="checkbox"/> Outbound
<input type="checkbox"/> Backend		Provide additional context to back-end services.	<input type="checkbox"/> Inbound

Reference:

<https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-policies> <https://docs.microsoft.com/en-us/azure/api-management/api-management-transformation-policies#forward-context-information-to-the-backend-service>

✉  **avocado0211**  1 year, 11 months ago

Inbound
Outbound
Backend
upvoted 33 times

✉  **g2000** 1 year, 3 months ago

last one is backend
<https://learn.microsoft.com/en-us/azure/api-management/forward-request-policy>
upvoted 4 times

✉  **uffuchsi** 1 year, 1 month ago

Last one has to be Inbound

Below is the sample snippet of the policy to send user context to backend service.

```
<policies>
<inbound>
```

```

<base />
<!-- Forward the name of the product associated with the subscription key in the request to the backend service. -->
<set-query-parameter name="x-product-name" exists-action="override">
<value>@(context.Product.Name)</value>
</set-query-parameter>
<!-- Forward the user id associated with the subscription key in the request as well as the region where the proxy processing the request is
hosted. -->
<set-header name="x-request-context-data" exists-action="override">
<value>@(context.UserId)</value>
<value>@(context.Deployment.Region)</value>
</set-header>
</inbound>

```

The configuration is part of inbound section.

<https://docs.microsoft.com/en-us/azure/api-management/policies/send-request-context-info-to-backend-service>
upvoted 1 times

✉  **uncledana** 9 months ago

Backend - <https://learn.microsoft.com/en-us/azure/api-management/api-management-howto-policies#-understanding-policy-configuration>

upvoted 1 times

✉  **srinz**  1 year, 7 months ago

Based on <https://docs.microsoft.com/en-us/azure/api-management/policies/send-request-context-info-to-backend-service>
Inbound is correct answer for the third question

```

<policies>
<inbound>
<base />
<!-- Forward the name of the product associated with the subscription key in the request to the backend service. -->
<set-query-parameter name="x-product-name" exists-action="override">
<value>@(context.Product.Name)</value>
</set-query-parameter>

```

```

<!-- Forward the user id associated with the subscription key in the request as well as the region where the proxy processing the request is
hosted. -->
<set-header name="x-request-context-data" exists-action="override">
<value>@(context.UserId)</value>
<value>@(context.Deployment.Region)</value>
</set-header>
</inbound>

```

upvoted 8 times

✉  **coffecold** 1 year, 6 months ago

It is adding a parameter and a header in the original request.

<backend>: statements to be applied before the request is forwarded to the backend service go here

<inbound> statements to be applied to the request go here.

So I think it can be both, the place that makes more sense is backend.

upvoted 1 times

✉  **OPT_001122** 1 year, 5 months ago

"statements to be applied before the request is forwarded to the backend service go here"

this is not mentioned anywhere in the question hence the confusion for inbound and backend.

upvoted 1 times

✉  **FeriAZ**  2 months, 2 weeks ago

Provide Additional Context to Back-end Services:

Policy Type: Backend.

Explanation: Backend policies are applied to the request after all inbound policies have been applied and just before the request is forwarded to the backend service. To provide additional context to your backend services, use a backend policy. This could include adding or transforming headers, setting context variables, or other modifications that enrich the request with additional information needed by the backend.

upvoted 1 times

✉  **130nk3r5** 4 months ago

Inbound policies are used to manipulate the request from the client before it is processed by the backend service. This can include supporting alternative input parameters.

Outbound policies are used to manipulate the response from the backend service before it is sent to the client. This can include removing formatting text from responses.

Backend policies are used to manipulate the request after it is processed by the inbound policy but before it is sent to the backend service. This can include providing additional context to back-end services.

So, for sure Inbound, Outbound, Backend

upvoted 2 times

✉  **nekkilodeon** 11 months, 4 weeks ago

The given answer is correct. Inbound, outbound, inbound
upvoted 1 times

✉ **adilkhan** 1 year ago

the provided answers are correct! Inbound
Outbound, Inbound <https://learn.microsoft.com/en-us/azure/api-management/set-backend-service-policy>
upvoted 1 times

✉ **hubekpeter** 1 year, 4 months ago

I think for the third question, you can use both inbound and backend. Context variable is available in both and it should not matter at what stage you'll add it. But in this article, they're using inbound policy. One newer knows how these exams are evaluated, therefore I'm sticting with inbound as an answer. <https://learn.microsoft.com/bs-latn-ba/azure/api-management/policies/send-request-context-info-to-backend-service>
upvoted 1 times

✉ **le129** 1 year, 7 months ago

Answer is correct. <https://docs.microsoft.com/en-us/azure/api-management/policies/send-request-context-info-to-backend-service>
upvoted 2 times

✉ **ivan0590** 1 year, 11 months ago

I'm not 100% sure, but based on the following statement I think that the third option should be backend:
"statements to be applied before the request is forwarded to the backend service go here"

Source:

<https://docs.microsoft.com/en-us/azure/api-management/api-management-howto-policies#-understanding-policy-configuration>

upvoted 4 times

✉ **sghaha** 1 year, 11 months ago

Korean Page:
<https://docs.microsoft.com/ko-kr/azure/api-management/api-management-transformation-policies#forward-context-information-to-the-backend-service>

upvoted 3 times

Get IT Certification

Unlock free, top-quality video courses on ExamTopics with a simple registration. Elevate your learning journey with our expertly curated content. Register now to access a diverse range of educational resources designed for your success. Start learning today with ExamTopics!

[Start Learning for free](#)